Optimization: Better Cost Estimation

R&G Chapter 15

(slides adapted from content by J.Gehrke, J.Shanmugasundaram, and/or C.Koch)

Recap

- 2 dimensions to search along for plans (or more)
  - What is the best access path? (π/σ equivs)
  - What is the best join order? (⋈ equivs)
- Consider the **cost** of each allowable plan.
  - Understanding how each operator's output size relates to its input size makes it possible to accurately estimate the cost of a plan.
  - Simplify the join order problem by exploring the cost of left-deep plans only.

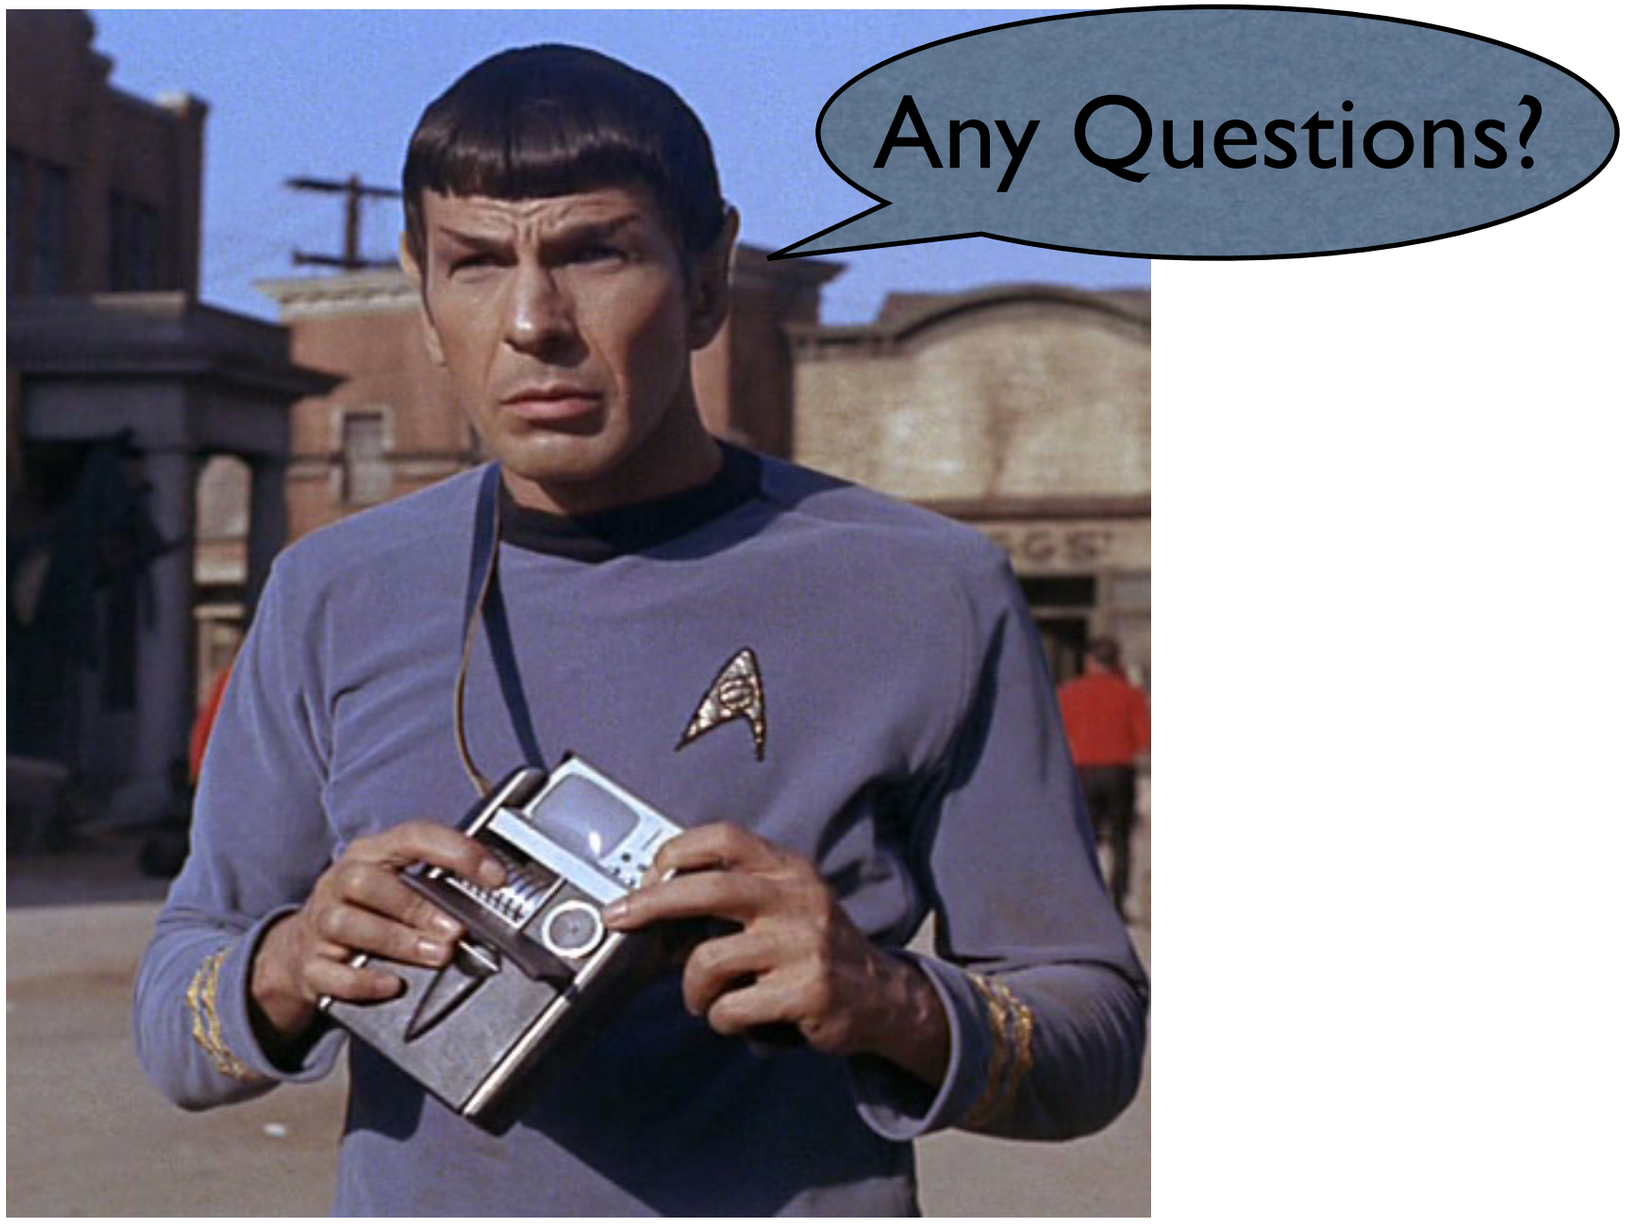

Recap: ISAM

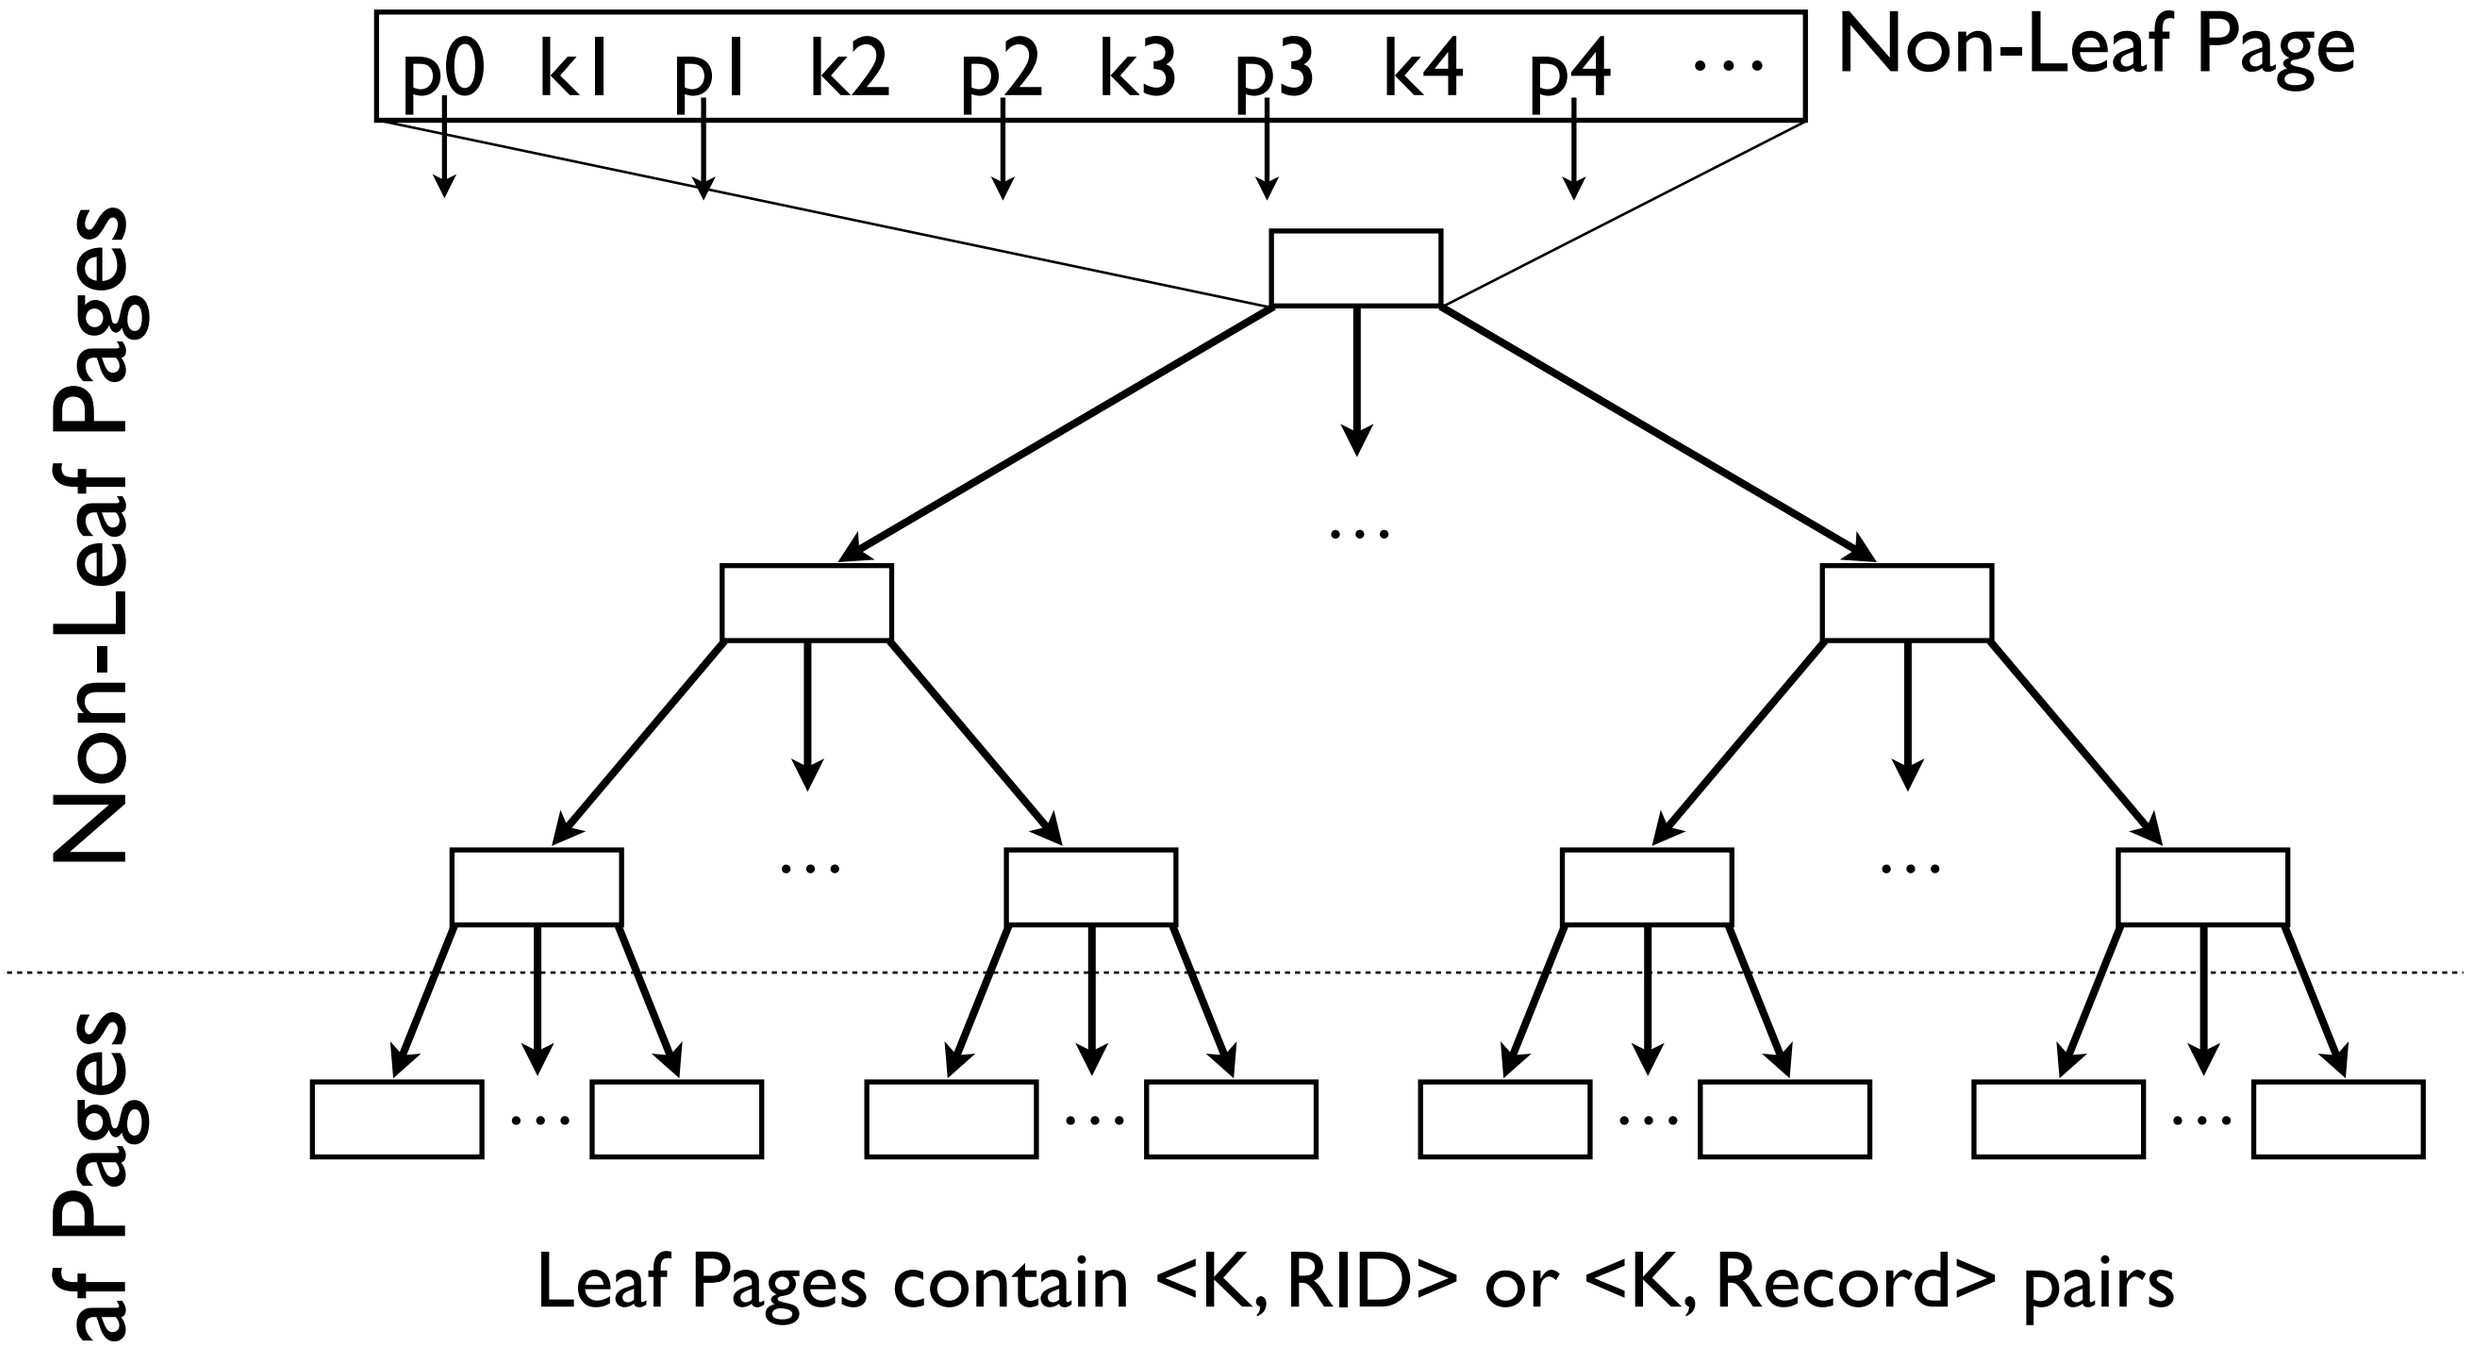

Friday, March 1, 13

ISAM: Index-structured access method

Recap: Static Hashing

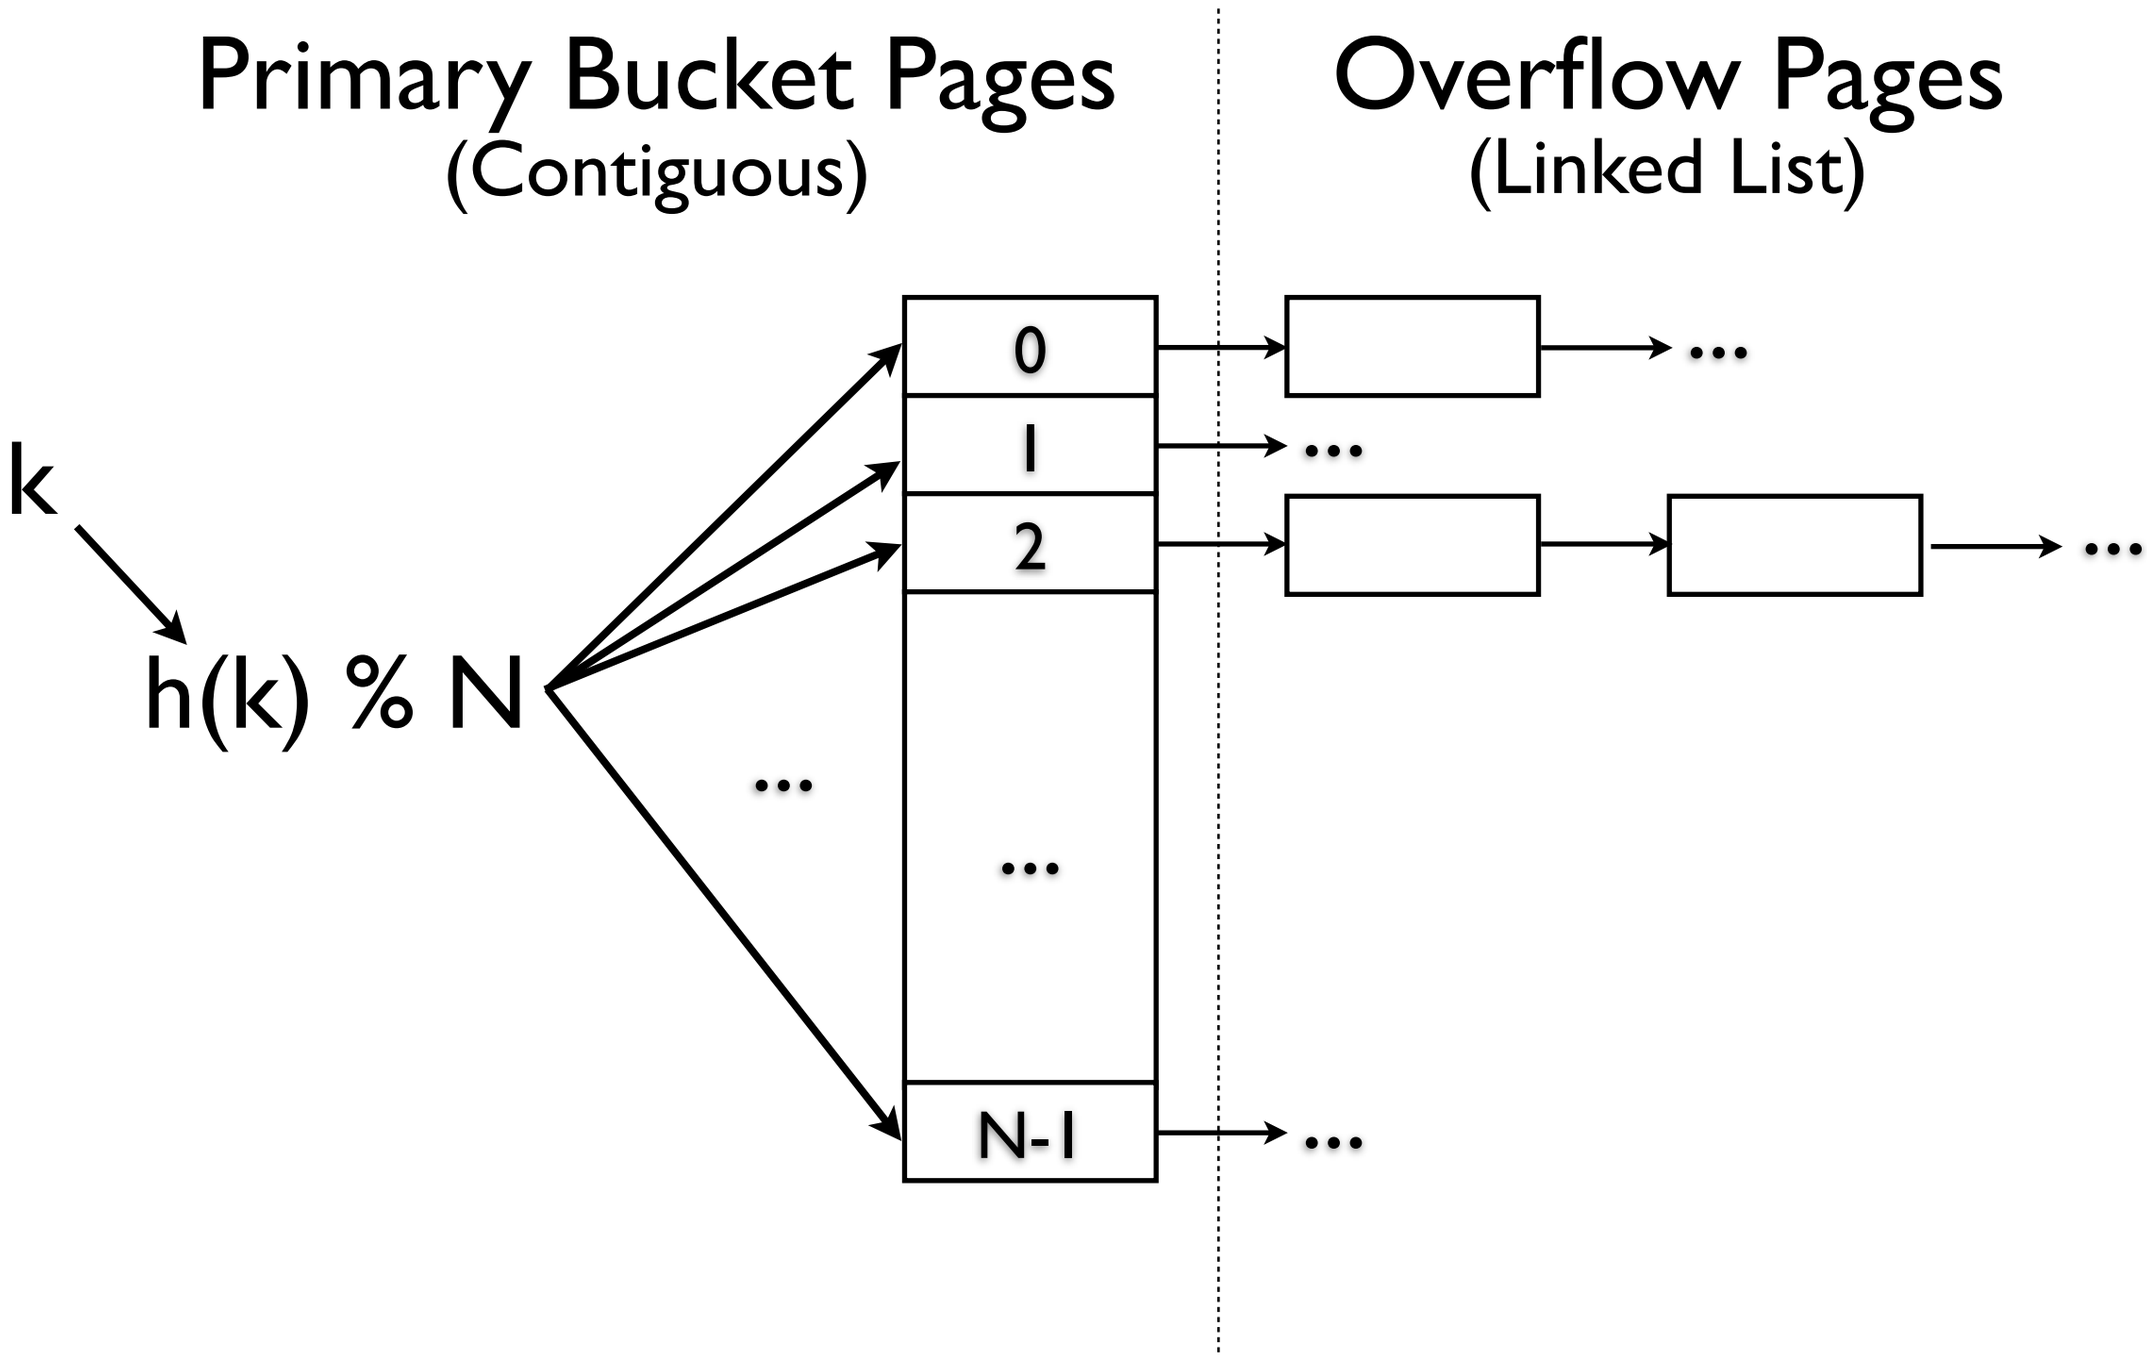

Project 2: sql.Index

Data Stream Generator: `sql.test.TestDataStream`

(`implements Iterator<Datum[]>`)

Index Class Tests Your Indexes

`index -hash hash.index`

Build a hash index in '`hash.index`'

`index -isam isam.index`

Build an ISAM index in '`isam.index`'

Project 2: sql.Index

`index -hash hash.index`

Build a hash index in '`hash.index`'

`index -hash hash.index -get 23`

Use '`hash.index`' to find `<23,...>`

Project 2: sql.Index

`index -isam isam.index`

Build an ISAM index in '`isam.index`'

`index -isam isam.index -scan`

Validate ISAM index scan

```
index -isam isam.index -scan
      -from 5 -to 23
```

Validate ISAM index scan over keys 5 to 23

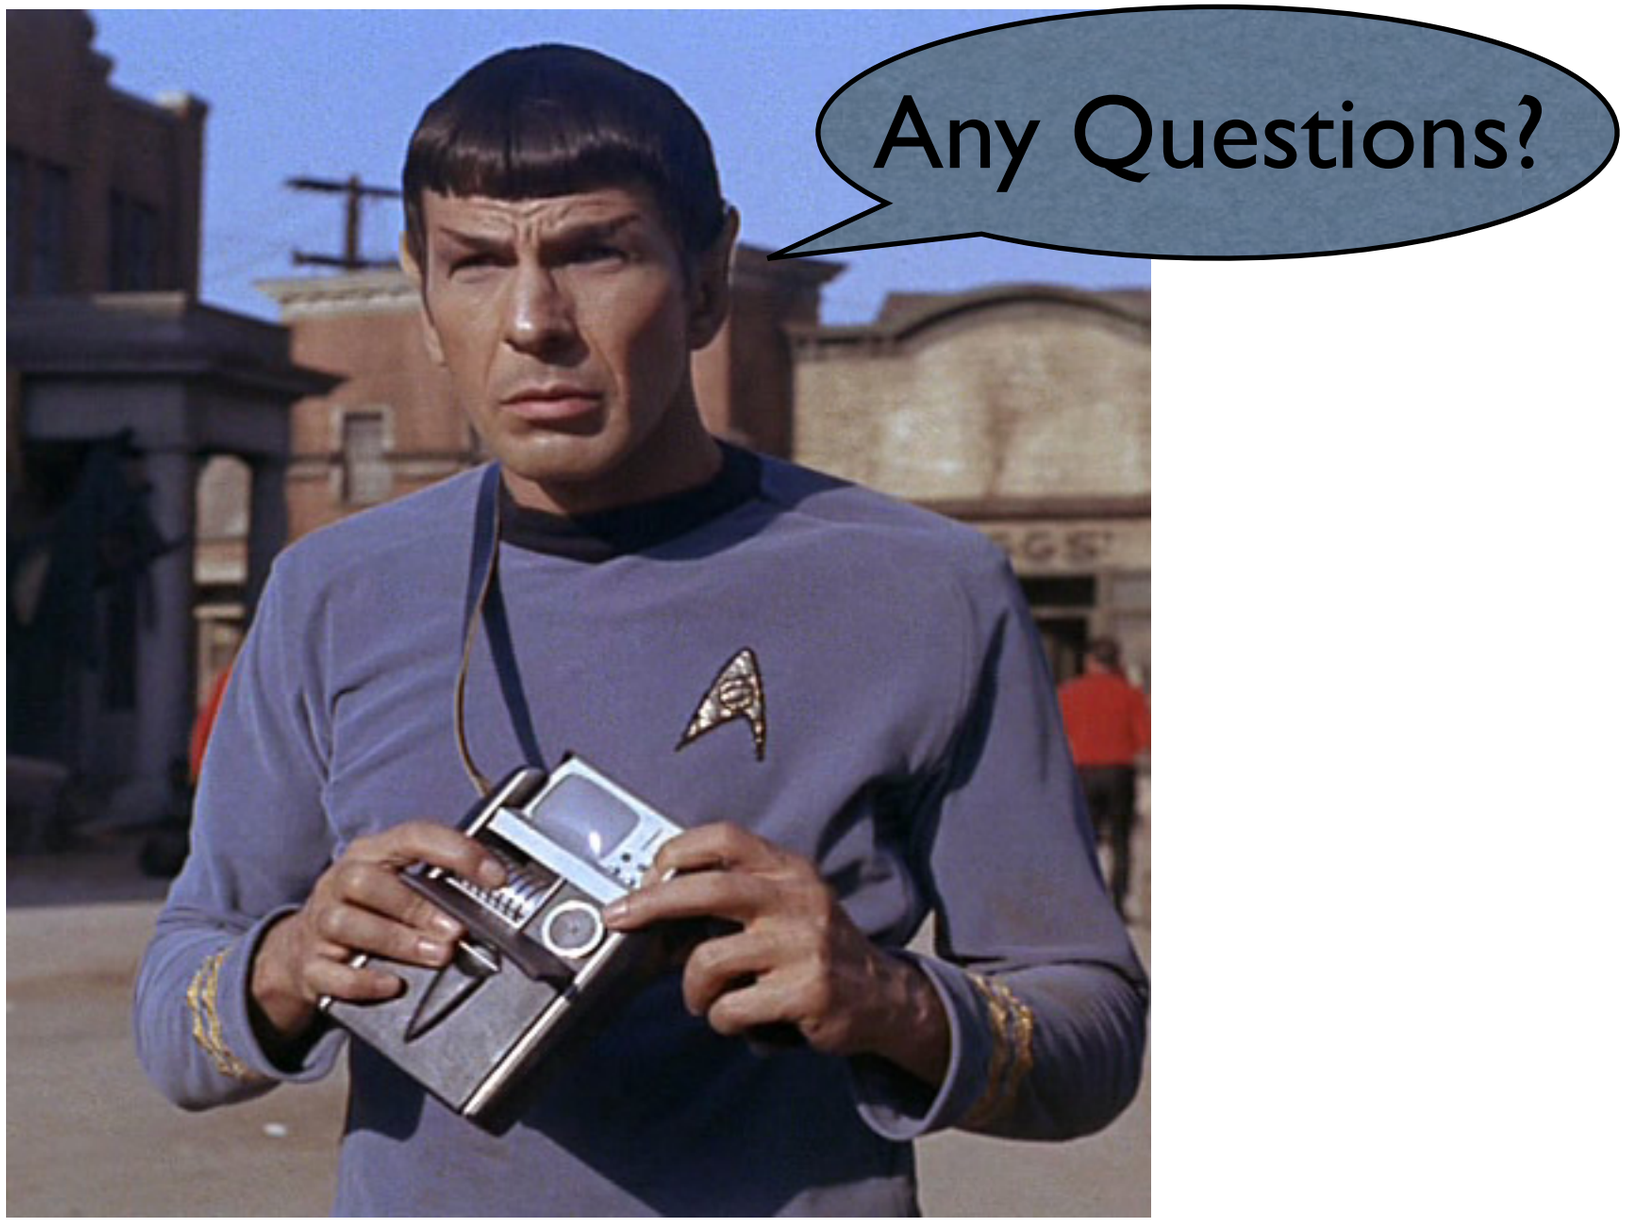

sql.buffer.BufferManager

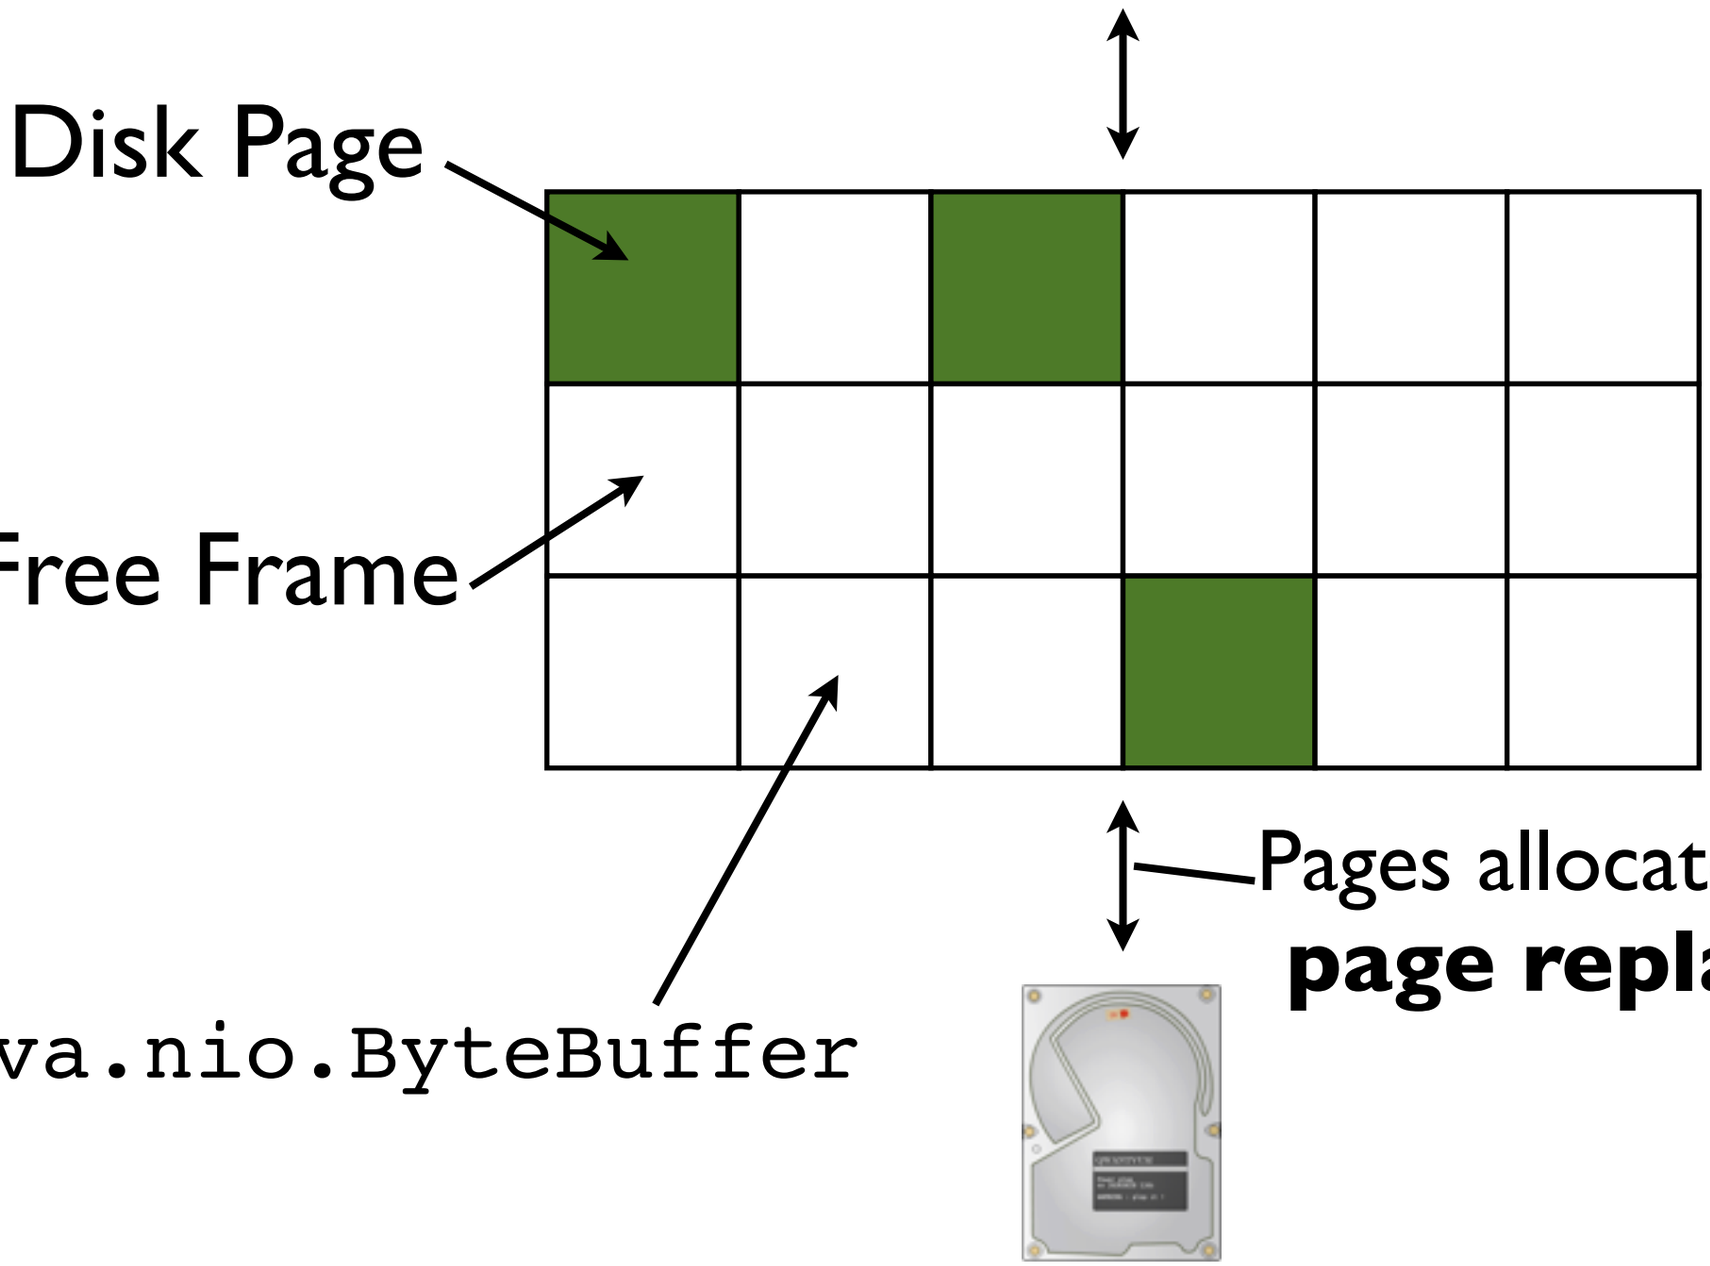

image credit: openclipart.org

Pinned Pages

- Pinning a page indicates that it is being used.
- The requestor must unpin the page when done.
  - The requestor must also indicate whether the page has been modified (with a ‘dirty’ bit)
    - Dirty pages must be written to disk
- Pages may be requested multiple times
  - Use a pin count (reference count) to keep track.

sql.buffer.ManagedFile

- Interface to the Buffer Manager
- Request ManagedFiles through FileManager
- Maps file pages to buffer frames
  - Pages Data In and Out of Memory
- getBuffer(), pin(), unpin(), dirty(), flush()

sql.data.DatumBuffer

- Self-Describing Records
- Wraps around Java's nio.ByteBuffer

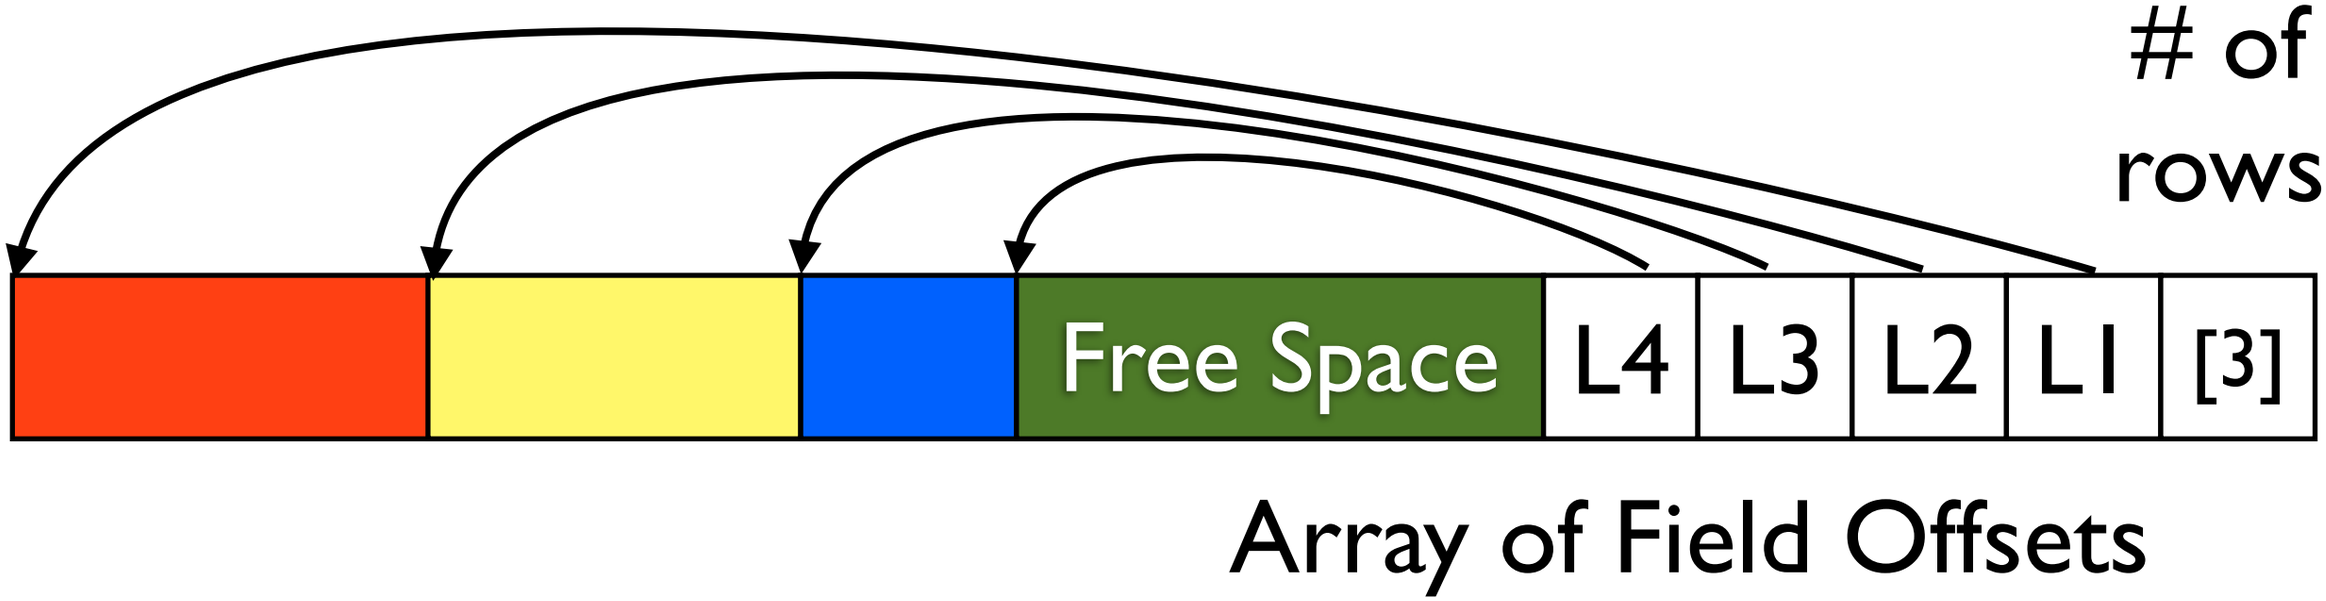

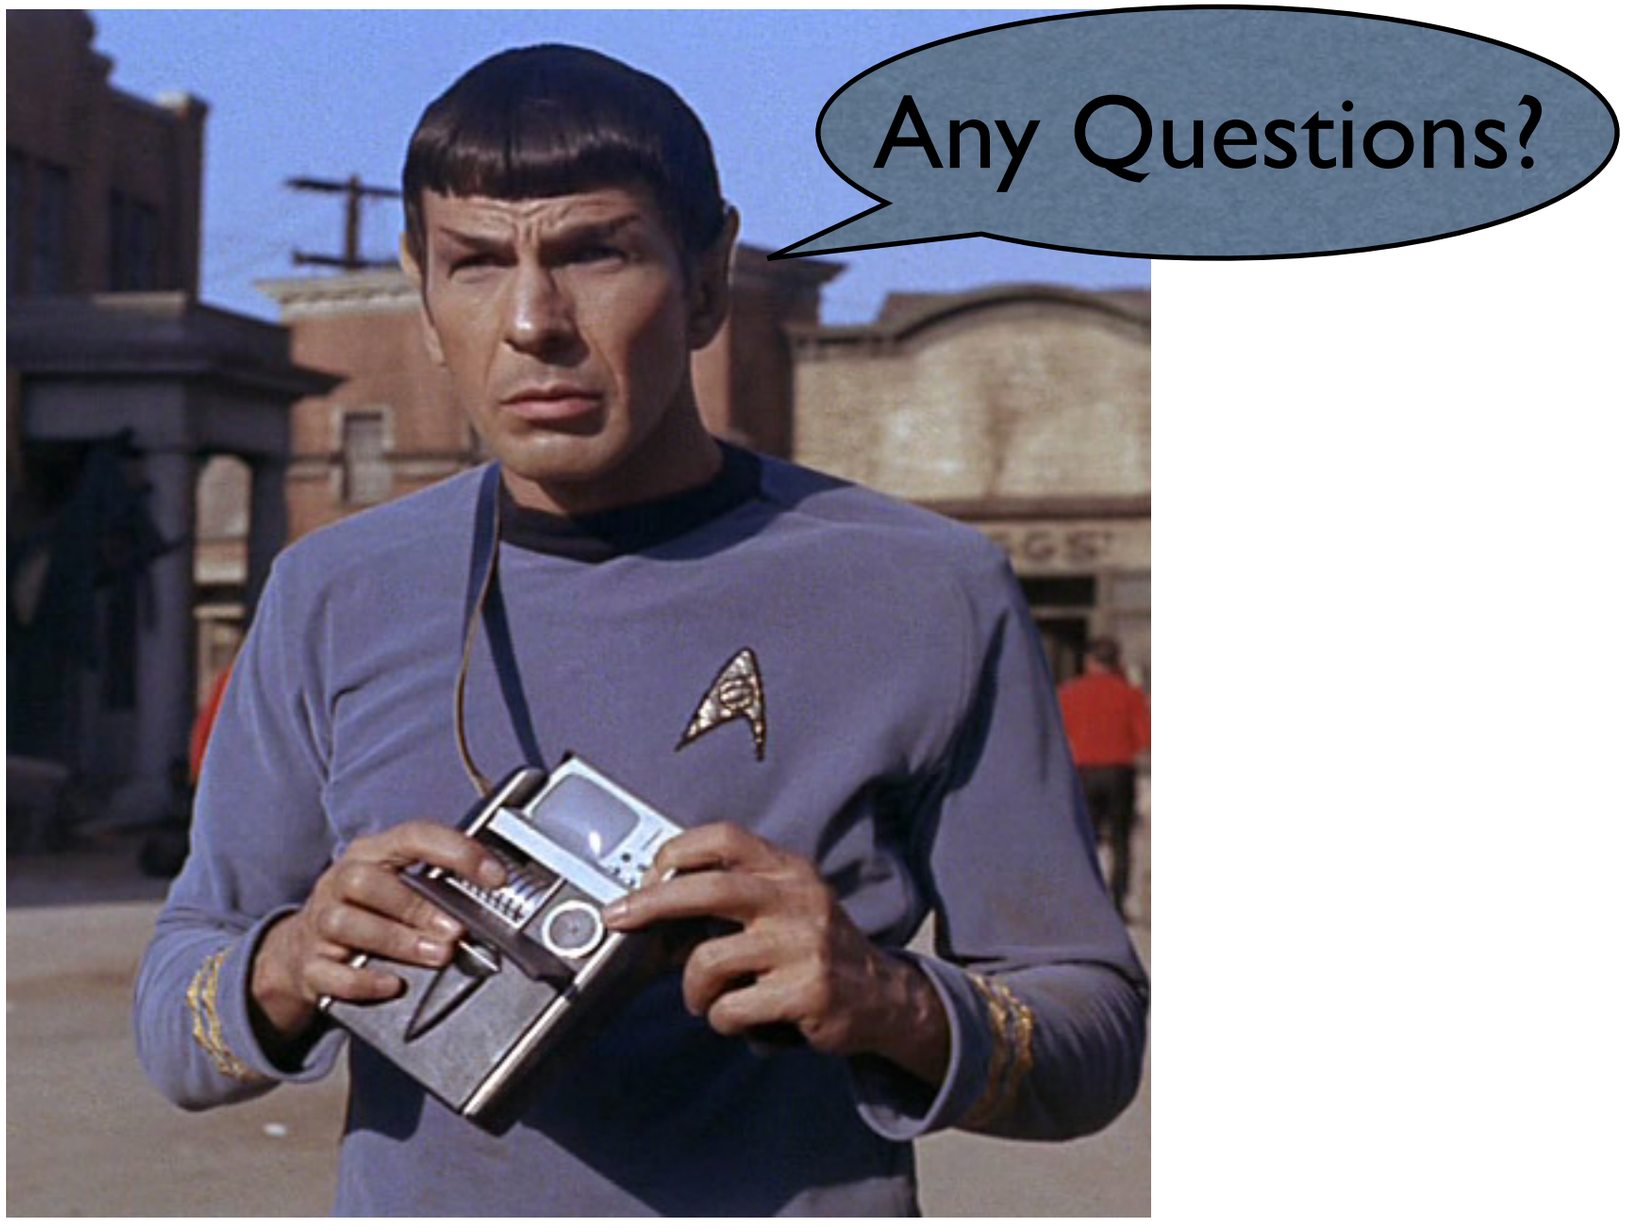

Onwards!

Better Cost Estimation

- What information about data can we use to get better estimates of cost?
  - Data-based statistics
  - Schema-based properties

Histograms

Uniform Distributions are a strong assumption!

(data is often skewed)

Histograms

People

Name	Age	Rank
<“Alice”,	21,	1 >
<“Bob”,	20,	2 >
<“Carol”,	21,	1 >
<“Dave”,	19,	3 >
<“Eve”,	20,	2 >
<“Fred”,	20,	3 >
<“Gwen”,	22,	1 >
<“Harry”,	20,	3 >

```
SELECT Name
FROM People
WHERE Rank = 3
   AND Age = 20
```

vs

```
…
   AND Age = 19
```

$RF_{Age} = 1/nkeys = 1/4$

$RF_{Rank} = 1/nkeys = 3/8$

Age is best!

Histograms

People

Name	Age	Rank
<“Alice”,	21,	1 >
<“Bob”,	20,	2 >
<“Carol”,	21,	1 >
<“Dave”,	19,	3 >
<“Eve”,	20,	2 >
<“Fred”,	20,	3 >
<“Gwen”,	22,	1 >
<“Harry”,	20,	3 >

```
SELECT Name
FROM People
WHERE Rank = 3
   AND Age = 20
```

vs

```
…
   AND Age = 19
```

$RF_{Age\text{-}20} = 1/nkeys = 1/2$

$RF_{Rank} = 1/nkeys = 3/8$

Age is worst!

Histograms

People

Name	Age	Rank
<“Alice”,	21,	1 >
<“Bob”,	20,	2 >
<“Carol”,	21,	1 >
<“Dave”,	19,	3 >
<“Eve”,	20,	2 >
<“Fred”,	20,	3 >
<“Gwen”,	22,	1 >
<“Harry”,	20,	3 >

```
SELECT Name
FROM People
WHERE Rank = 3
   AND Age = 20
```

vs

```
…
   AND Age = 19
```

$RF_{Age\text{-}19} = {}^{1}/_{nkeys} = {}^{1}/_{8}$

$RF_{Rank} = {}^{1}/_{nkeys} = {}^{3}/_{8}$

Age is best!

Histograms

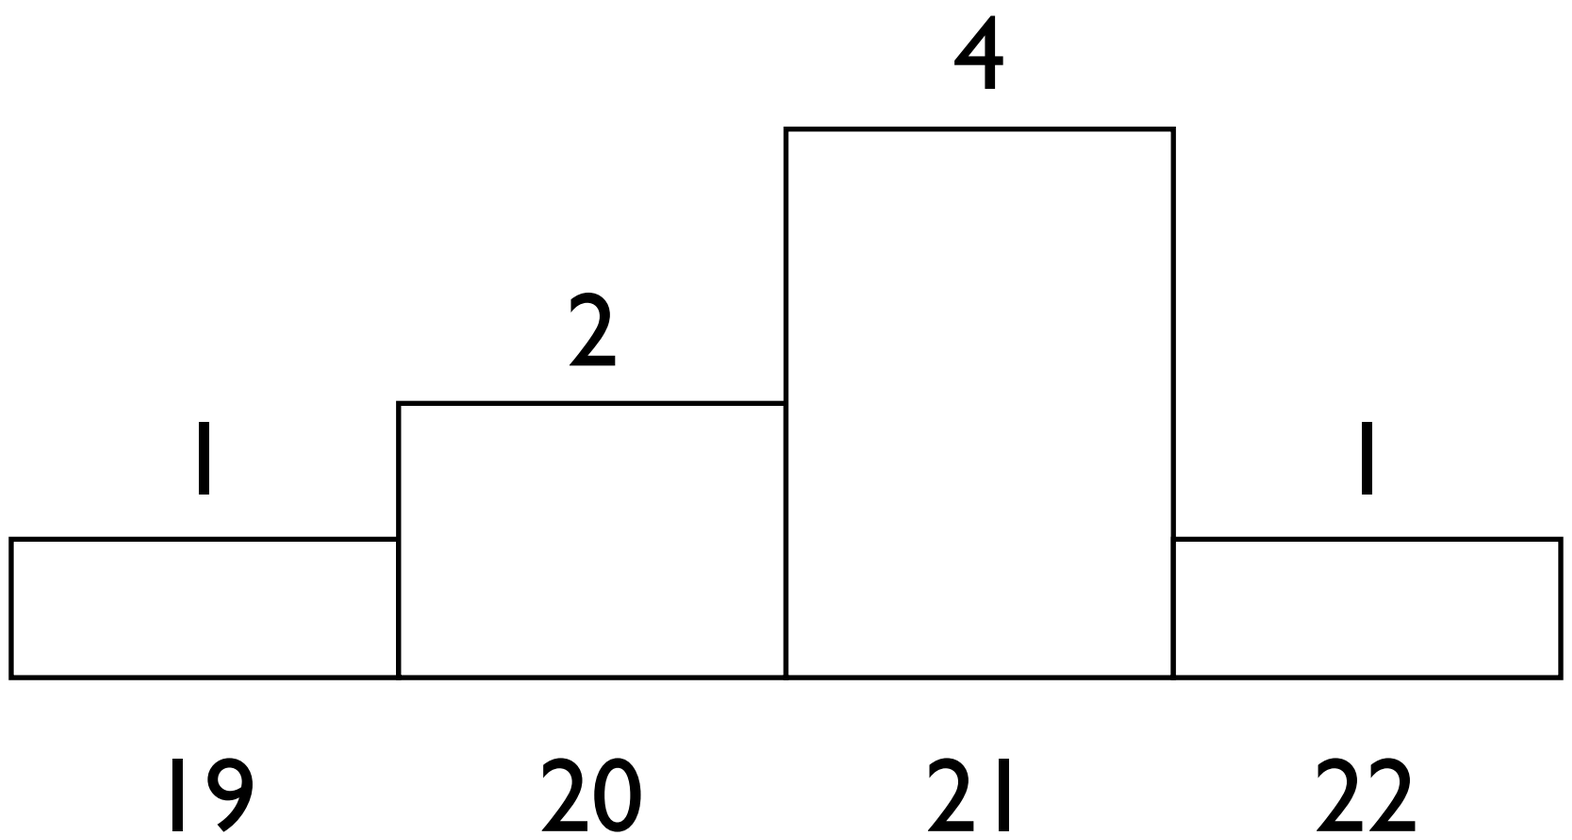

Histograms

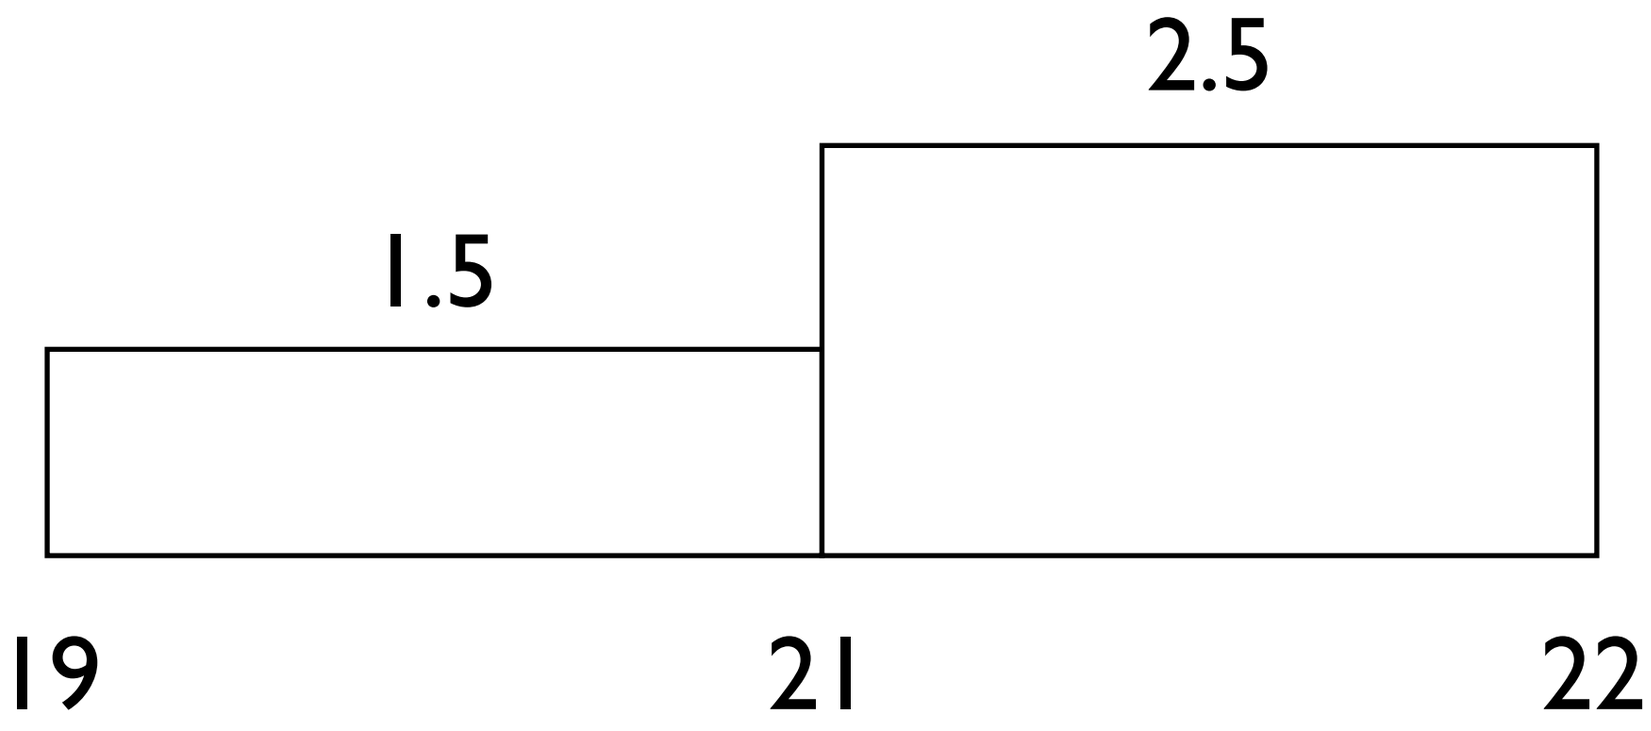

Histograms

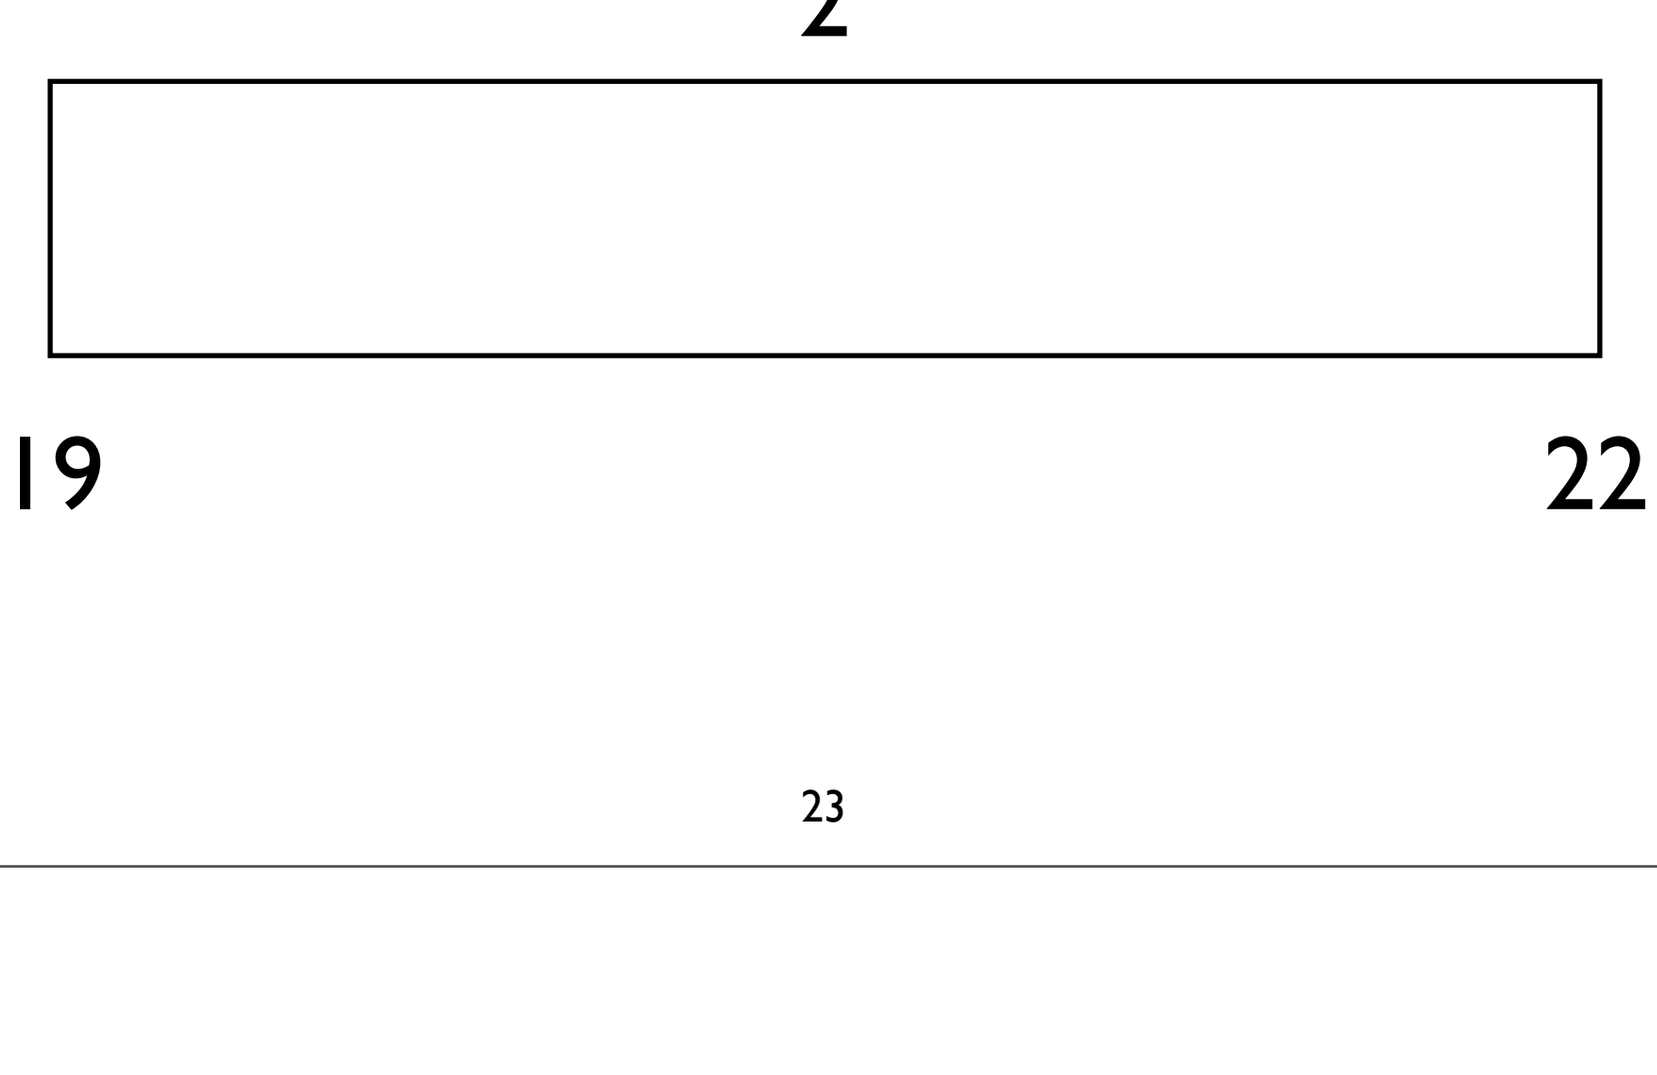

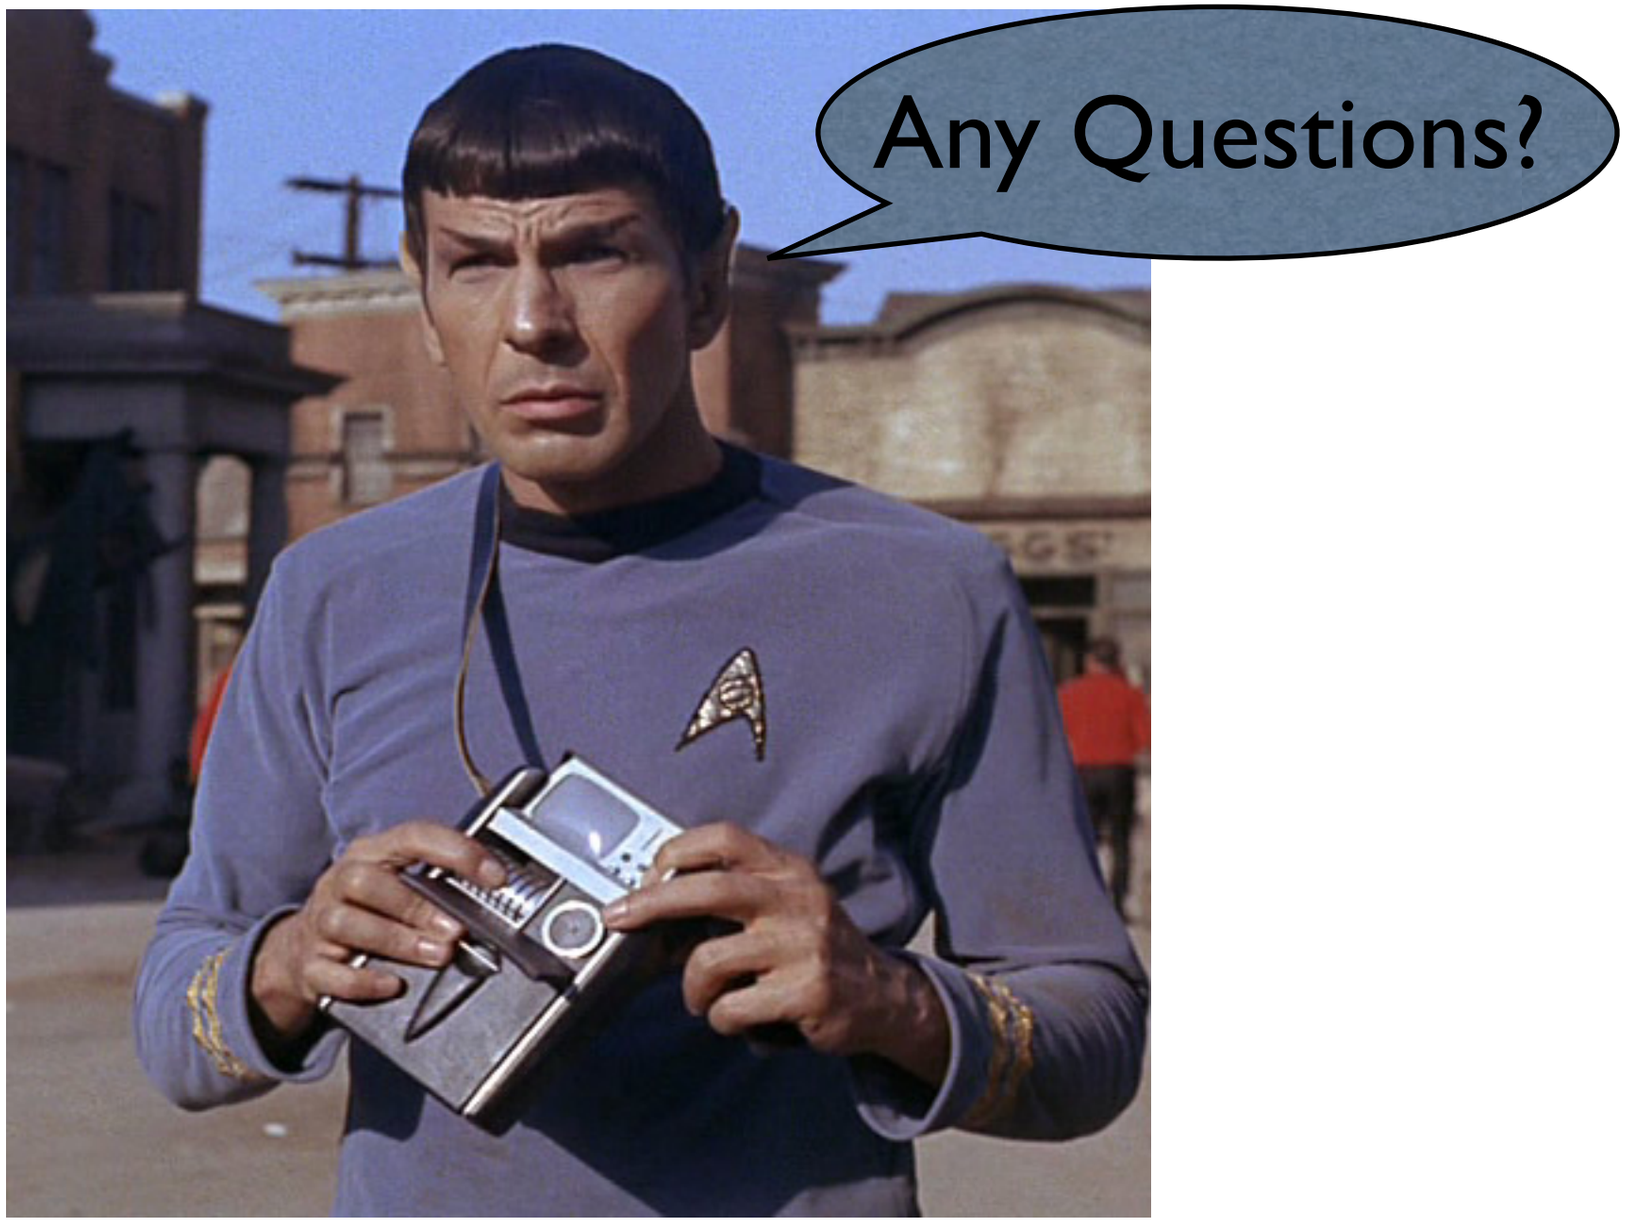

Histograms

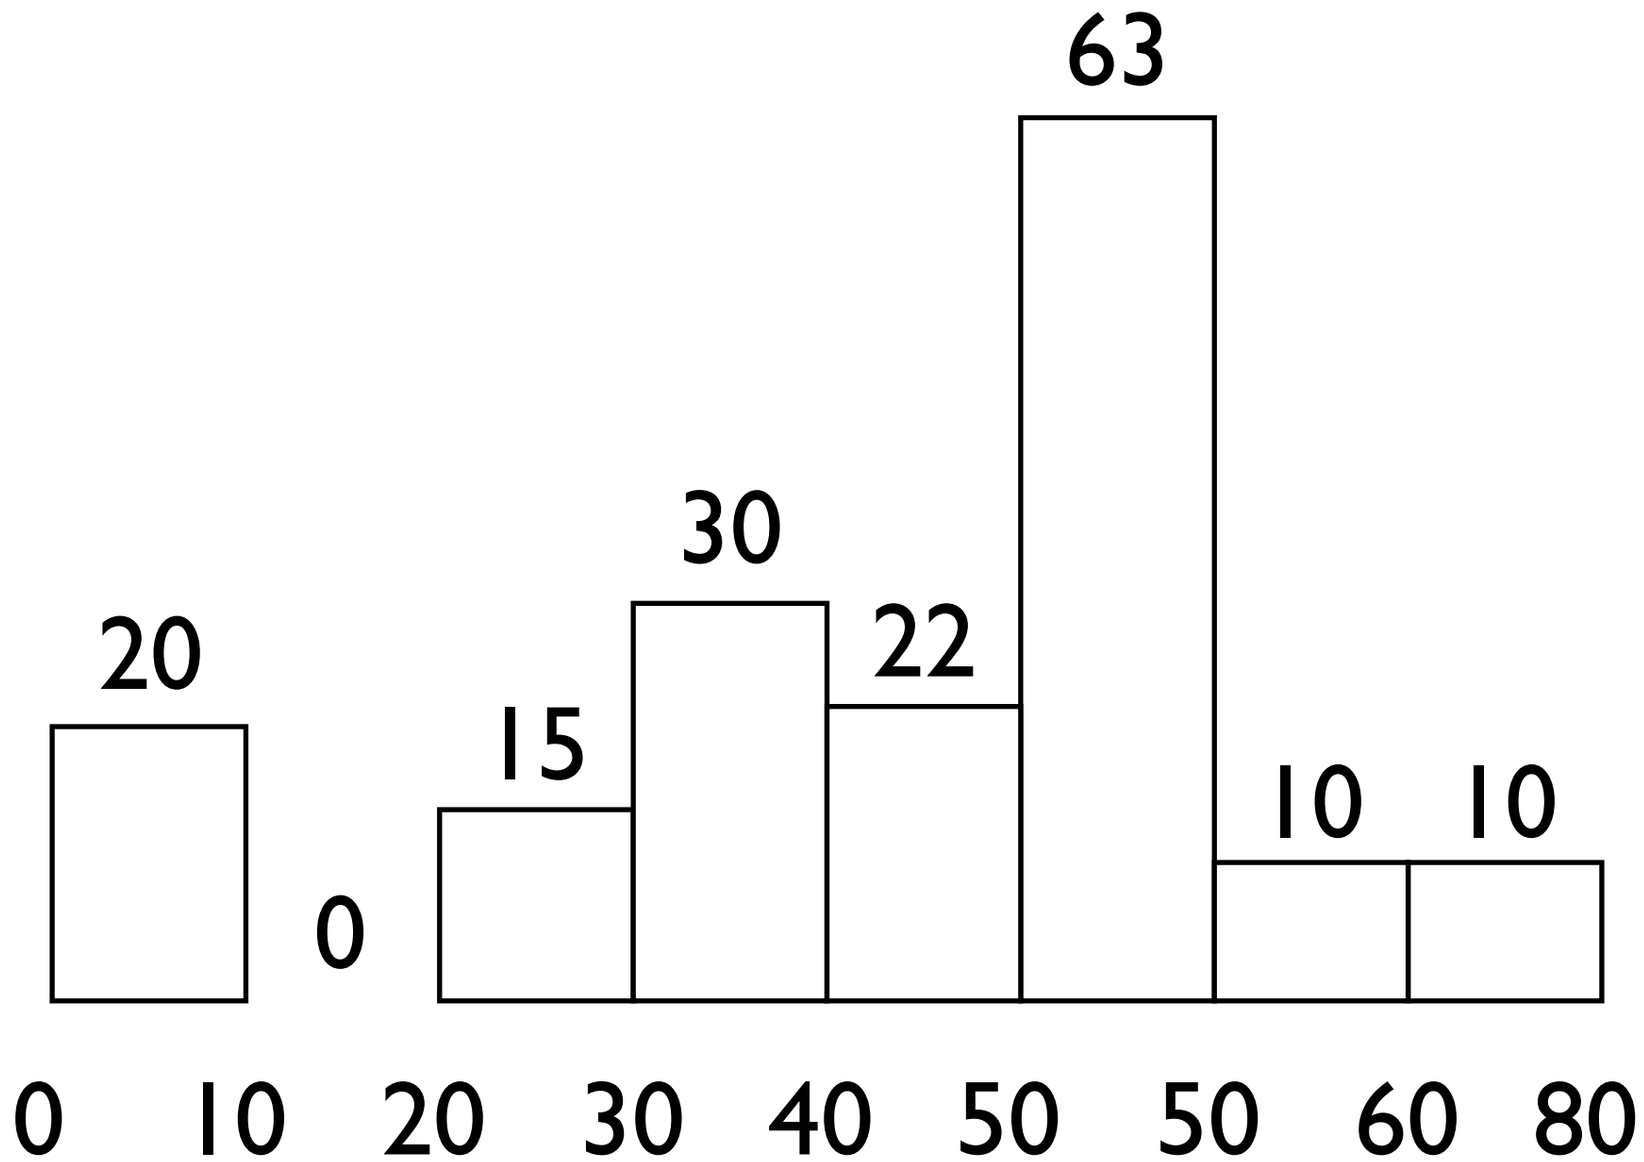

SELECT … WHERE A = 33

Friday, March 1, 13

Total number of keys: 20+15+30+22+63+10+10 = 170
(30/10)/170 = 3/170 = 3/180 = 0.0176 = 1.8%

Histograms

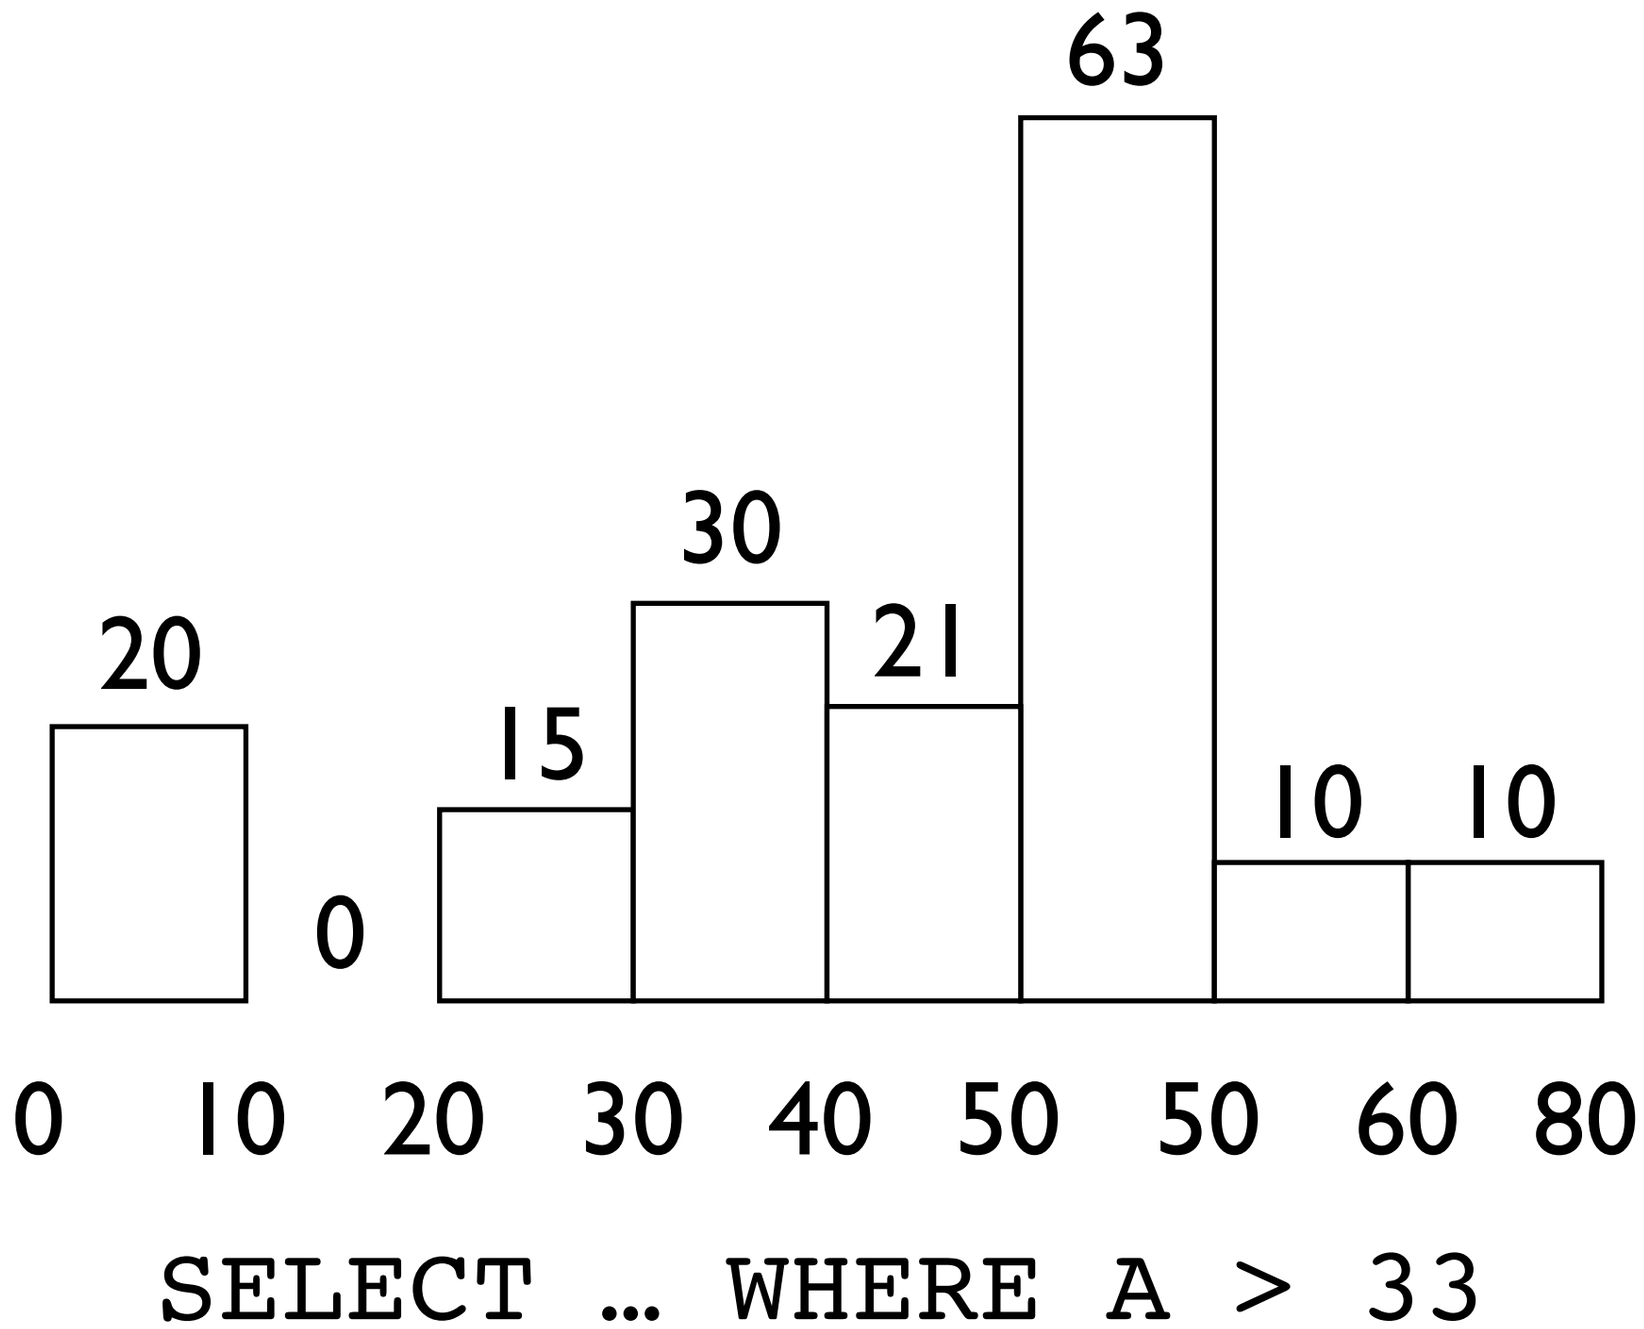

Friday, March 1, 13

Total number of keys: 20+15+30+22+63+10+10 = 170
Keys matching predicate: ((40-33)/10)*30+21+63+10+10 = 21+104 = 125
125/170 = 0.735 = 74%

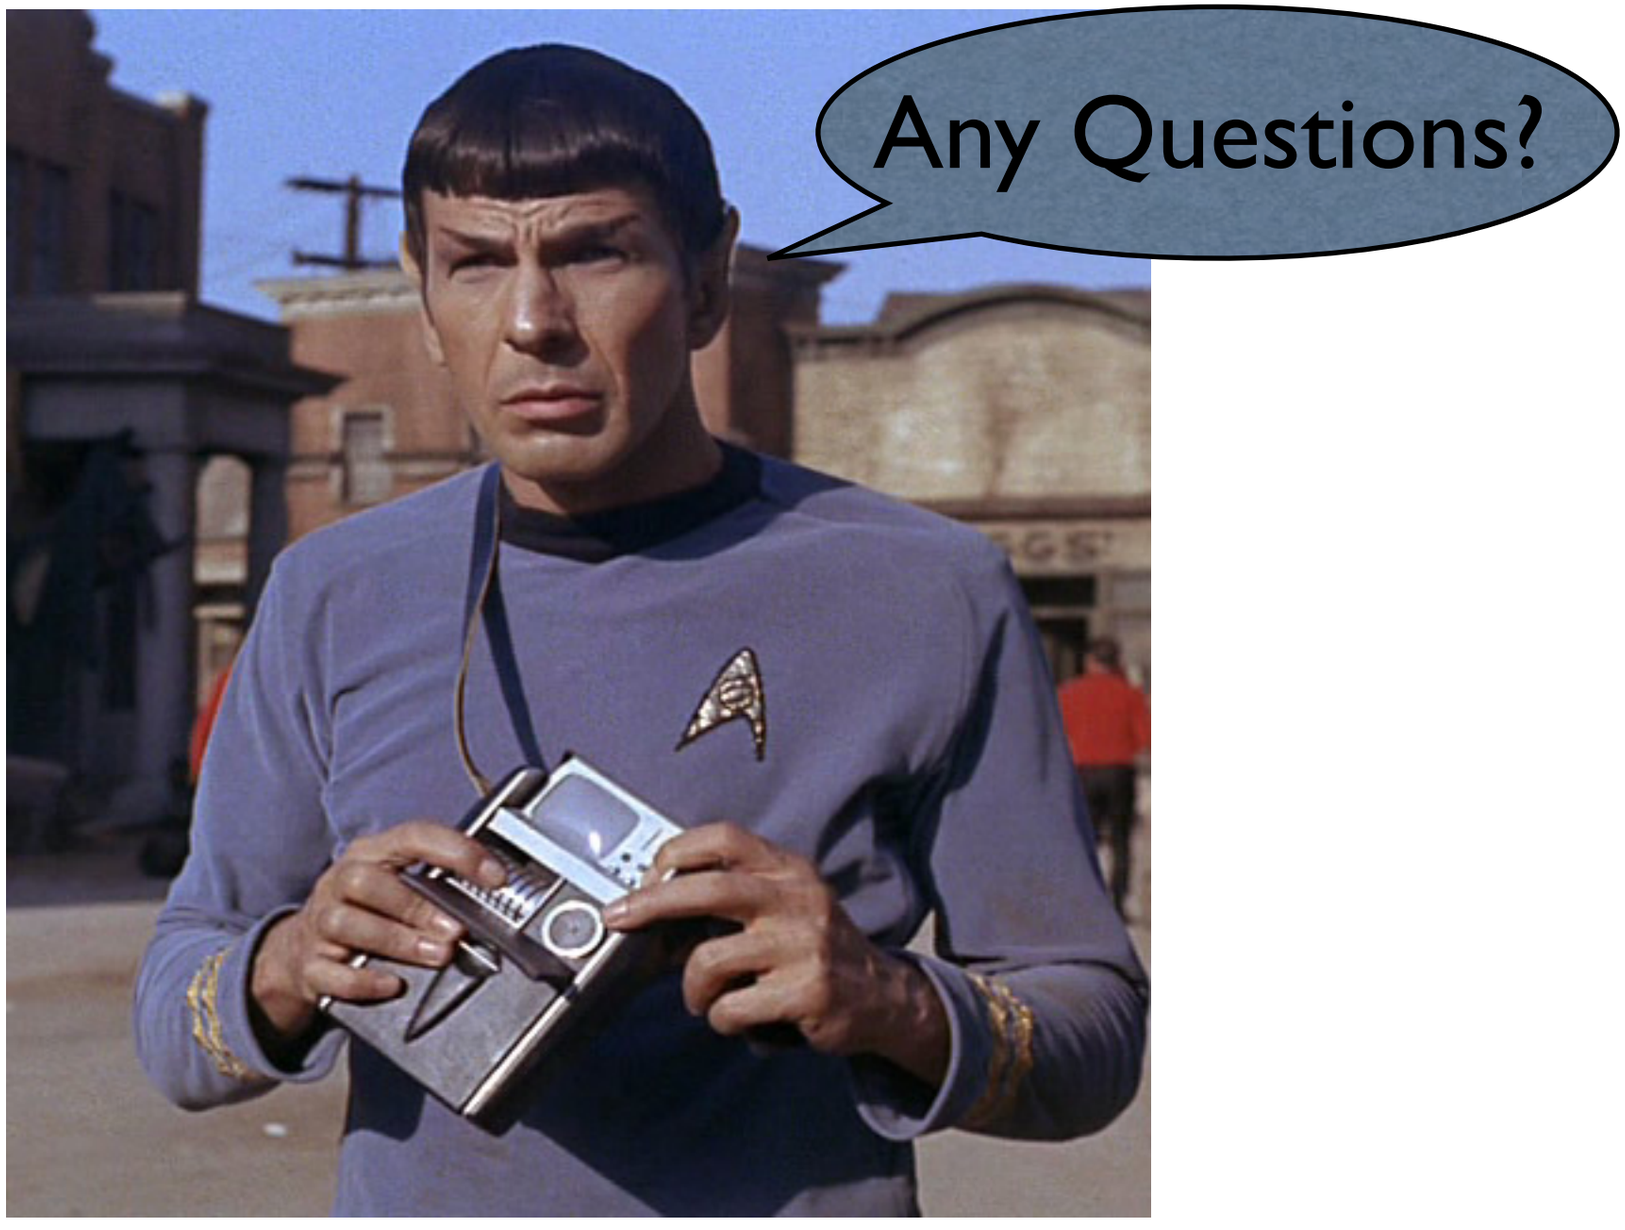

Using Constraints

- A Key attribute has one distinct value per row (equality selects exactly one row)
- Foreign Key joins generate one row for each row in the **referencing** relation.
- Cascade relationship guarantees EXACTLY one row per reference.