

# Transactions

## R&G Chapter 16

(slides adapted from content by J.Gehrke, J.Shanmugasundaram, and/or C.Koch)

Project I Graded - Grades available by Wednesday.  
(along with HW1, HW2. 3-4 Graded Soon)

1- DO NOT MODIFY TEST CASE FILES  
(please see me if you have, to avoid getting a 0)

2- Submit using tar/gz/zip. Please do not use rar.

3- Testing done with JVM/SDK 1.6  
Please do not use 1.7/1.8 features  
(Generics Type-Inference, Strings in Switches)

# Sample TEAM file

(Makes grading scripts happy)

```
kirkj  
spock32  
mccoyle
```

One UBIT per line.  
Nothing else!

# Grading

Code Didn't Compile: Tentative 0  
(See me)

Ran all provided RA Tests Successfully: 55 pts

Ran GB-Agg RA Tests Successfully: 5 pts

Ran all provided SQL Tests Successfully: 40 pts

Ran GB-Agg SQL Tests Successfully: +5 pts

Submitted to Original Deadline: +20 pts

## Of projects that successfully compiled

**Min:** 35.0

**Max:** 119.2

**Avg:** 91.4 (+/- 24.9)

Reminder: Midterm on Monday

Covers Material Up to and Including Optimization

See Syllabus for Exact Chapter/Section #s

# Time to get up into the guts of a database

## How do databases keep data moving at warp speed?



# Database Workloads

- Two major interactive database workloads:
  - OLAP: Online Analytical Processing
    - Big queries about large, static data.
  - OLTP: Online Transaction Processing
    - Many queries and updates on small data.



# OLTP Workloads

- Small data manipulations: Lots of IOs.
- Keep the disk active by running many tasks in parallel.
- Keep latencies low by making progress on multiple tasks at once.
- It is the DBMS's responsibility to interleave parallel tasks correctly.

# Transactions

What does it mean for a database operation to be correct?

The database doesn't know anything about the business logic. It doesn't know how the user/calling application is using data. All it knows is what values have been read and which values have been written.

We need to develop a notion of “correctness” relative to this view of the data -- reads and writes.

# Transactions

What does it mean for a database operation to be correct?

How does a database interact with its users?

The database doesn't know anything about the business logic. It doesn't know how the user/calling application is using data. All it knows is what values have been read and which values have been written.

We need to develop a notion of “correctness” relative to this view of the data -- reads and writes.

# Transactions


- Users group sequences of interactions with the DBMS into a Transaction.
- *Guarantee*: From the user's perspective, transactions execute fully and on their own.
- *Guarantee*: Transactions preserve data consistency (as per Integrity Constraints).
- The DBMS interleaves DB operations while respecting the above guarantees.

# Transactions

- Challenge 1: Interleaving
  - How do we negate the effects of interleaving operations from two parallel transactions?
- Challenge 2: Crashes
  - How do we avoid leaving the system in an inconsistent state after a crash?

# Atomicity

- A transaction completes by committing, or terminates by aborting.
- Logging is used to undo aborted transactions.
- **Atomicity:** A transaction is (or appears as if it were) applied in one 'step', independent of other transactions.
- All ops in a transaction commit or abort together.
- All other transactions commit 'logically' in a fixed order. (Serializability)



# Example

T1: BEGIN  $A=A+100$ ,  $B=B-100$  END

T2: BEGIN  $A=1.06*A$ ,  $B=1.06*B$  END

- Intuitively, T1 transfers \$100 from A to B and T2 credits both accounts with interest.
- What are possible interleaving errors?



# Example: Schedule

Time

T1

T2

$A = A + 100$

$A = 1.06 * A$

$B = B - 100$

$B = 1.06 * B$



# Example: Schedule

Time

T1

T2

$A = A + 100$

$A = 1.06 * A$

$B = B - 100$

$B = 1.06 * B$

OK!

16

# Example: Schedule

Time

T1

T2

$$A = A + 100$$

$$A = 1.06 * A$$

$$B = 1.06 * B$$

$$B = B - 100$$



# Example: Schedule

Time

T1

T2

$A = A + 100$

$A = 1.06 * A$

$B = 1.06 * B$

$B = B - 100$

Not OK!

# Example: The DBMS's View

Time

T1

T2

R ( A )

W ( A )

R ( A )

W ( A )

R ( B )

W ( B )

R ( B )

W ( B )



# Example: The DBMS's View

Time

T1

T2

R ( A )

W ( A )

R ( A )

W ( A )


R ( B )

W ( B )

R ( B )

W ( B )

Not OK!



# Scheduling

- **Serial Schedule:** A schedule that does not interleave the actions of different transactions.
- (Two) **Equivalent Schedules:** For any database state, the effect (on the set of objects in the database) of executing the first schedule is identical to the effect of executing the second.
- **Serializable Schedule:** A schedule that is equivalent to some serial execution of the transactions.
- If the transactions preserve consistency, every serializable schedule does too.



# Interleaving Anomalies

- Reading uncommitted data (write-read/WR conflicts; aka “dirty reads”)

T1: R(A), W(A), R(B), W(B), ABRT  
T2: R(A), W(A), CMT,

- Unrepeatable Reads (read-write/RW conflicts)


T1: R(A), R(A), W(A), CMT  
T2: R(A), W(A), CMT,

# Interleaving Anomalies

- Overwriting Uncommitted Data (write-write/WW Conflicts)

T1 :  $W(A)$  ,  $W(B)$  , CMT

T2 :  $W(A)$  ,  $W(B)$  , CMT ,




# Lock-Based Concurrency

- Strict Two-Phase Locking
  - Before reading, obtain a S(hared) lock
  - Before writing, obtain a X(clusive) lock
  - Hold locks until transaction completes
- Why two-phase locking?

# Lock-Based Concurrency

- Non-Strict 2PL Variant
  - Release locks anytime, but can't acquire locks after first release.
- General 2PL issues
  - What can go wrong with locking?
  - What are some ways around these issues?



# Aborting Transactions

- If a transaction  $T_i$  is aborted, all of its actions have to be undone!
- What if  $T_j$  reads a value modified by  $T_i$ ?
- How can we prevent this from happening?
- The DBMS maintains a log of every write, which it uses to undo aborted xacts.
- The log also assists in crash recovery!


# The Log

- The log records:
  - $T_i$  writes: Both old and new values of object.
    - Log record must make it to disk before the changed data page!
  - $T_i$  commits/aborts: A record of the event.
- Log records are chained (stored as a linked list) for each xact, so it's easy to undo an xact.



# The Log

- DBMSes often store redundant duplicates of the log on stable storage.
- All log-related activities (and locking/deadlock detection) are handled transparently by the DBMS.



# Transaction Variants

- Cross-server transactions.
  - Multiple DBMS servers participate.
- Non-interactive transactions.
  - Specify full transactions programmatically.
- Functionally-limited transactions.
  - Read-only, No side-effects, Limited writes

# Crash Recovery Preview

- The ARIES recovery algorithm (3 Phases)
  - **Analysis:** Scan through the log to identify all xacts active at time of crash and all dirty pages in the buffer pool.
  - **Redo:** Redo all updates to dirty pages in the buffer pool to ensure that logged updates are carried out.
  - **Undo:** Use 'before' value from log to cancel out the writes of all xacts that were active at the time of crash. (need to guard vs crashing during recovery)

# Summary

- Concurrency control is one of the most important features provided by a DBMS.
- Users do not need to worry about concurrency issues.
  - DBMS automatically inserts locking/scheduling requests as needed.
- Write-ahead logging is used to undo the actions of aborted transactions.
- Consistent State: Only the effects of committed xacts are seen by other xacts.