

# Concurrency Control I

R&G Chapter 17

(slides adapted from content by J.Gehrke, J.Shanmugasundaram, and/or C.Koch)

# Reminders

- Midterm on Monday.
  - Questions similar to the homeworks.
  - Closed book.
  - Homework 1-4 solutions posted tonight.
- No homework this week.

# Recap: Transactions

- OLTP Workloads: Lots of data updates.
- Transactions group database operations.
  - ... is a sequence of reads and writes.
  - ... is executed atomically.
  - ... preserves integrity constraints.
    - ... even if the system crashes.

# Crash Recovery Preview

- The ARIES recovery algorithm (3 Phases)
  - **Analysis:** Scan through the log to identify all xacts active at time of crash and all dirty pages in the buffer pool.
  - **Redo:** Redo all updates to dirty pages in the buffer pool to ensure that logged updates are carried out.
  - **Undo:** Use 'before' value from log to cancel out the writes of all xacts that were active at the time of crash. (need to guard vs crashing during recovery)

# Schedules

- An order in which a sequence of operations are executed.
- A schedule for a set of transactions is an interleaving of all operations in the transaction.
- The only restriction is that operations within a transaction are executed in order.

# Recap: Scheduling

- **Serial Schedule:** A schedule that does not interleave the actions of different transactions.
- **(Two) Equivalent Schedules:** For any database state, the effect (on the set of objects in the database) of executing the first schedule is identical to the effect of executing the second.
- **Serializable Schedule:** A schedule that is equivalent to some serial execution of the transactions.
- If the transactions preserve consistency, every serializable schedule does too.

# Ensuring Serializable Schedules

- **Option 1:** Use reader/writer locks.
  - ... latency/throughput cost (pessimistic).
- **Option 2:** Abort conflicting transactions.
  - ... fairness/throughput cost (optimistic).
- **Option 3:** Predeclare reads/writes, and schedule operations before execution.
  - ... flexibility cost.

# Ensuring Serializable Schedules

- **Option 1:** Use reader/writer locks.
  - ... latency/throughput cost (pessimistic).
- **Option 2:** Abort conflicting transactions.
  - ... fairness/throughput cost (optimistic).
- **Option 3:** Predeclare reads/writes, and schedule operations before execution.
  - ... flexibility cost.





Any Questions?

How do we define correctness precisely?

# Conflict Serializability

- Two schedules are conflict equivalent if:
  - ... they involve the same set of operations from the same set of transactions.
  - ... every pair of conflicting actions between two transactions is ordered in the same way.
- Not quite the same as regular equivalence (it's stronger).
  - Defined over operations rather than possible effects.
- A schedule is conflict serializable if it is conflict-equivalent to some serial schedule.

# Example

Time

T1

T2

R ( A )

W ( A )

R ( A )

W ( A )

R ( B )


W ( B )

R ( B )

W ( B )


||

# Example



I2

# Example



(Dependency Graph)

# Example



(Dependency Graph)

The cycle in the graph reveals a problem:  
T1's output depends on T2 and visa versa.

# Dependency Graphs

- One node per transaction
- One edge (from  $T_i$  to  $T_k$ ) if  $T_k$  reads/writes an object most recently written by  $T_i$ .
- **Claim:** A schedule is conflict serializable if and only if its dependency graph is acyclic.





Any Questions?

# View Serializability

- Schedules in terms of pre- and post-conditions.
- Schedules  $S1$  and  $S2$  are view-equivalent if:
  - ...  $T_i$  reads the initial value of  $A$  in  $S1$ , then it reads the initial value of  $A$  in  $S2$ .
  - ...  $T_i$  reads a value of  $A$  written by  $T_k$  in  $S1$  then it reads the value of  $A$  written by  $T_k$  in  $S2$ .
  - ...  $T_i$  writes the final value of  $A$  in  $S1$  then  $T_i$  writes the final value of  $A$  in  $S2$ .

# Example

Time

T1

T2

T3

R ( A )

W ( A )

W ( A )

W ( A )



# Example

Time

T1

T2

T3

R ( A )

W ( A )

W ( A )

W ( A )



# Example

Time

T1

T2

T3

R ( A )

Write order irrelevant  
(T3 overwrites either way)

W ( A )

W ( A )

W ( A )



Any Questions?

# Implementing Locks

# Take 1: 2-Phase Locking

- To Recap:
  - Obtain reader/writer locks on objects.
  - Once an xact releases a lock, it can no longer acquire any new locks.
- **Claim:** 2-phase locking allows only conflict-serializable schedules.
  - If xact A modifies a value, no other xact can read/modify that value until A 'completes'.





Any Questions?

# Lock Management

- Lock Manager (sometimes part of the Transaction manager) handles lock/unlock requests.
- Lock table entry:
  - # of transactions currently holding a lock
  - Type: Shared (reader) or Exclusive (writer) Lock
  - Pointer to queue of lock requests
- Locking/Unlocking implemented as atomic ops
- Lock Upgrades: Convert Shared Lock into Exclusive.

# Deadlocks

- Deadlock: A cycle of transactions waiting on each other's locks
  - Problem in 2PL; xact can't release a lock until it completes
- How do we handle deadlocks?
  - **Anticipate**: Prevent deadlocks before they happen.
  - **Detect**: Identify deadlock situations and abort one of the deadlocked xacts.

# Deadlock Detection

- **Baseline:** If a lock request can not be satisfied, the transaction is blocked and must wait until the resource is available.
- Create a waits-for graph:
  - Nodes are transactions
  - Edge from  $T_i$  to  $T_k$  if  $T_i$  is waiting for  $T_k$  to release a lock.
- Periodically check for cycles in the graph.

# Example

Time



T1

T2

T3

T4

S ( A )  
R ( A )

X ( B )  
W ( B )

S ( B )


S ( C )  
R ( C )

X ( C )

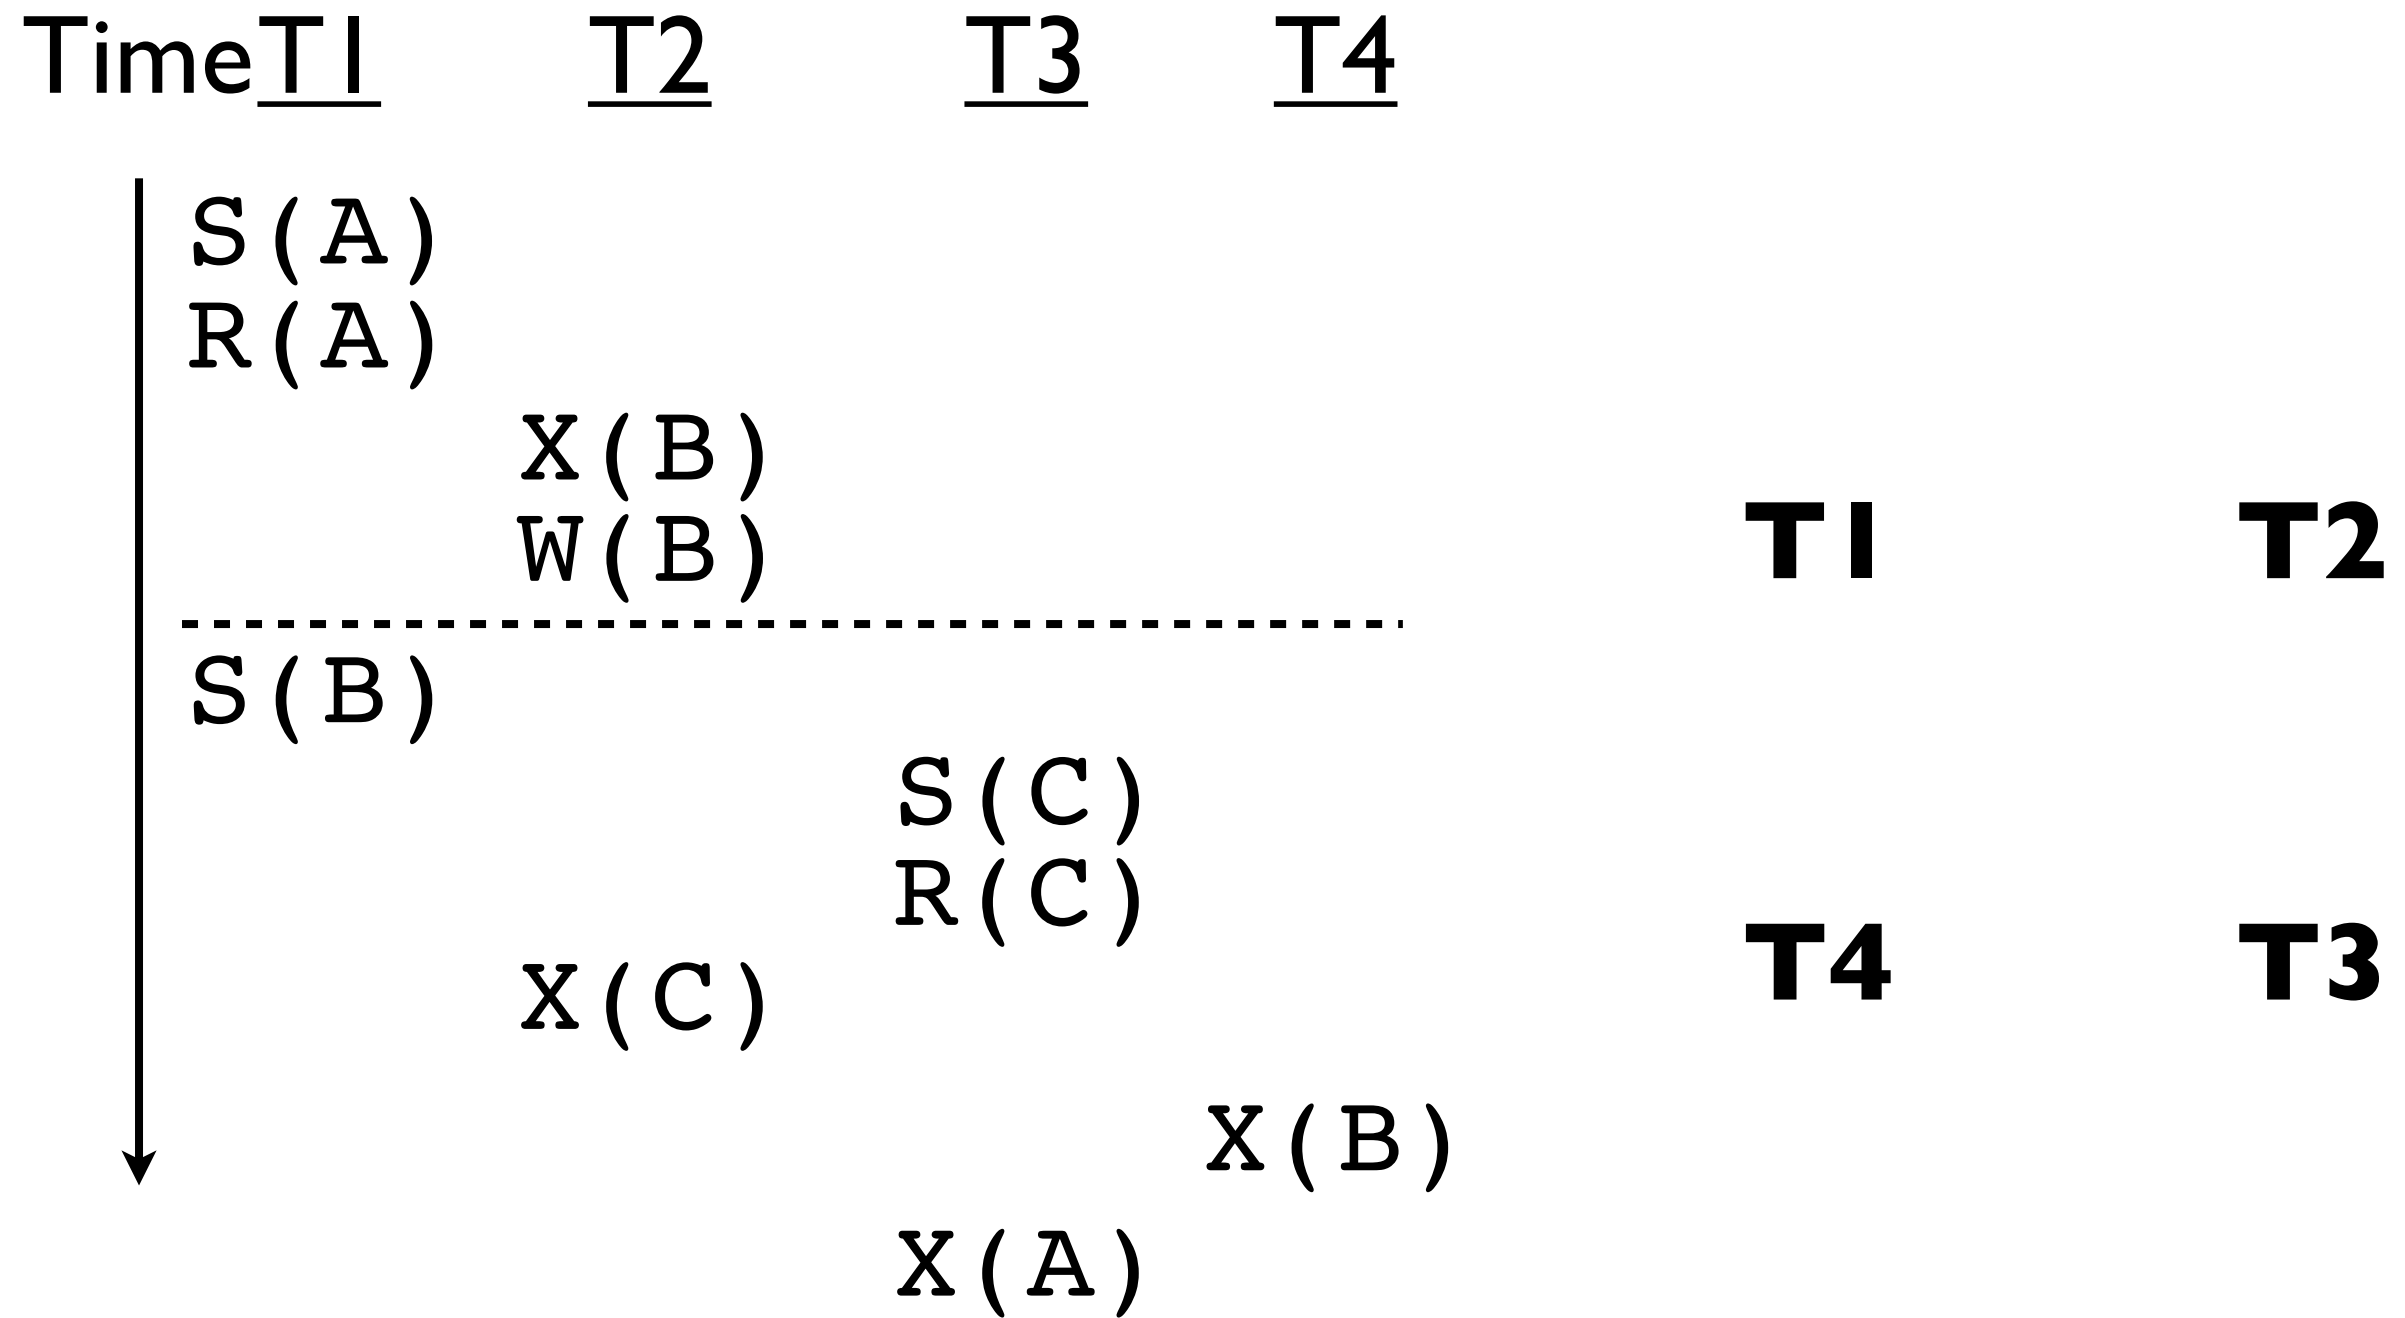
X ( A )

X ( B )

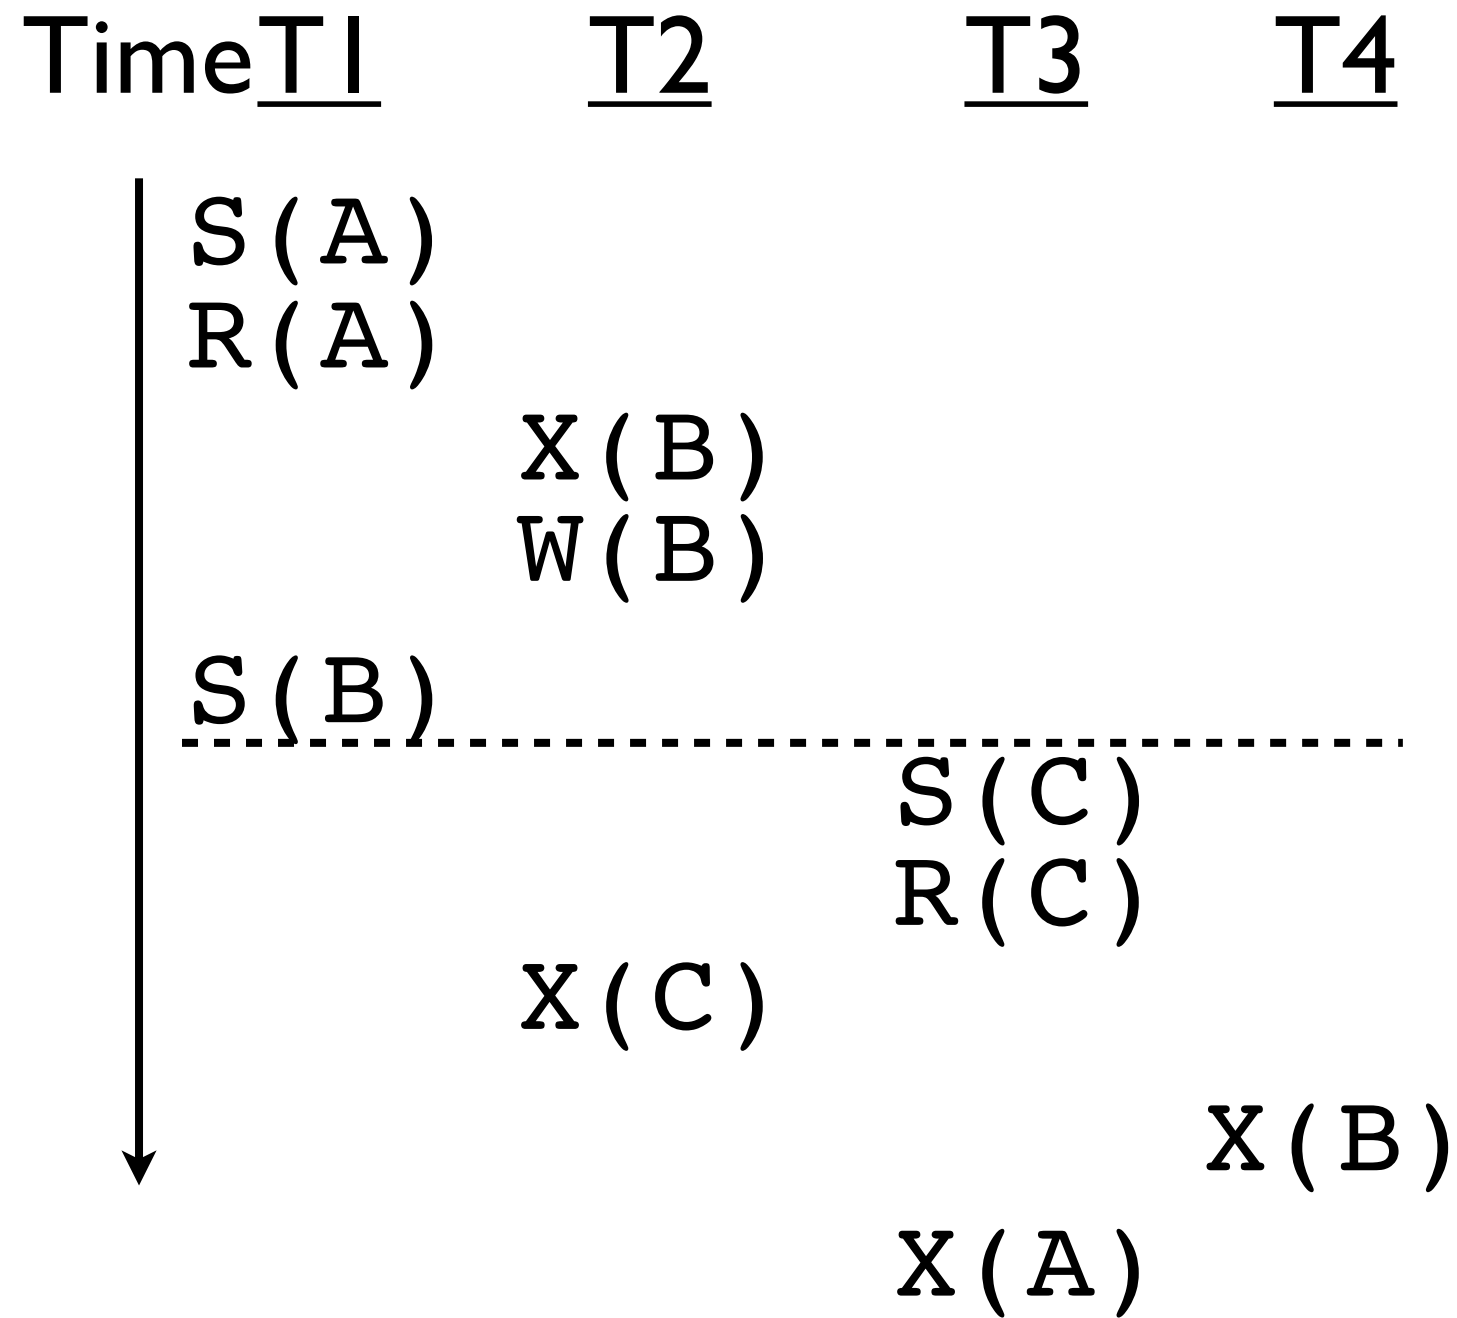
# Example



# Example



# Example





**T1** → **T2**

**T4**



**T3**





# Example





# Example



# Example



# Example





Any Questions?

# Deadlock Prevention

- Ensure that dependencies are monotonic (and consequently acyclic)
- Assign each transaction a priority based on the timestamp at which it starts.
- When a transaction fails to acquire a lock:
  - Wait if monotonicity would be preserved.
  - Kill one transaction otherwise.

# Deadlock Prevention

- Policy 1 (Wait-Die): If  $T_i$  has a higher priority, wait for  $T_k$ , otherwise  $T_i$  aborts.
- Policy 2 (Wait-Wound): If  $T_i$  has a higher priority,  $T_k$  aborts, otherwise  $T_i$  waits.
- Protect fairness by restarting the aborted transaction with its original timestamp.



Any Questions?