

Concurrency Control 2

R&G Chapter 17

(slides adapted from content by J.Gehrke, J.Shanmugasundaram, and/or C.Koch)

Reminders

- Midterm on Monday.
 - Questions similar to the homeworks.
 - Closed book.
 - Homework 1-4 solutions posted tonight.
- No homework this week.

Recap-Equivalence

- Schedule Equivalence (when are 2 scheds the same)
 - Conflict Equivalent
 - All conflicting ops are in the same order
 - View Equivalent
 - All reads read the same value (initial or from the same xact)
 - All final writes come from the same xact.
- (Conflict | View) Serializable
 - = (Conflict | View) Equivalent to a Serial schedule.

Recap-Deadlock Prevention


- 1) Track which xacts are waiting for which xacts (waits-for graph)
 - If a cycle exists, kill one of the xacts.
- 2) Prioritize transactions
 - a) Lower priority xacts die when they try to take a lock held by a higher priority xact.
 - b) Higher priority xacts kill any xact holding a lock that they need.

What is an Object (that we can lock)?

What do we lock?


- The entire database?
- A table in the database?
- Individual pages a table is stored on?
- Individual tuples in a table?

What do we lock?



What do we lock?

- We don't need to decide what to store!
- Data containers are nested!



What do we lock?

Is it safe to allow some transactions to lock tables while other transactions to lock tuples?

New Lock Modes

- 'Intent' to lock (a child)
 - Intent-to-Lock Shared (IS)
 - Intent-to-Lock Exclusive (IX)
- Actual lock (on the object)
 - Lock-Shared (S)
 - Lock-Exclusive (x)
- Lock Shared + Intent-to-Lock Exclusive (SIX)

New Lock Modes

Lock Mode(s) Currently Held By Other Xacts

Lock Mode Desired

	None	IS	IX	S	X
None	valid	valid	valid	valid	valid
IS	valid	valid	valid	valid	fail
IX	valid	valid	valid	fail	fail
S	valid	valid	fail	valid	fail
X	valid	fail	fail	fail	fail

Hierarchical Locks

- Lock Objects Top-Down
 - Before acquiring a lock on an object, an xact must have at least an intention lock on its parent!
- For example:
 - To acquire a S on an object, an xact must have an IS, IX on the object's parent (why not S, SIX, or X?)
 - To acquire an X (or SIX) on an object, an xact must have a SIX, or IX on the object's parent.

Hierarchical Locks

	None	IS	IX	S	X
None	valid	valid	valid	valid	valid
IS	valid	valid	valid	valid	fail
IX	valid	valid	valid	fail	fail
S	valid	valid	fail	valid	fail
X	valid	fail	fail	fail	fail

T1: Holds a S on a table

T2: Wants an X on a row of the table

Hierarchical Locks

	None	IS	IX	S	X
None	valid	valid	valid	valid	valid
IS	valid	valid	valid	valid	fail
IX	valid	valid	valid	fail	fail
S	valid	valid	fail	valid	fail
X	valid	fail	fail	fail	fail

T1: Holds a X on a table

T2: Wants an S on a row of the table

Hierarchical Locks

	None	IS	IX	S	X
None	valid	valid	valid	valid	valid
IS	valid	valid	valid	valid	fail
IX	valid	valid	valid	fail	fail
S	valid	valid	fail	valid	fail
X	valid	fail	fail	fail	fail

T1: Holds a S on a row of the table

T2: Wants a S on the table

Hierarchical Locks

	None	IS	IX	S	X
None	valid	valid	valid	valid	valid
IS	valid	valid	valid	valid	fail
IX	valid	valid	valid	fail	fail
S	valid	valid	fail	valid	fail
X	valid	fail	fail	fail	fail

T1: Holds a S on a row of the table

T2: Wants an X on the table

Hierarchical Locks

	None	IS	IX	S	X
None	valid	valid	valid	valid	valid
IS	valid	valid	valid	valid	fail
IX	valid	valid	valid	fail	fail
S	valid	valid	fail	valid	fail
X	valid	fail	fail	fail	fail

T1: Holds an X on a row of the table
T2: Wants an X on the table

Hierarchical Locks

- Release locks bottom up (child before parent)
- This protocol is equivalent to acquiring all child locks when the parent is acquired.
- S requires IS/IX on parent (conflicts with X on parent)
- X requires IX/SIX on parent (conflicts with S, X on parent)

Example

- T1 scans R and updates several tuples
- T1 acquires SIX on R
- **Scan:** T1 already has S on R.
- **Update:** T1 acquires X on individual tuples as needed (needs IX on pages).

Example

- T2 uses an index to read part of R
- T2 gets an IS lock on R
- **Scan:** T2 repeatedly acquires S locks on pages of R as needed.

Example

- T3 scans all of R
 - T3 gets an S lock on R
- OR
- T3 could behave like T2
(IS+S on each page or tuple)
- Use lock escalation to decide which

Problem:
Locking assumes that we can lock all objects!
(What happens if we insert objects?)

T1

T2

```
DELETE FROM Officers
WHERE rank = 1
AND age =
  (SELECT MAX(age)
   FROM Officers WHERE rank=1)
LIMIT 1;
```

Time



T1

T2

```
DELETE FROM Officers
WHERE rank = 1
      AND age = (71)
      (SELECT MAX(age)
       FROM Officers WHERE rank=1)
LIMIT 1;
```

Time



T1

T2

DELETE FROM Officers
WHERE rank = 1
AND age = (71)
(SELECT MAX(age)
FROM Officers WHERE rank=1)
LIMIT 1;

INSERT INTO Officers(rank,age)
VALUES (1, 96);

Time



Time



T1

```
DELETE FROM Officers
WHERE rank = 1
  AND age = (71)
  (SELECT MAX(age)
   FROM Officers WHERE rank=1)
LIMIT 1;
```

T2

```
INSERT INTO Officers(rank,age)
          VALUES (1, 96);

DELETE FROM Officers
WHERE rank = 2
  AND age = (80)
  (SELECT MAX(age)
   FROM Officers WHERE rank=2)
LIMIT 1;
```

T1

T2

DELETE FROM Officers
WHERE rank = 1
AND age = (71)
(SELECT MAX(age)
FROM Officers WHERE rank=1)
LIMIT 1;

INSERT INTO Officers(rank,age)
VALUES (1, 96);

DELETE FROM Officers
WHERE rank = 2
AND age = (80)
(SELECT MAX(age)
FROM Officers WHERE rank=2)
LIMIT 1;

SELECT MAX(age)
FROM Officers(rank,age)
WHERE rank = 2 (63)

Time
↓

T1

T2

DELETE FROM Officers
WHERE rank = 1
AND age = (71)
(SELECT MAX(age)
FROM Officers WHERE rank=1)
LIMIT 1;

INSERT INTO Officers(rank,age)
VALUES (1, 96);

DELETE FROM Officers
WHERE rank = 2
AND age = (80)
(SELECT MAX(age)
FROM Officers WHERE rank=2)
LIMIT 1;

SELECT MAX(age)
FROM Officers(rank,age)
WHERE rank = 2 (63)

Time
↓

[INCONSISTENT]

The Problem

- T1 assumes that it has locked all sailor records with rating = 1
- Solution 1: Lock entire table (expensive!)
- Solution 2: Lock a **predicate**.
(e.g., rating = 1)

(To be continued next week)