

Concurrency Control 3

R&G Chapter 17

(slides adapted from content by J.Gehrke, J.Shanmugasundaram, and/or C.Koch)


Reminders

- Homework due Monday after break
(Posted tonight)

Recap: New Lock Modes


- 'Intent' to lock (a child)
 - Intent-to-Lock Shared (IS)
 - Intent-to-Lock Exclusive (IX)
- Actual lock (on the object)
 - Lock-Shared (S)
 - Lock-Exclusive (x)
- Lock Shared + Intent-to-Lock Exclusive (SIX)

What do we lock?



What do we lock?

- Objects Locked Top-Down
- Before acquiring a lock on an object, a transaction must have at least an intention lock on its parent!



Recap: New Lock Modes

Lock Mode(s) Currently Held By Other Xacts

Lock Mode Desired

	None	IS	IX	S	X
None	valid	valid	valid	valid	valid
IS	valid	valid	valid	valid	fail
IX	valid	valid	valid	fail	fail
S	valid	valid	fail	valid	fail
X	valid	fail	fail	fail	fail

Problem:
Locking assumes that we can lock all objects!
(What happens if we insert objects?)

T1

T2

```
DELETE FROM Officers
WHERE rank = 1
      AND age =
      (SELECT MAX(age)
       FROM Officers WHERE rank=1)
LIMIT 1;
```

Time



T1

T2

```
DELETE FROM Officers
WHERE rank = 1
      AND age = (71)
      (SELECT MAX(age)
       FROM Officers WHERE rank=1)
LIMIT 1;
```

Time



T1

T2

DELETE FROM Officers
WHERE rank = 1
AND age = (71)
(SELECT MAX(age)
FROM Officers WHERE rank=1)
LIMIT 1;

INSERT INTO Officers(rank,age)
VALUES (1, 96);

Time



Time



T1

```
DELETE FROM Officers
WHERE rank = 1
  AND age = (71)
  (SELECT MAX(age)
   FROM Officers WHERE rank=1)
LIMIT 1;
```

T2

```
INSERT INTO Officers(rank,age)
          VALUES (1, 96);

DELETE FROM Officers
WHERE rank = 2
  AND age = (80)
  (SELECT MAX(age)
   FROM Officers WHERE rank=2)
LIMIT 1;
```

T1

T2

DELETE FROM Officers
WHERE rank = 1
AND age = (71)
(SELECT MAX(age)
FROM Officers WHERE rank=1)
LIMIT 1;

INSERT INTO Officers(rank,age)
VALUES (1, 96);

DELETE FROM Officers
WHERE rank = 2
AND age = (80)
(SELECT MAX(age)
FROM Officers WHERE rank=2)
LIMIT 1;

SELECT MAX(age)
FROM Officers(rank,age)
WHERE rank = 2 (63)

Time
↓

T1

T2

DELETE FROM Officers
WHERE rank = 1
AND age = (71)
(SELECT MAX(age)
FROM Officers WHERE rank=1)
LIMIT 1;

INSERT INTO Officers(rank,age)
VALUES (1, 96);

DELETE FROM Officers
WHERE rank = 2
AND age = (80)
(SELECT MAX(age)
FROM Officers WHERE rank=2)
LIMIT 1;

SELECT MAX(age)
FROM Officers(rank,age)
WHERE rank = 2 (63)

Time
↓

[INCONSISTENT]

The Problem

- T1 assumes that it has locked all sailor records with rating = 1
- Solution 1: Lock entire table (expensive!)
- Solution 2: Lock a **predicate**.
(e.g., rank = 1)
- Solution 3: Lock **index page(s)**.
(equivalent to a range/predicate lock)

Index Locking

- If there is an index on r rank, T_I locks the index page(s) with r rank = 1.
- If no such entries exist, lock the page where an entry would go.
- Need to stop other indexes from being locked
 - 1) Have a single primary index at any time.
 - 2) Conflict detection a'la predicate locking.

Predicate Locking

- Grant a lock on all records that satisfy some logical predicate (e.g., $\text{age} > 2 * \text{rank}$)
- Need a way to compute predicate intersections: Is it possible for a record to satisfy both predicates?
 - No analytical solution in the worst case.
 - Use a conservative approximation.

Predicate Locking

- Index locking is a special case of predicate locking when an index exists for the predicate.
- Can easily identify (superset) of records that satisfy multiple predicates based on pages.
- Easy to lock pages. Harder to lock predicates.
 - Predicate locking is rarely used.

Locking in B+ Trees

- How can we efficiently lock specific leaf pages?
 - **Note:** This is not quite multiple granularity locking
 - We may need to modify tree nodes (IX wrong).
 - Modifications to leaves may not affect the tree.
- **Solution I:** Ignore the tree and lock **every** page accessed (S) or modified (X) (both tree and leaf).
- Horrible performance! Why?

Locking in B+ Trees

- Higher levels of the tree are often only used to direct searches for leaf pages.
- We only need an X on a tree page if a child of the page could possibly split/merge on an insert/delete.
- We can exploit these observations to design a locking protocol that guarantees serializability, even though it violates 2PL.

Simple Tree Locking Algorithm


- **Scan:** Start at the root and descend.
 - Repeatedly S lock node, then unlock parent.
- **Update:** Start at the root and descend
 - Repeatedly X lock node.
 - If all children of a node are locked and safe, release the parent lock
 - Safe node: A node that will not propagate changes.

Simple Tree Locking Algorithm

- **Scan:** Start at the root and descend.
 - Repeatedly S lock node, then unlock parent.
- **Update:** Start at the root and descend
 - Repeatedly X lock node.
 - If all children of a node are locked and safe, release the parent lock
 - Safe node: A node that will not propagate changes.

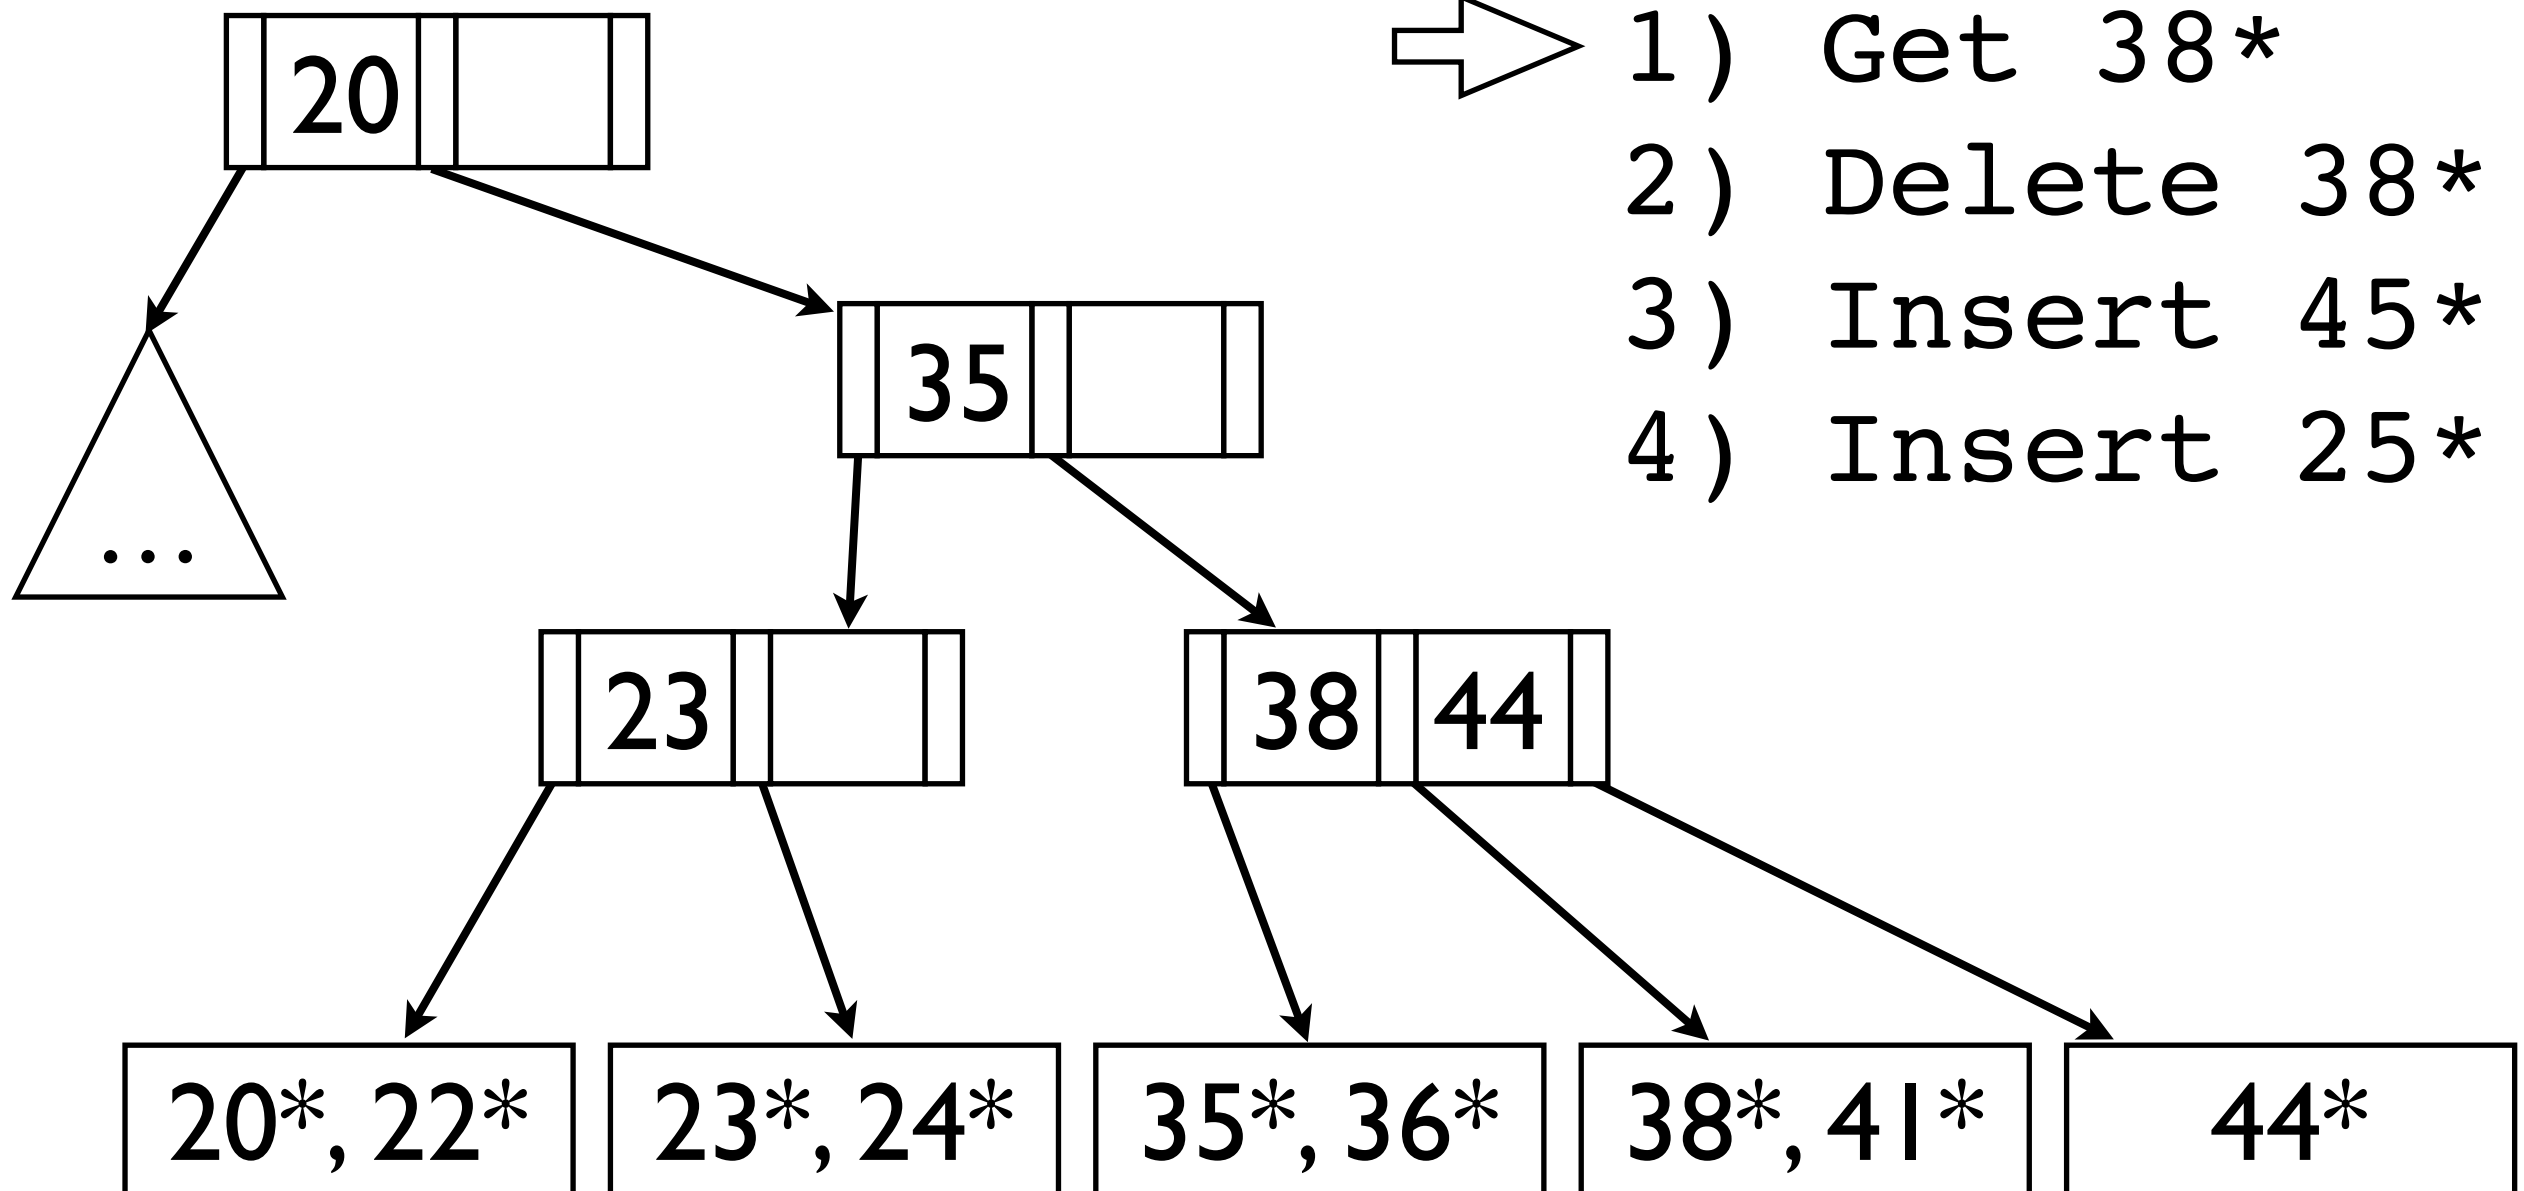
When is a node safe for single inserts? single deletes?

Examples



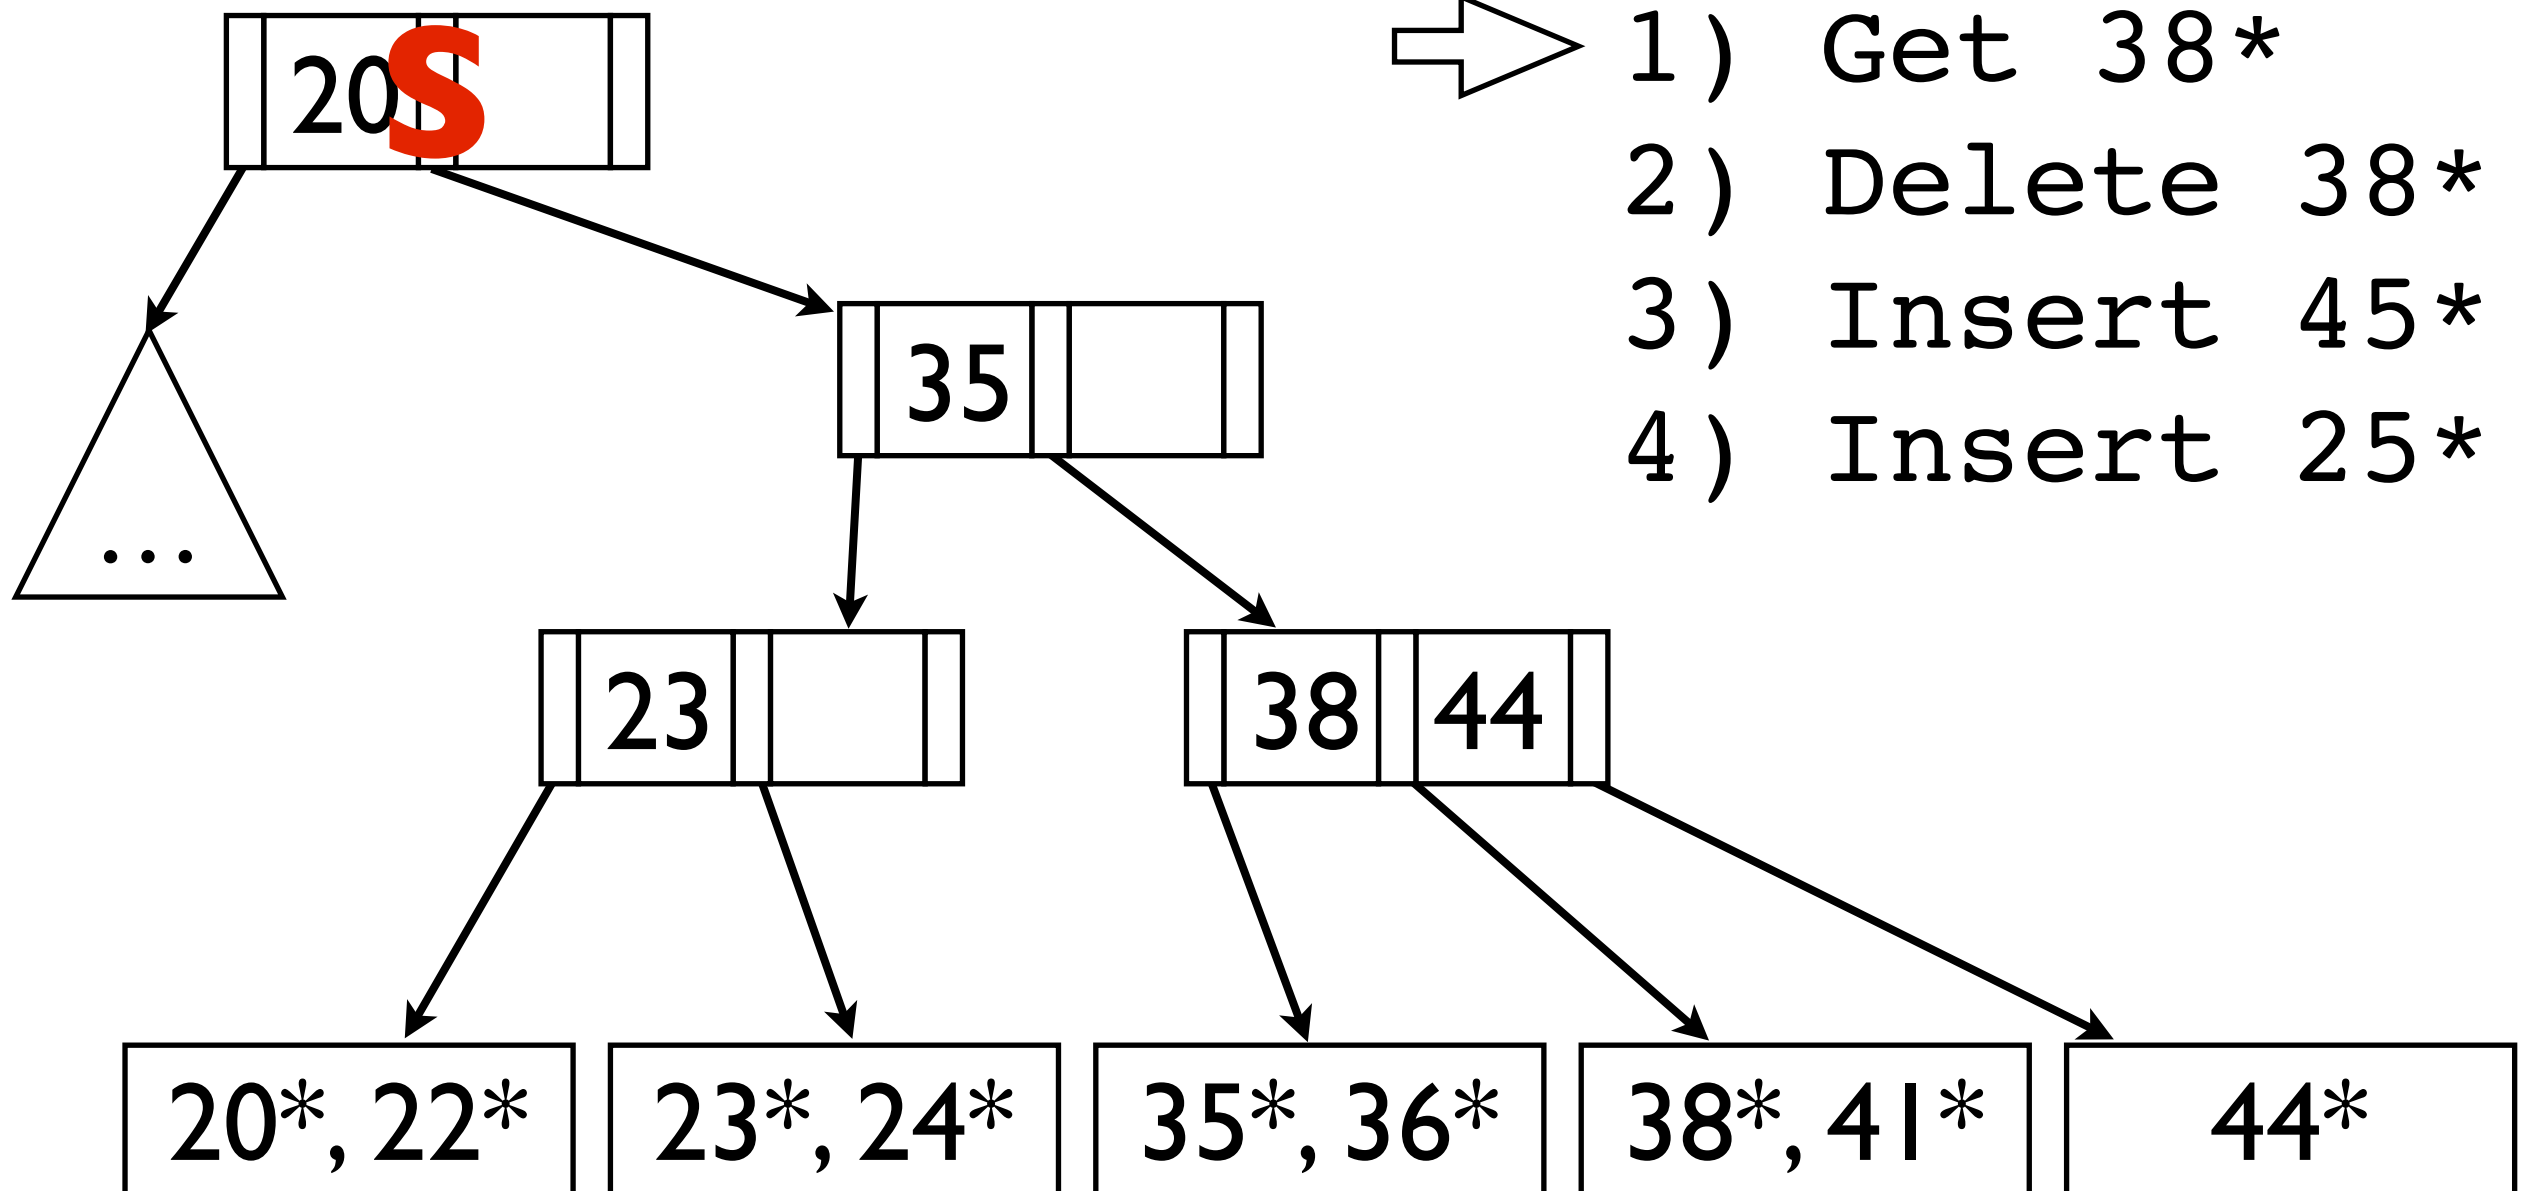
Examples

- ⇒
- 1) Get 38*
 - 2) Delete 38*
 - 3) Insert 45*
 - 4) Insert 25*



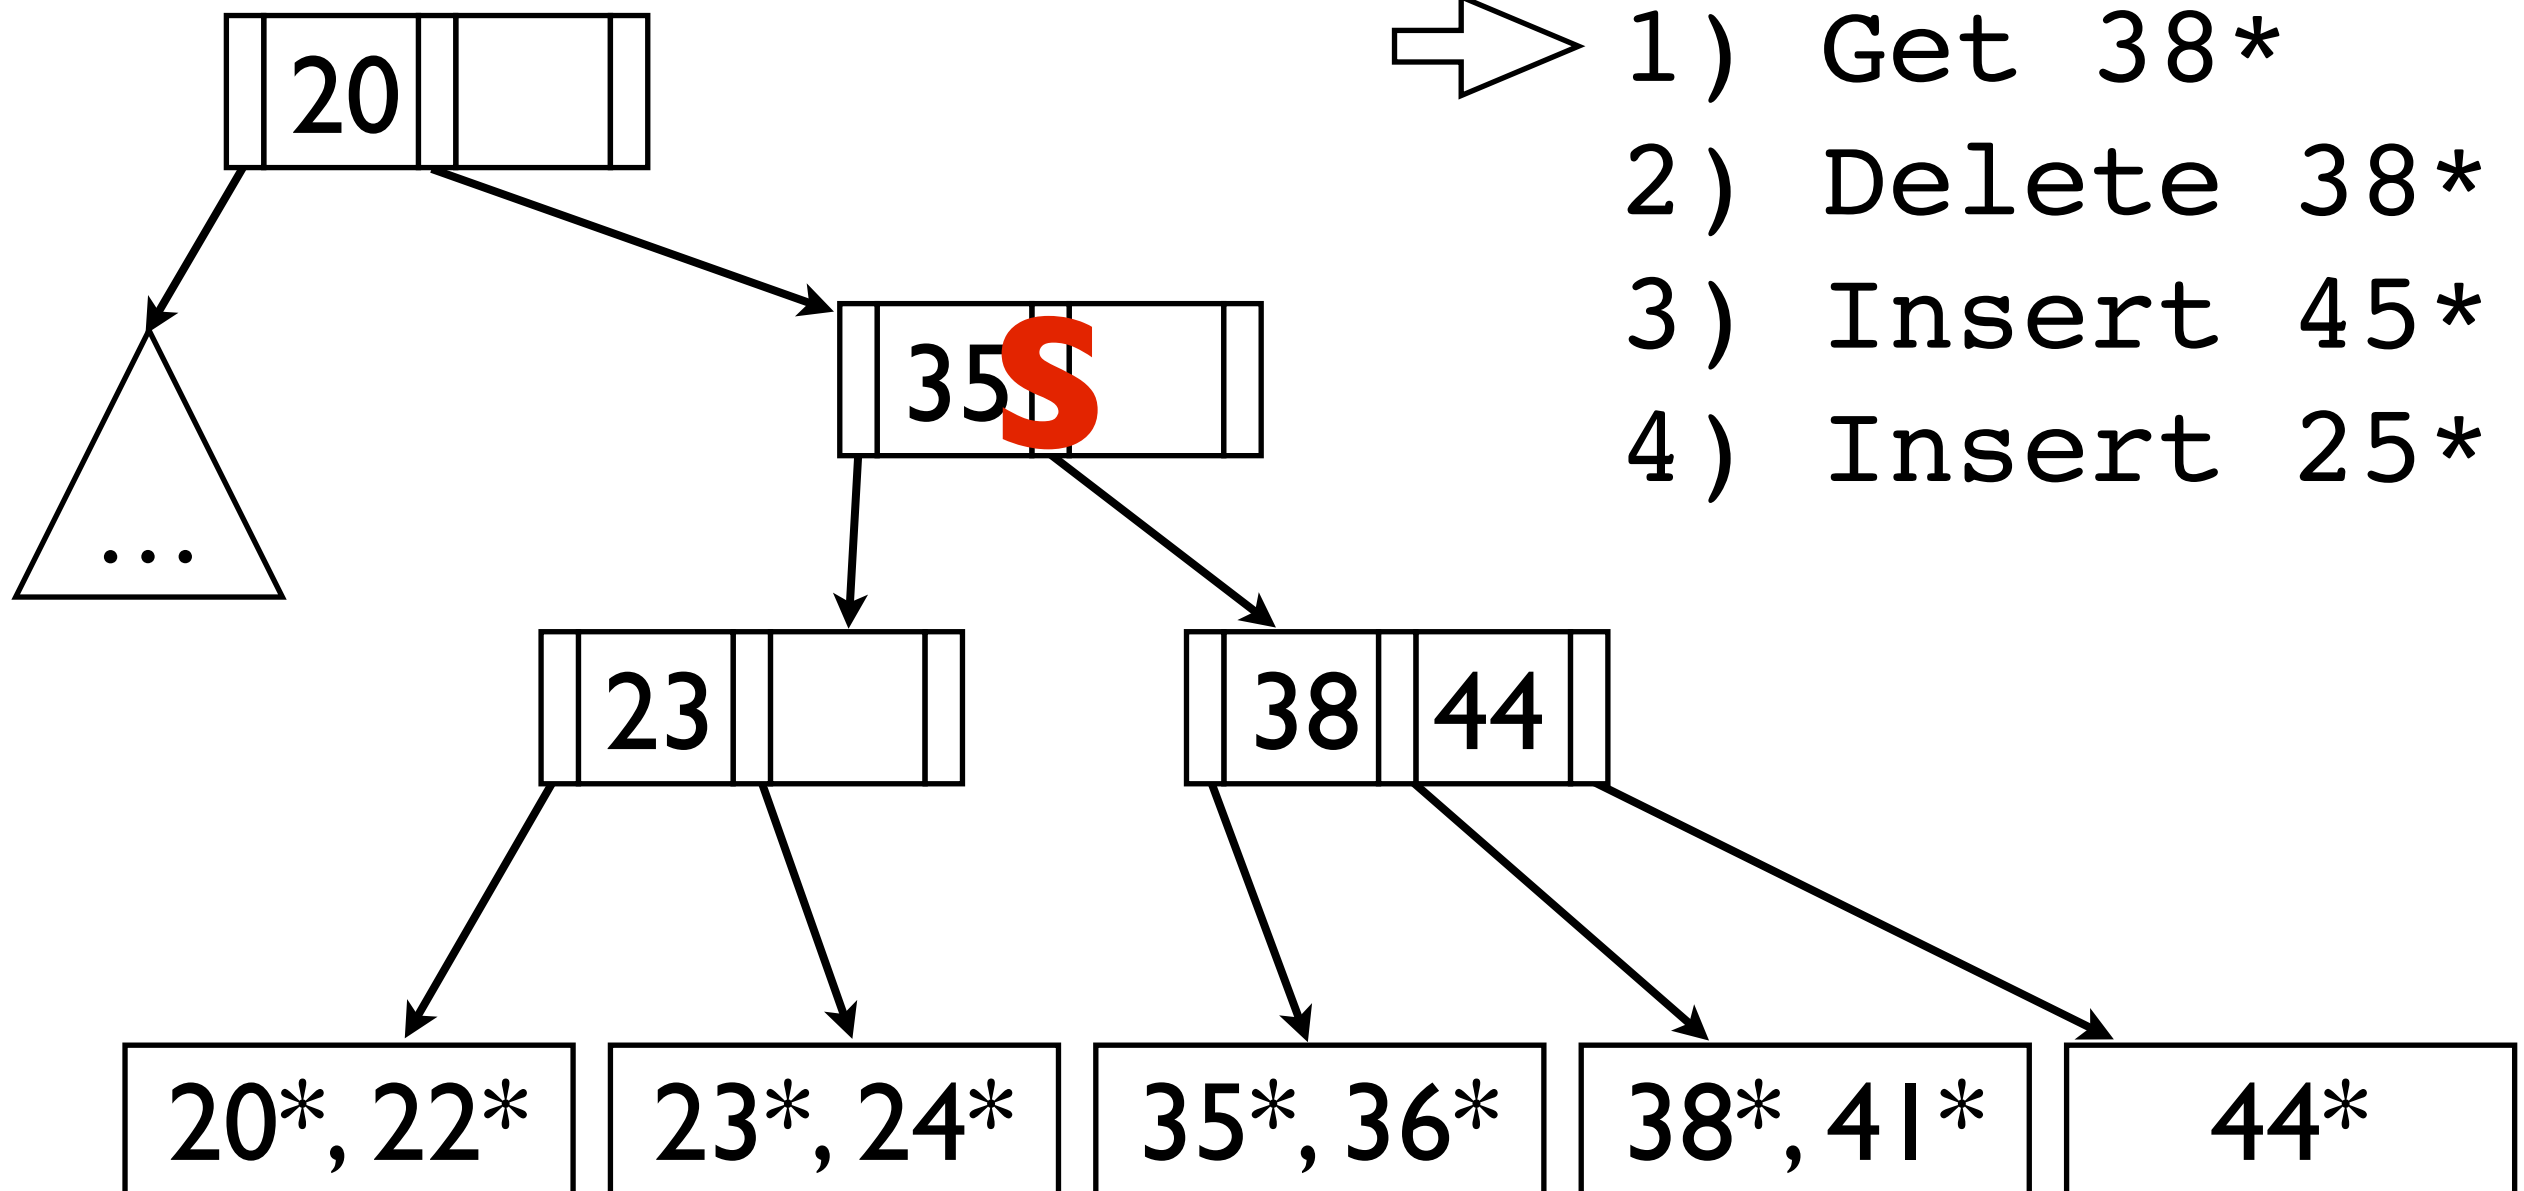
Examples

- ⇒
- 1) Get 38*
 - 2) Delete 38*
 - 3) Insert 45*
 - 4) Insert 25*




Examples

- ⇒
- 1) Get 38*
 - 2) Delete 38*
 - 3) Insert 45*
 - 4) Insert 25*

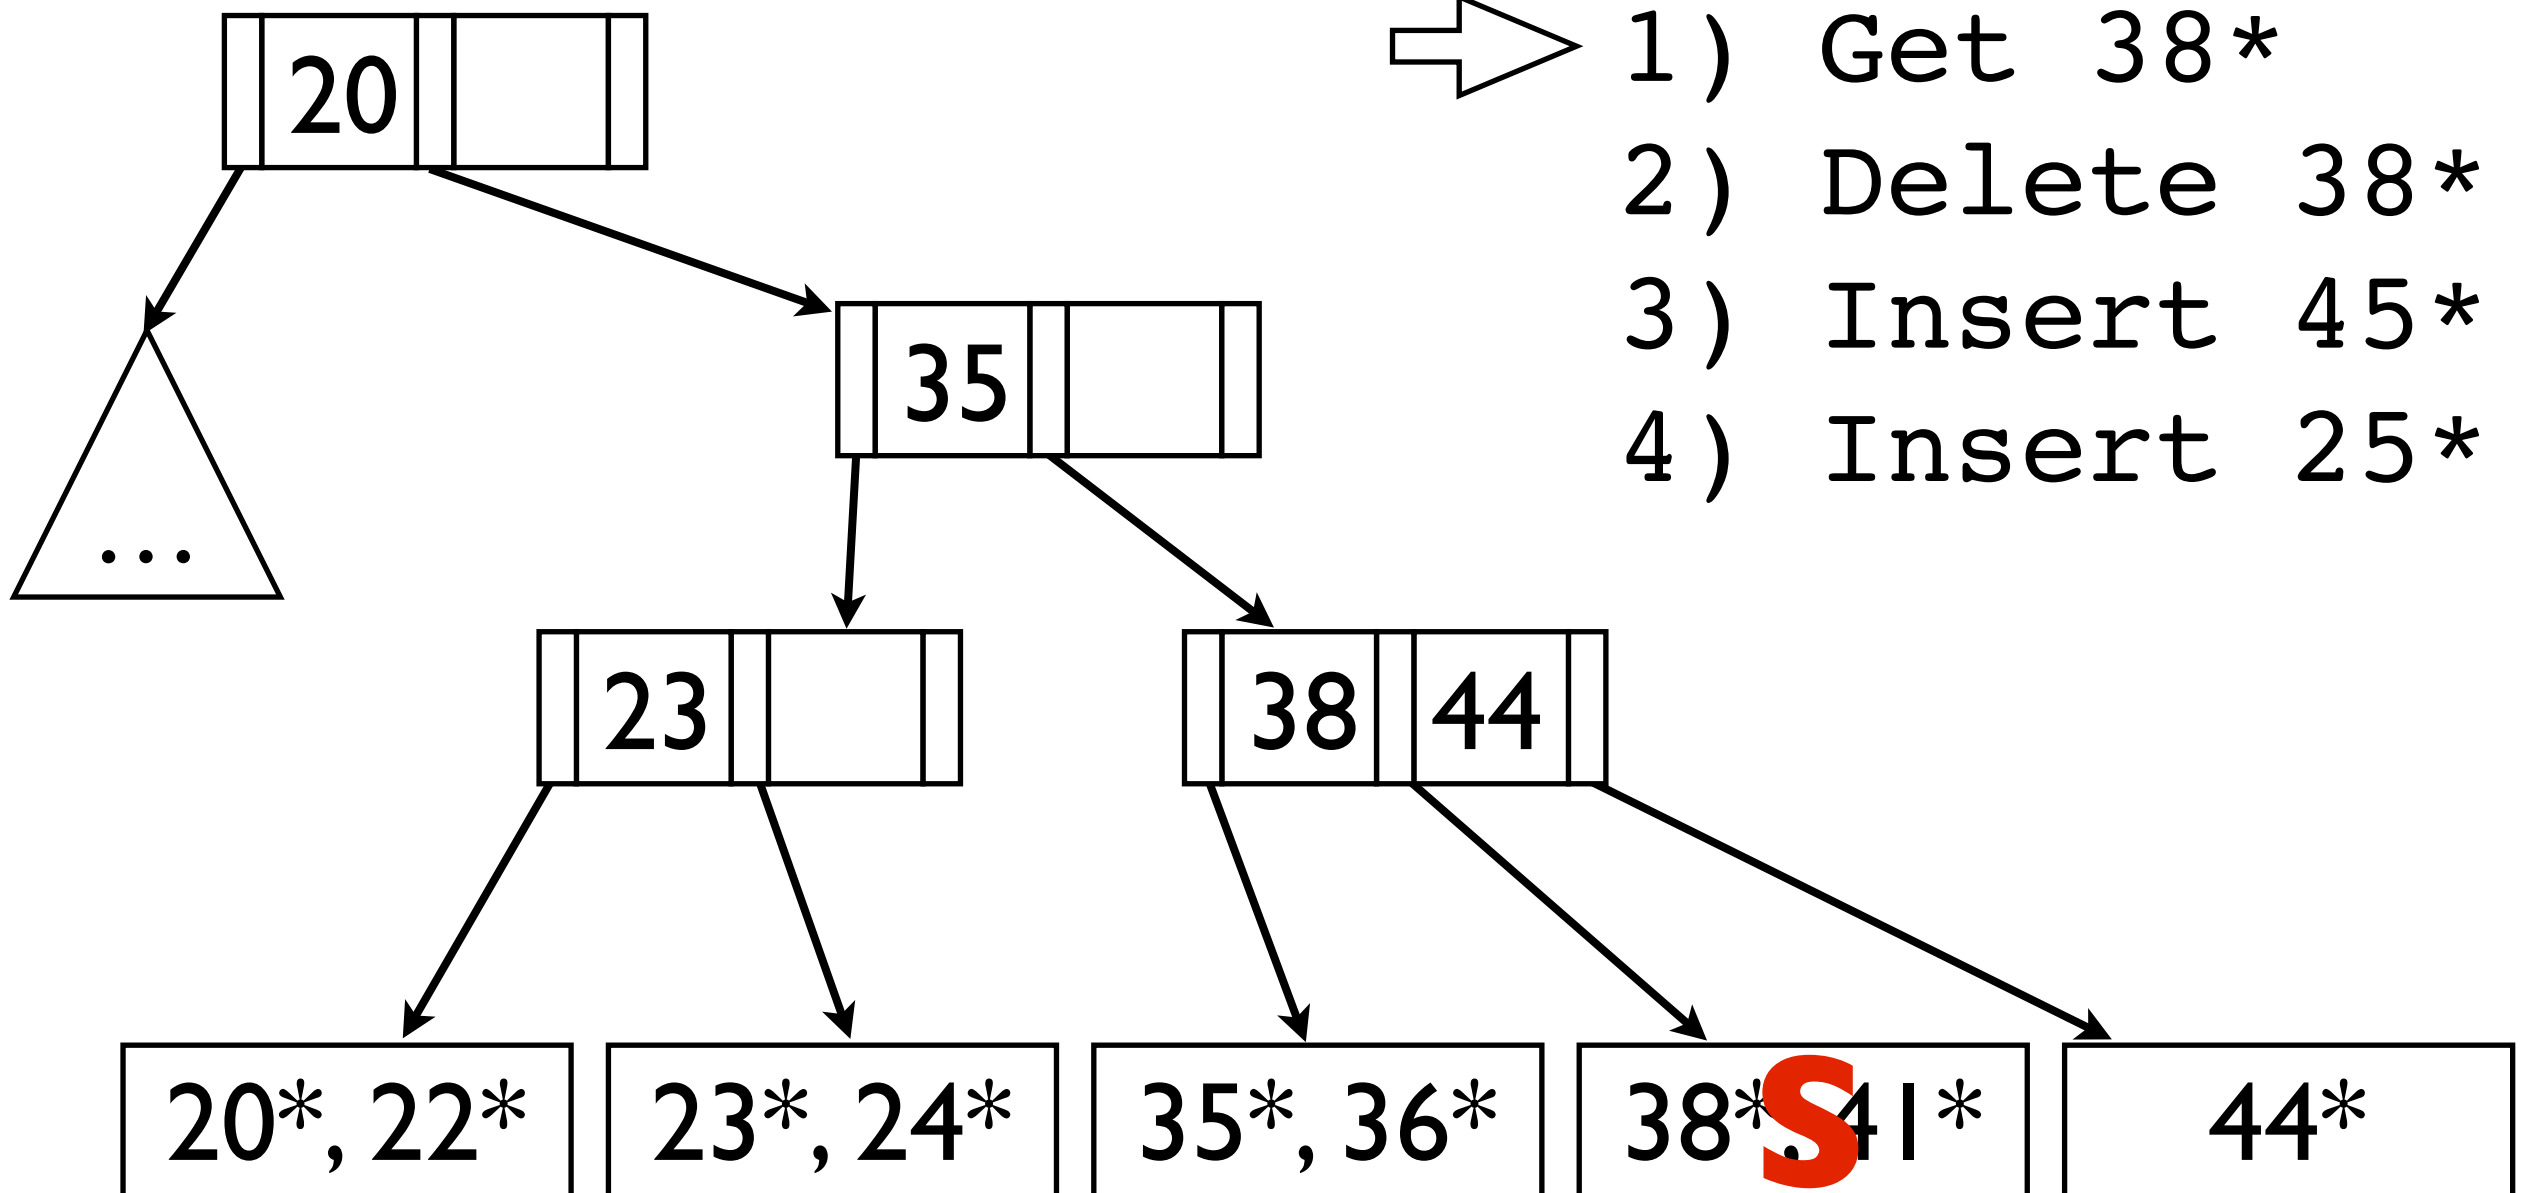


Examples


- ⇒
- 1) Get 38*
 - 2) Delete 38*
 - 3) Insert 45*
 - 4) Insert 25*




Examples



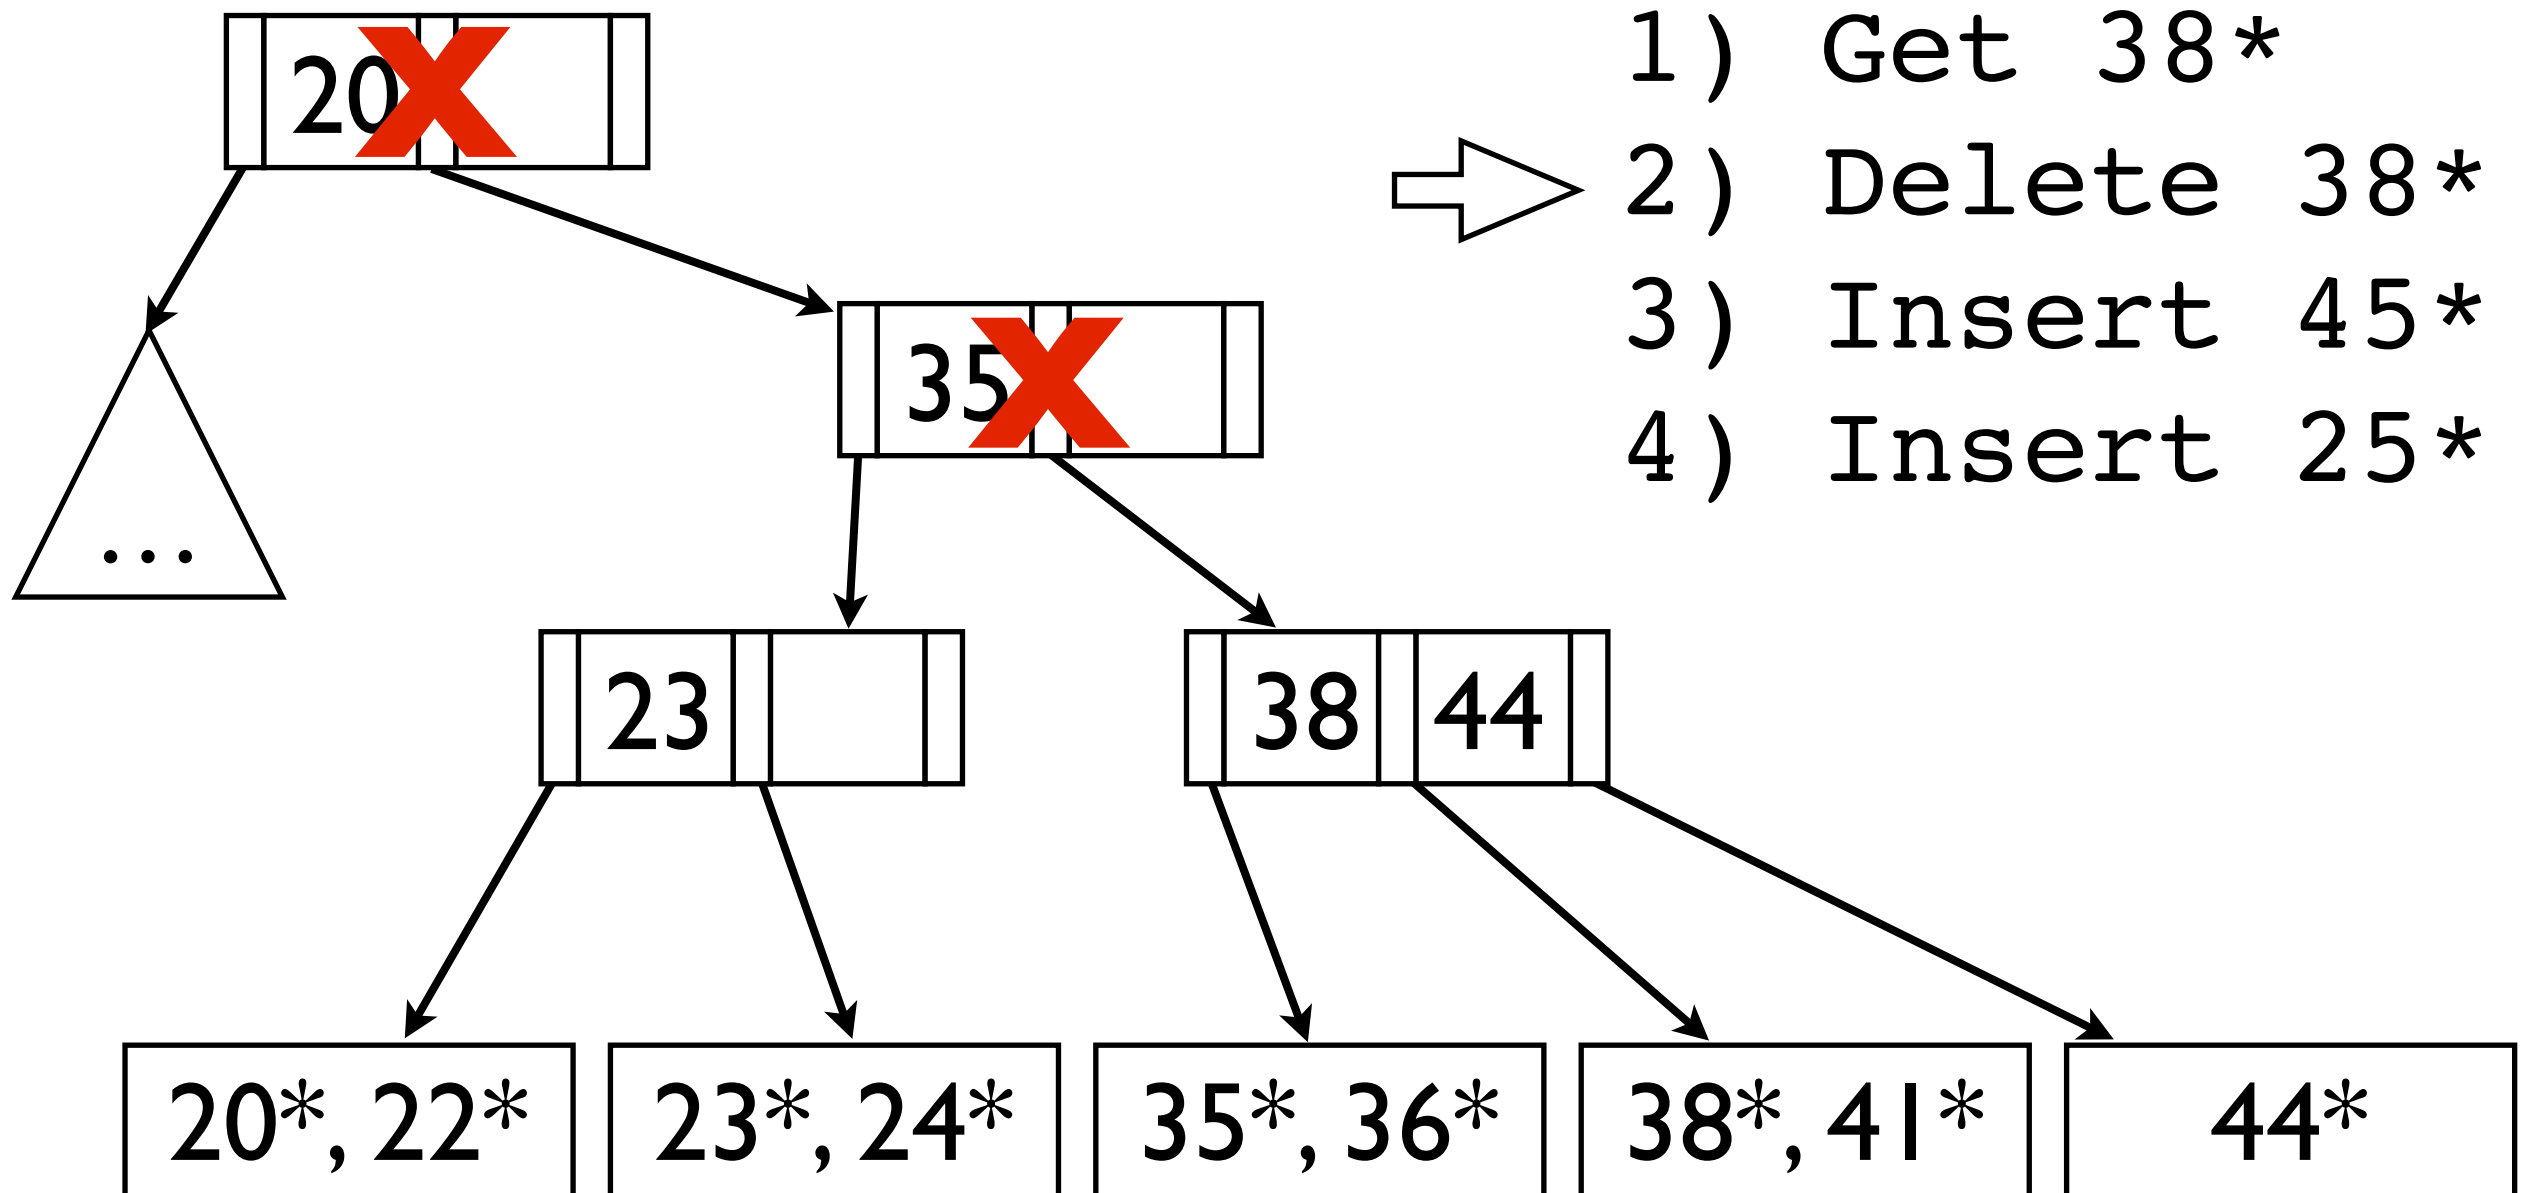
Examples




Examples



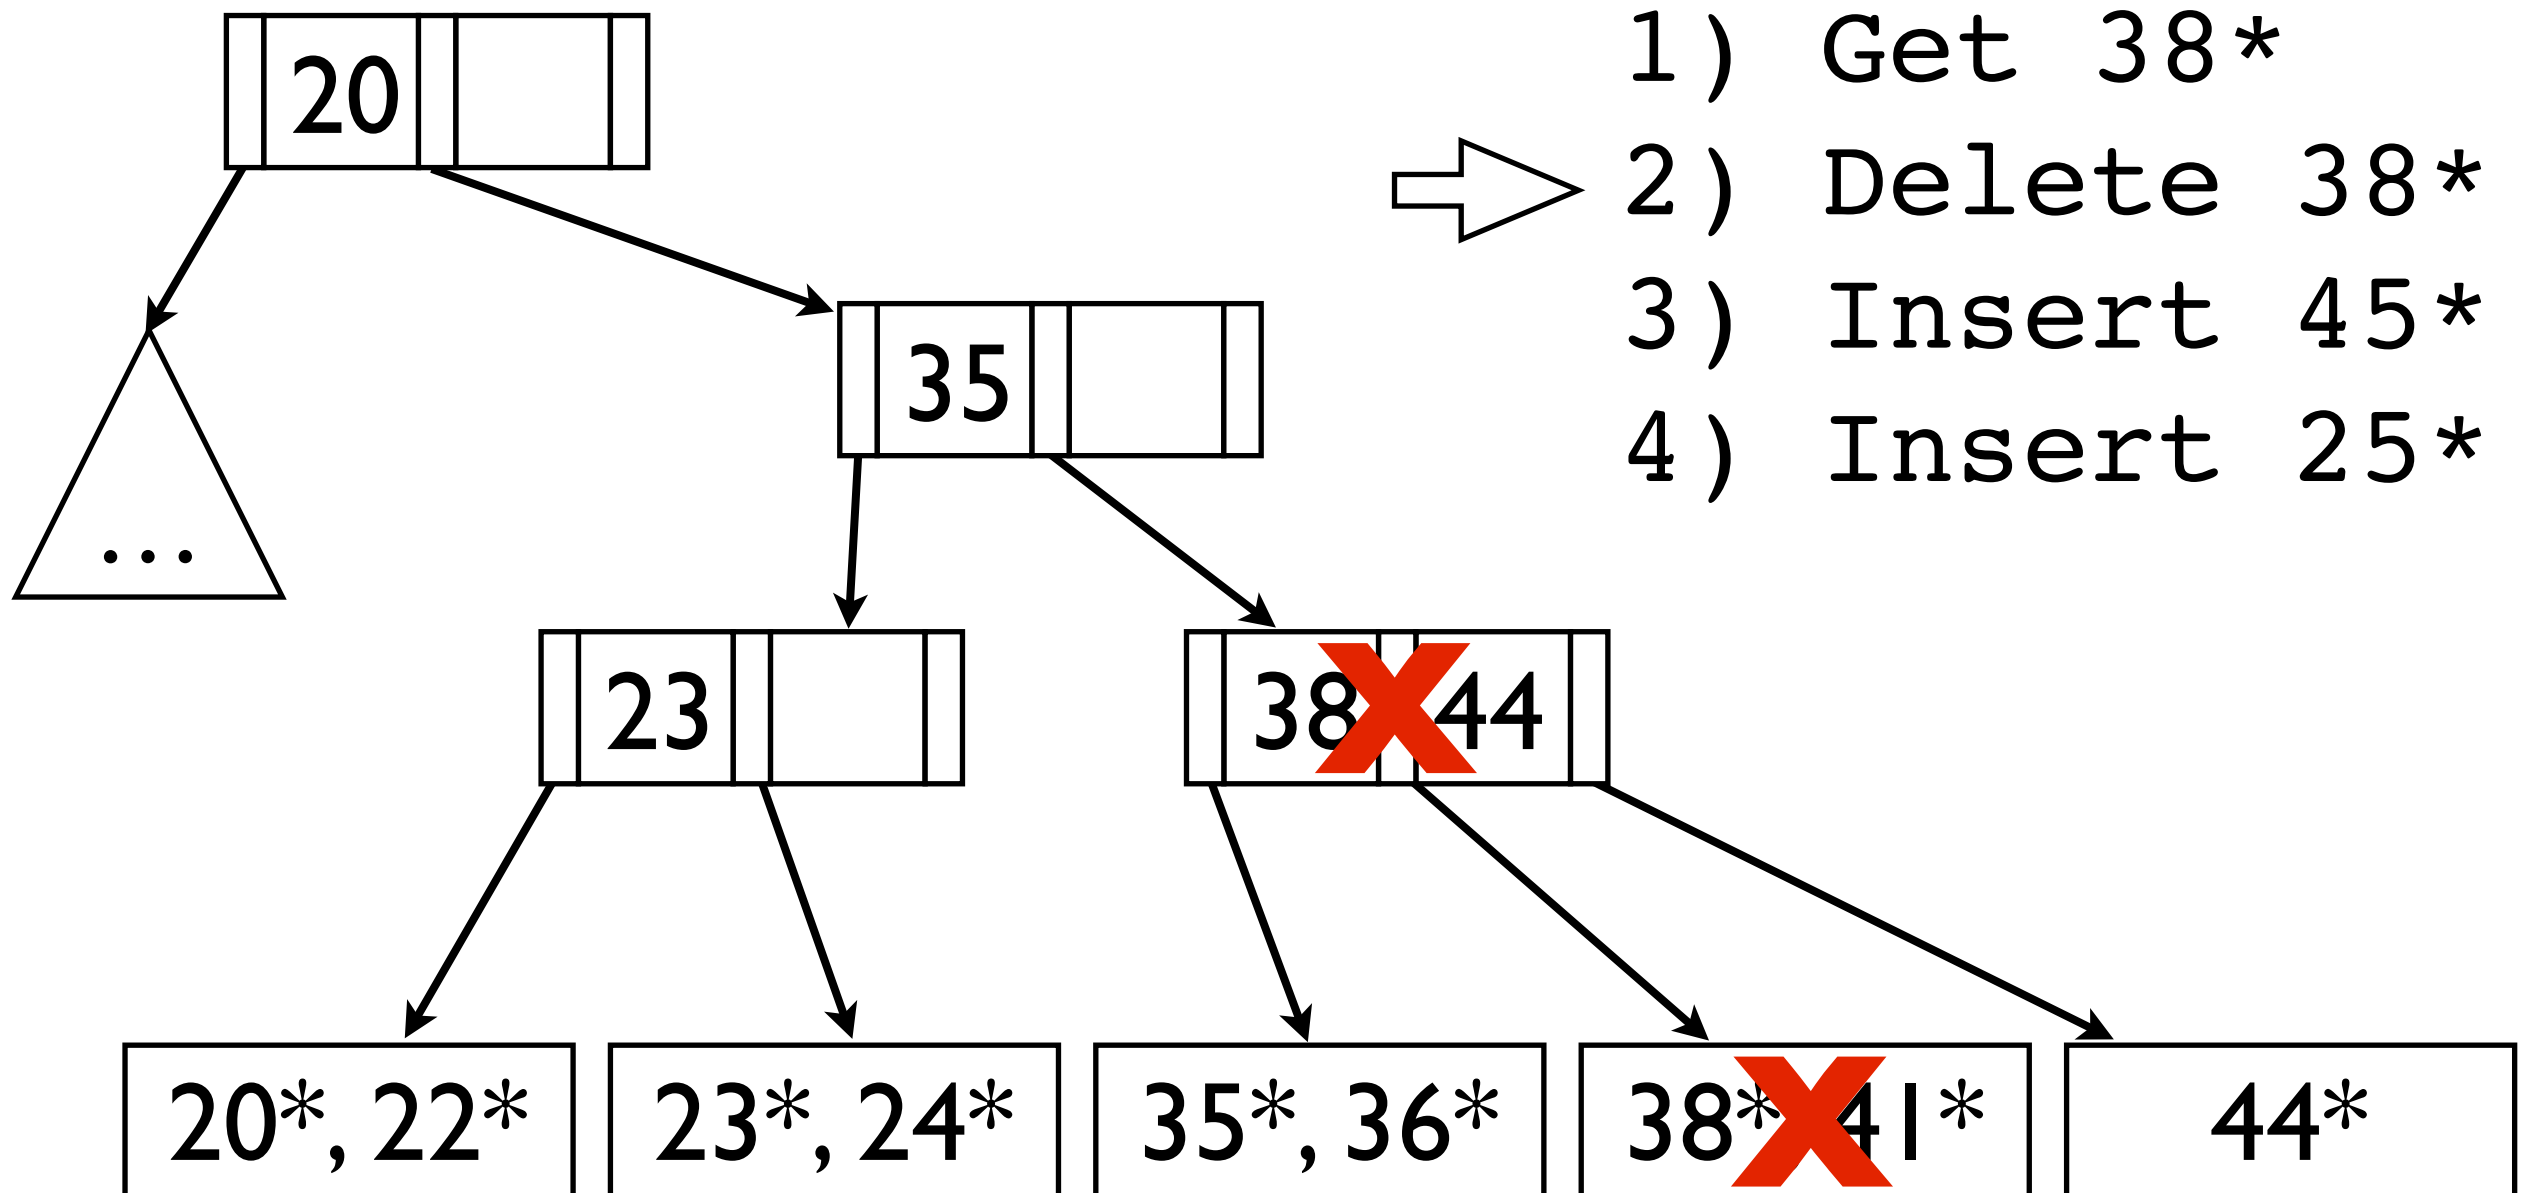
Examples



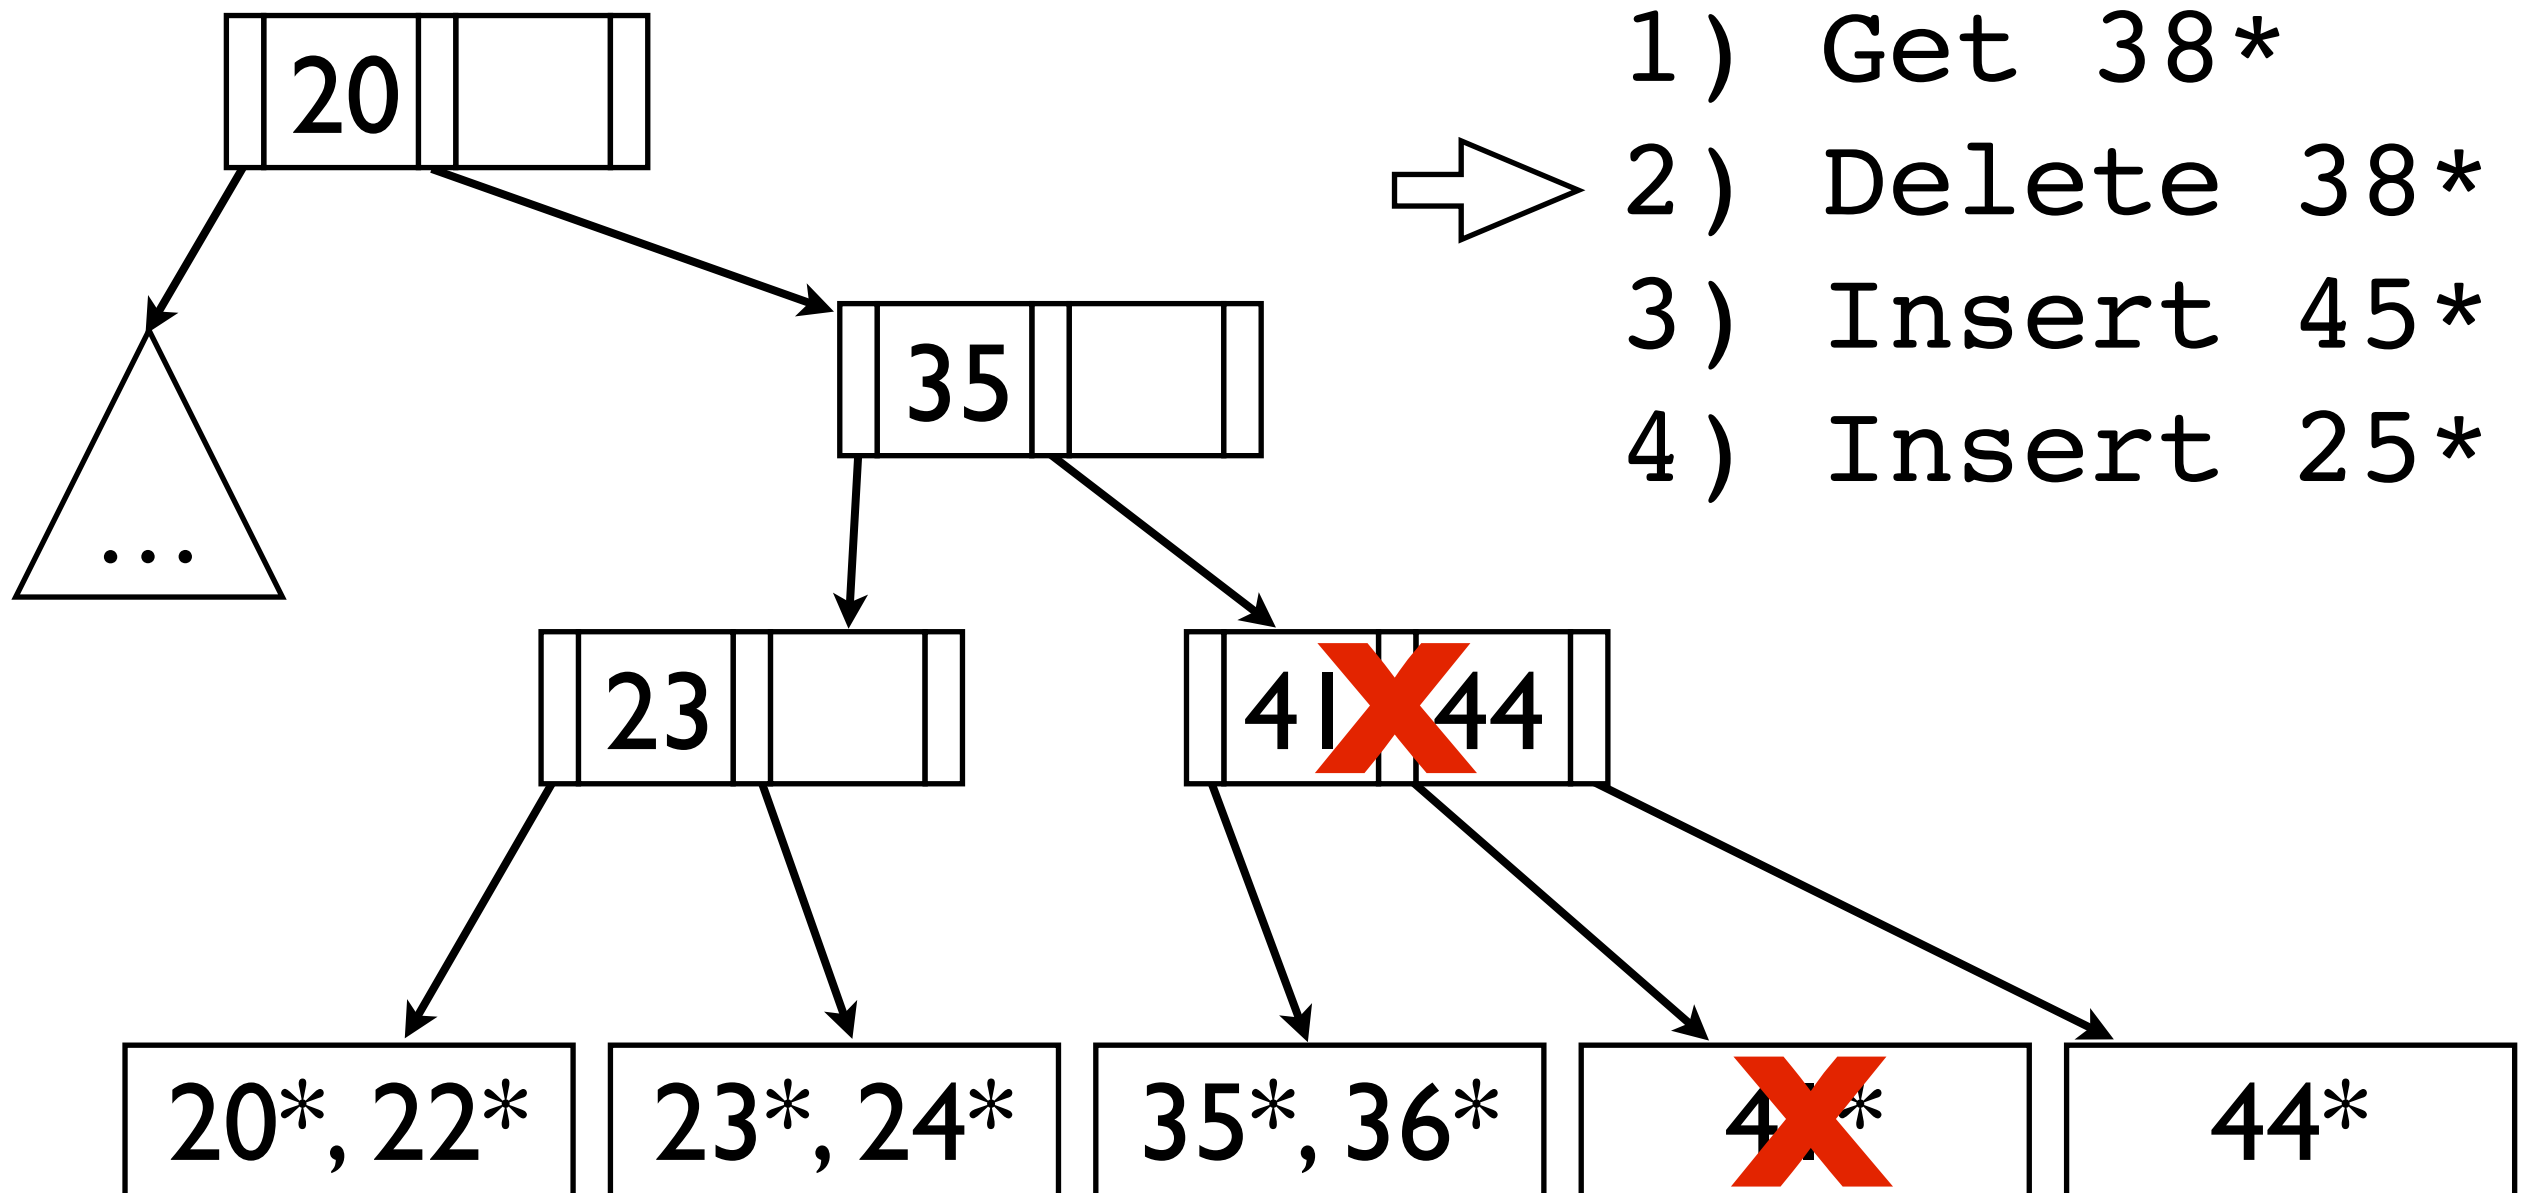
Examples



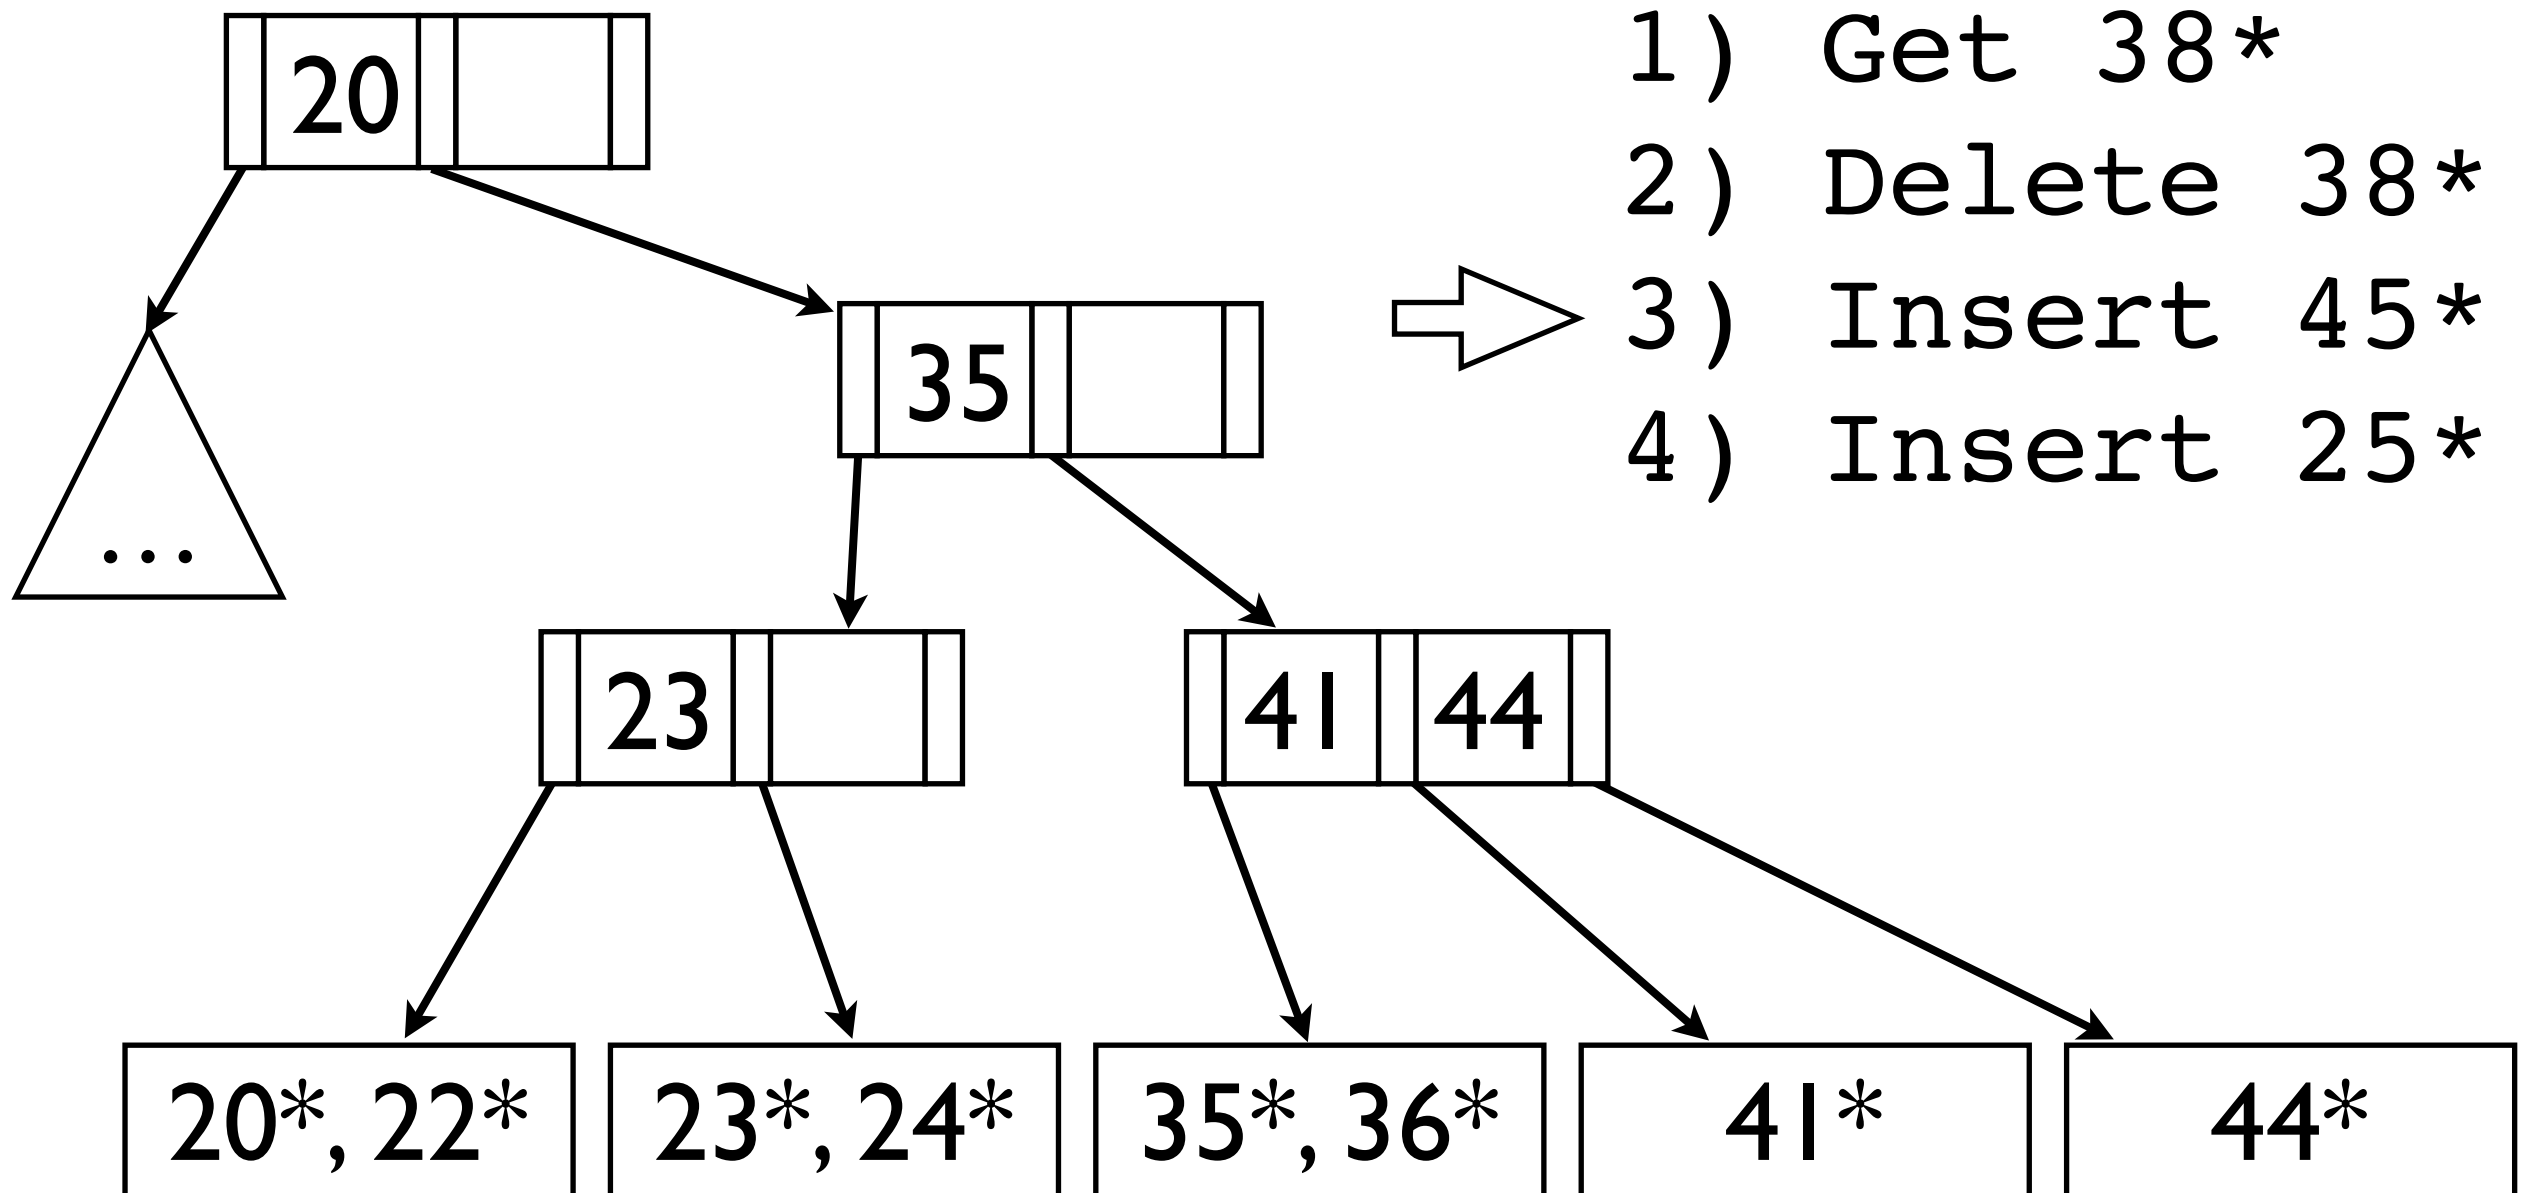
Examples



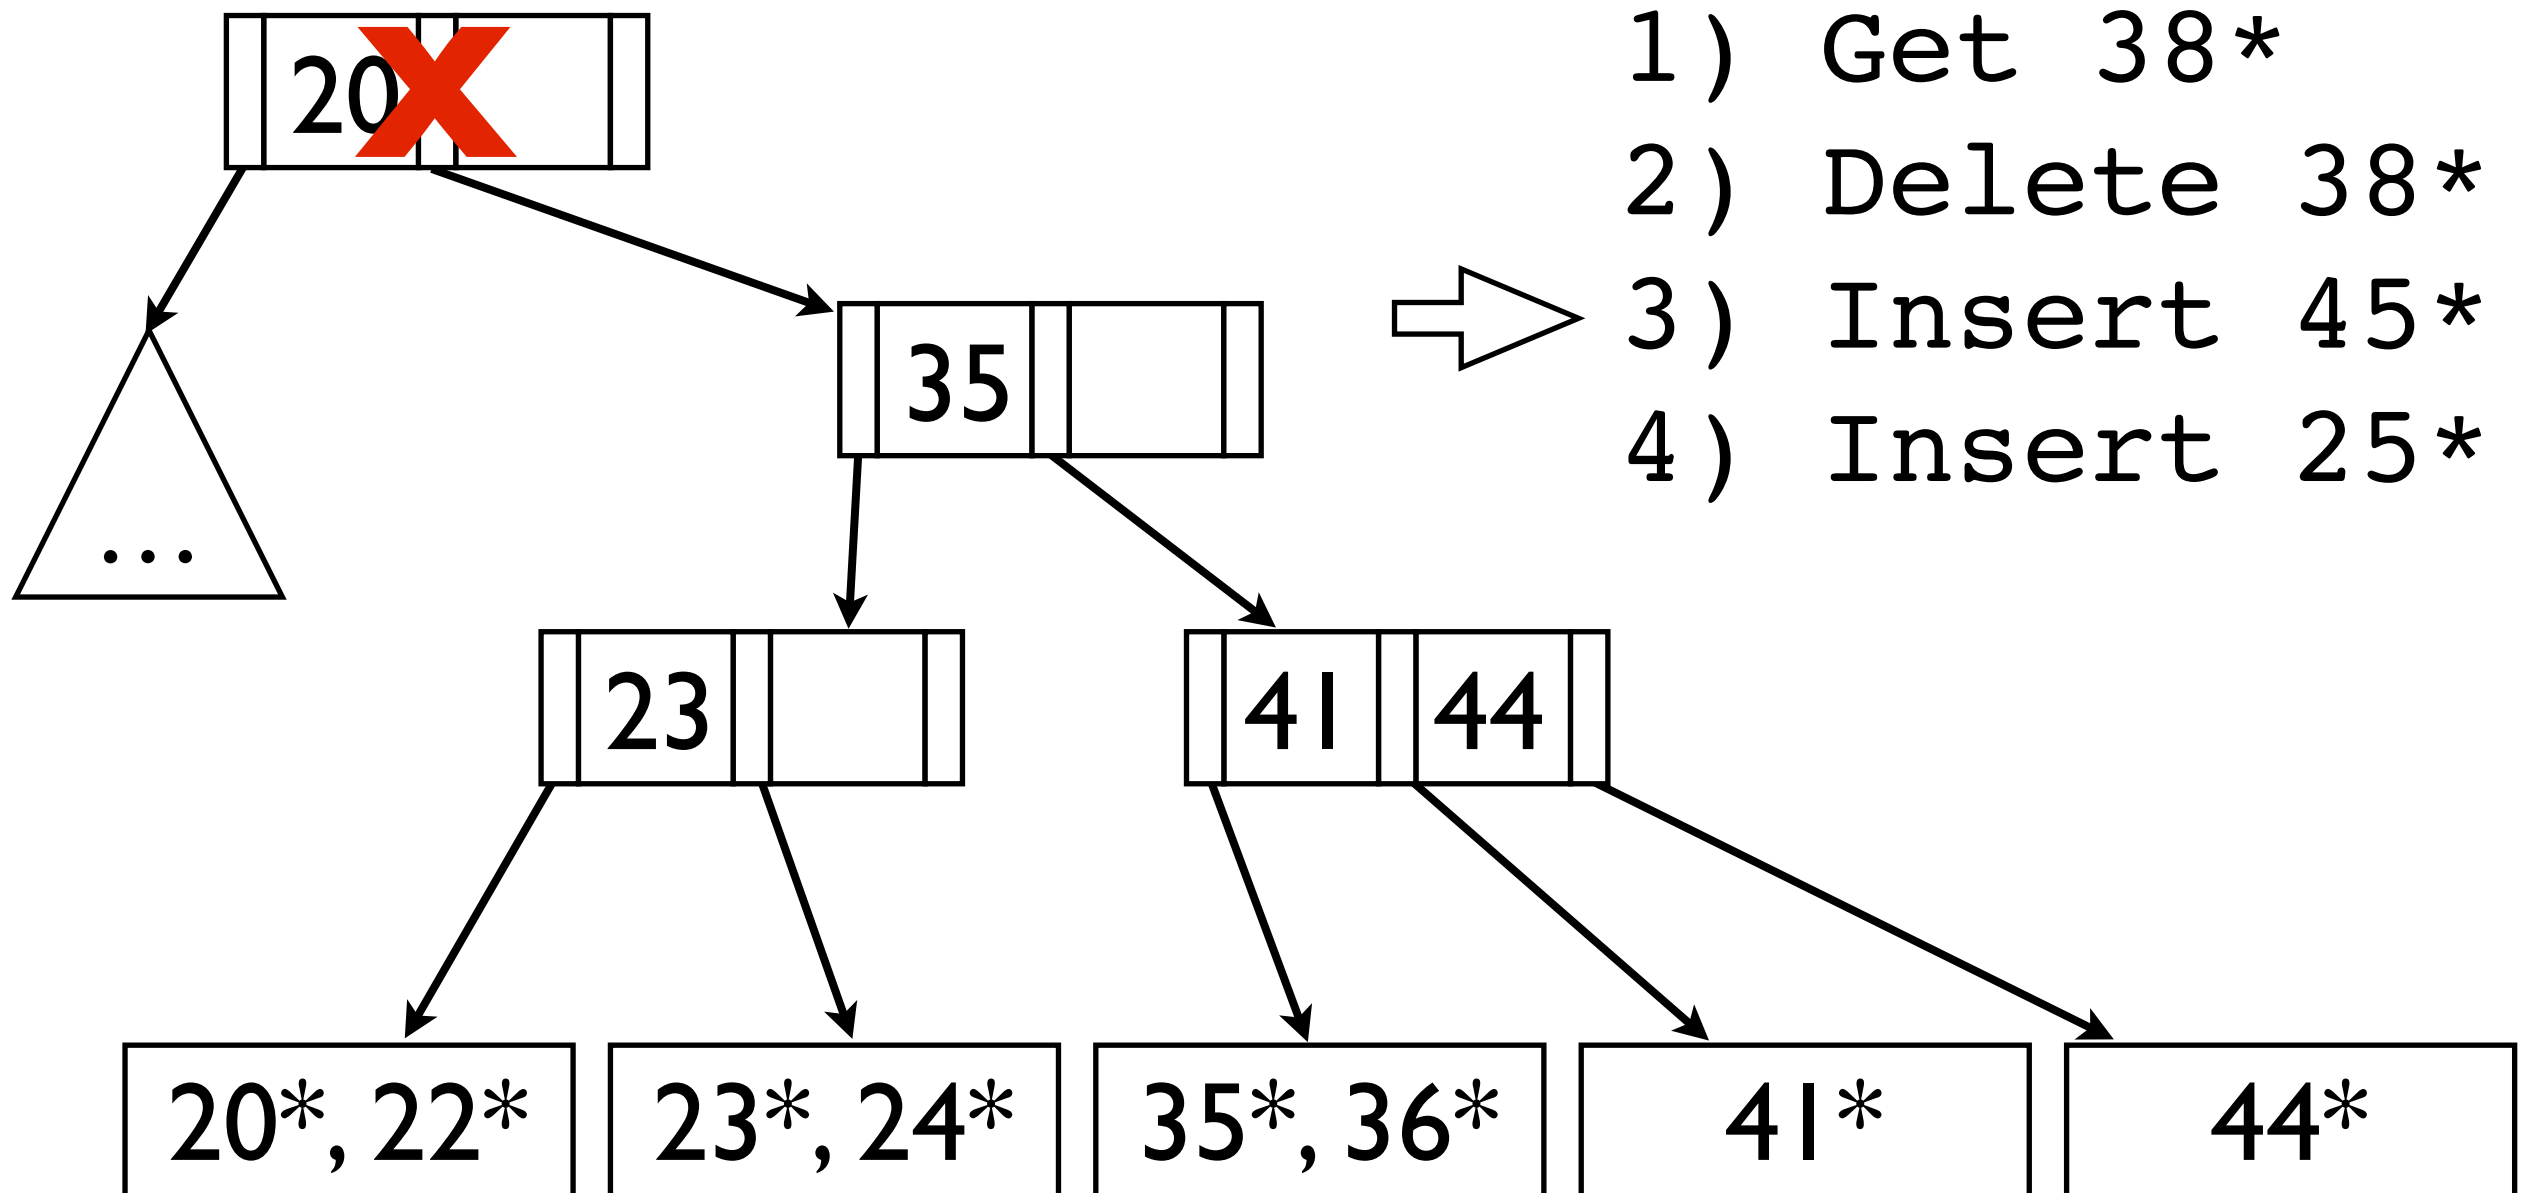
Examples



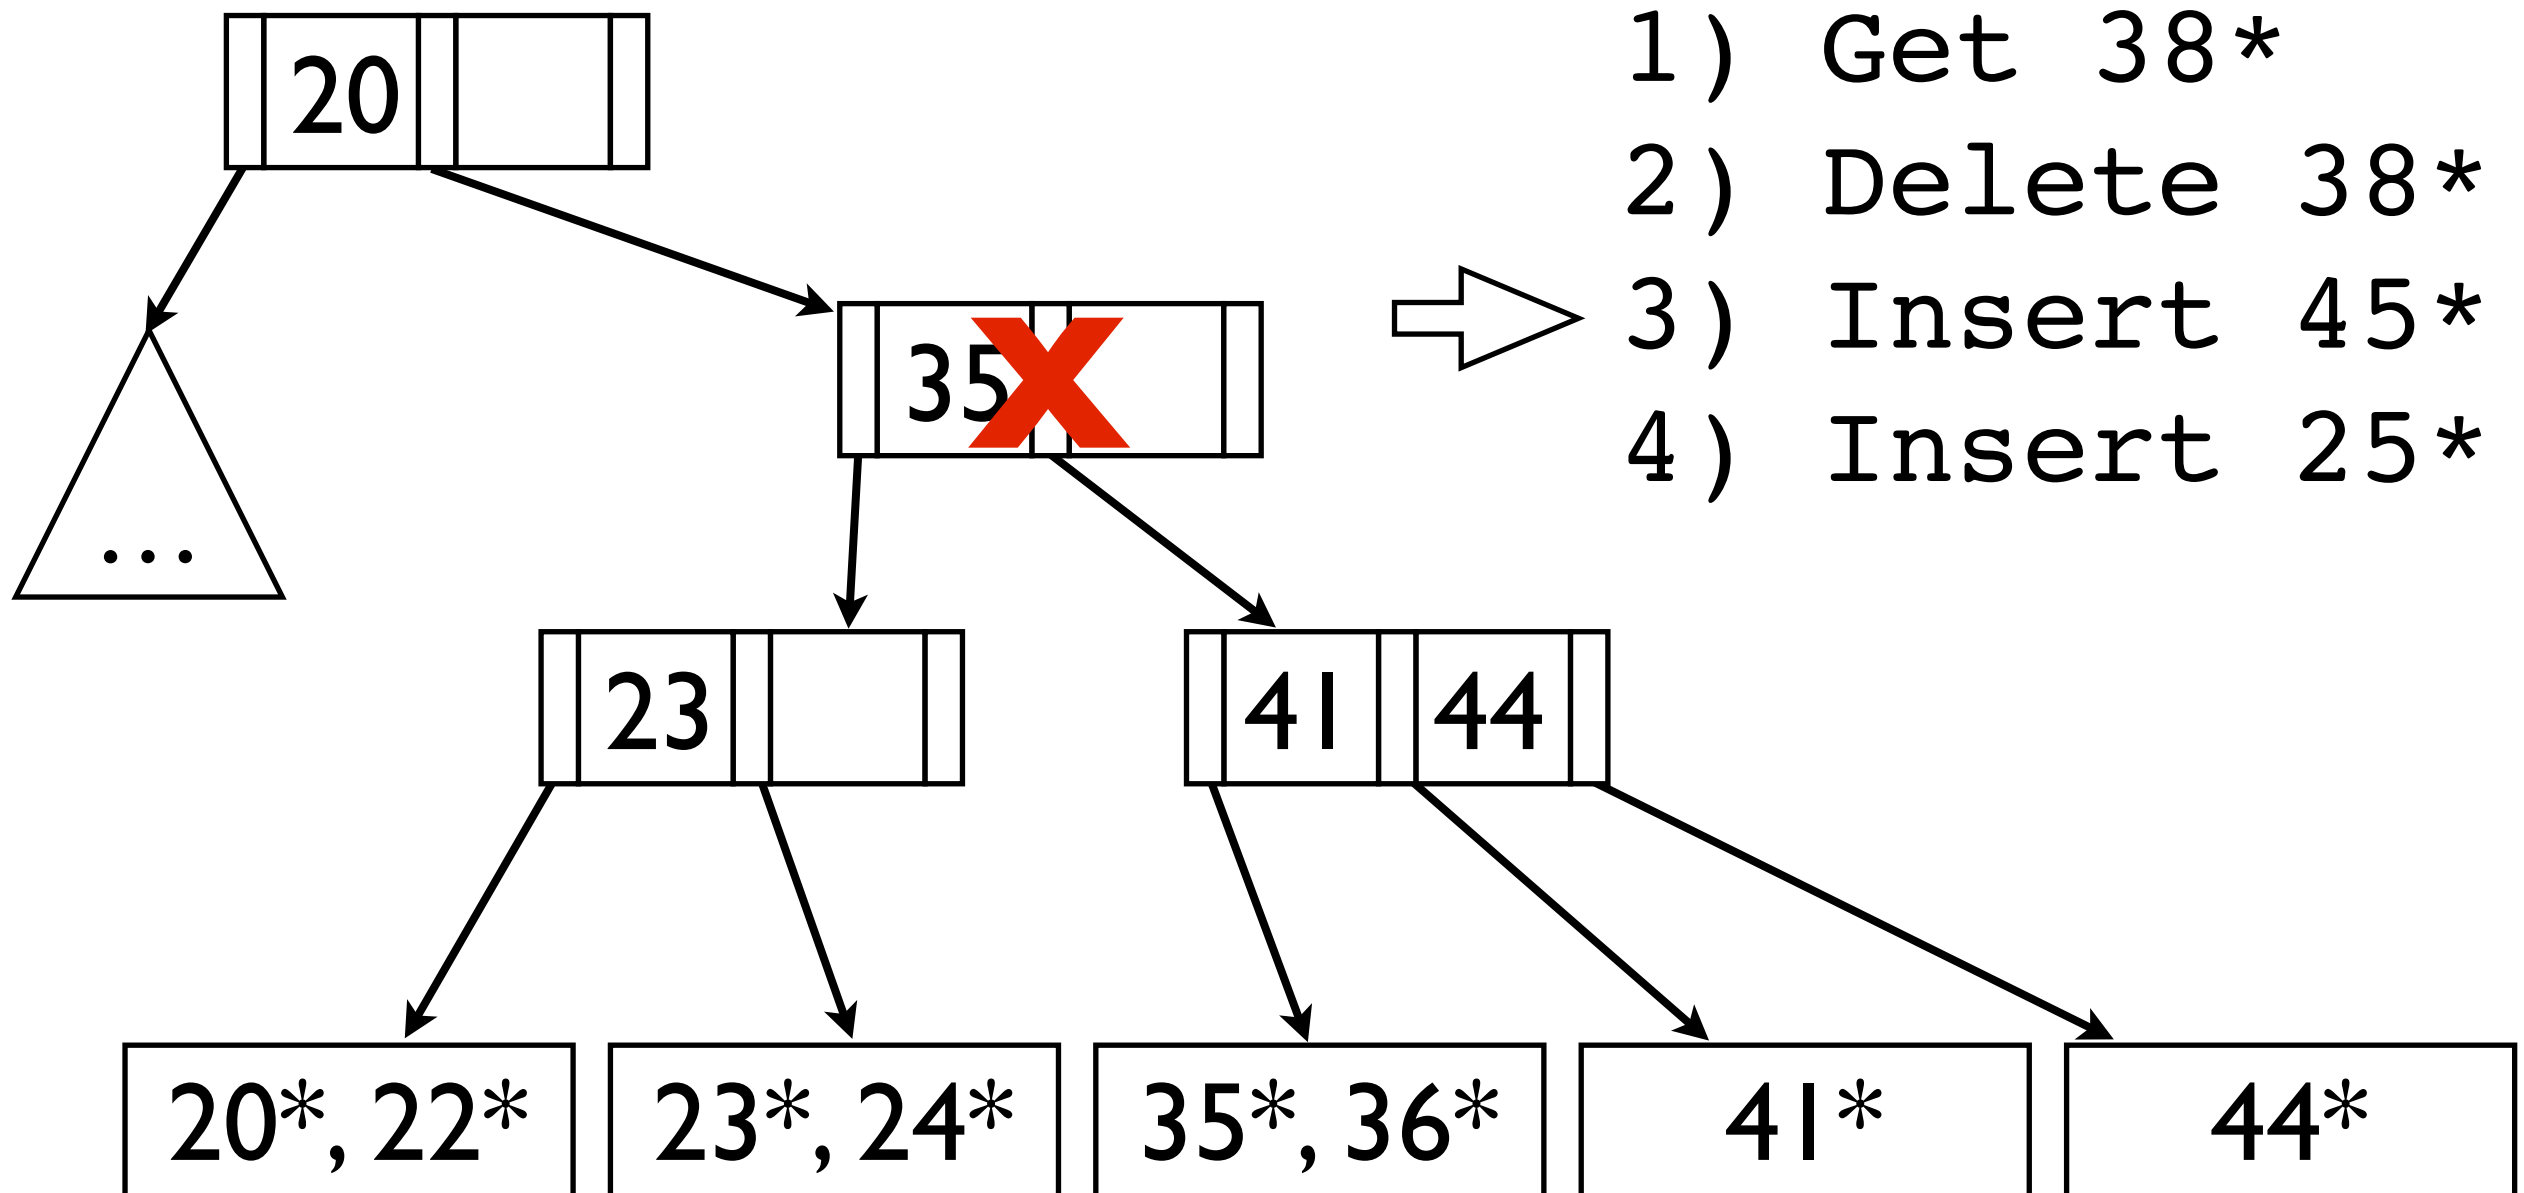
Examples



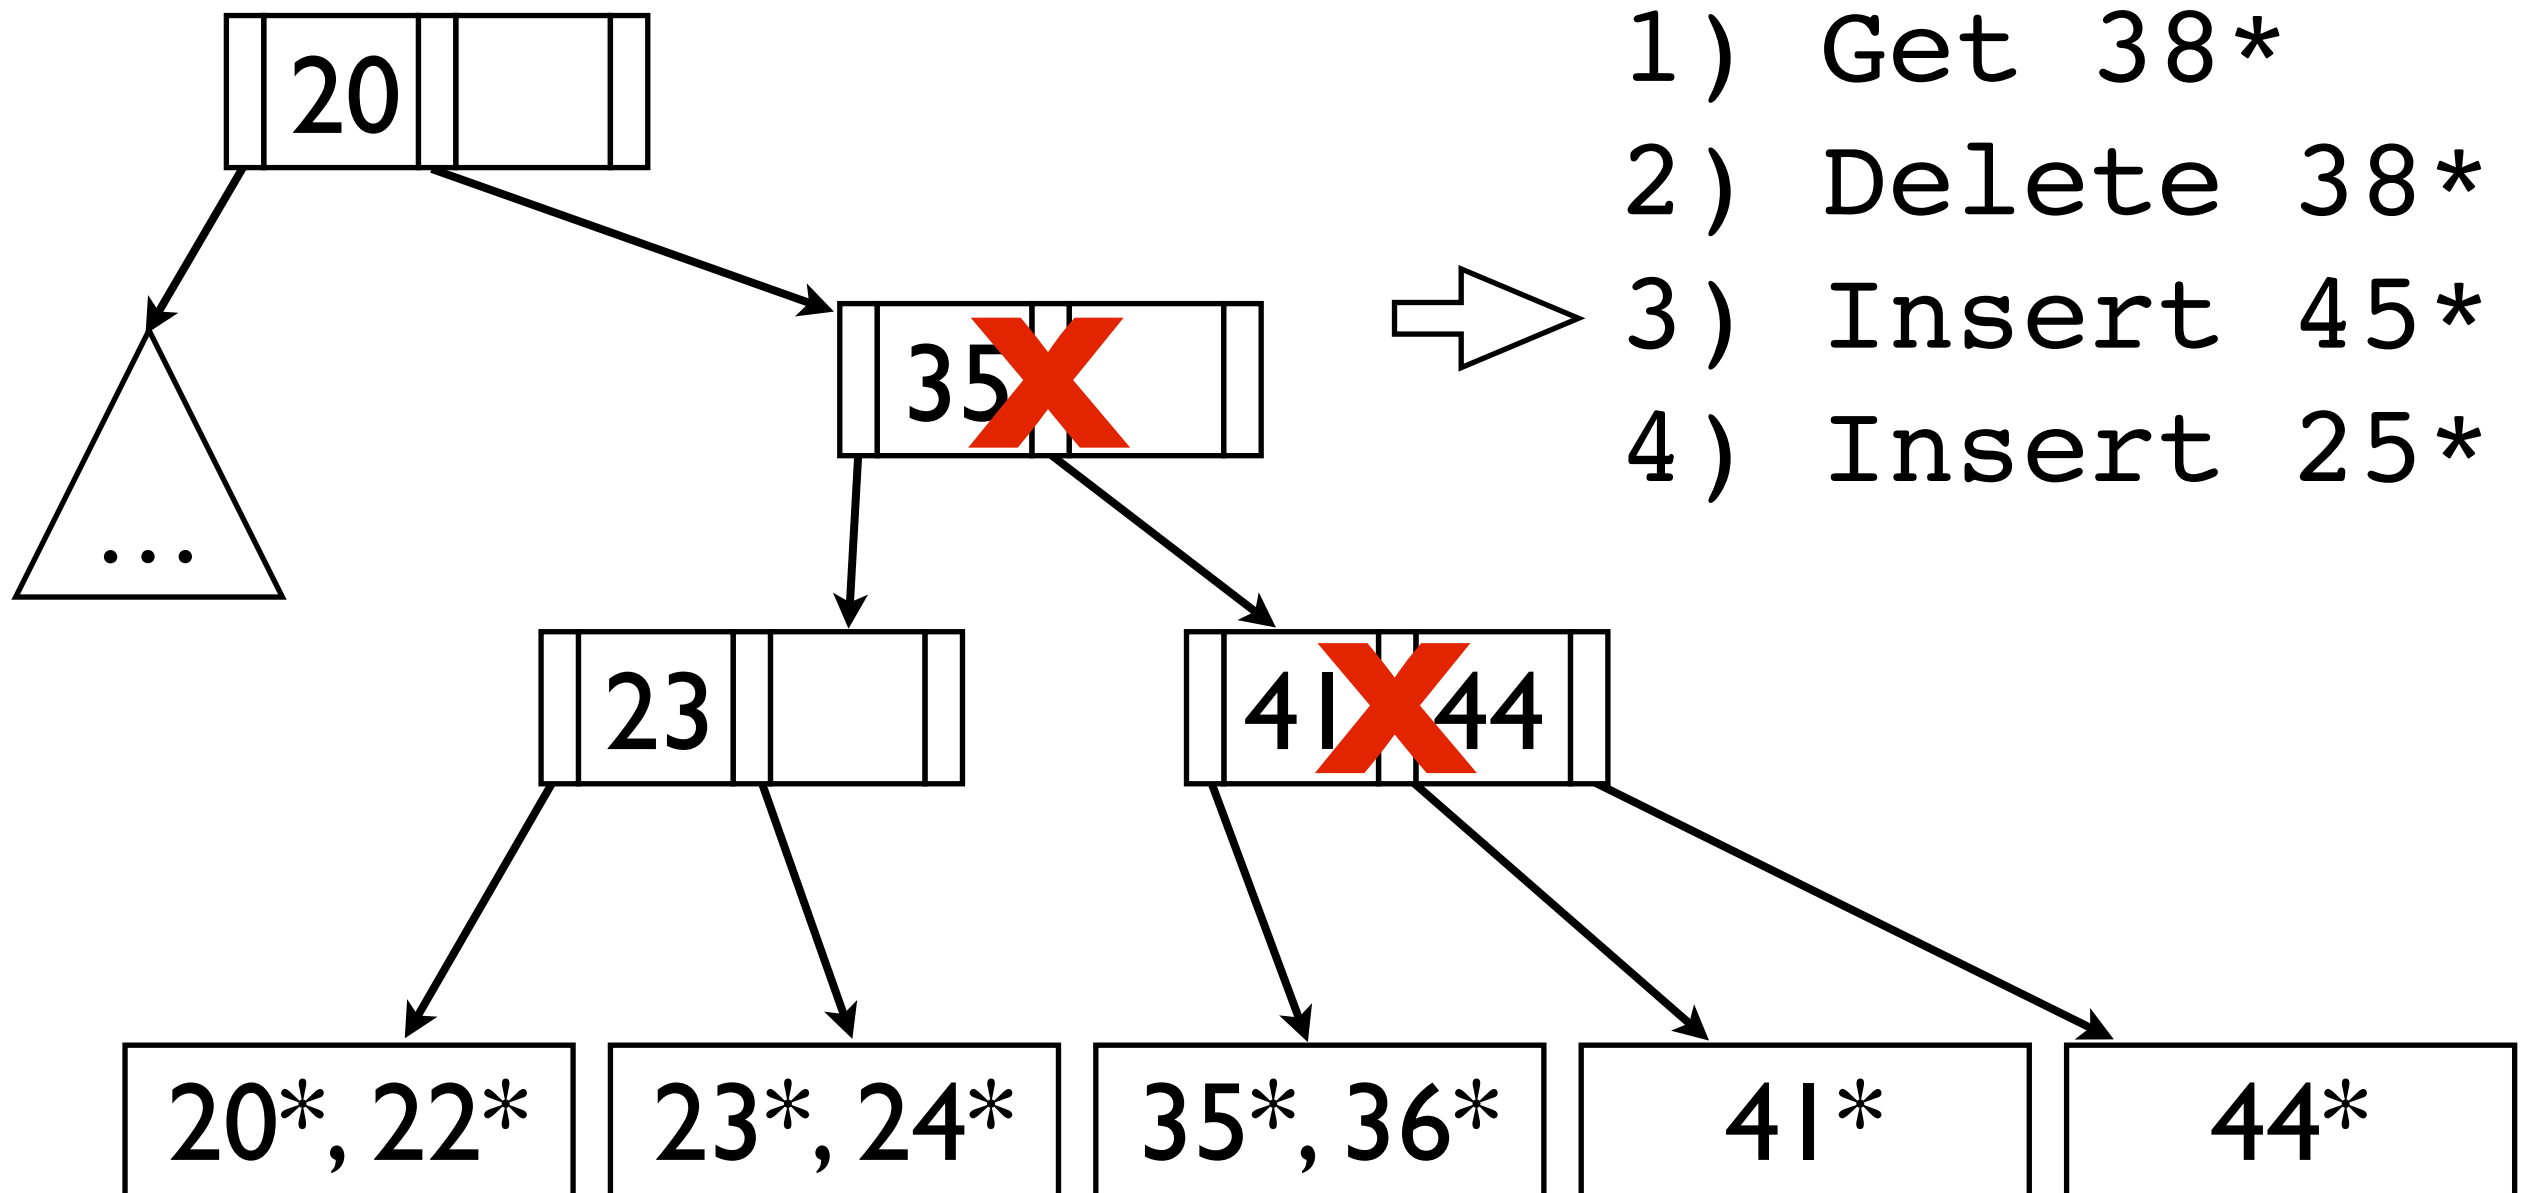
Examples



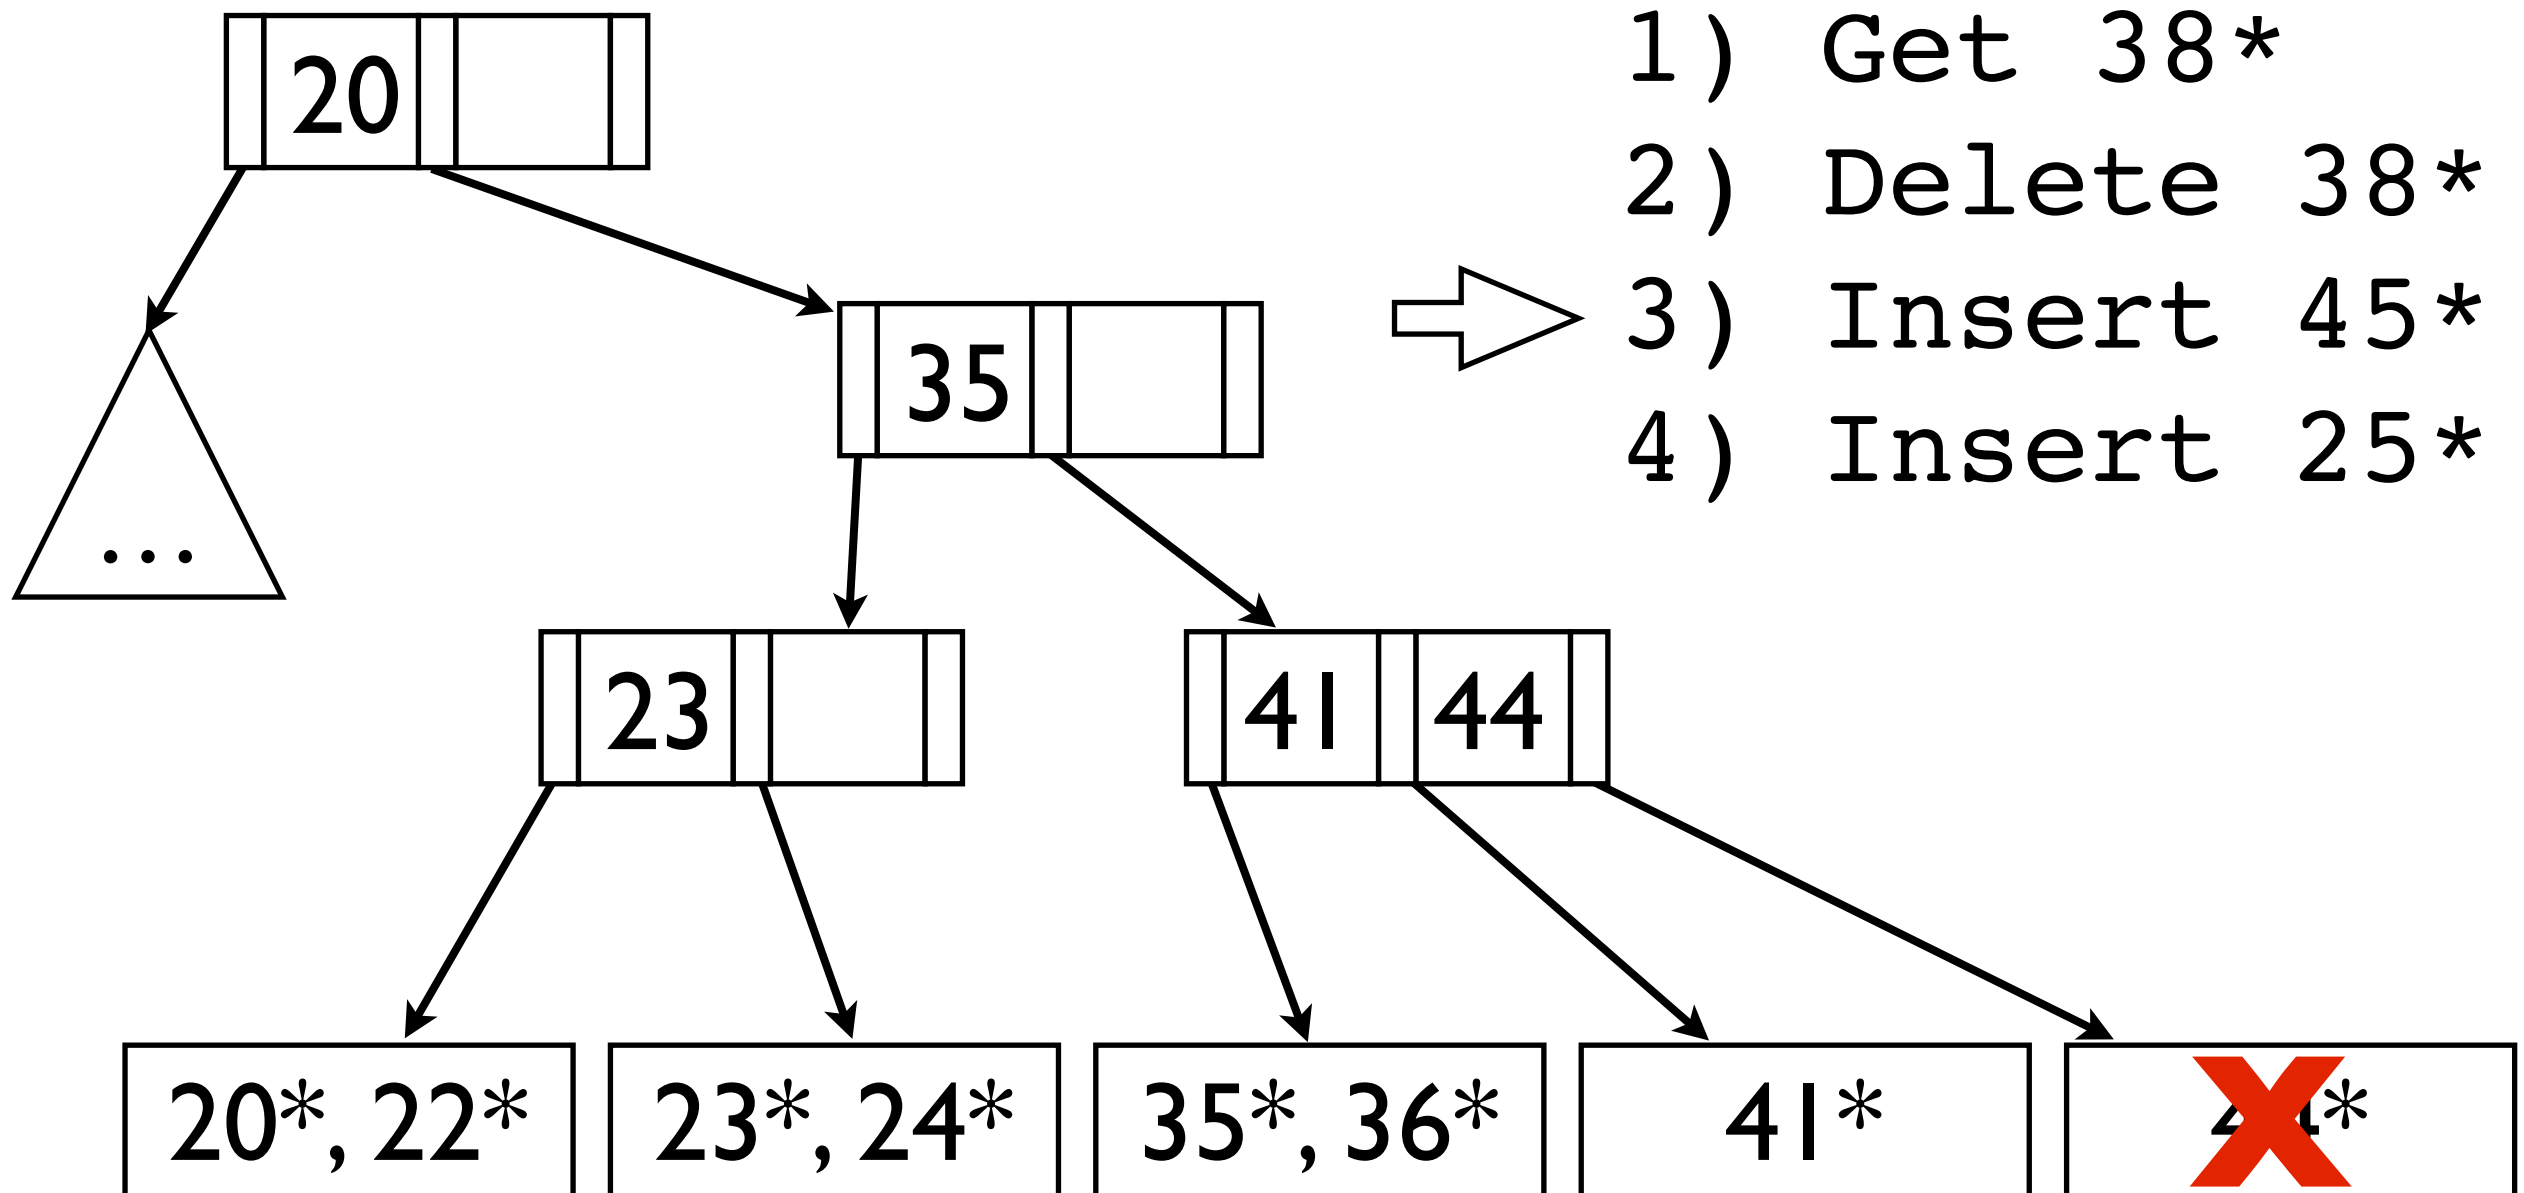
Examples




Examples




Examples




Examples



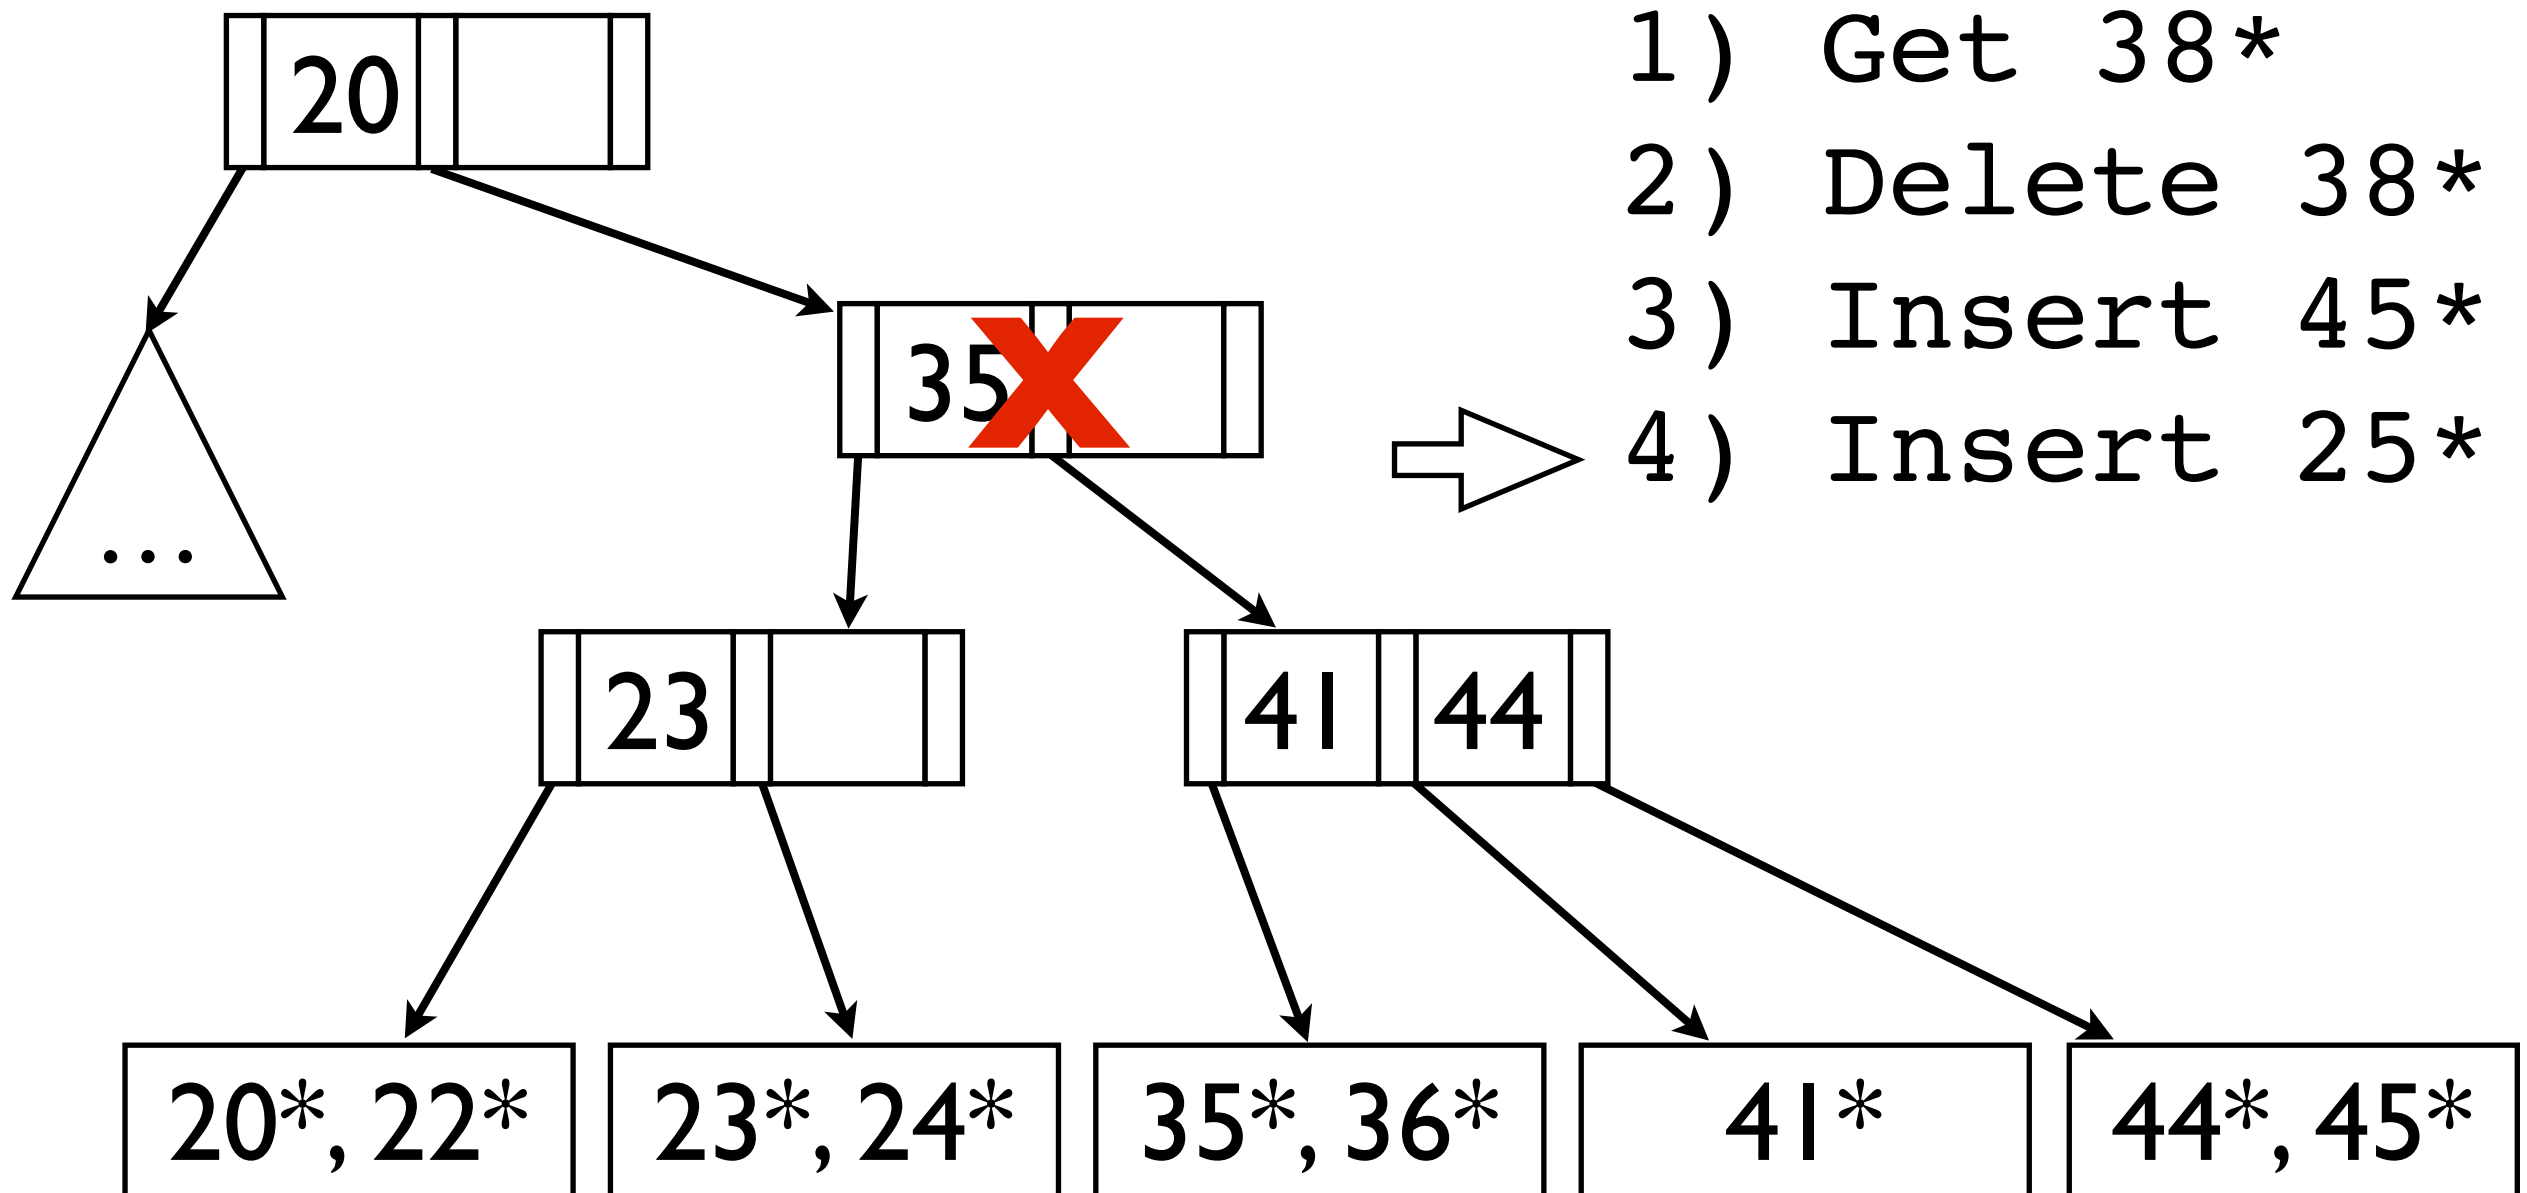
Examples



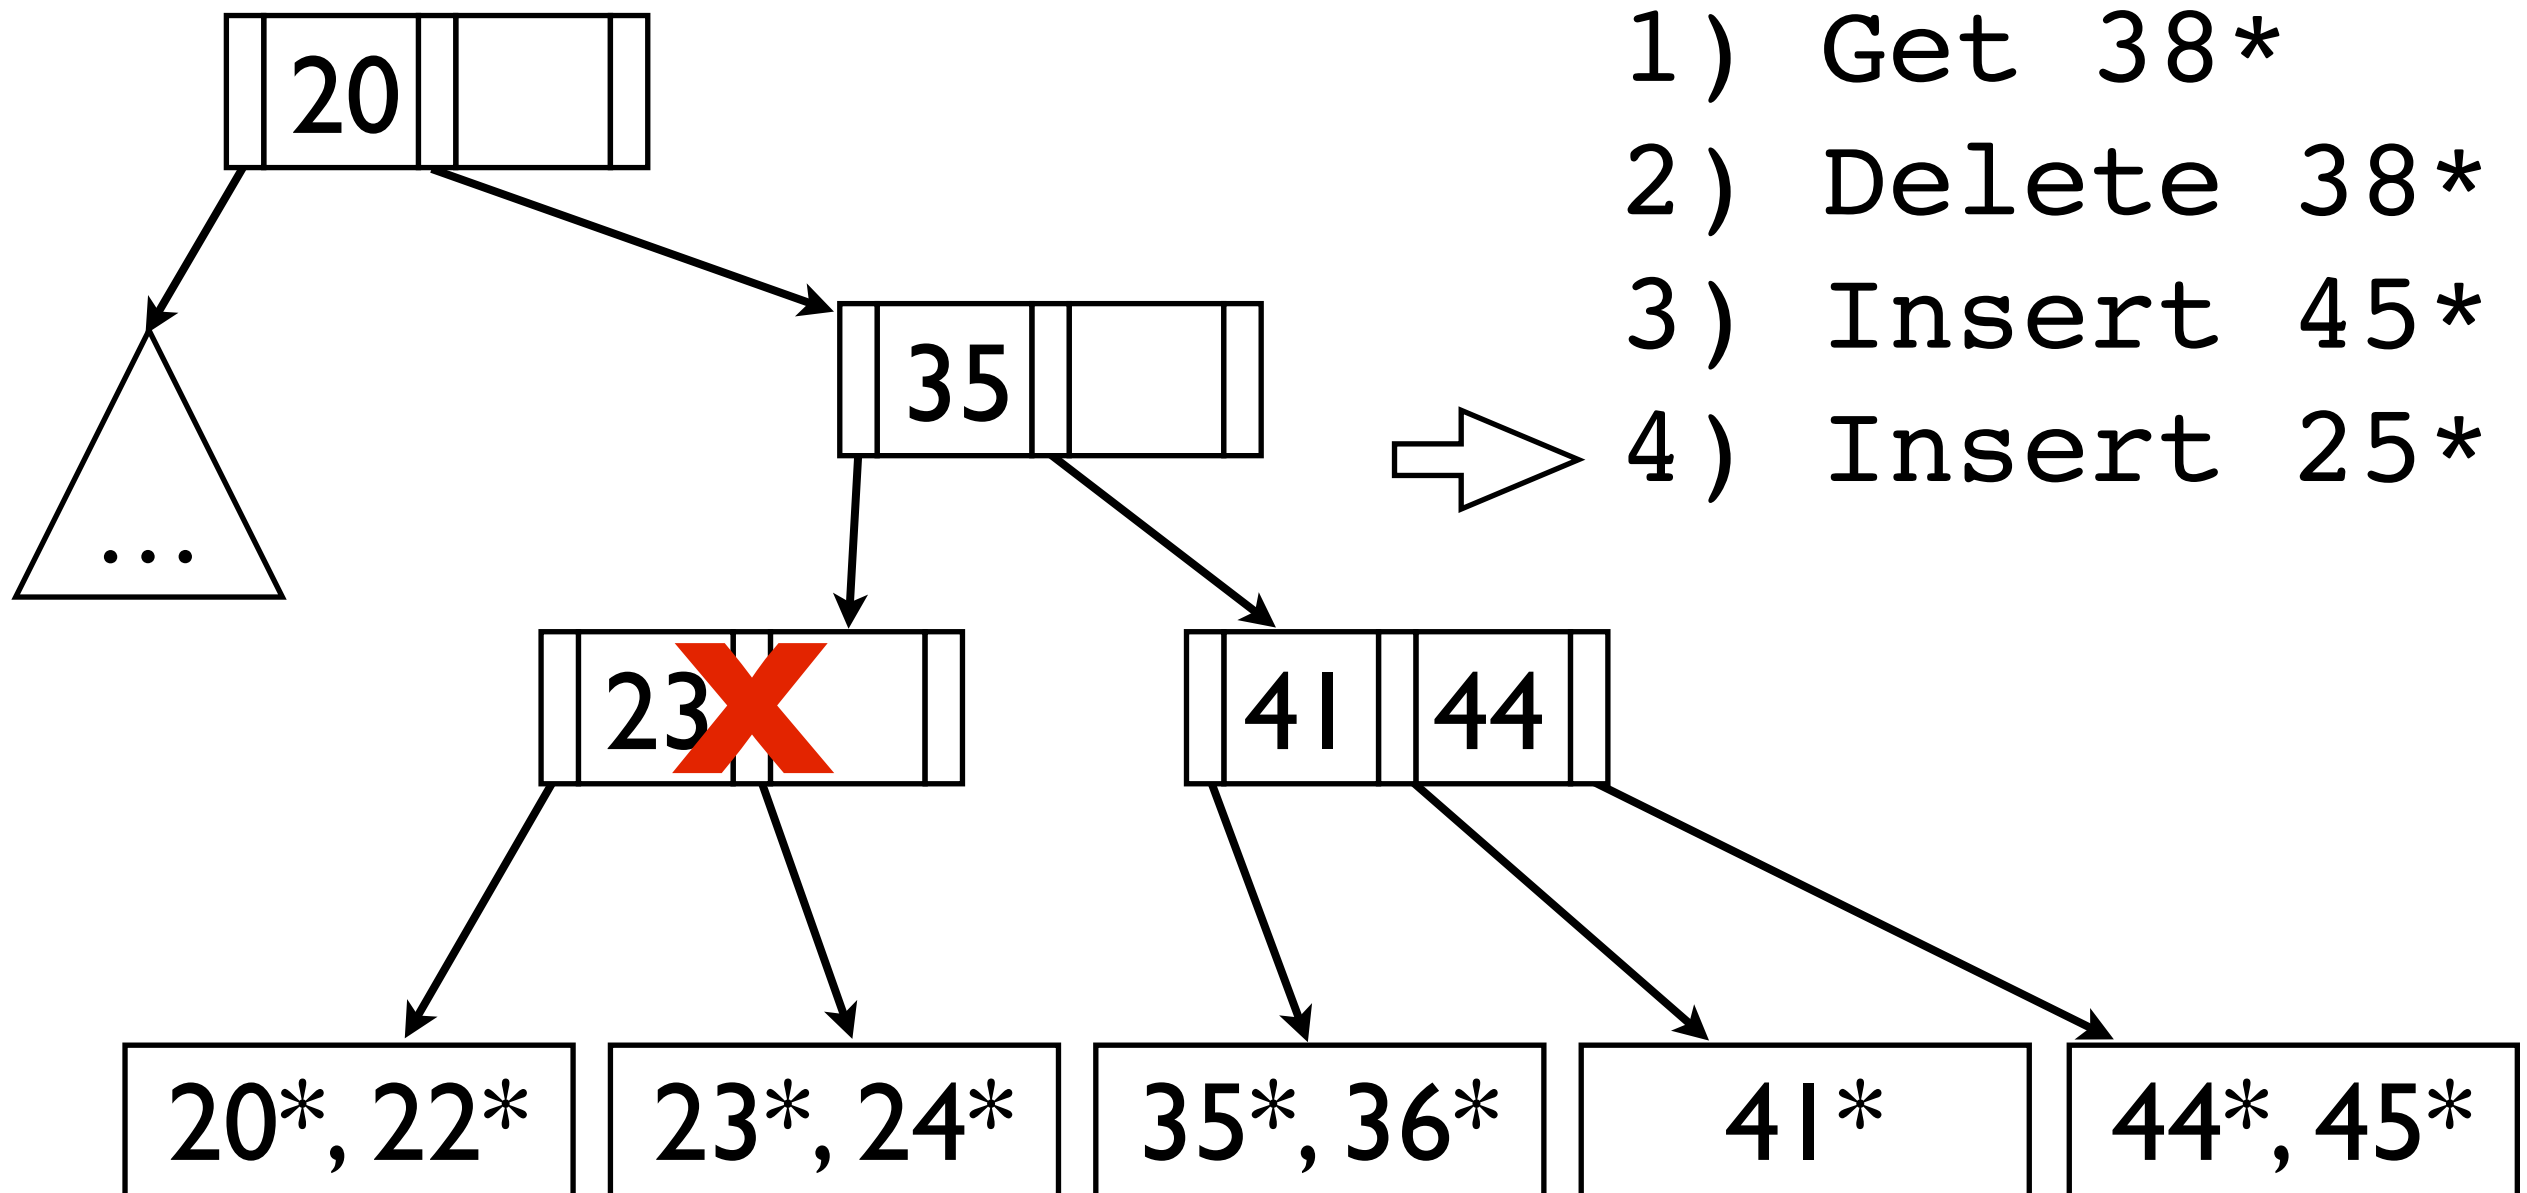
Examples



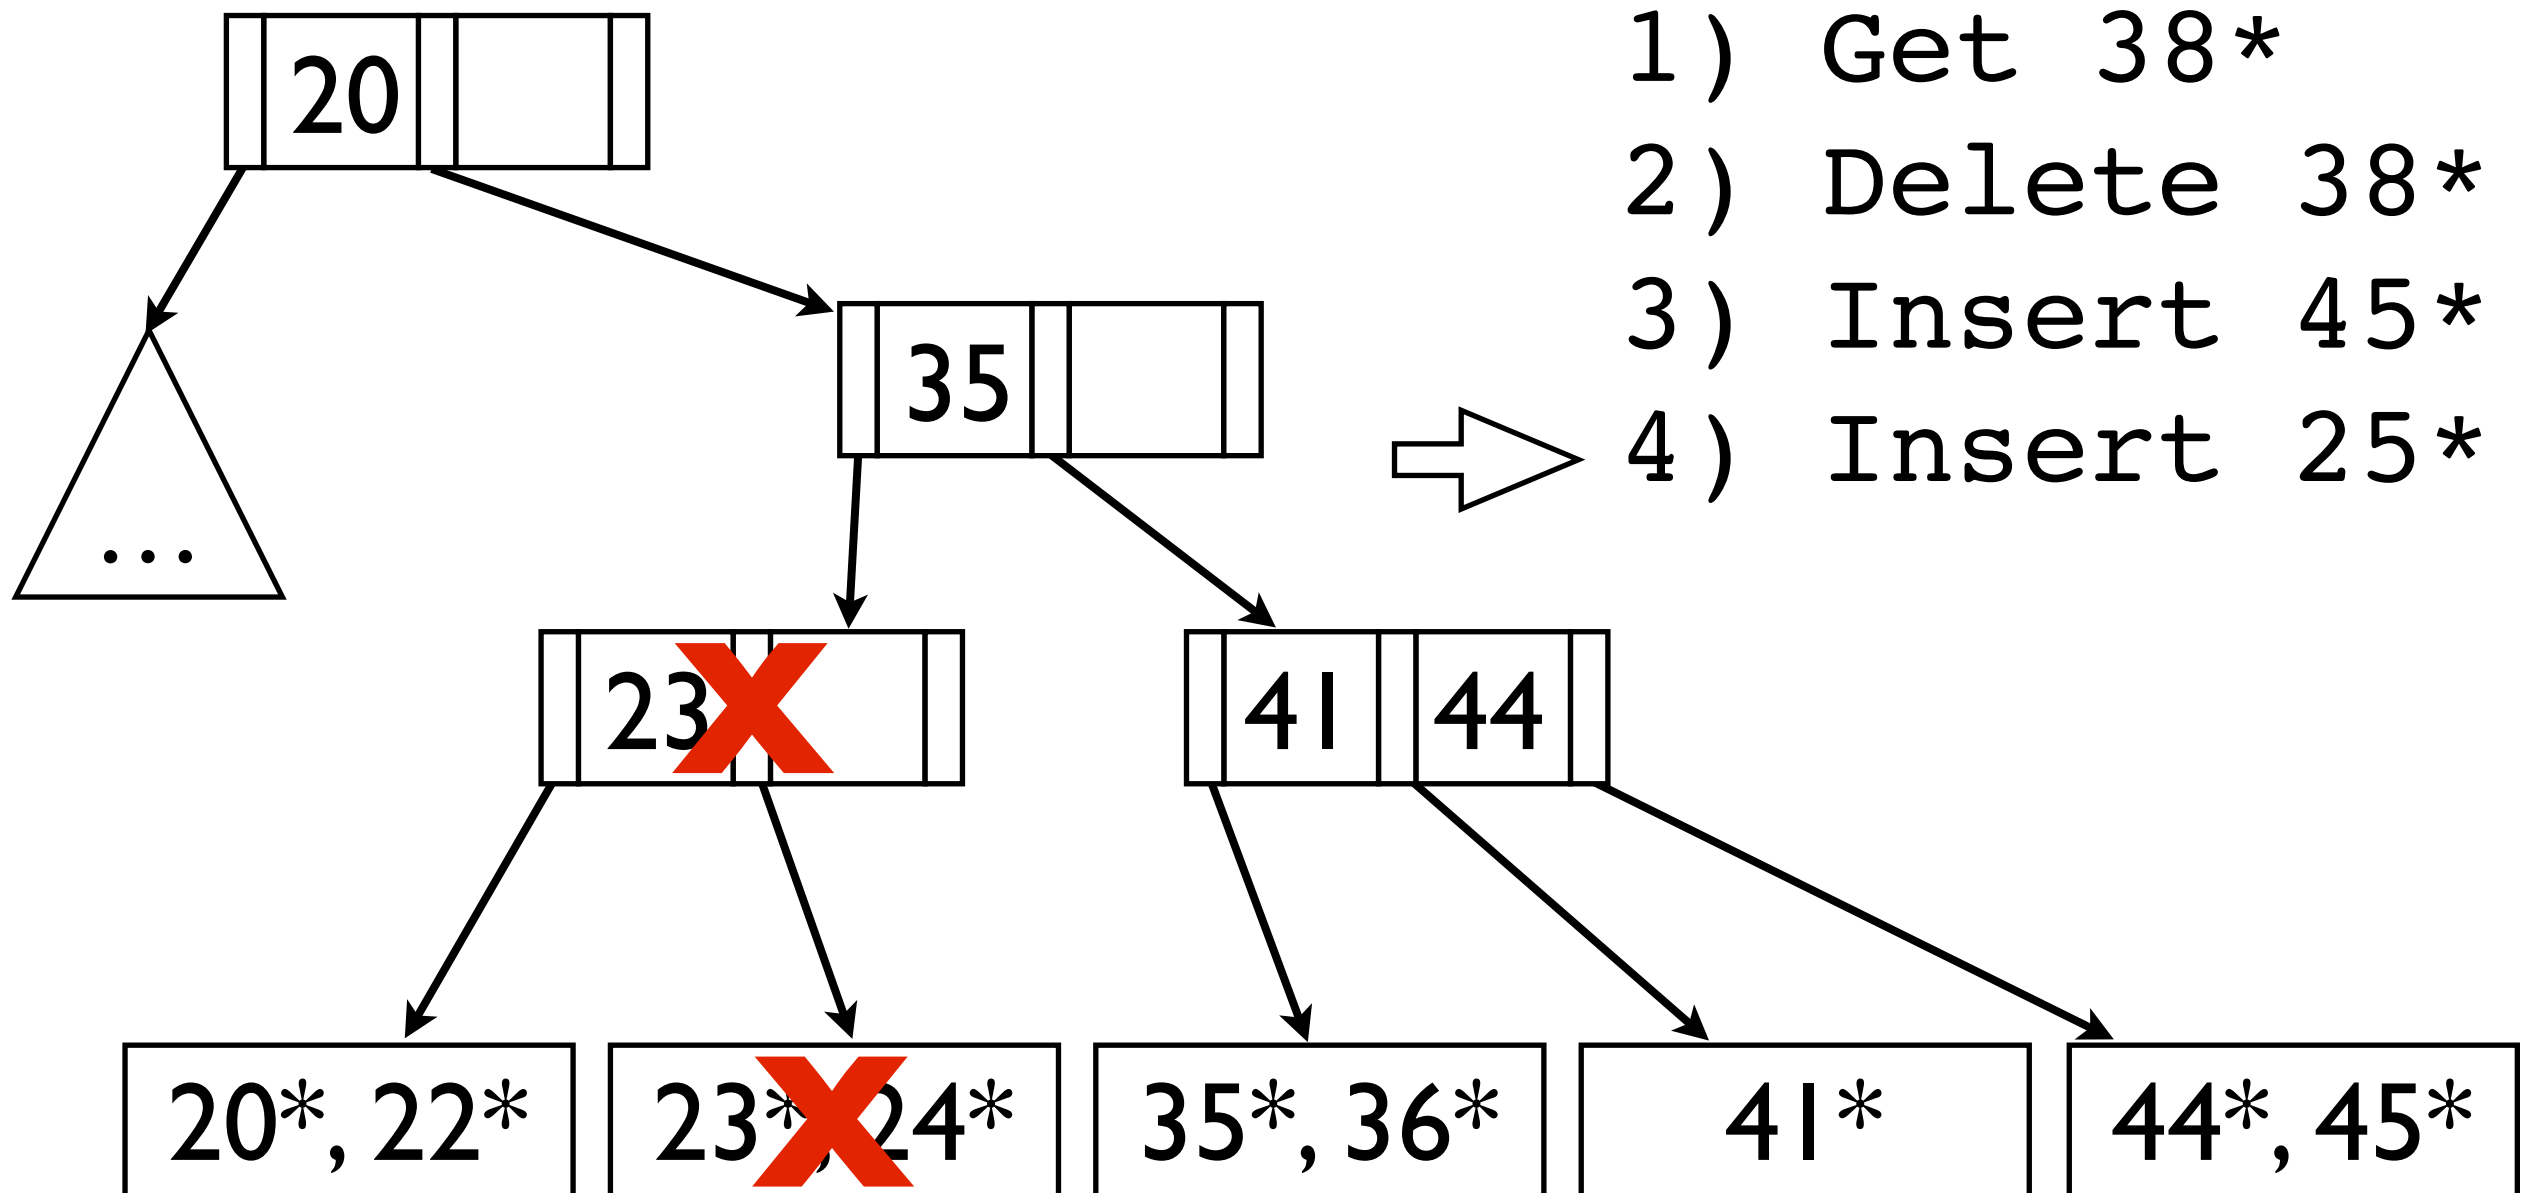
Examples



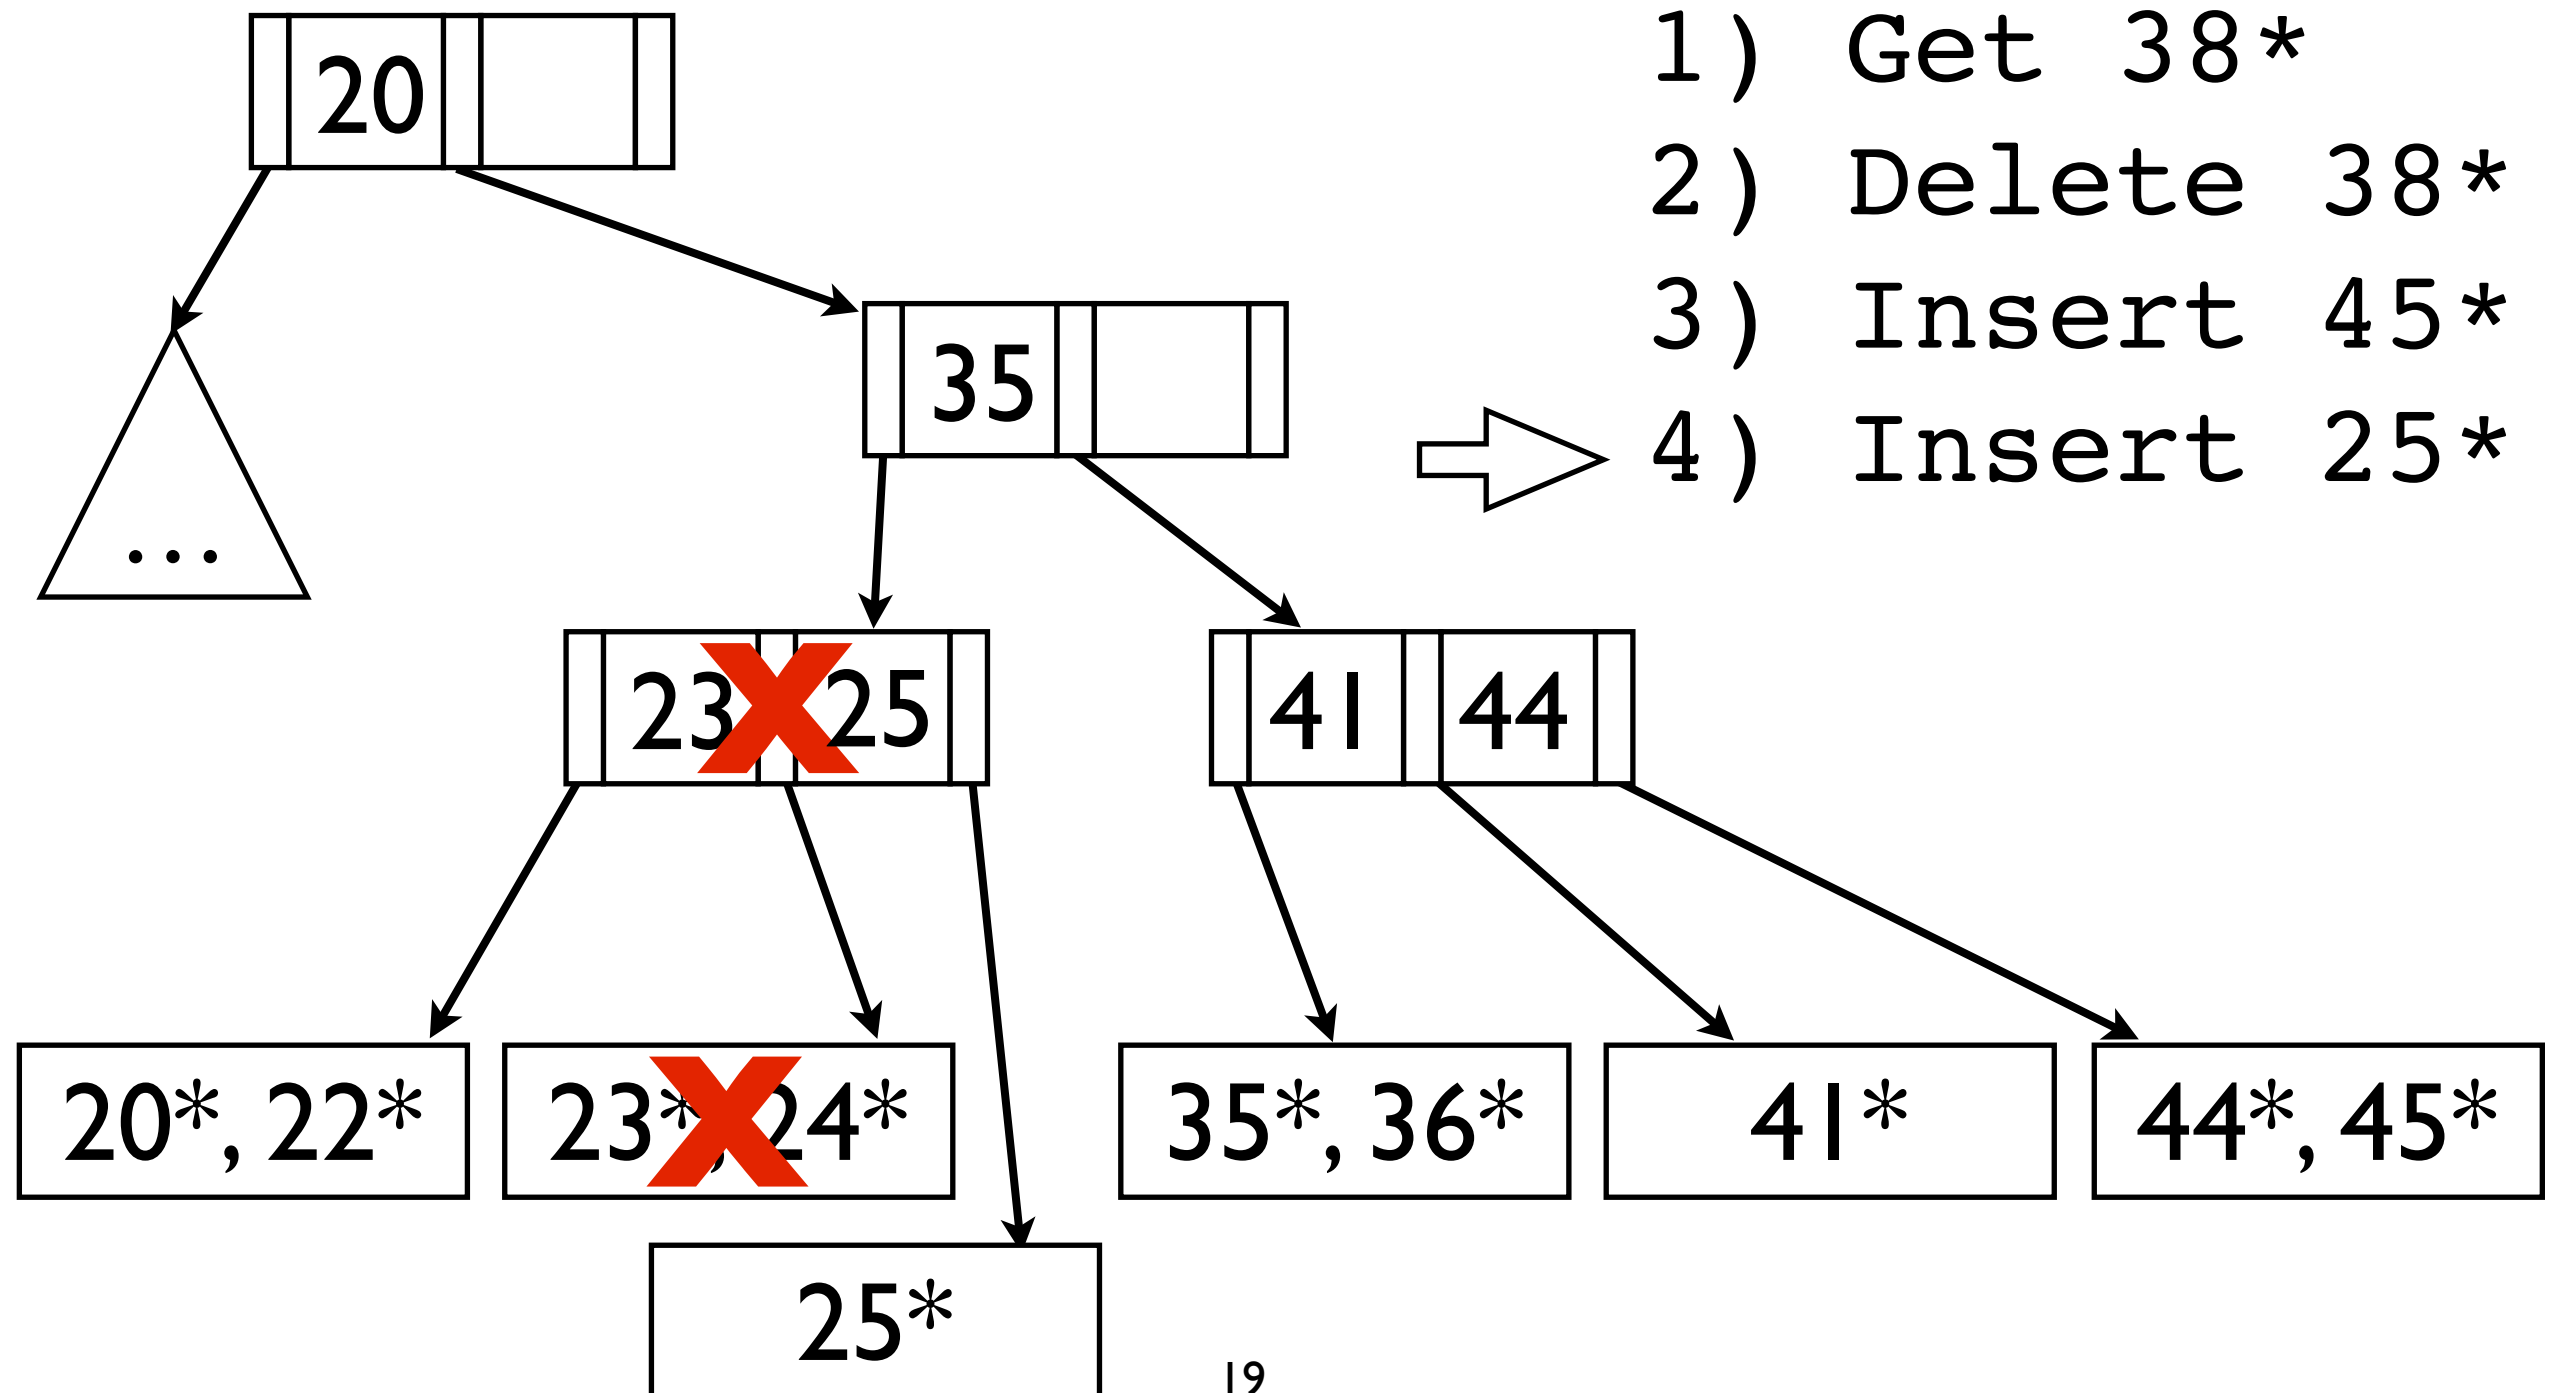
Examples



Examples




Examples



Better Tree Locking Algorithm (Bayer Schkolnick)


- **Scan:** As before
- **Update:** Set locks as if for search (using S)
 - Acquire X lock on the leaf.
 - If leaf is not safe, release all locks and restart using the simple algorithm.
- Gambles that only the leaf node will be modified.
 - S locks set on first pass are wasted otherwise.
 - In practice, better than the simple algorithm.

Examples




Examples

- ⇒
- 1) Get 38*
 - 2) Delete 38*
 - 3) Insert 45*
 - 4) Insert 25*




Examples

- ⇒
- 1) Get 38*
 - 2) Delete 38*
 - 3) Insert 45*
 - 4) Insert 25*



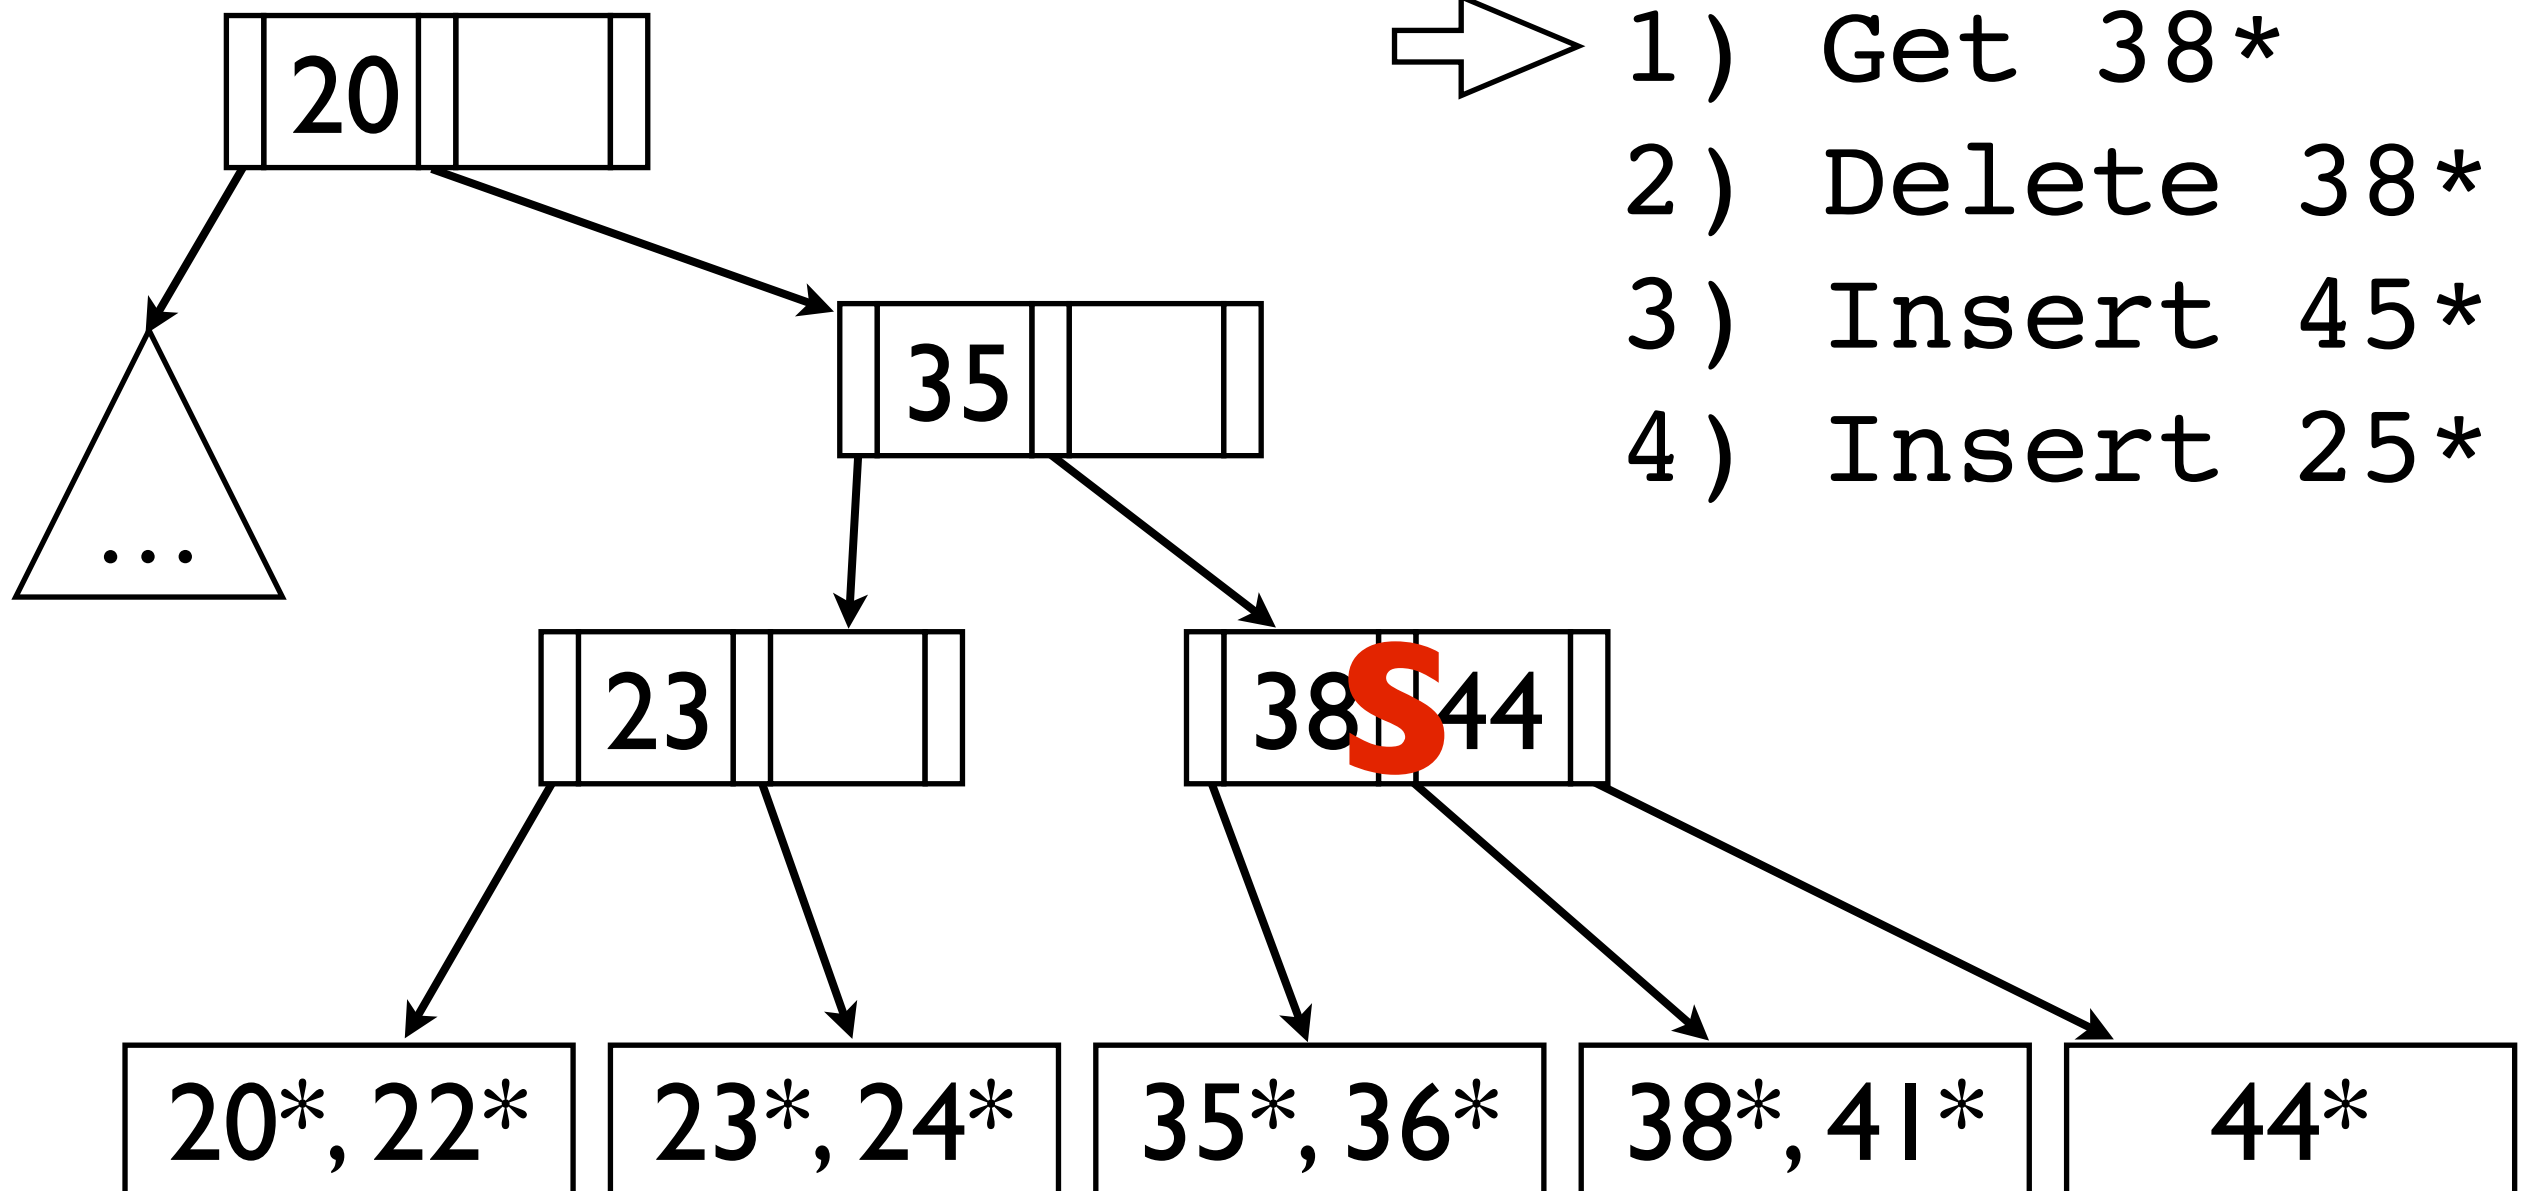
Examples

- ⇒
- 1) Get 38*
 - 2) Delete 38*
 - 3) Insert 45*
 - 4) Insert 25*



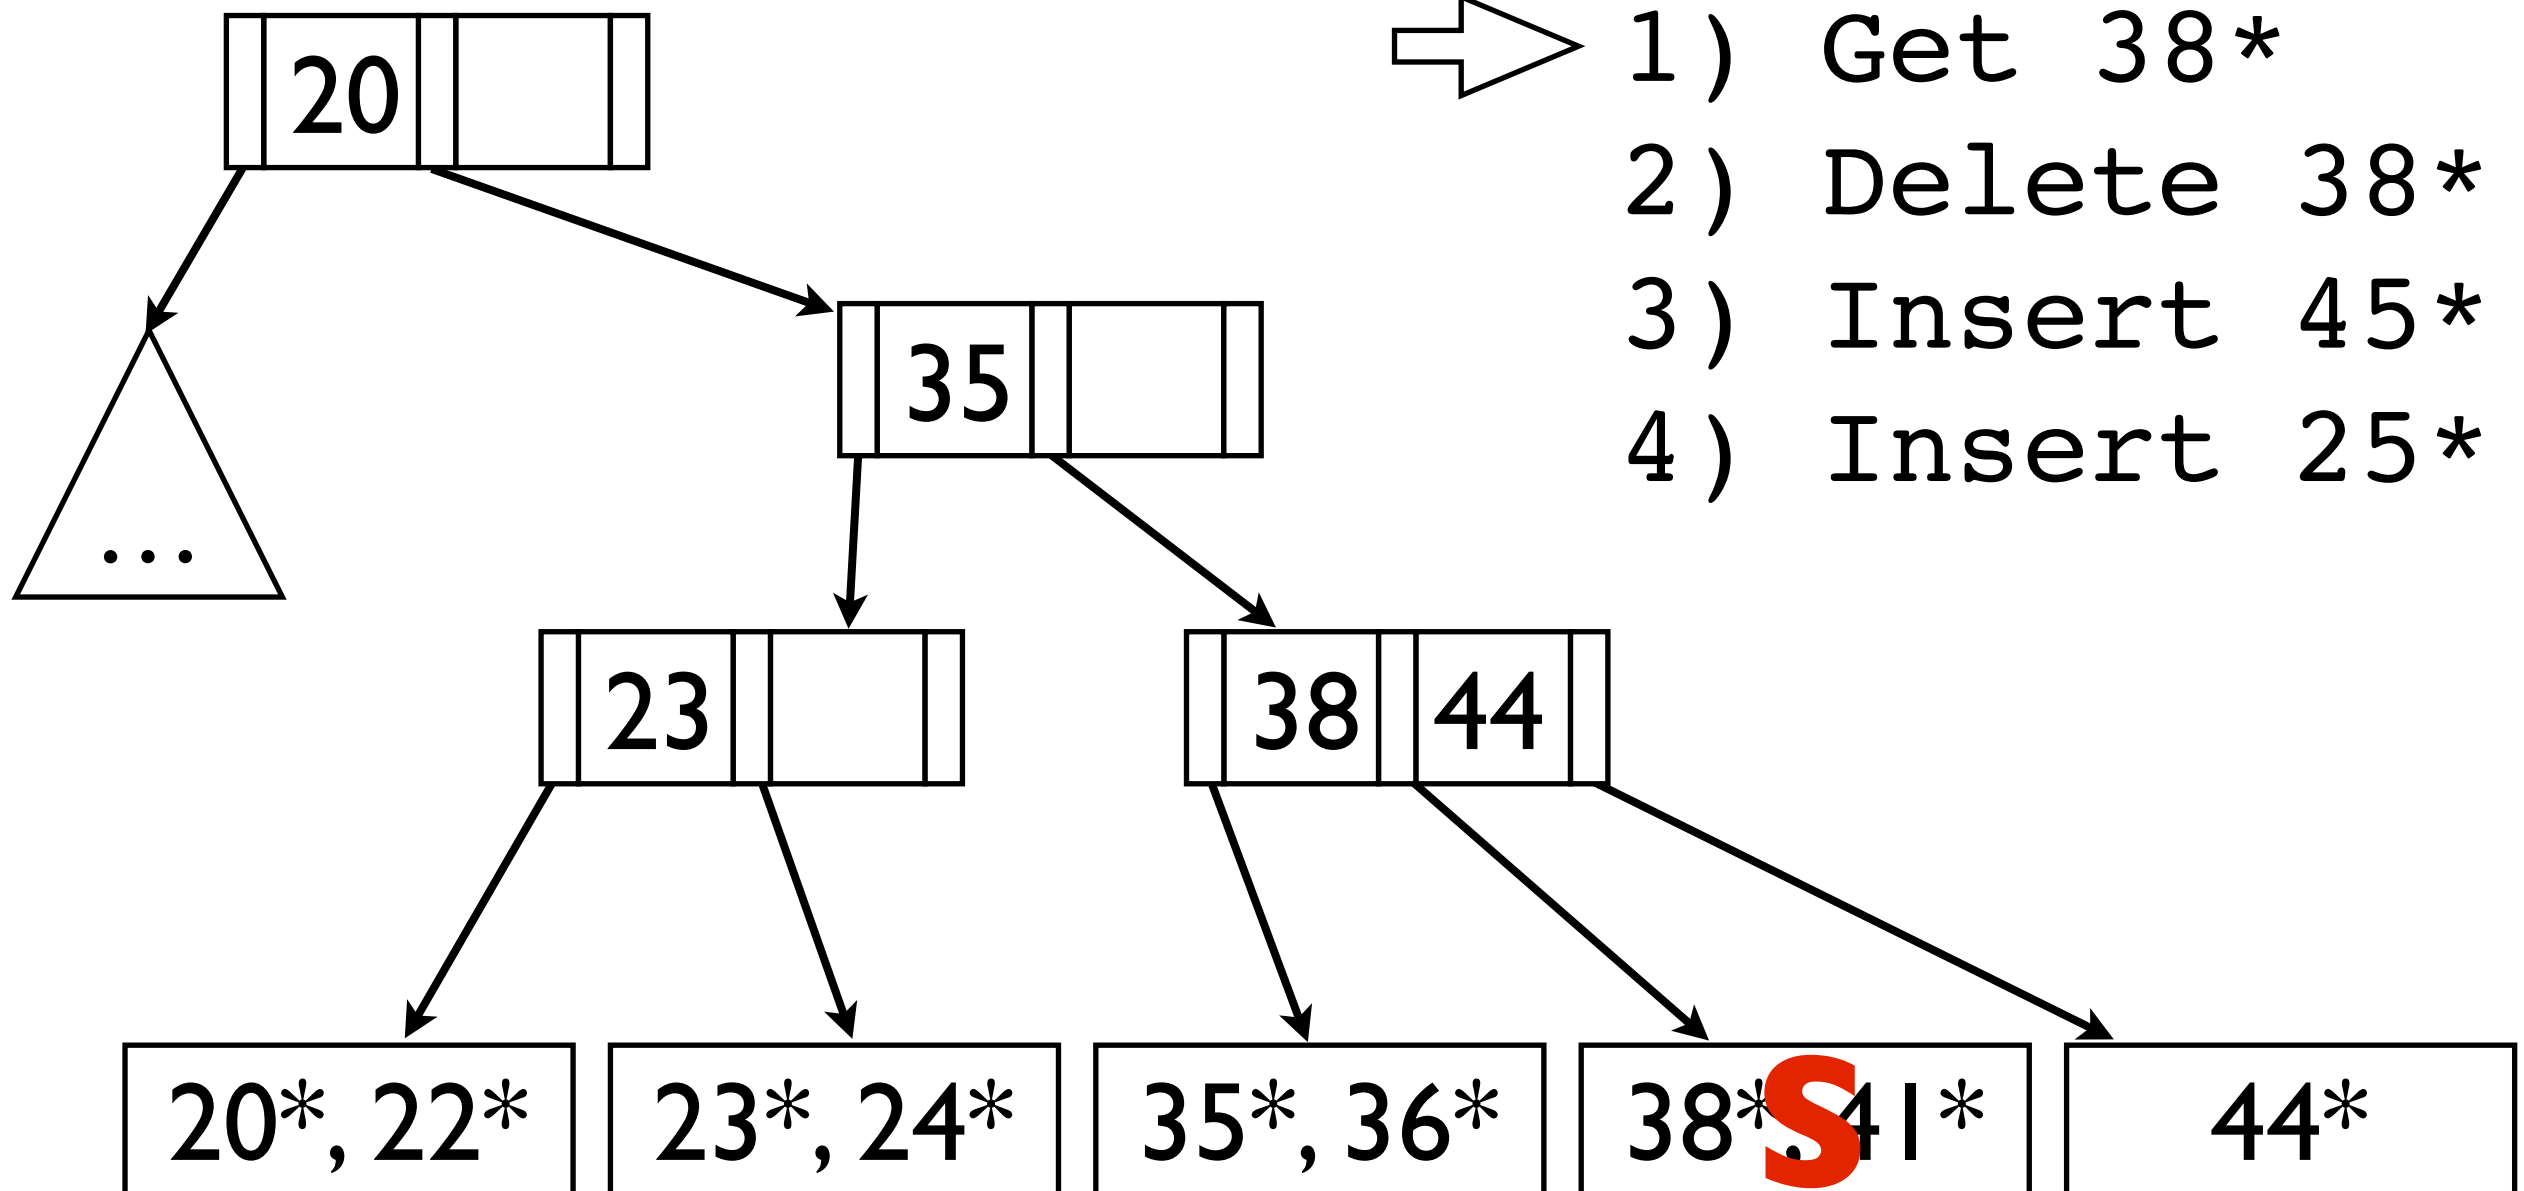
Examples

- ⇒
- 1) Get 38*
 - 2) Delete 38*
 - 3) Insert 45*
 - 4) Insert 25*




Examples


- ⇒
- 1) Get 38*
 - 2) Delete 38*
 - 3) Insert 45*
 - 4) Insert 25*



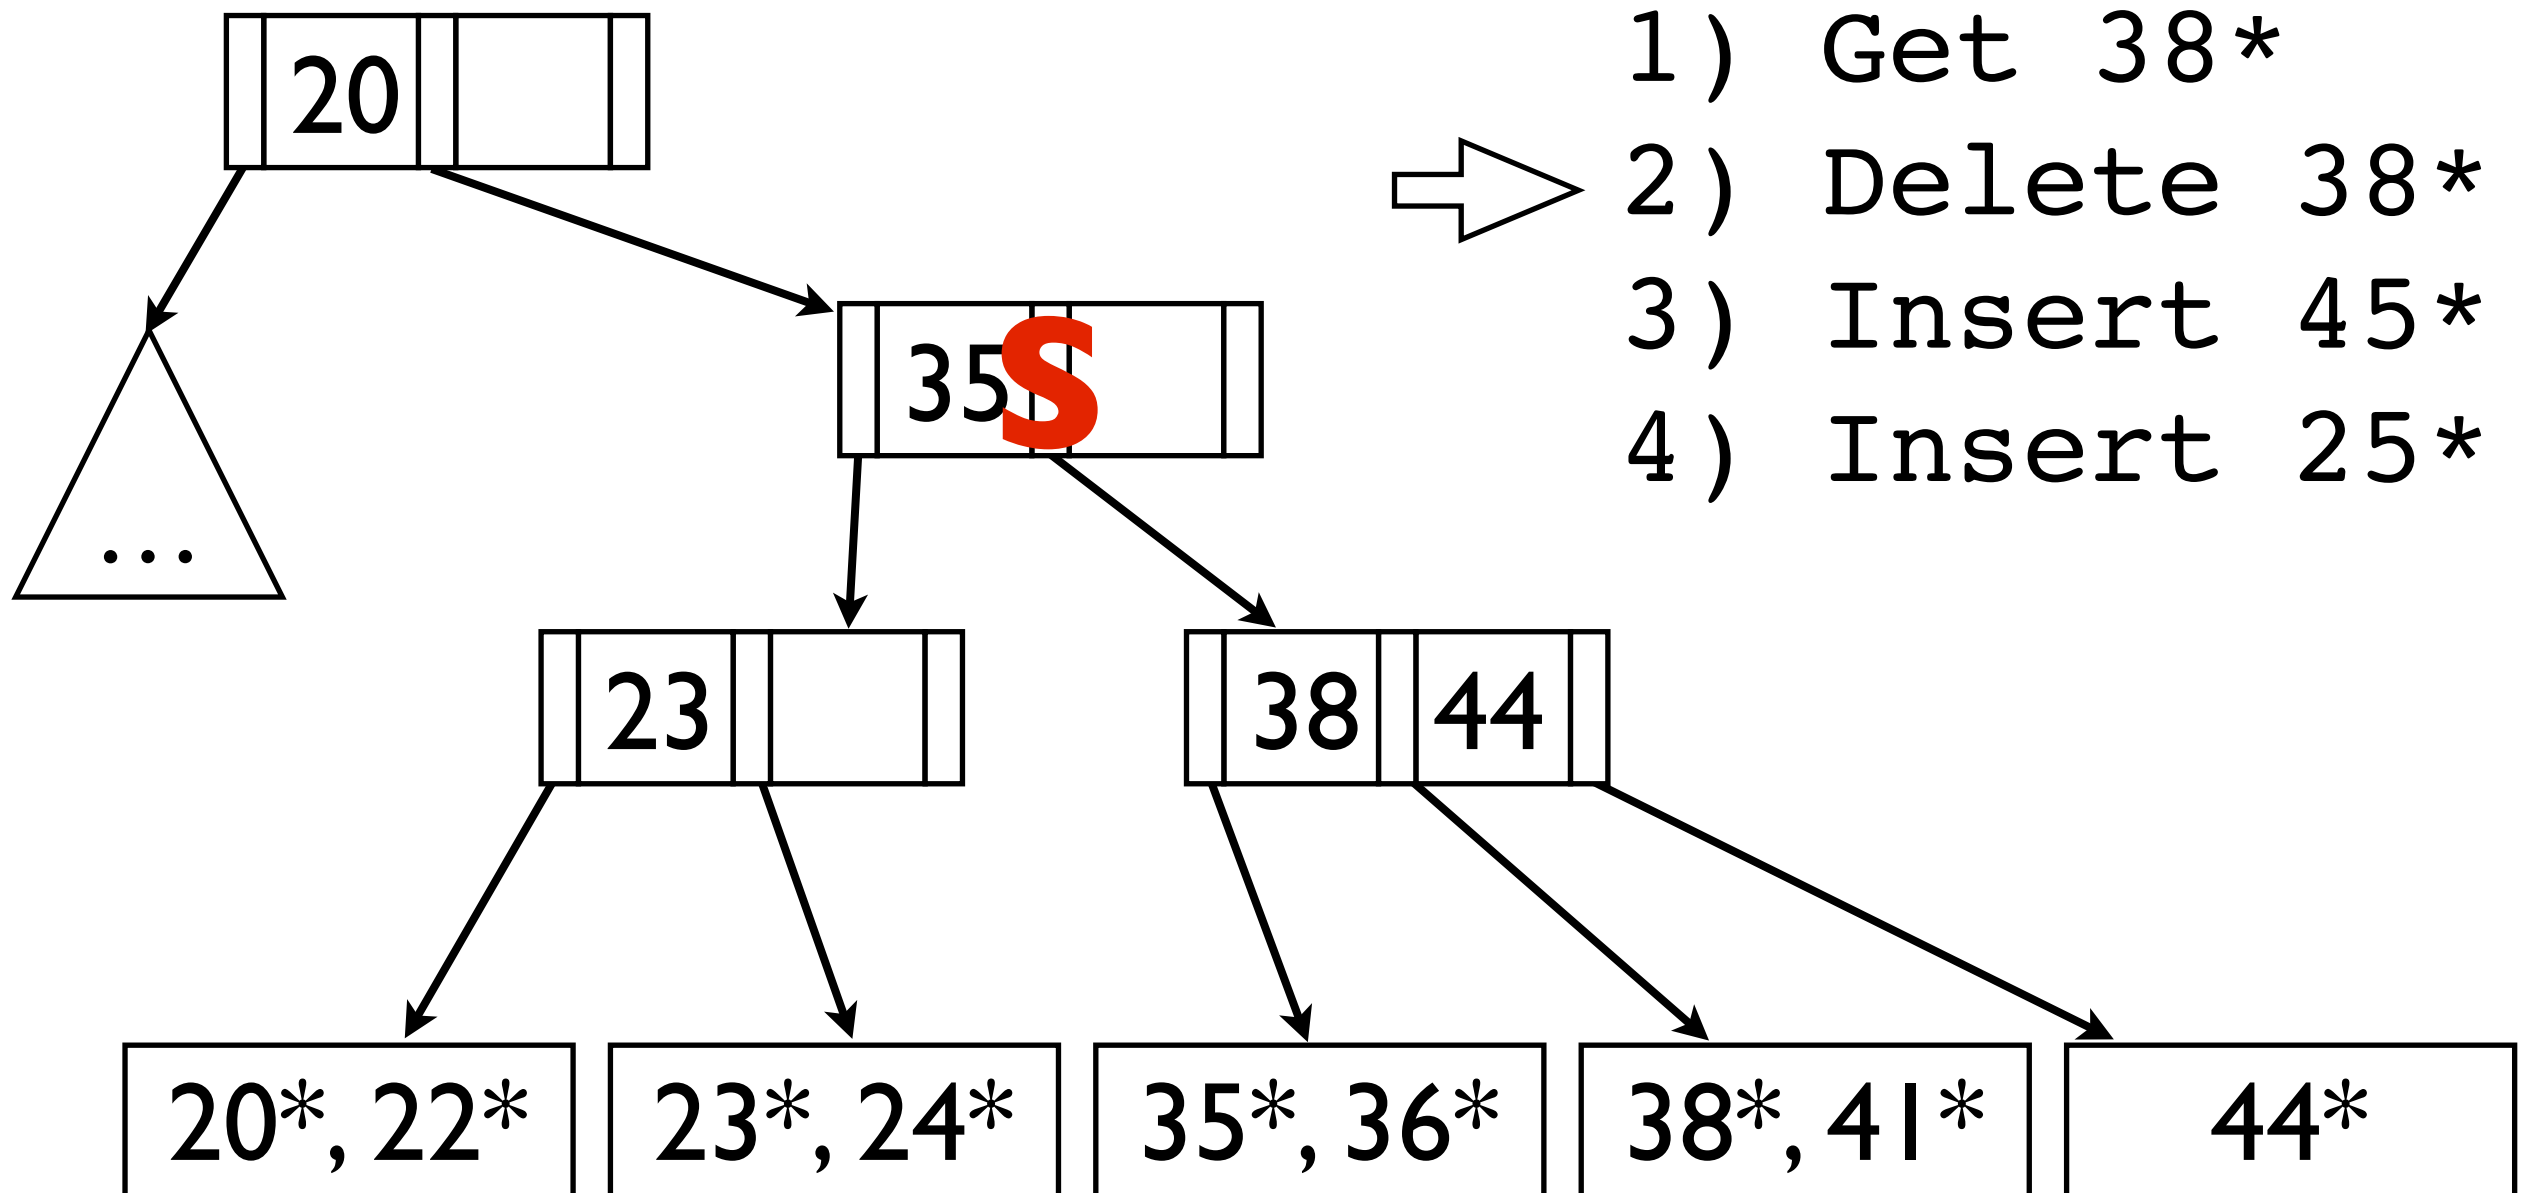
Examples



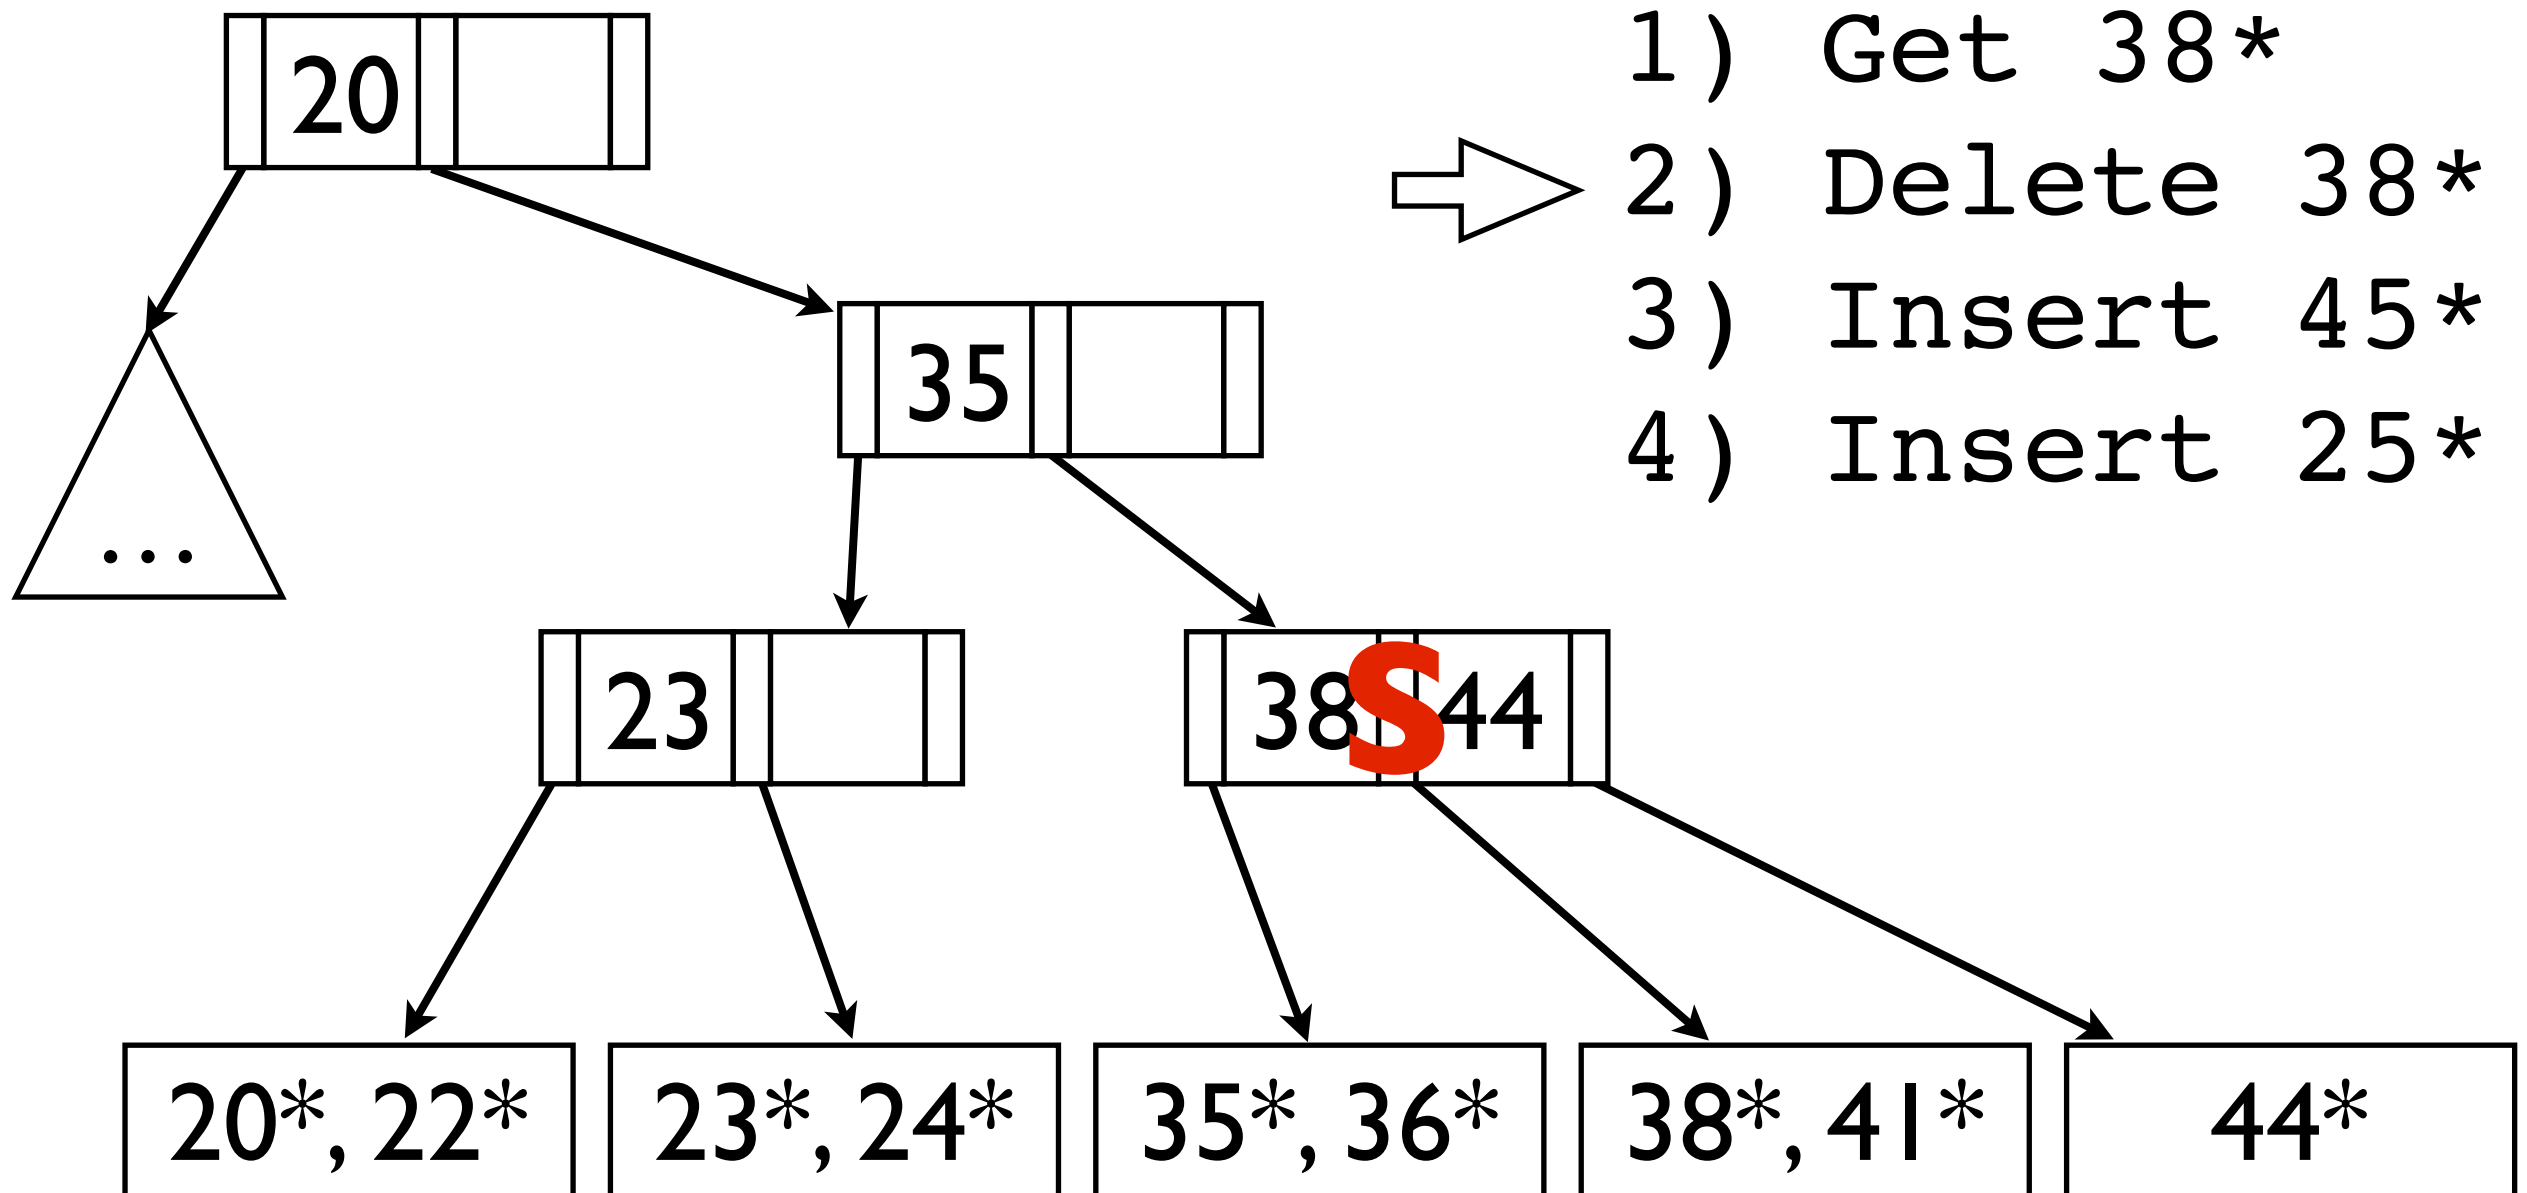
Examples



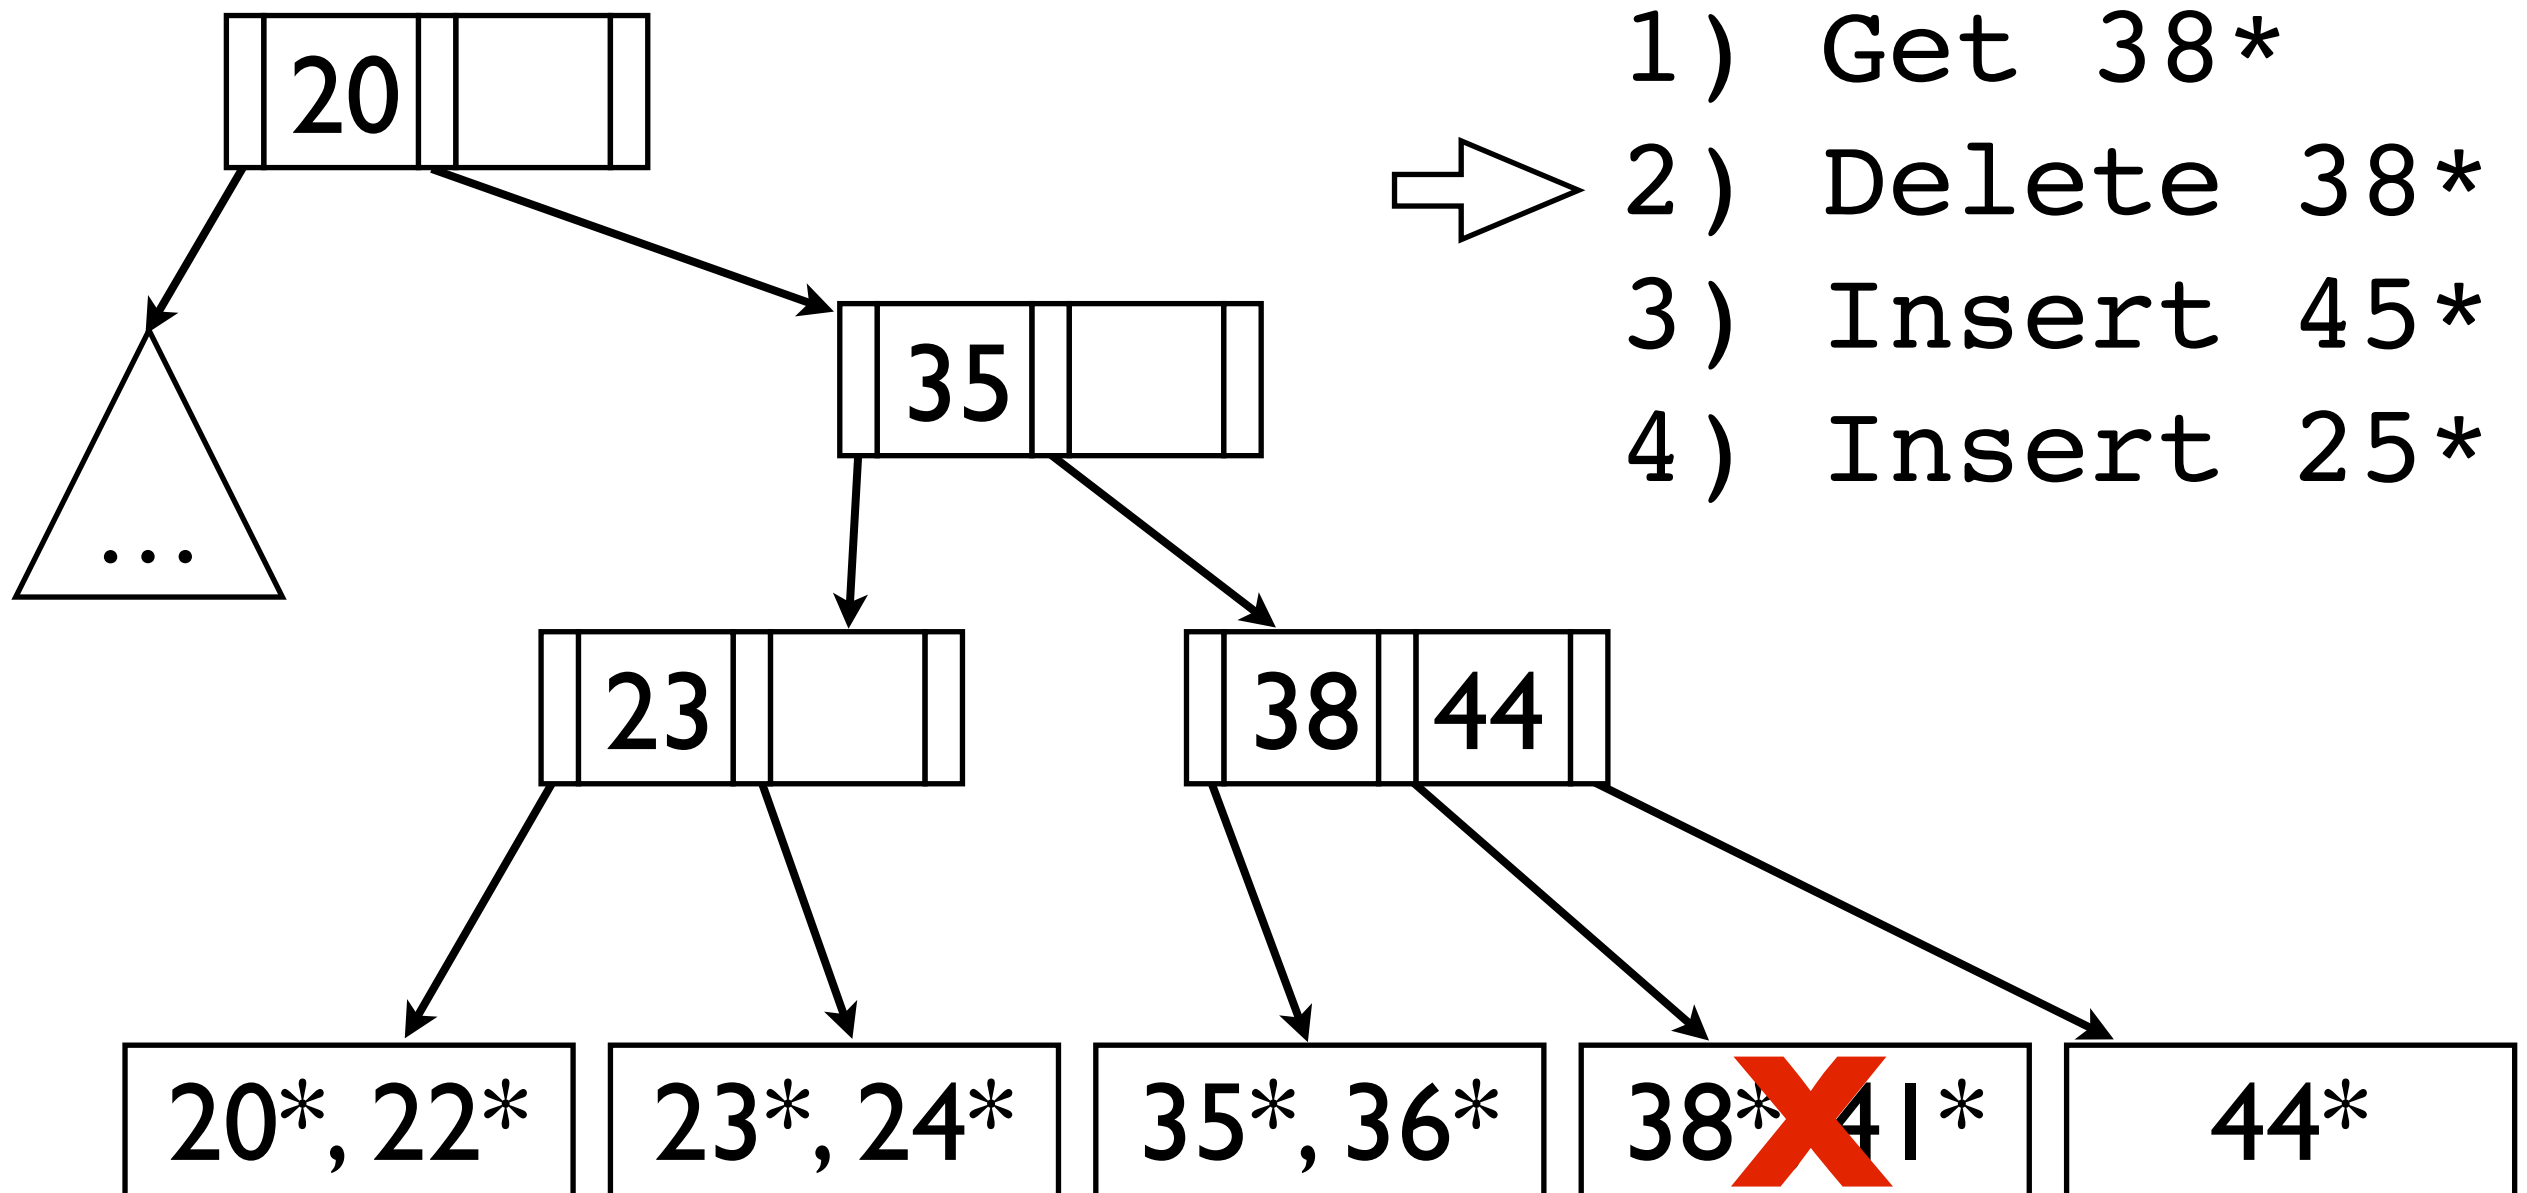
Examples



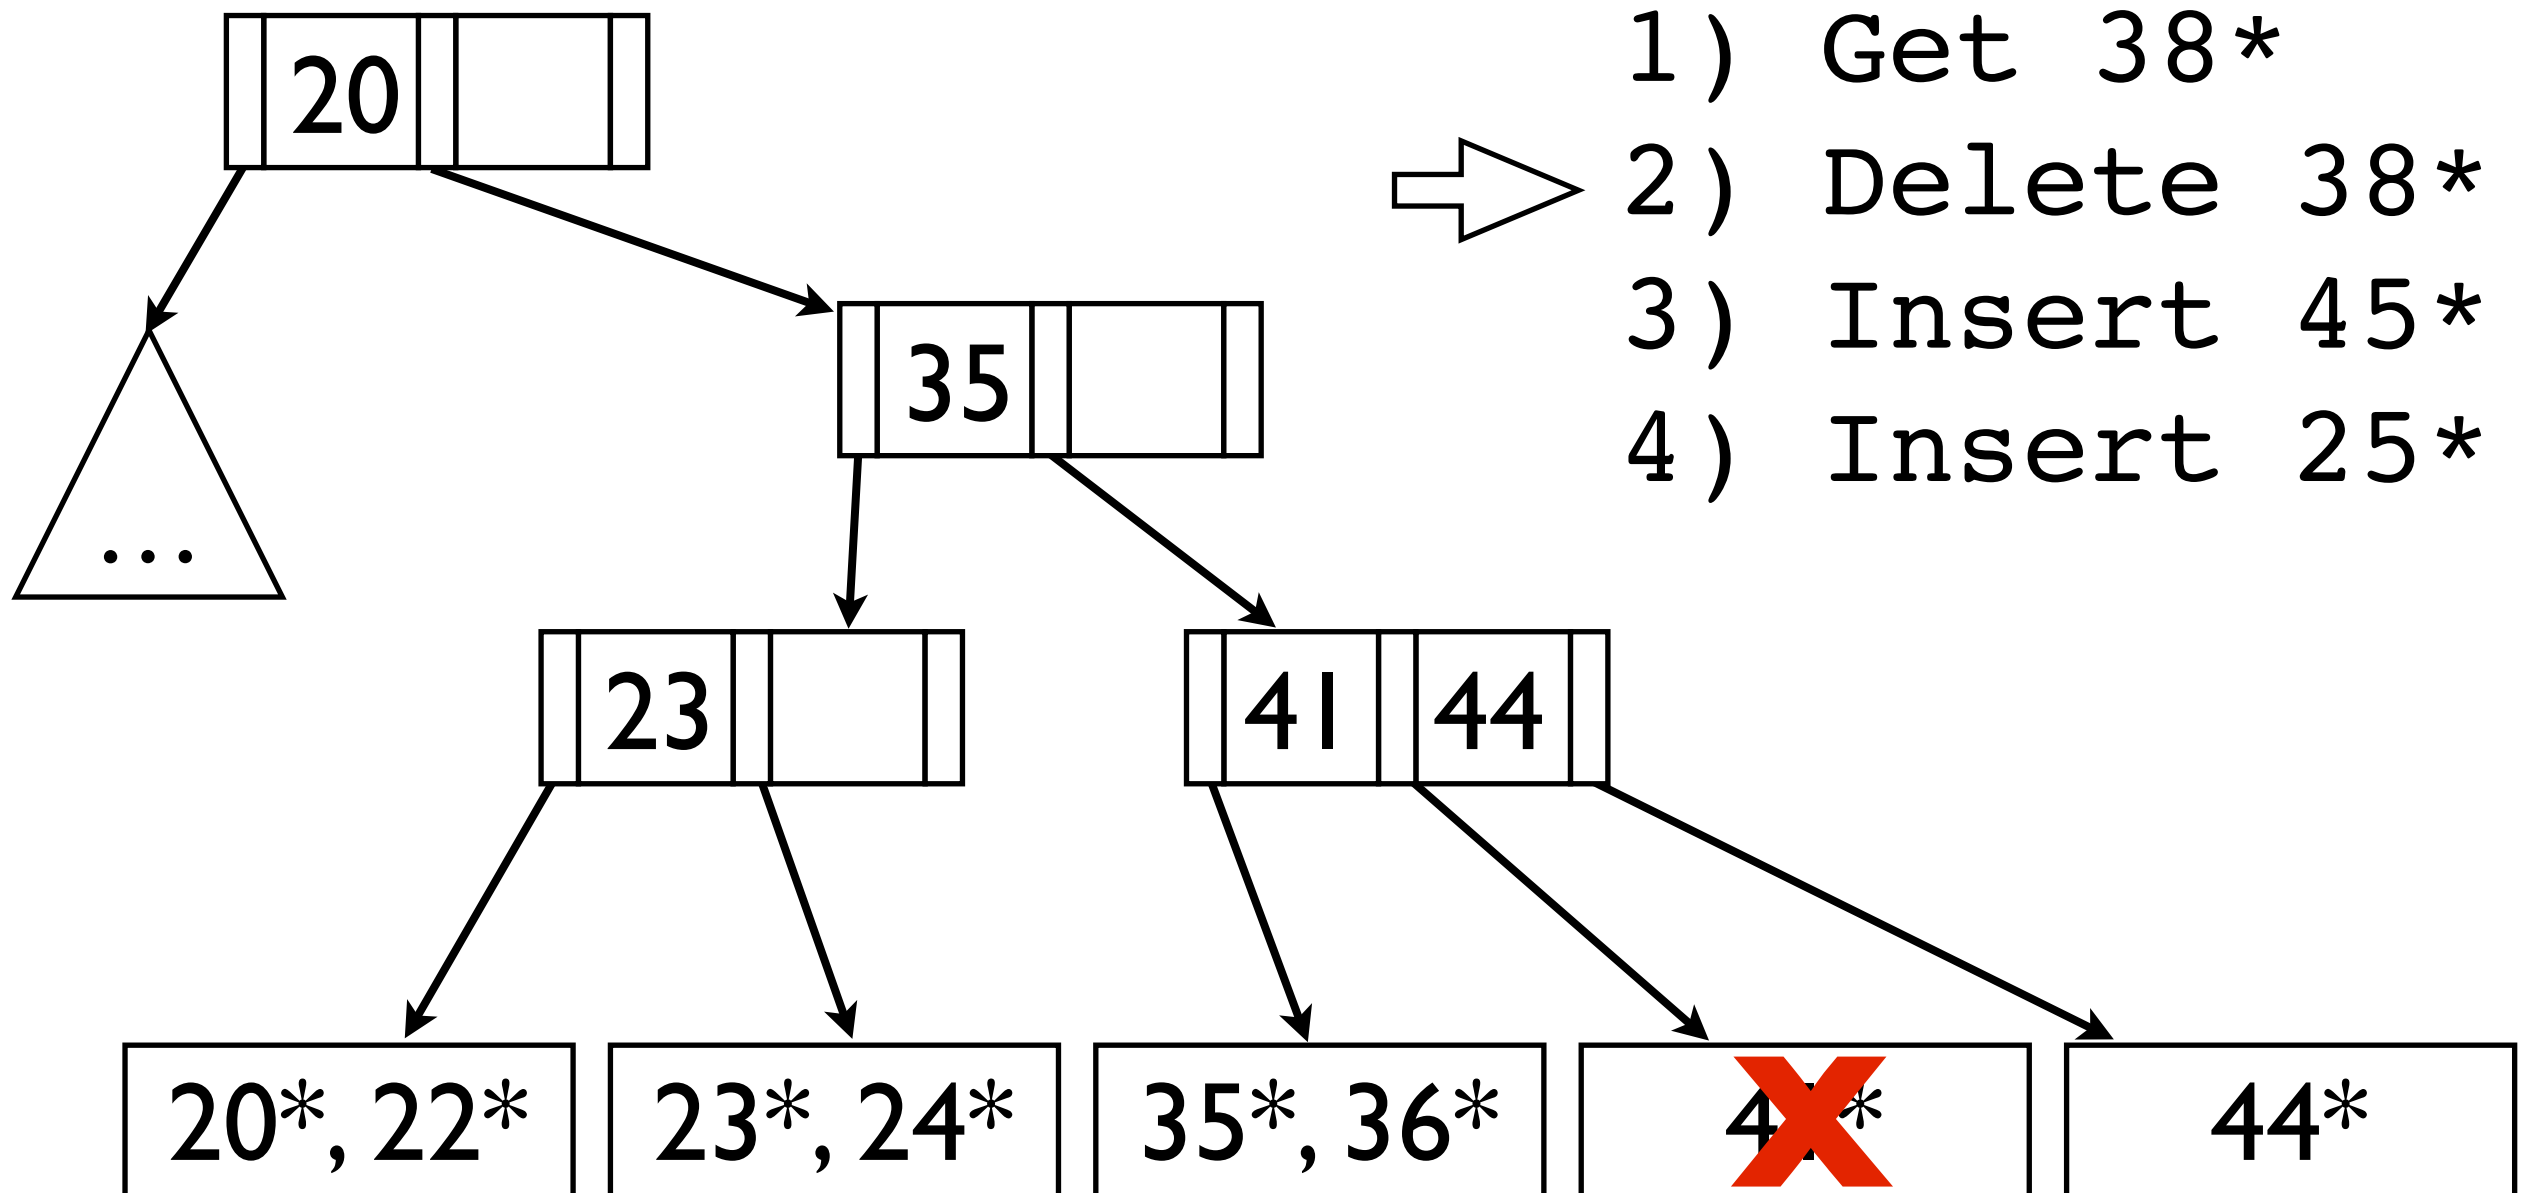
Examples



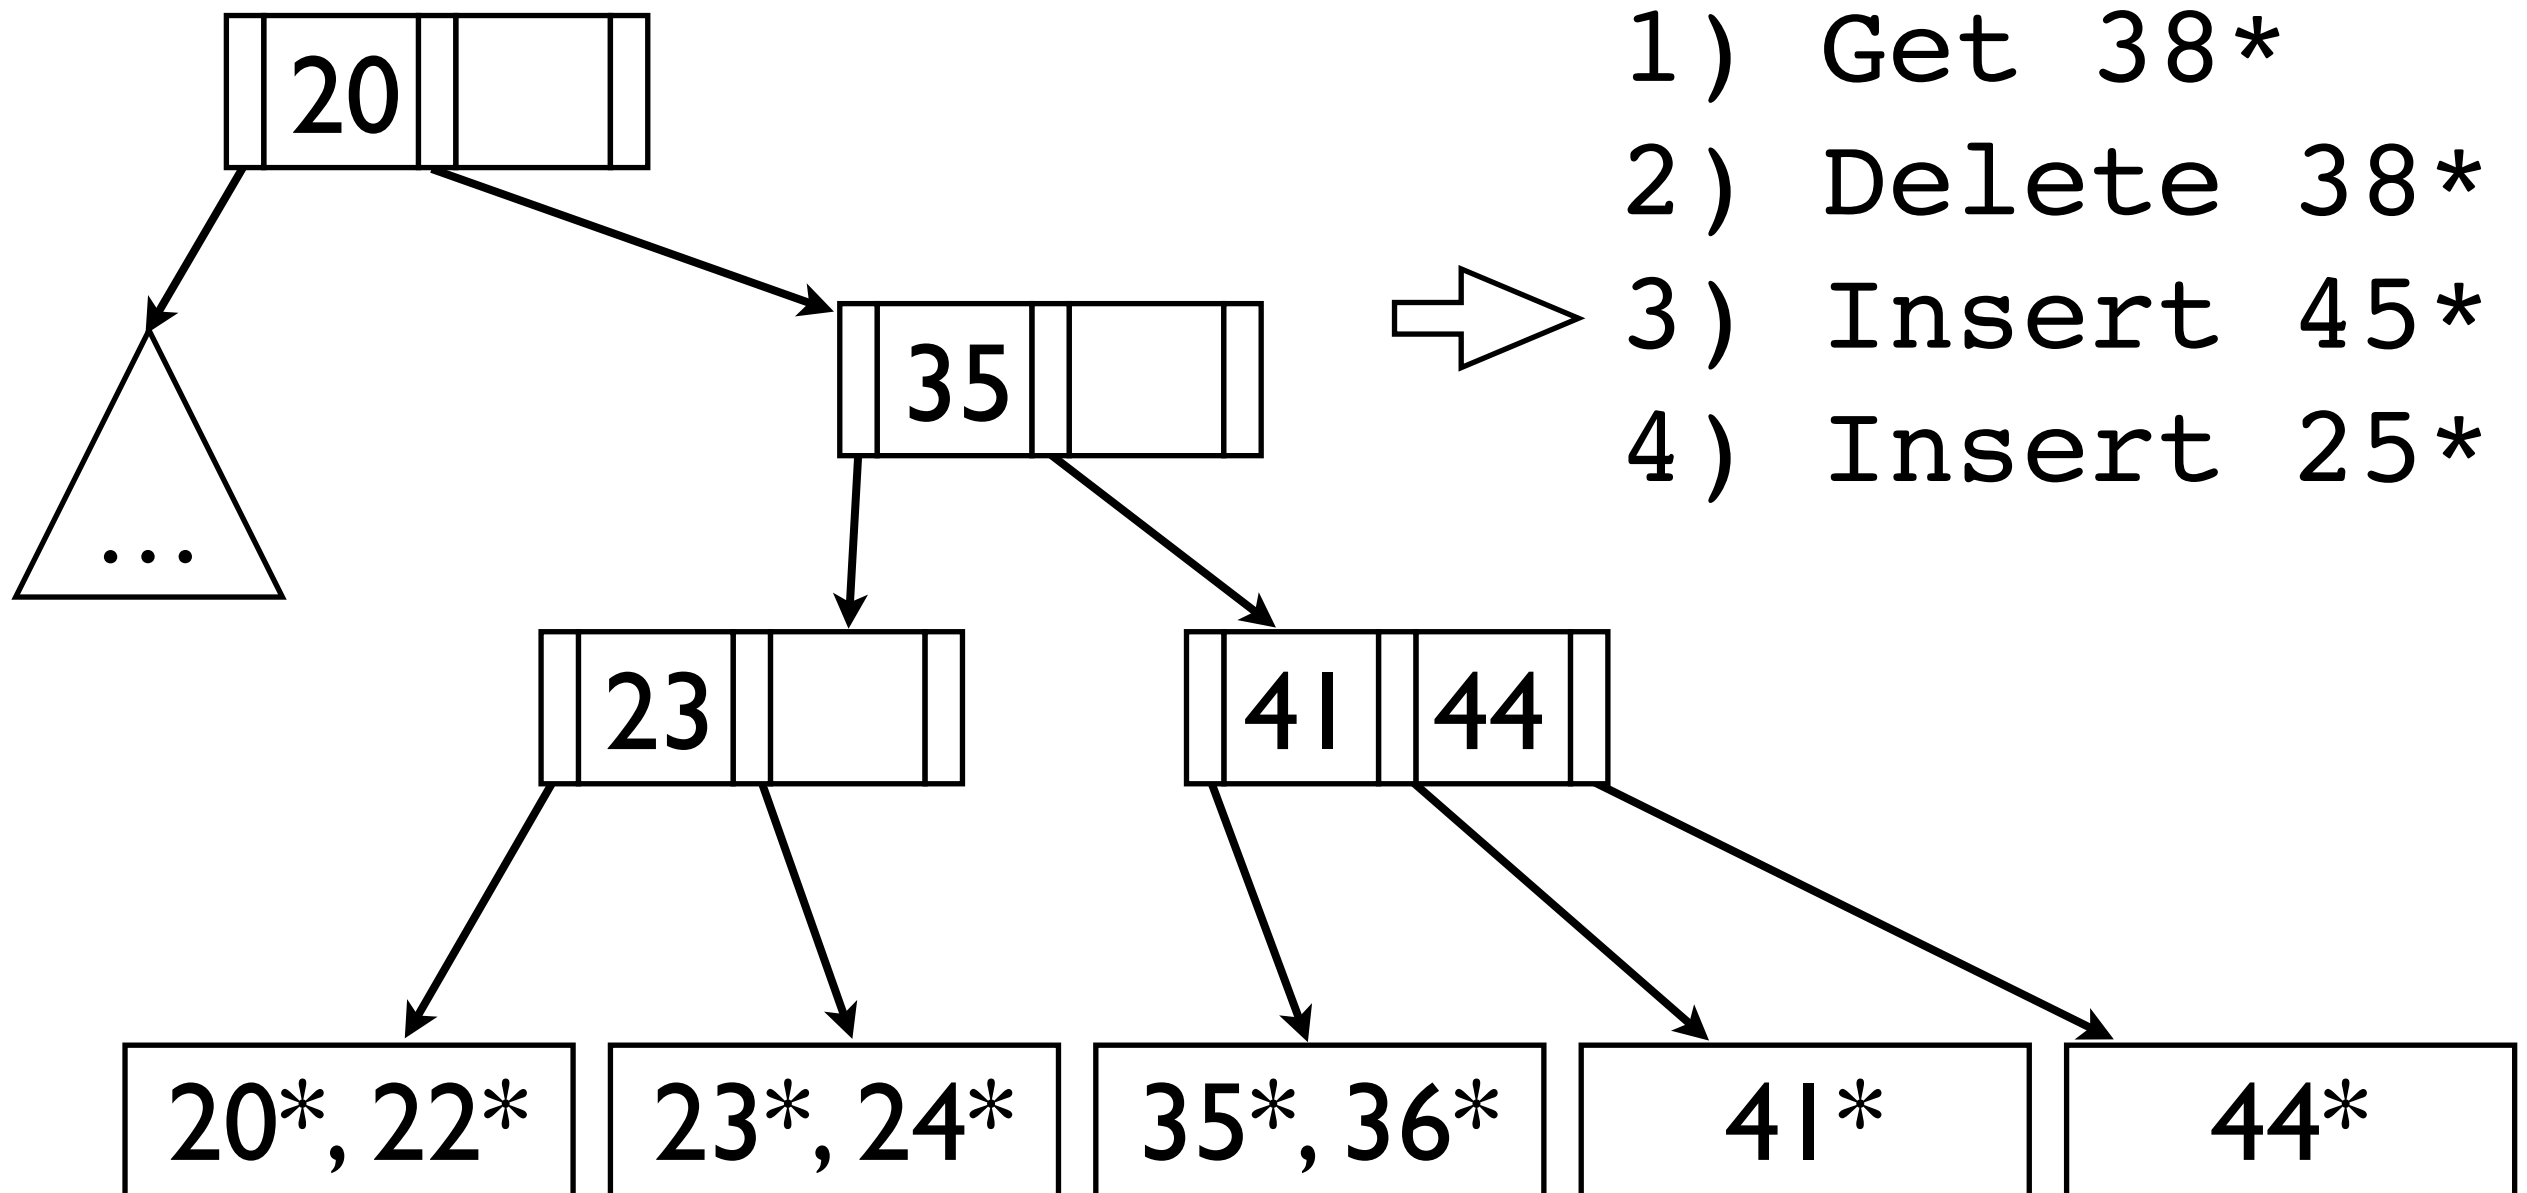
Examples



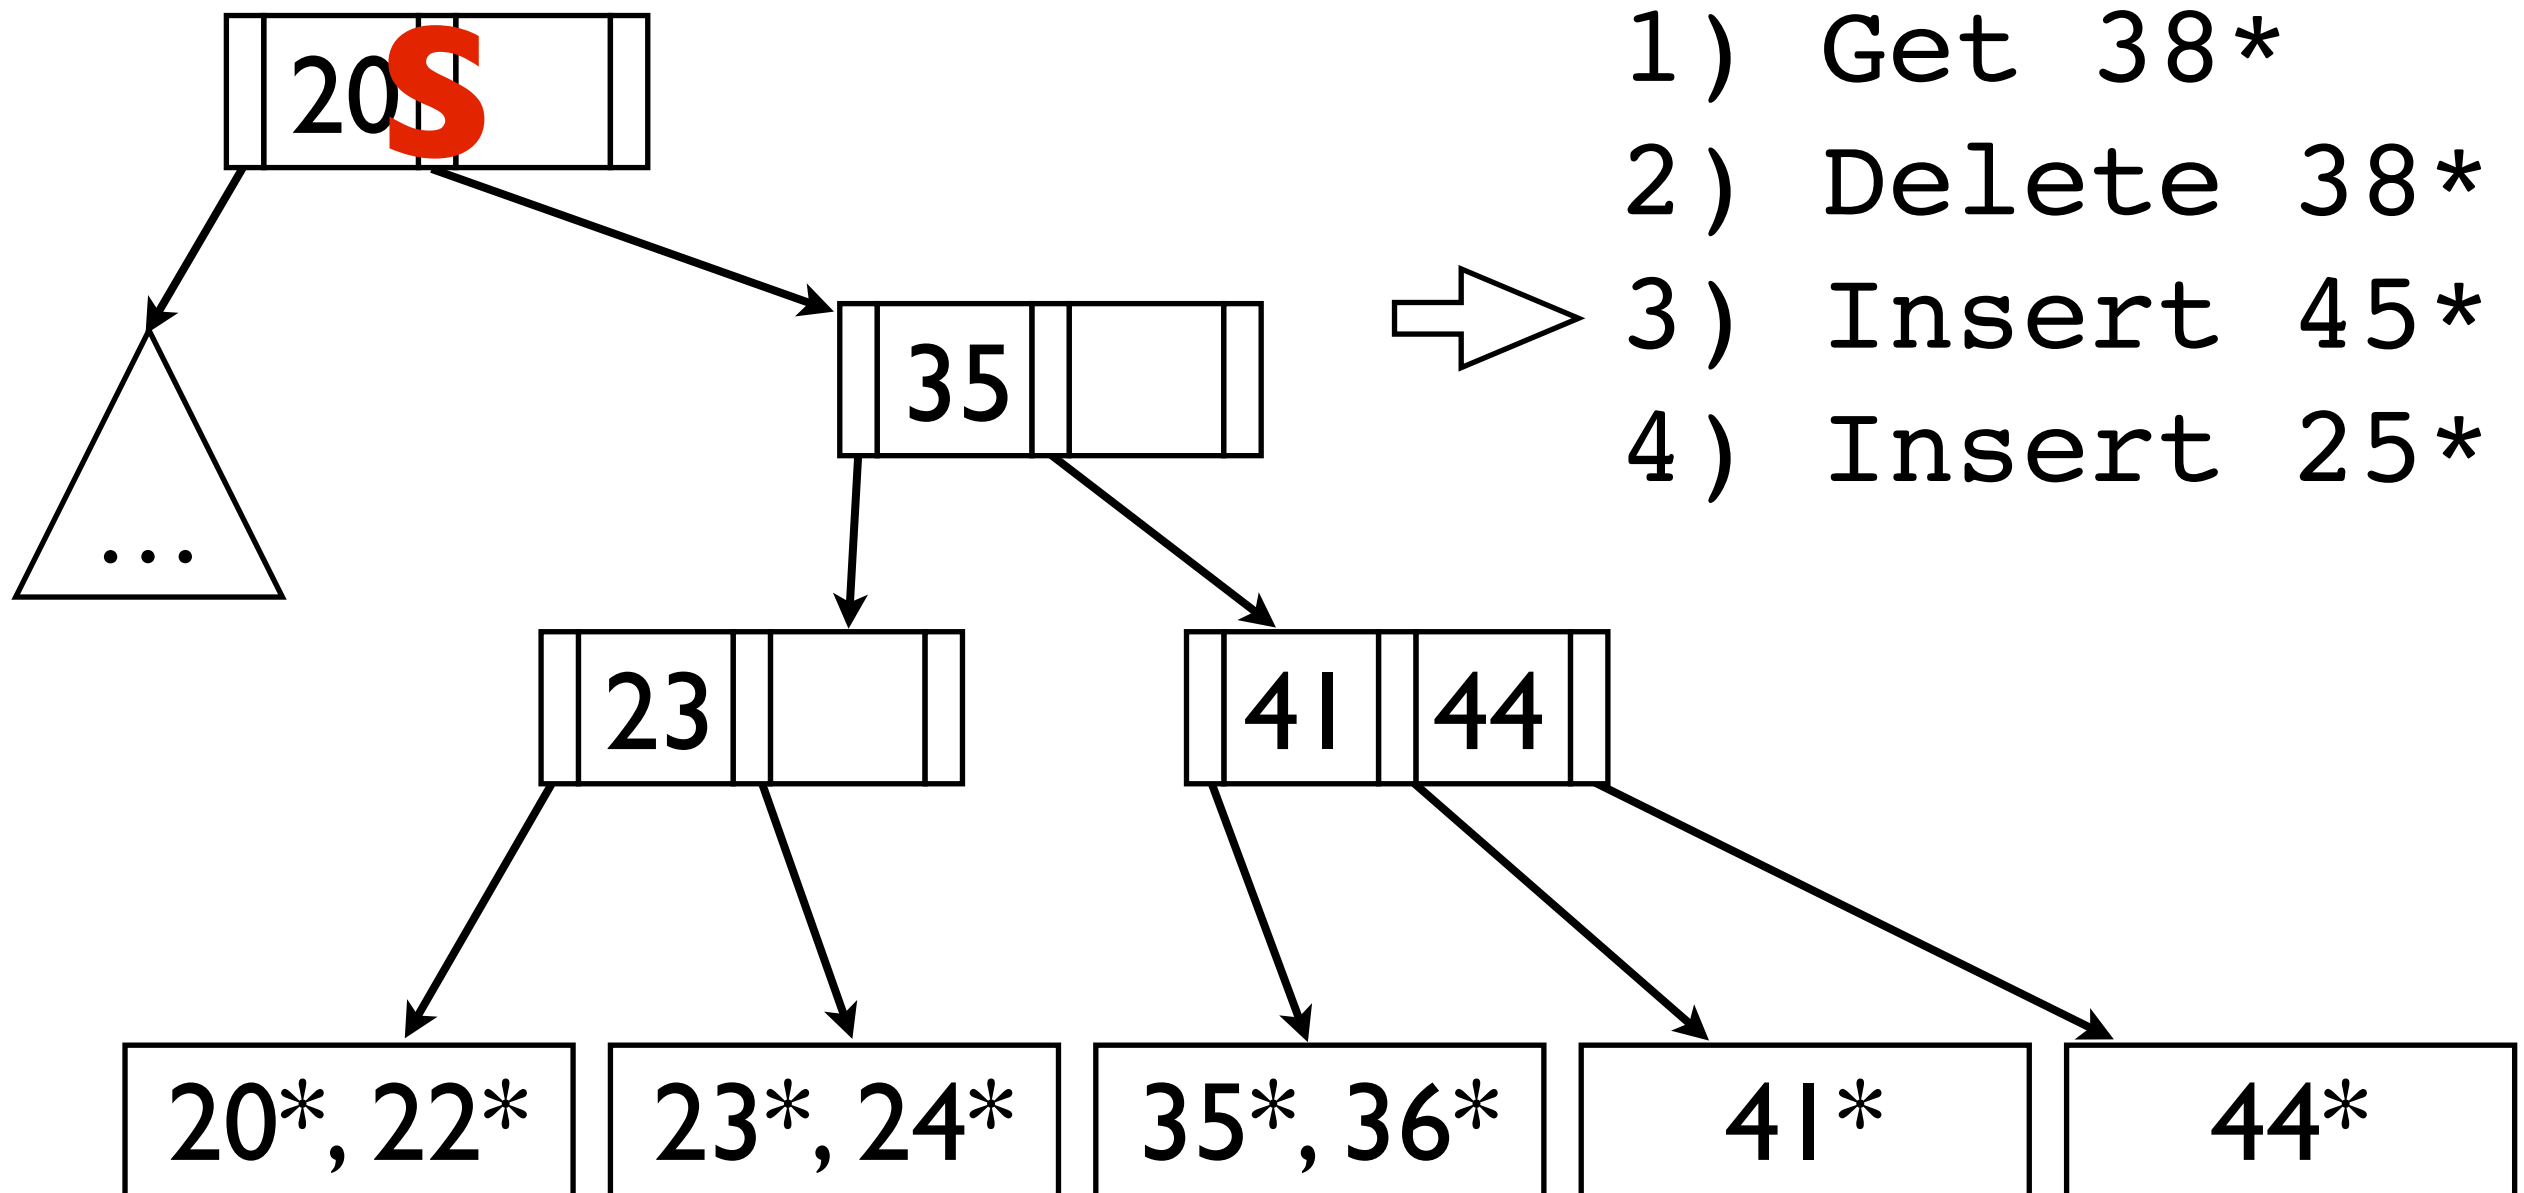
Examples



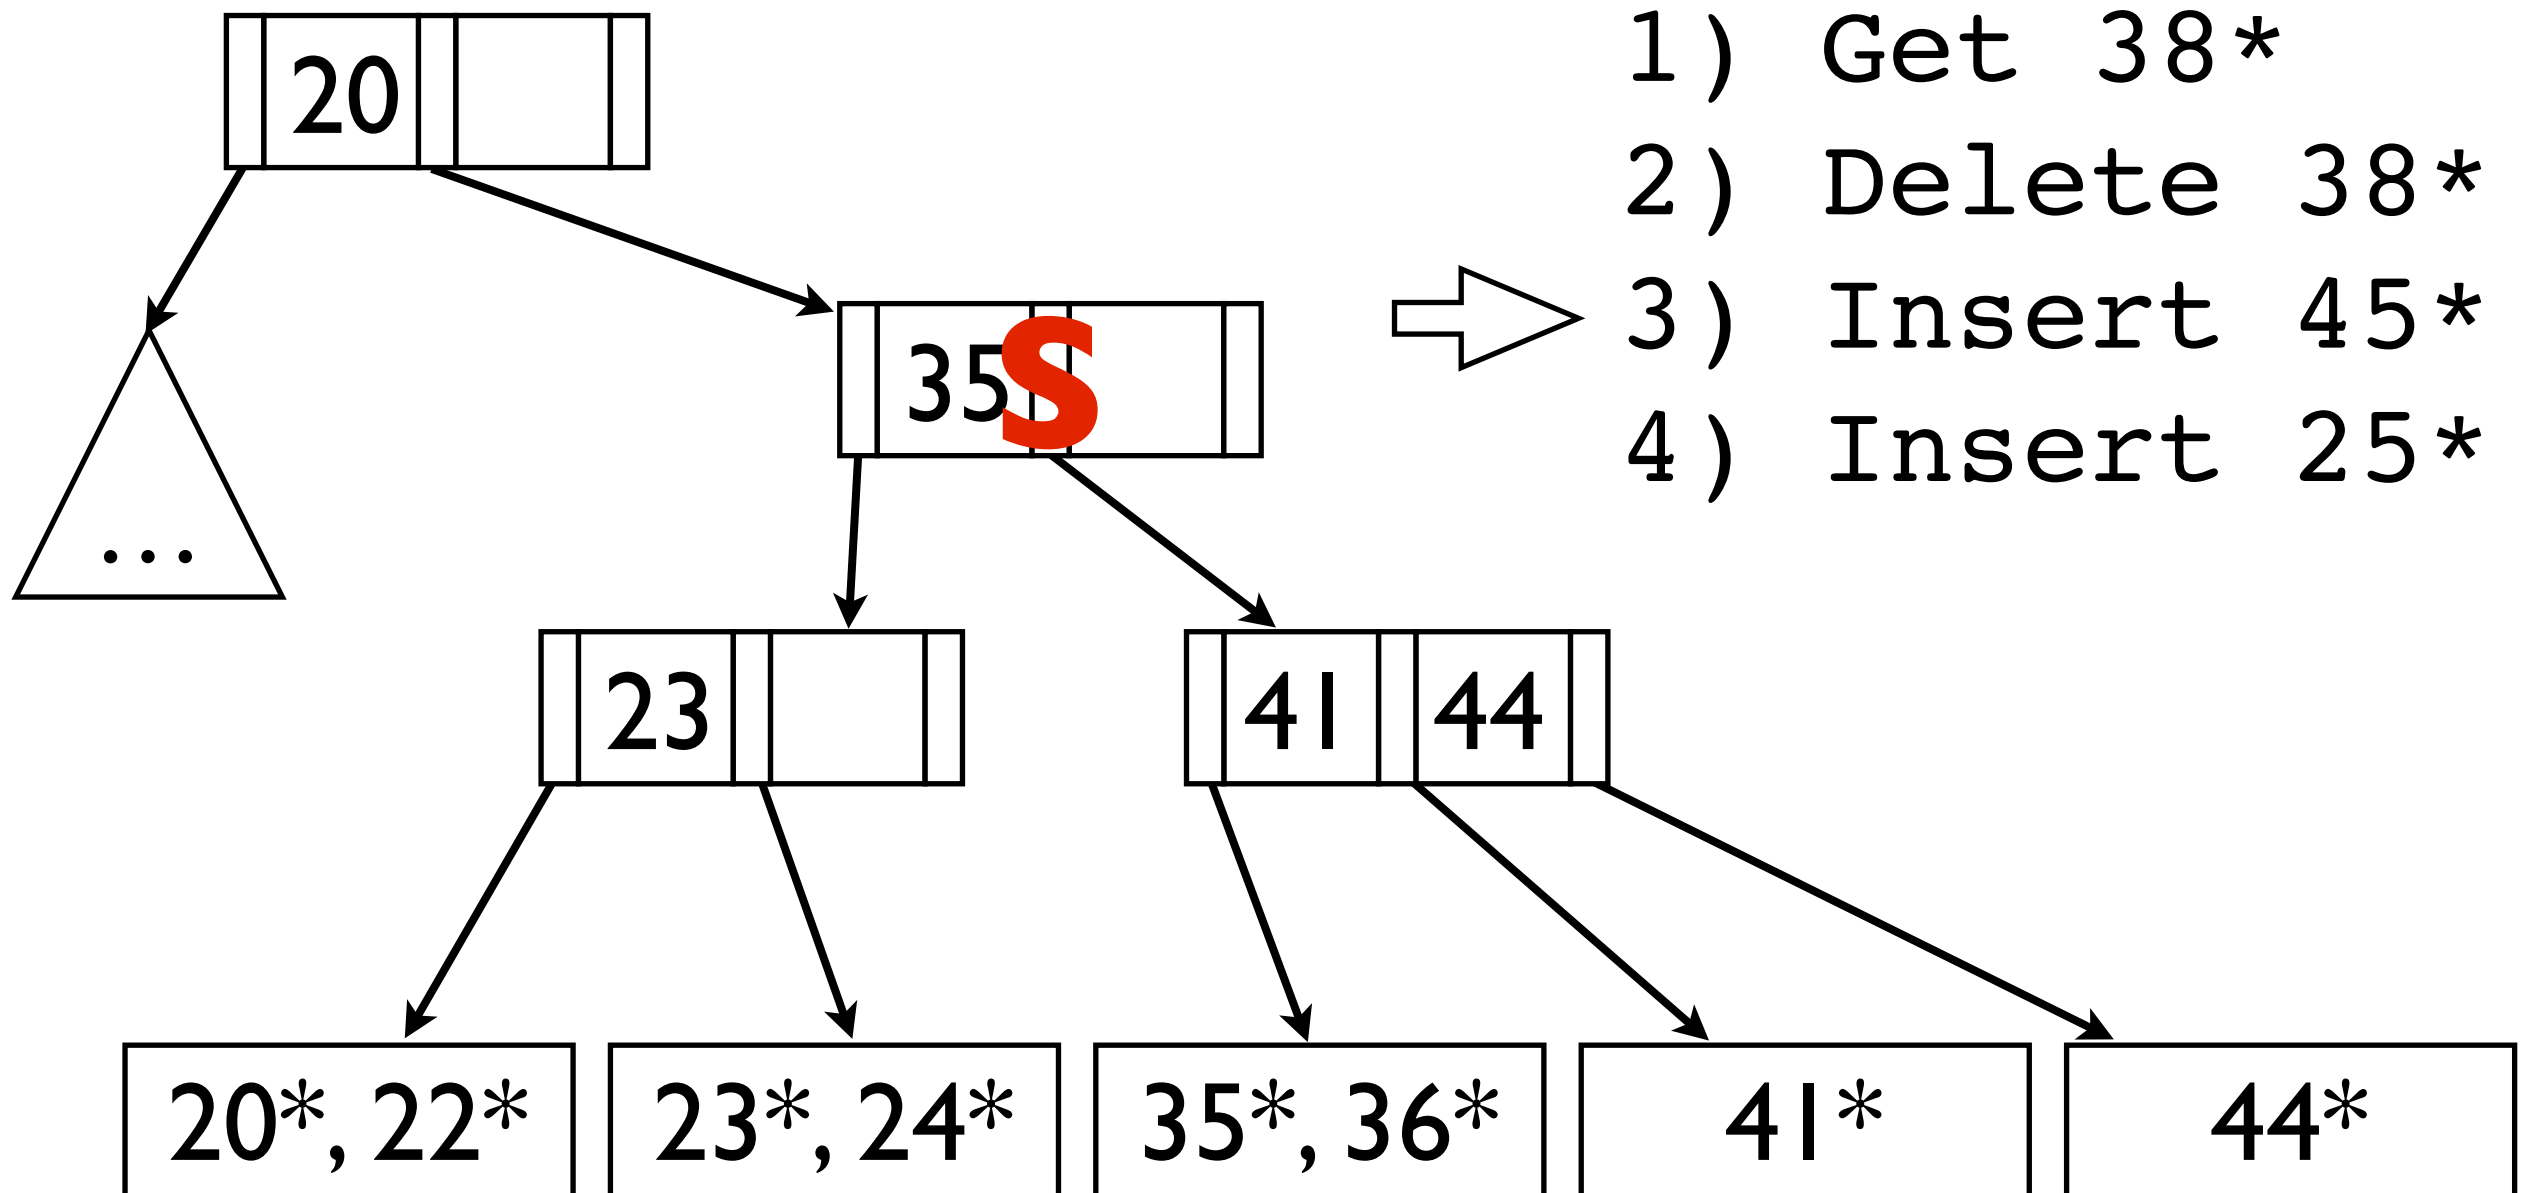
Examples



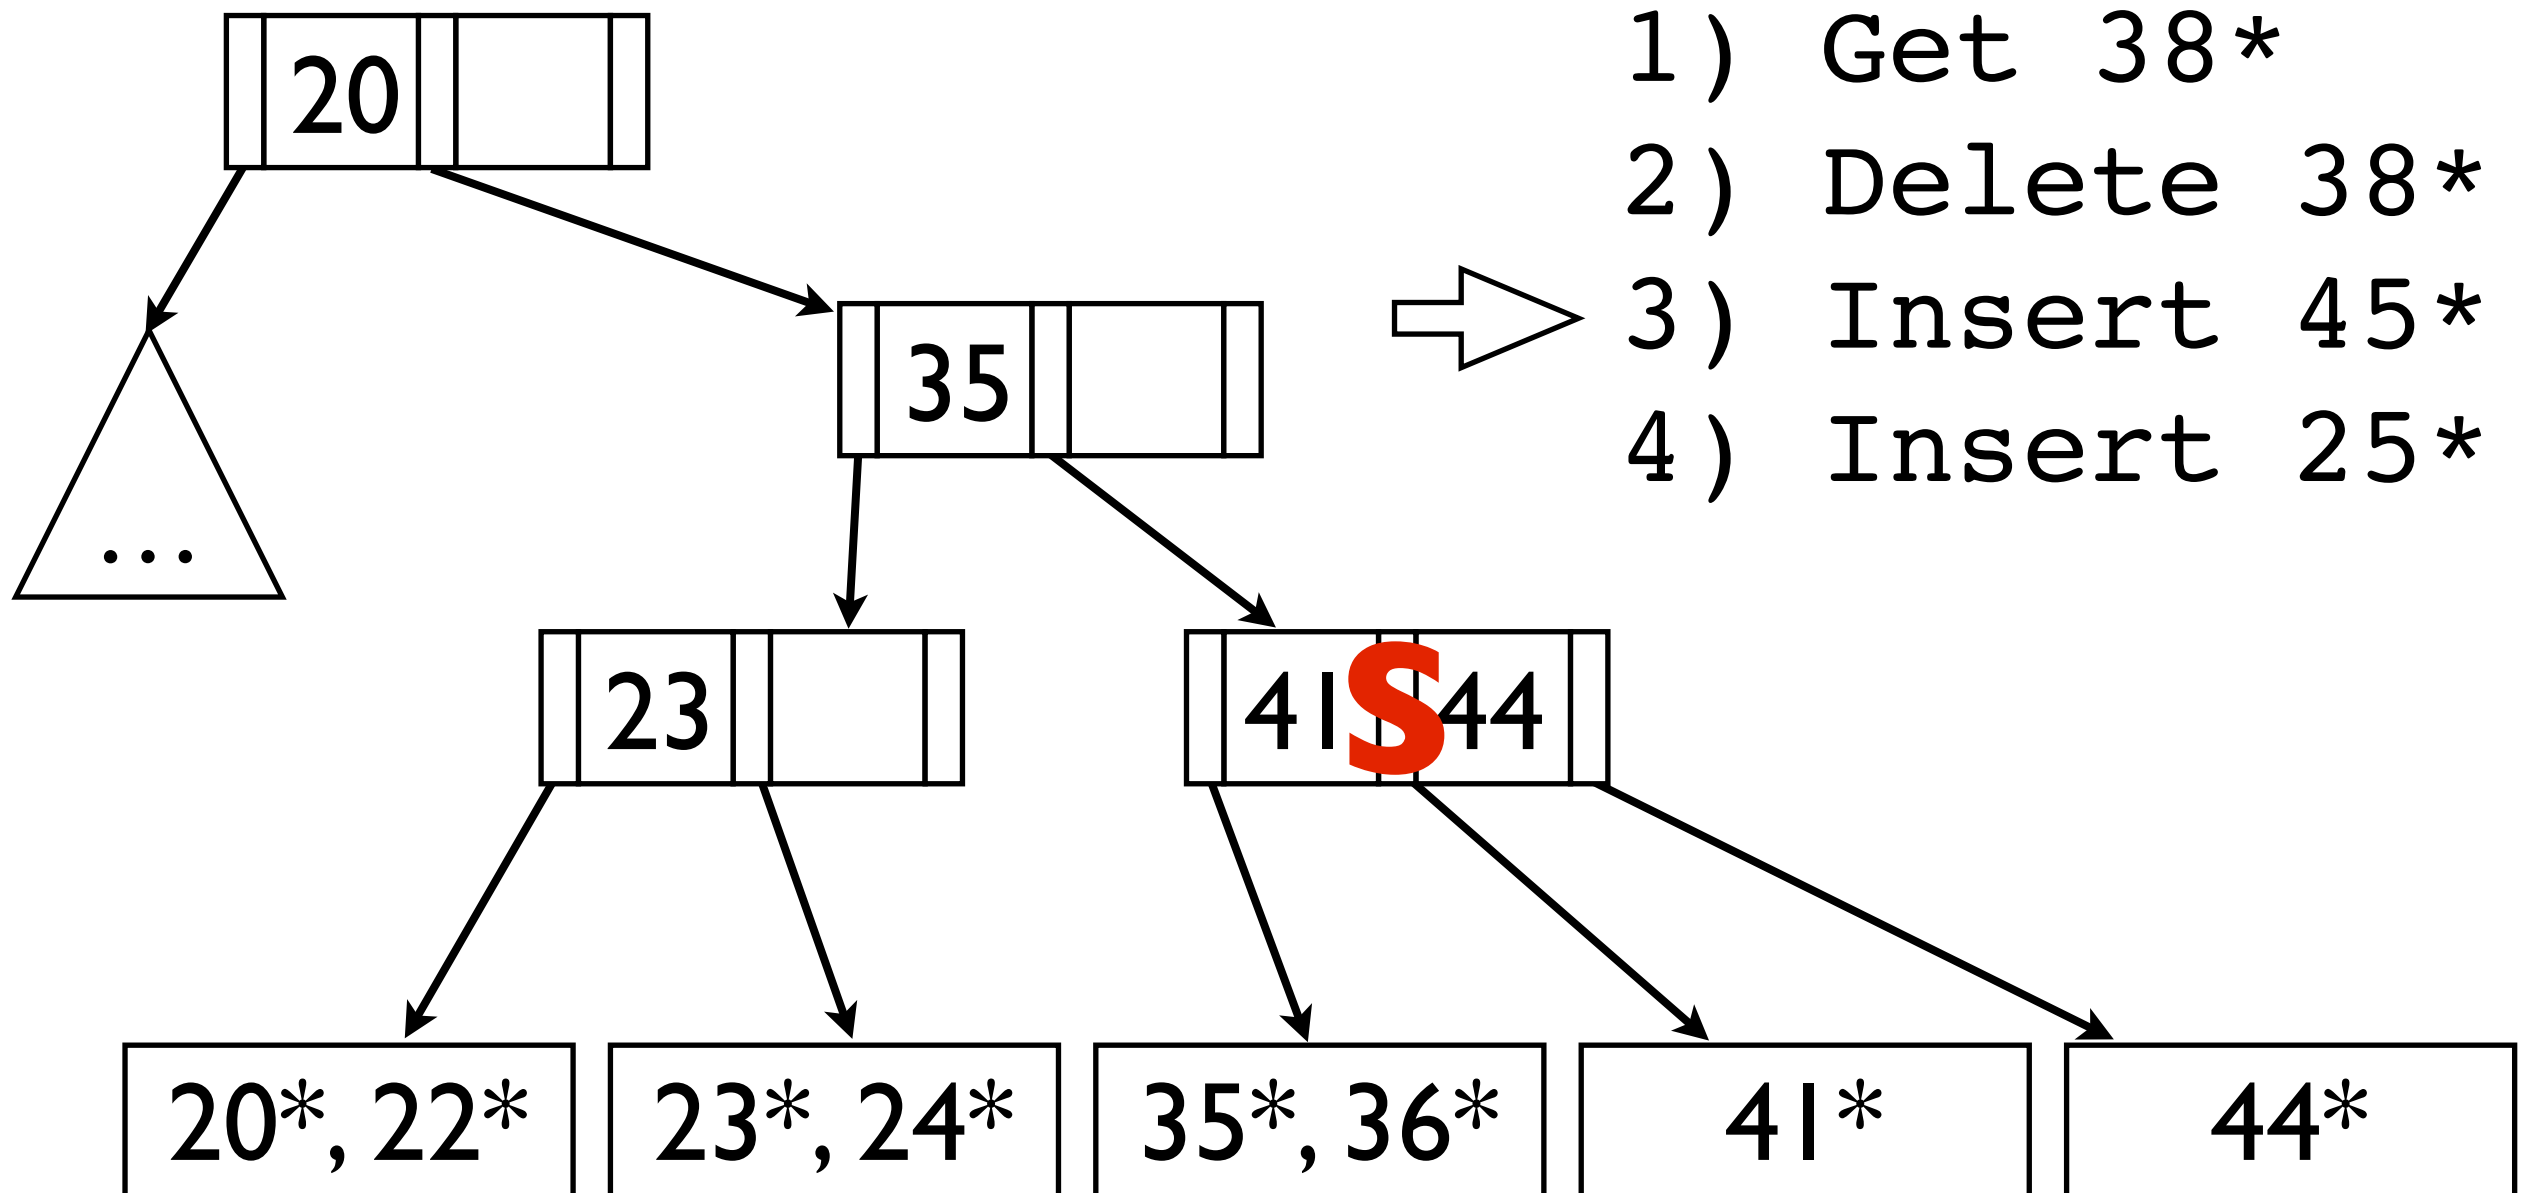
Examples




Examples



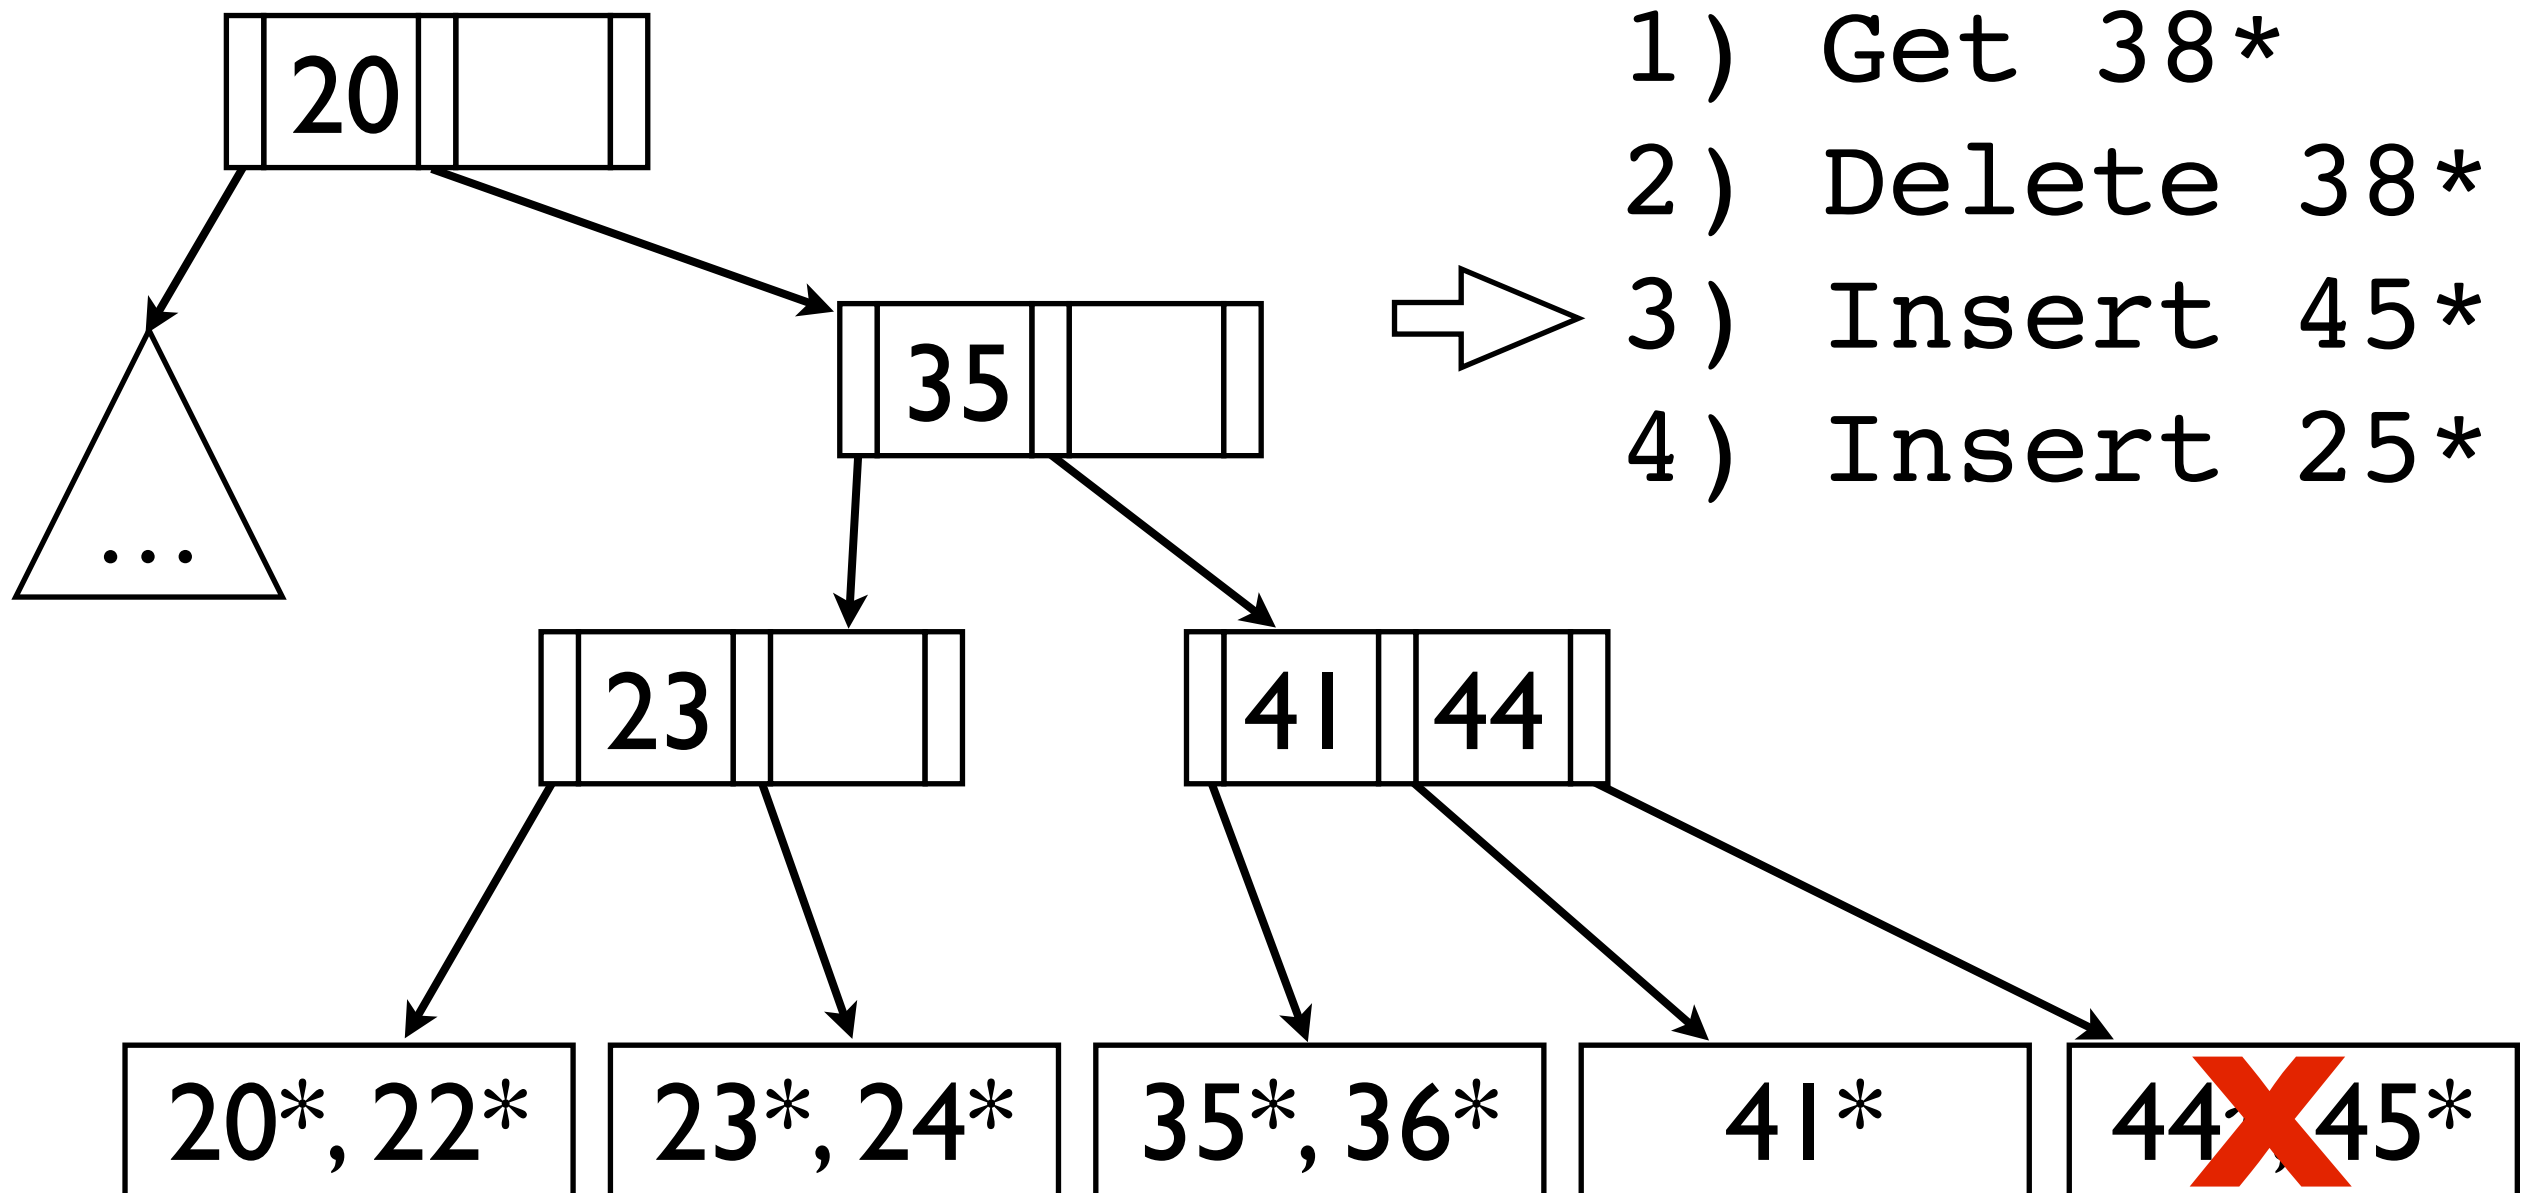
Examples




Examples




Examples



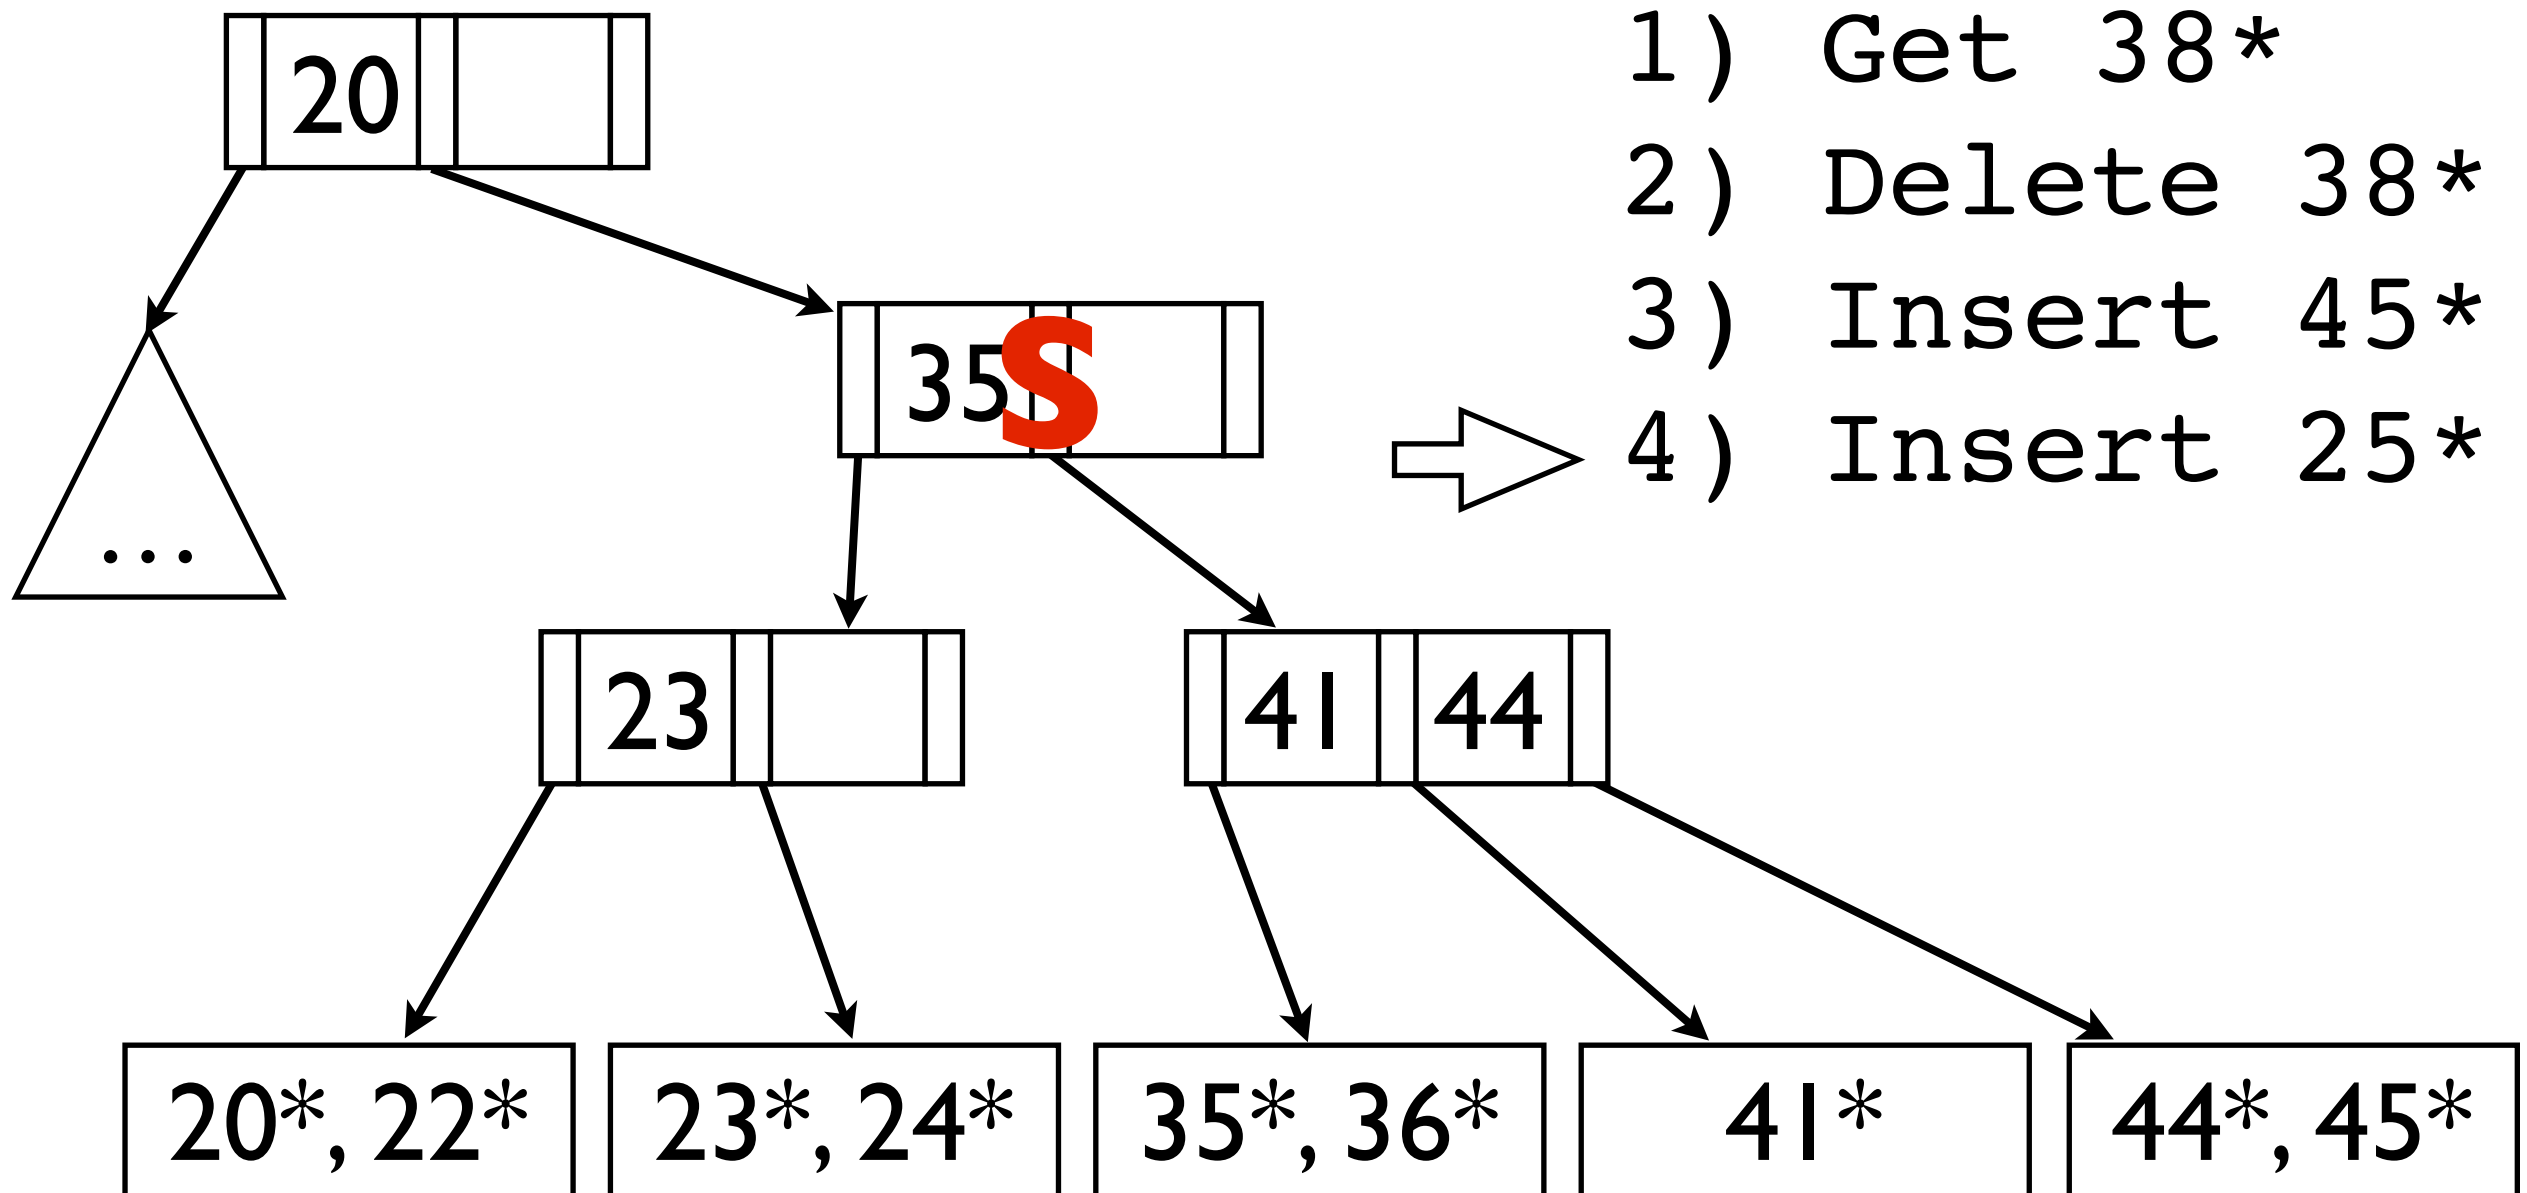
Examples




Examples



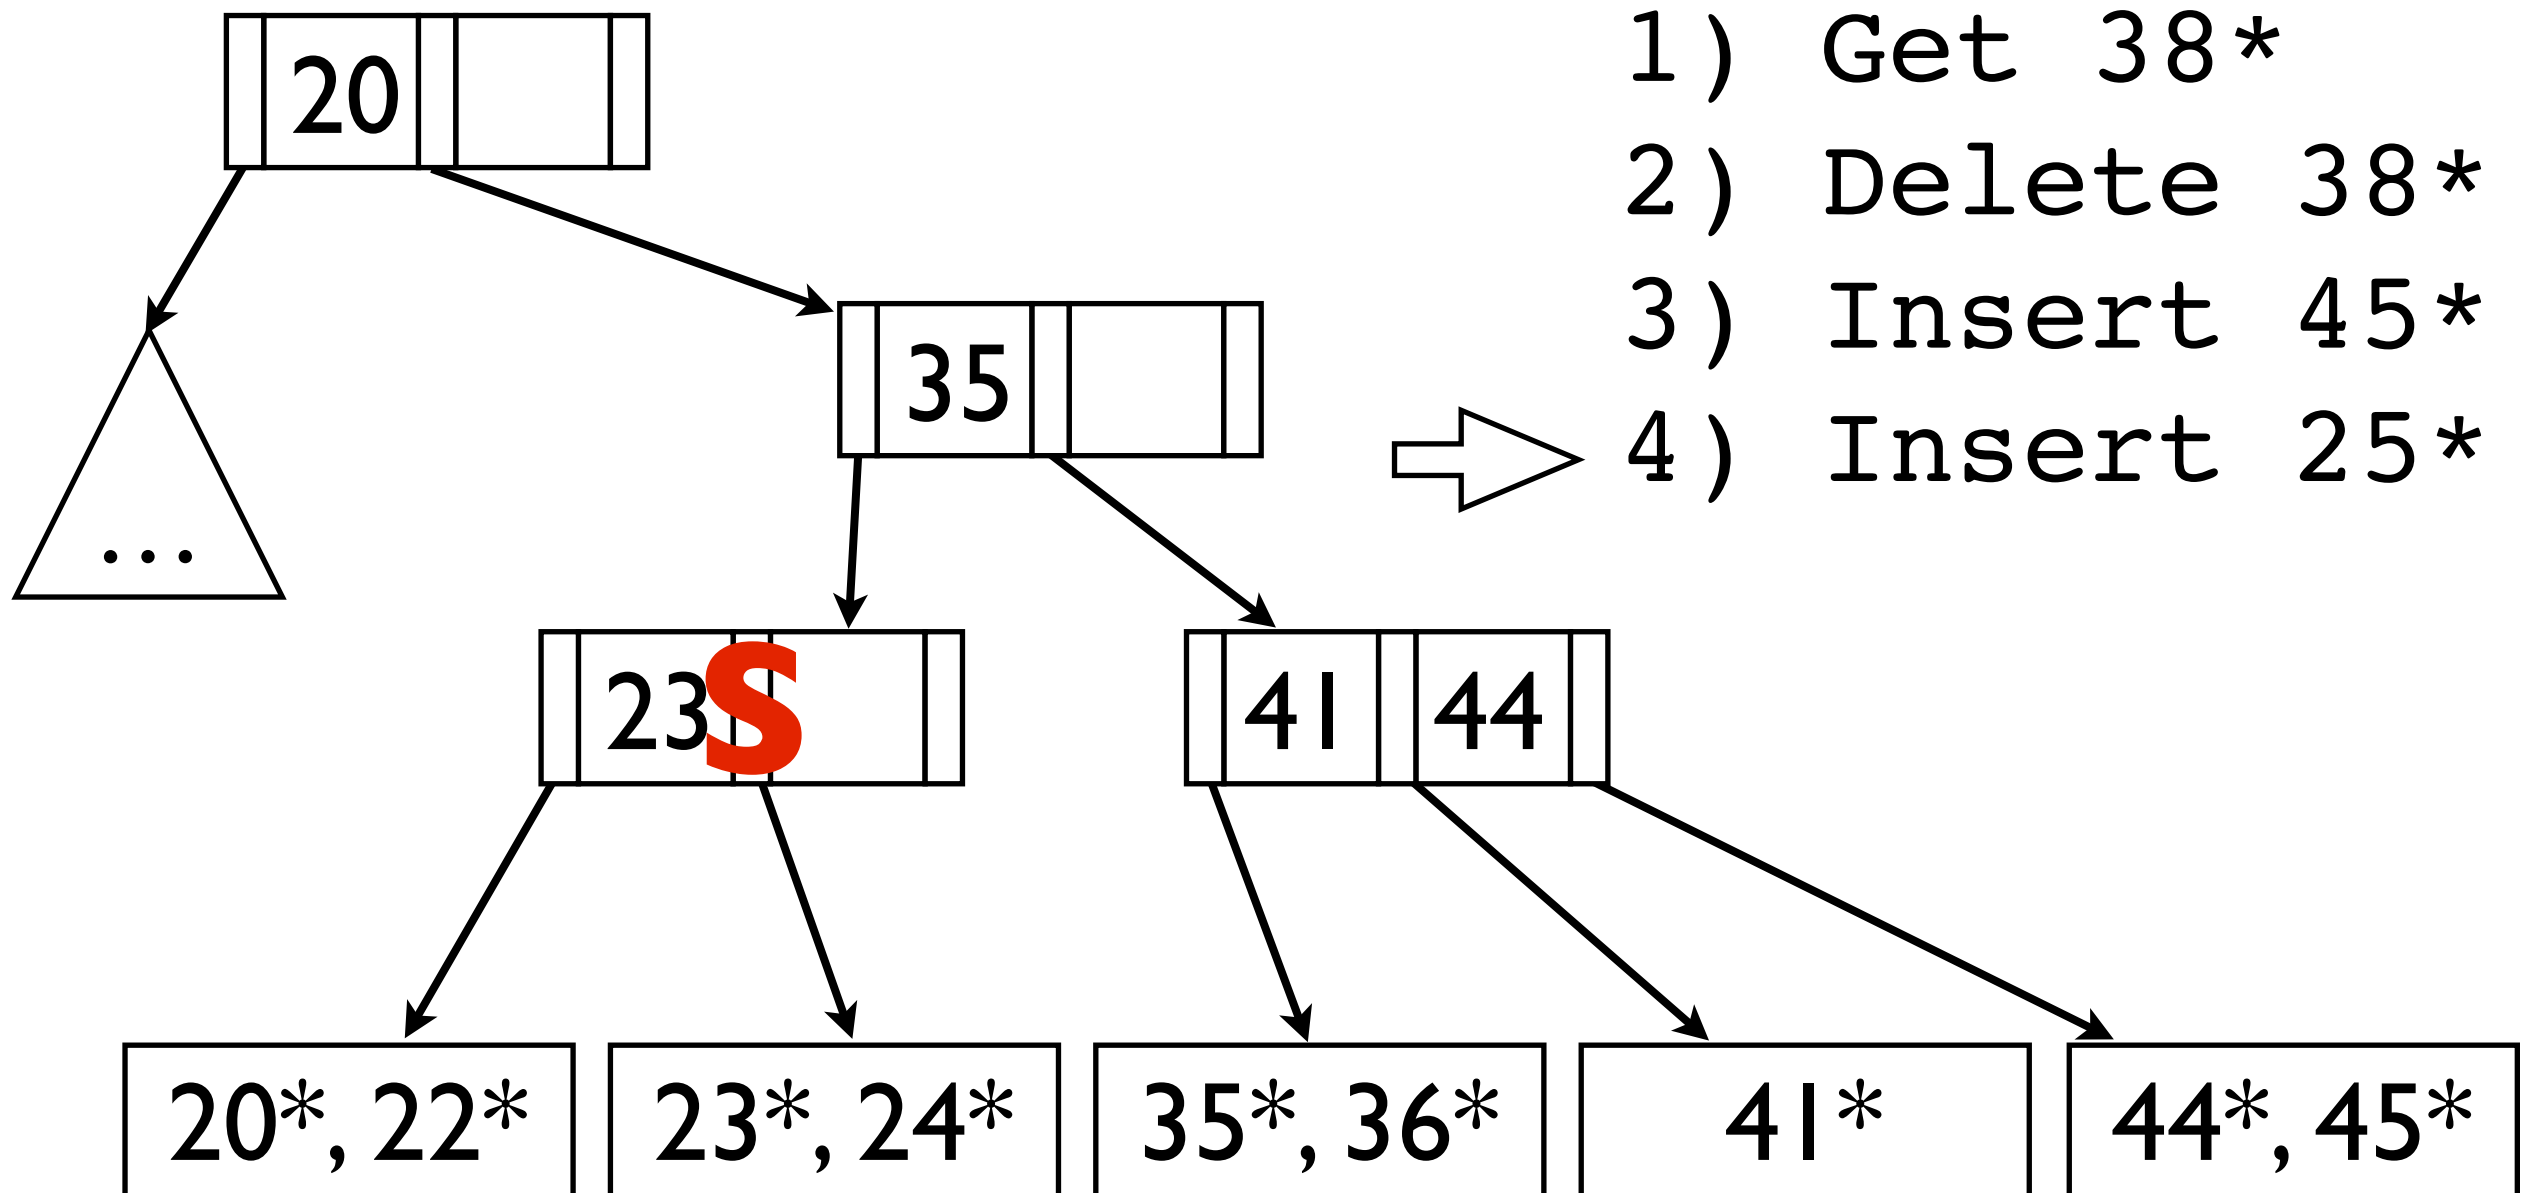
Examples



Examples



Examples



Restart!

Even Better Algorithm

- **Scan:** As before
- **Update:** Acquire IX locks on tree nodes
 - Acquire X lock on the leaf.
 - Convert IX to X locks as needed. Release rest.
 - Do conversion top-down to limit deadlock
- Contrast IX locks here with their use in multiple-granularity locking.

Hybrid Algorithm

- The likelihood that we need an X lock decreases as we move up the tree!

