

Parallel Databases

R&G Chapter 22

(slides adapted from content by J.Gehrke, J.Shanmugasundaram, and/or C.Koch)

Creating Parallelism

- Division of Labor:
 - Stages of Execution: Pipeline Parallelism
 - Parts of Data: Partition Parallelism
- Benefits: Speed Up vs Scale-Up
- Architectures:
 - Shared Memory, Shared Disk, Shared Nothing

Creating Parallelism

- Inter-Operator Parallelism
 - (Mostly) Easy: Partition operators between nodes.
- Inter-Query Parallelism
 - Hard: Need to deal with xacts (next lecture)
- Intra-Operator Parallelism
 - Our focus for today.

Partitioning the Data

- Which machine does a data value end up at?
- Range Partitioning:
 - Split data up into N ranges.
- Hash Partitioning
 - Split data up based on hash value.
- Round Robin
 - Value 1 to machine 1, 2 to machine 2, etc...

Parallel Sorts

- Record as of March 2013: TritonSort (UCSD)
 - 0.725 TB/Min using 52 nodes
 - 1 Trillion Records in ~2.5 min
- <http://sortbenchmark.org>
- http://sortbenchmark.org/2011_06_tritonsort.pdf

Parallel Sorts

6

7

5

1

2


4

8

3

Parallel Sorts

Partition



Parallel Sorts

Partition

Local Sort

6

7

5

1

2

4

8

3

6

7

1

5

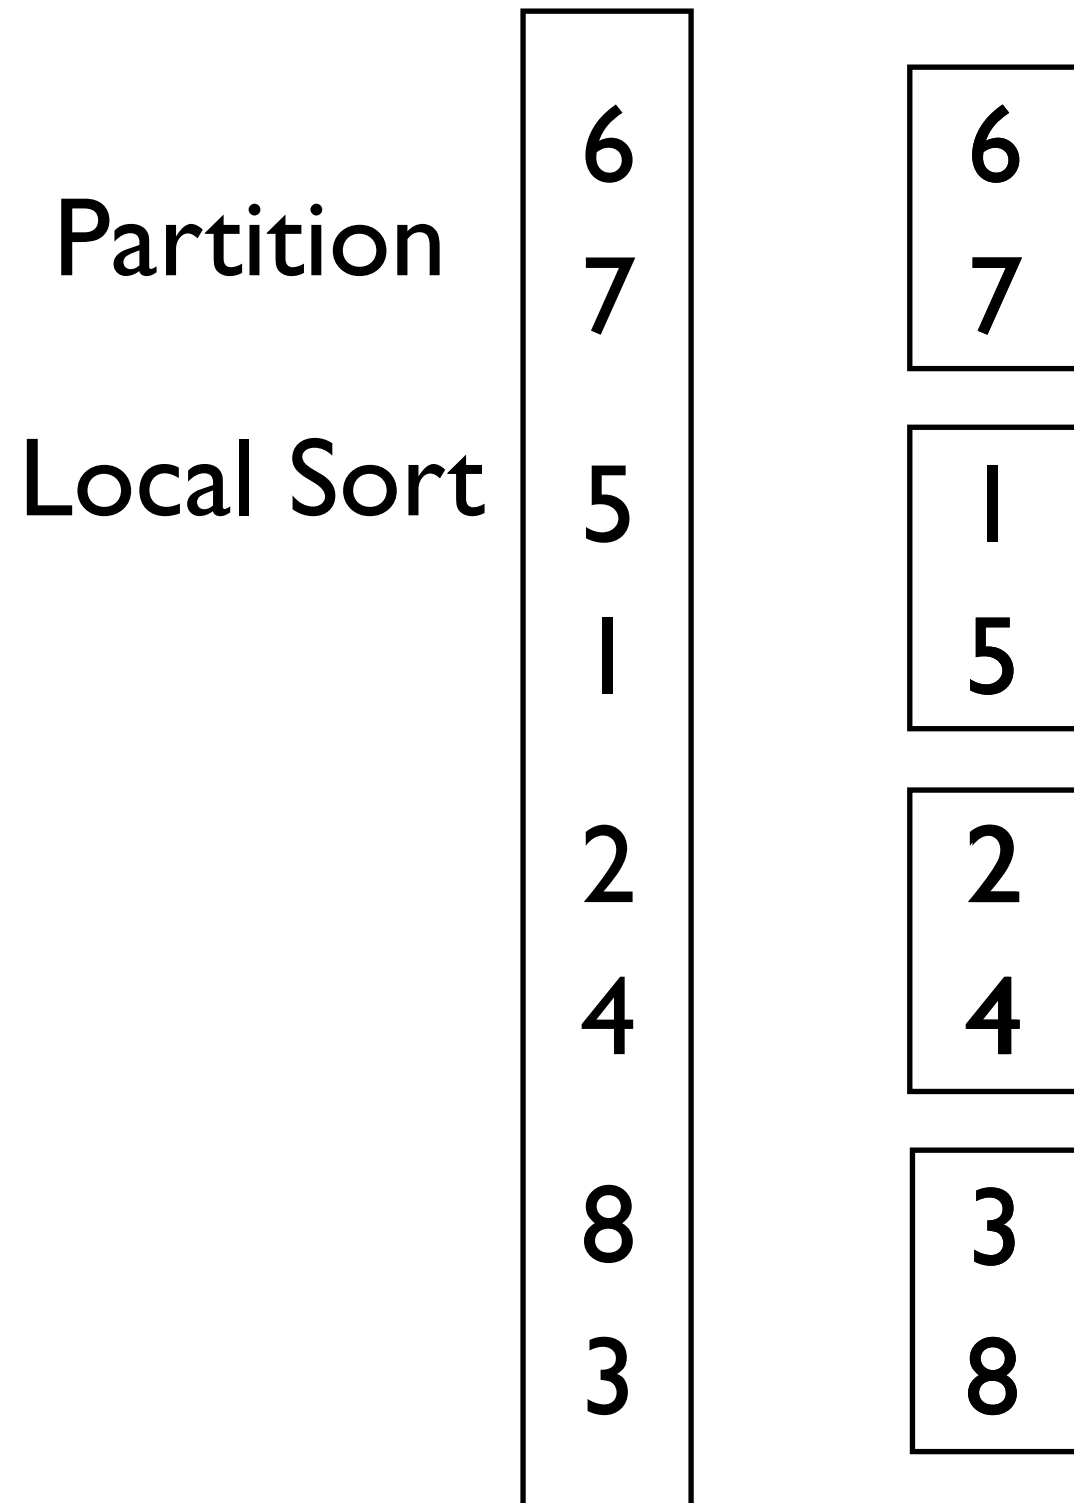
2

4

3

8

Parallel Sorts



Parallel Sorts

Partition

6

7

6

7

1

6

Local Sort

5

1

5

1

Merge Sort

2

4

8

3

2

4

3

8

2

3

Parallel Sorts

Partition

6

7

6

7

1

5

Local Sort

5

1

6

1

5

7

Merge Sort

2

2

2

4

4

3

8

3

4

3

8

8

Parallel Sorts

Partition

6

7

6

7

1

5

Local Sort

5

1

6

1

5

7

Merge Sort

2

2

2

4

4

3

8

3

4

3

8

8

Parallel Sorts

Partition

6

7

6

7

1

5

1

2

Local Sort

5

1

6

3

1

5

7

4

Merge Sort

2

2

2

5

Recur

4

4

3

6

8

3

4

7

3

8

8

8

Parallel Sorts

- General Strategy
 - Scan in parallel
 - Range partition to other nodes
 - As tuples come in, sort locally
- Problem: Skew!
 - How do we pick partition boundaries?

Parallel Scans

- Scan the data separately on each machine
- Merge the results together
 - E.g., compute $\text{SUM}(X)$ on each machine
 - SUM all computed $\text{SUM}(X)$
- May not need all machines to participate.
- We can build an index at each site.

Parallel Aggregates

- Aggregate Functions:
 - Distributive: $F(A,B,C,D) = F(F(A,B),F(C,D))$
 - Algebraic: $F(A,B,C,D) = G(H(A,B),H(C,D))$
 - Holistic: Can't decompose F.
- What are some examples of each?
- How do we exploit this for ||ism?

Parallel Group By

- Compute partial aggregates locally for each group.
- Use hash group-by columns to partition groups.
- Send partial aggregate to a site for each group.
- Each group's site merges partial aggregates.


Parallel Joins-NLJ

- Analogous to Block-NLJ
- One site computes each block-pair.




Parallel Joins-NLJ

- Analogous to Block-NLJ
- One site computes each block-pair.




Parallel Joins-NLJ

- Analogous to Block-NLJ
- One site computes each block-pair.




Parallel Joins-NLJ

- **Optimization:** Skip blocks that are guaranteed to be empty




Parallel Joins-NLJ

- **Optimization:** Skip blocks that are guaranteed to be empty



Parallel Joins-NLJ

- **Optimization:** Skip blocks that are guaranteed to be empty



Parallel Joins-NLJ

- **Limitation:** Requires a large amount of network communication.
 - Why does this happen?
 - Why is this a problem?
 - What can be done about it?

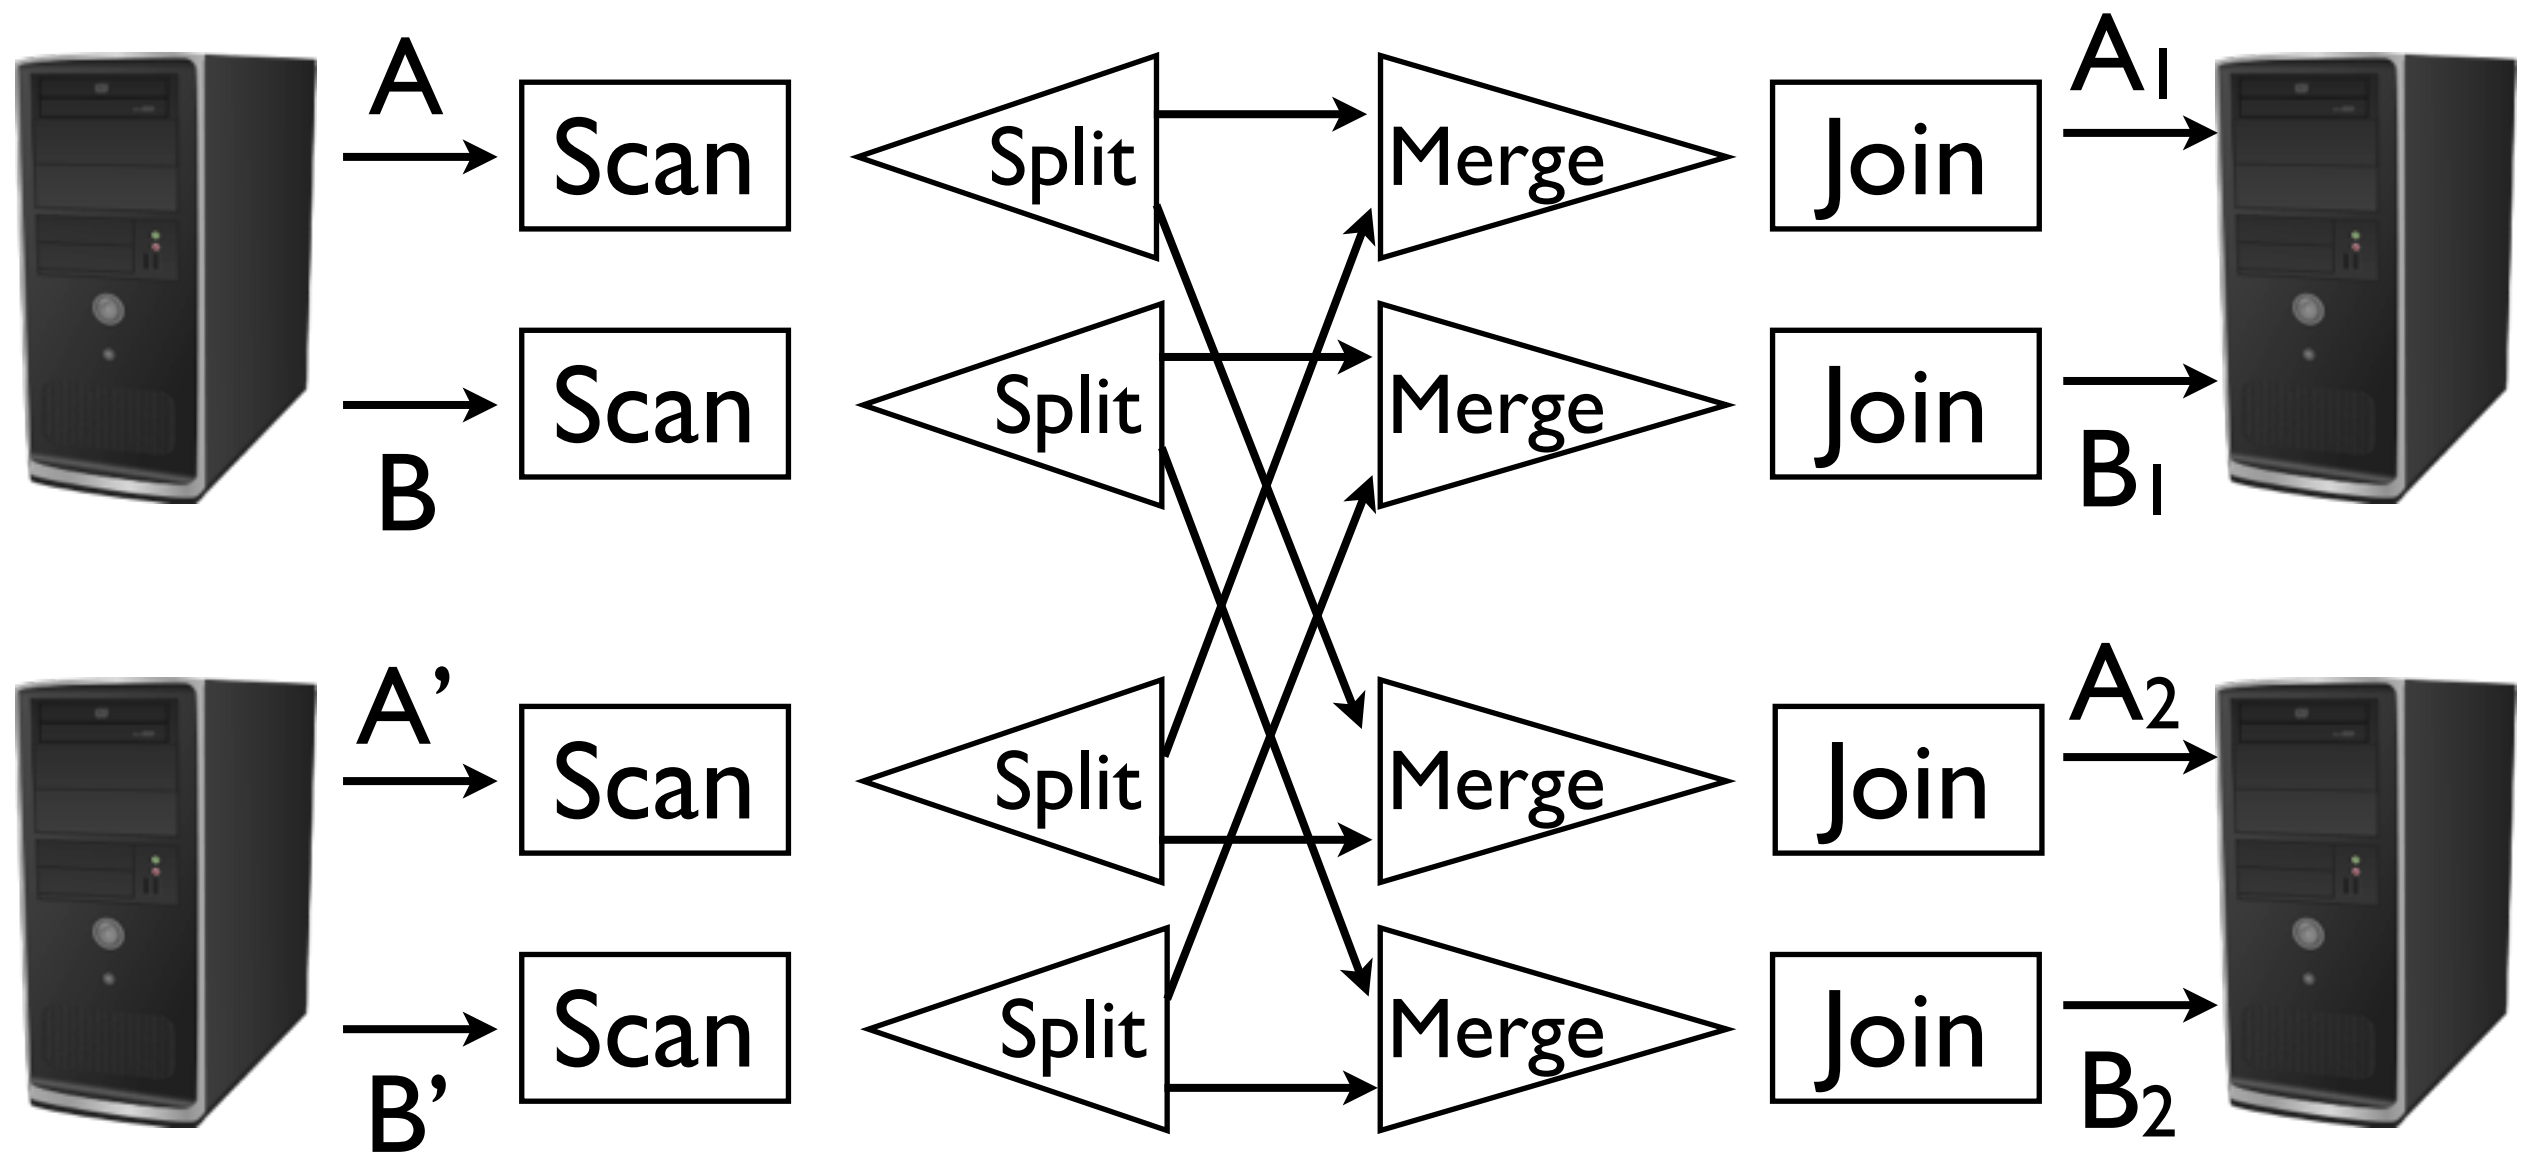
Parallel Joins-Merge Join

- Case 1: Identical Range Partitioning
 - Devolves to Single-Site Sort/Merge
- Case 2: Different Partitioning Schemes
 - Sorting gives range partitions! (what about skew?)
- Case 3: Similar Range Partitioning
 - How would you handle this case?


Parallel Joins-Hash

- Phase 1: Partition/Split the data by hash
 - Each partition goes to a different site.
 - (should produce roughly equivalent workloads)
- Phase 2: Join
 - Perform Joins locally
- Almost always the fastest for Equi-Join
 - Common algorithm for equi-joins in Map/Reduce.

Parallel Joins-Hash




Multi-Stage Query Plans



The Partitioning Phase (Phase I) of Hash Join is Blocking!

How do we get Operator-Level Parallelism?

Multi-Stage Query Plans



The Partitioning Phase (Phase I) of Hash Join is Blocking!

How do we get Operator-Level Parallelism?

Optimizing || Query Plans

- Building a fast parallel query execution engine is (relatively) easy.
- Optimizing queries for ||ism is much harder.
 - Even more variables to consider.
 - Optimizer complexity is high!
 - Lots of research being done in this area.

Optimizing || Query Plans

- Common Approach: 2 Phase Optimizer
 - Phase 1: Pick best sequential plan.
 - Phase 2: Assign operators to sites.
 - “Decorate” query tree as above.
- What’s wrong with this?

Example

- `SELECT * FROM Officers WHERE Name < "Spock"`
- 'Officers' partitioned on name with secondary index on name as well.
- Sites 1-5, all < 'Spock' : Table Scan Best Access Path
- Site 6: 'Scott' through 'Uhura' : Index Scan!