

Project 3

Optimized Query Execution

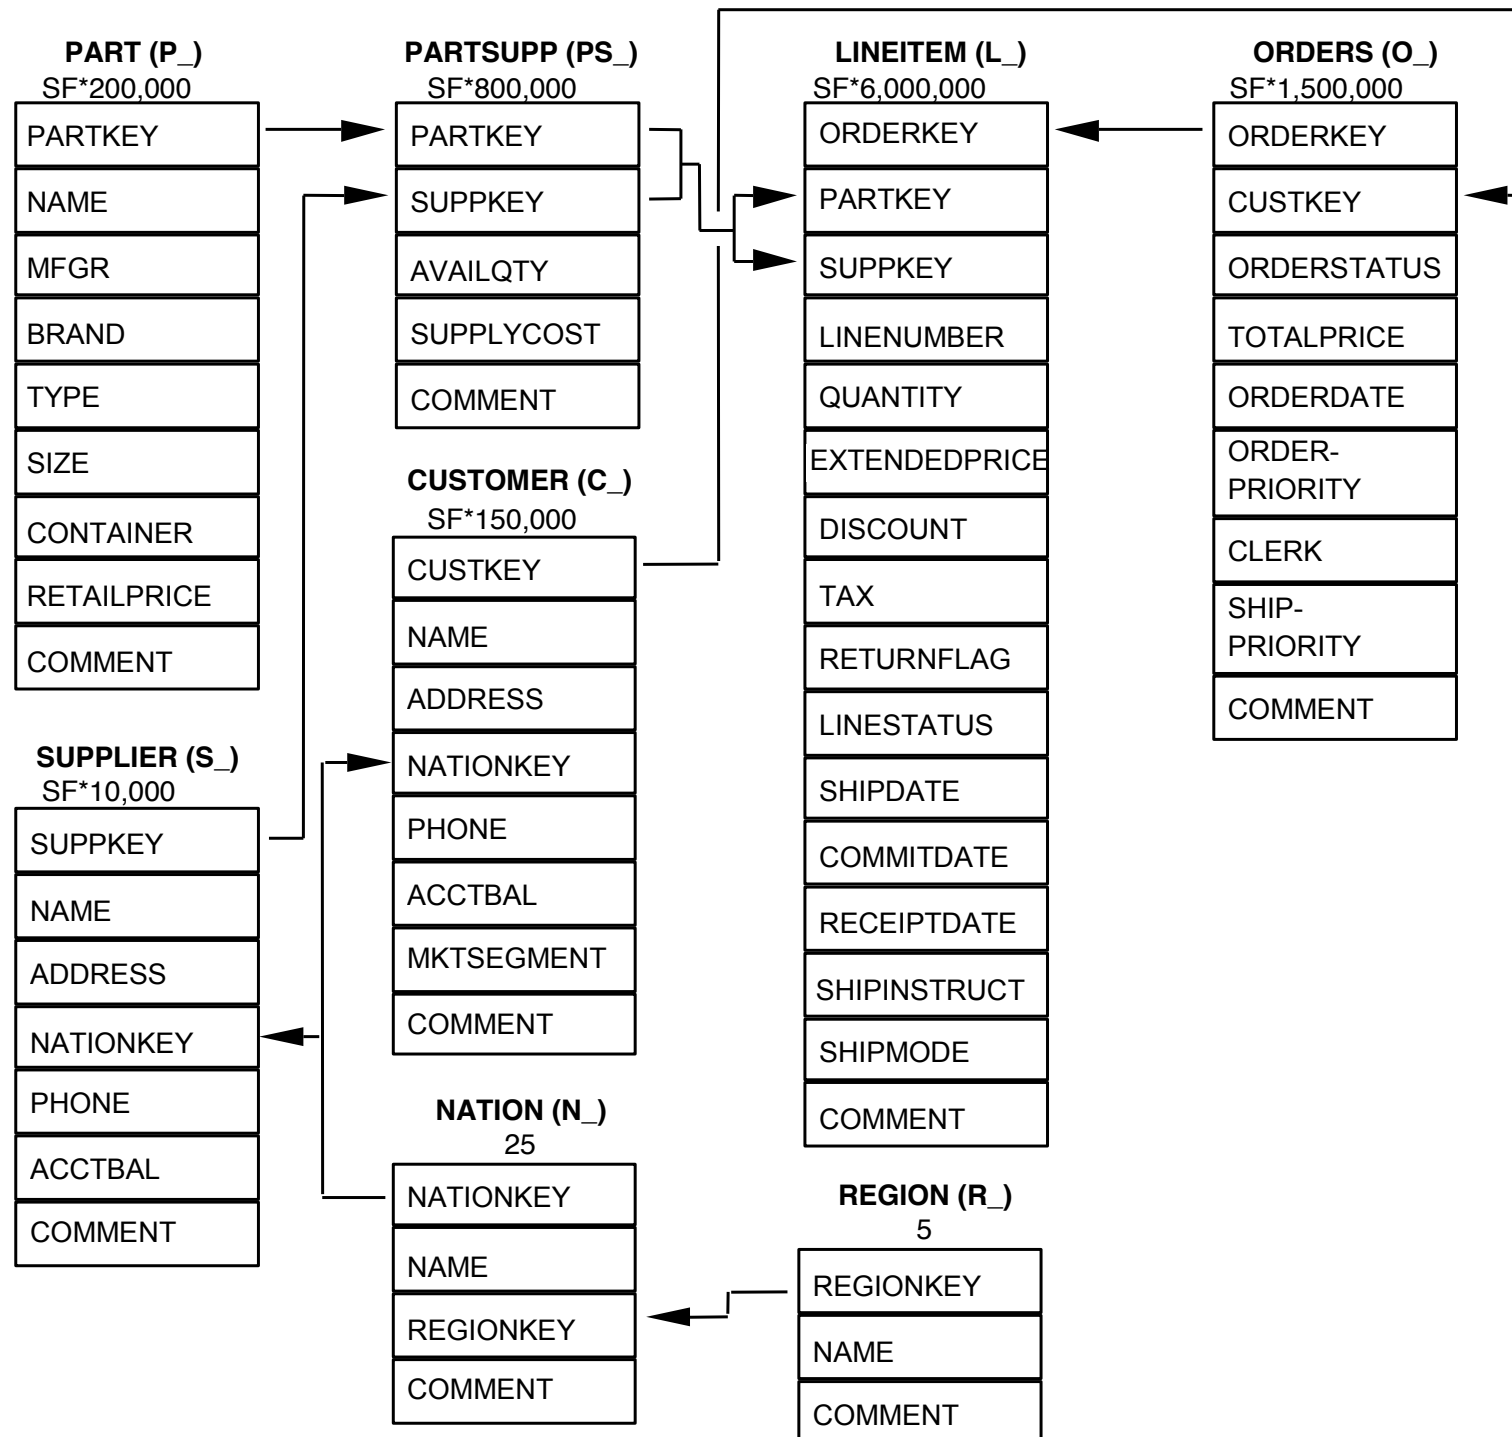
Overview

1. Get your query processor to work with TPC-H.
2. Implement Hybrid Hash or Sort/Merge Join and a Relational Algebra Query Optimizer.
3. Implement Index Scan and Index-Nested Loop Join, and integrate them into your Optimizer.

TPC-H

- TPC - The Transaction Processing Council.
 - ... defines database benchmarks.
- Benchmark H: Decision Support
 - <http://www.tpc.org/tpch/>
 - A major benchmark for database systems.
 - 22 query templates and a data generator.

TPC-H Schema



Part I: DBGen

- Compiling DBGen: See instructions in project writeup.
 - Copy makefile.suite and modify it
 - Special instructions for OSX
- `dbgen -s [size in GB]`
 - Creates one .tbl file for each table.

DBGGen Data File

```
1|3691|O|194029.55|1996-01-02|5-LOW|Clerk#000000951|0|nstructions sleep furiously among |
2|7801|O|60951.63|1996-12-01|1-URGENT|Clerk#000000880|0| foxes. pending accounts at the pending, silent asymptot|
3|12332|F|247296.05|1993-10-14|5-LOW|Clerk#000000955|0|sly final accounts boost. carefully regular ideas cajole carefully. depos|
4|13678|O|53829.87|1995-10-11|5-LOW|Clerk#000000124|0|sits. slyly regular warthogs cajole. regular, regular theodolites acro|
5|4450|F|139660.54|1994-07-30|5-LOW|Clerk#000000925|0|quickly. bold deposits sleep slyly. packages use slyly|
6|5563|F|65843.52|1992-02-21|4-NOT SPECIFIED|Clerk#000000058|0|ggle. special, final requests are against the furiously specia|
7|3914|O|231037.28|1996-01-10|2-HIGH|Clerk#000000470|0|ly special requests |
32|13006|O|166802.63|1995-07-16|2-HIGH|Clerk#000000616|0|ise blithely bold, regular requests. quickly unusual dep|
33|6697|F|118518.56|1993-10-27|3-MEDIUM|Clerk#000000409|0|uriously. furiously final request|
34|6101|O|75662.77|1998-07-21|3-MEDIUM|Clerk#000000223|0|ly final packages. fluffily final deposits wake blithely ideas. spe|
35|12760|O|192885.43|1995-10-23|4-NOT SPECIFIED|Clerk#000000259|0|zzle. carefully enticing deposits nag furio|
36|11527|O|72196.43|1995-11-03|1-URGENT|Clerk#000000358|0| quick packages are blithely. slyly silent accounts wake qu|
37|8612|F|156440.15|1992-06-03|3-MEDIUM|Clerk#000000456|0|kly regular pinto beans. carefully unusual waters cajole never|
38|12484|O|64695.26|1996-08-21|4-NOT SPECIFIED|Clerk#000000604|0|haggle blithely. furiously express ideas haggle blithely furiously regular re|
39|8177|O|307811.89|1996-09-20|3-MEDIUM|Clerk#000000659|0|ole express, ironic requests: ir|
64|3212|F|30616.90|1994-07-16|3-MEDIUM|Clerk#000000661|0|wake fluffily. sometimes ironic pinto beans about the dolphin|
65|1627|P|99763.79|1995-03-18|1-URGENT|Clerk#000000632|0|ular requests are blithely pending orbits-- even requests against the deposit|
66|12920|F|100991.26|1994-01-20|5-LOW|Clerk#000000743|0|y pending requests integrate|
67|5662|O|167270.36|1996-12-19|4-NOT SPECIFIED|Clerk#000000547|0|symptotes haggle slyly around the furiously iron|
68|2855|O|305815.83|1998-04-18|3-MEDIUM|Clerk#000000440|0| pinto beans sleep carefully. blithely ironic deposits haggle furiously acro|
69|8449|F|228015.94|1994-06-04|4-NOT SPECIFIED|Clerk#000000330|0| depths atop the slyly thin deposits detect among the furiously silent accou|
70|6434|F|133507.10|1993-12-18|5-LOW|Clerk#000000322|0| carefully ironic request|
71|338|O|244449.86|1998-01-24|4-NOT SPECIFIED|Clerk#000000271|0| express deposits along the blithely regul|
96|10778|F|72504.36|1994-04-17|2-HIGH|Clerk#000000395|0|oost furiously. pinto|
97|2107|F|128590.11|1993-01-29|3-MEDIUM|Clerk#000000547|0|hang blithely along the regular accounts. furiously even ideas after the|
98|10448|F|62956.90|1994-09-25|1-URGENT|Clerk#000000448|0|c asymptotes. quickly regular packages should have to nag re|
99|8891|F|136624.34|1994-03-13|4-NOT SPECIFIED|Clerk#000000973|0|e carefully ironic packages. pending|
100|14701|O|204408.59|1998-02-28|4-NOT SPECIFIED|Clerk#000000577|0|heodolites detect slyly alongside of the ent|
101|2800|O|142434.13|1996-03-17|3-MEDIUM|Clerk#000000419|0|ding accounts above the slyly final asymptote|
102|73|O|172239.95|1997-05-09|2-HIGH|Clerk#000000596|0| slyly according to the asymptotes. carefully final packages integrate furious|
103|2911|O|147675.81|1996-06-20|4-NOT SPECIFIED|Clerk#000000090|0|ges. carefully unusual instructions haggle quickly regular f|
128|7396|F|57495.50|1992-06-15|1-URGENT|Clerk#000000385|0|ns integrate fluffily. ironic asymptotes after the regular excuses nag around |
129|7114|F|273469.52|1992-11-19|5-LOW|Clerk#000000859|0|ing tithes. carefully pending deposits boost about the silently express |
130|3697|F|169065.69|1992-05-08|2-HIGH|Clerk#000000036|0|le slyly unusual, regular packages? express deposits det|
131|9275|F|143898.91|1994-06-08|3-MEDIUM|Clerk#000000625|0|after the fluffily special foxes integrate s|
132|2641|F|180897.80|1993-06-11|3-MEDIUM|Clerk#000000488|0|sits are daringly accounts. carefully regular foxes sleep slyly about the|
133|4400|O|120535.39|1997-11-29|1-URGENT|Clerk#000000738|0|usly final asymptotes |
134|620|F|203218.79|1992-05-01|4-NOT SPECIFIED|Clerk#000000711|0|lar theodolites boos|
135|6049|O|280793.15|1995-10-21|4-NOT SPECIFIED|Clerk#000000804|0|l platelets use according t|
```

DBGen Data File

- ' | ' -Separated Fields
- Datatypes
 - Integers
 - Decimals (Floats)
 - Strings (Varchar, Char)
 - Dates (YYYY-MM-DD)
 - Represent as Integer:YYYYMMDD

Todo?


- Add a main function
 - `./sql [filename]`
- Generate a 0.1 GB Database with DBGen
- Add support for parsing DBGen files.
 - Add support for parsing date types
- Test TPC-H Queries 1, 3, 5, 6, 10 and 19

Part 2: Optimization

$$\pi_{\vec{a}}(\sigma_{c \wedge c'}(\pi_{\vec{a}'}(R) \times (S \times T)))$$


Part 2: Optimization

$$\pi_{\vec{a}}(\sigma_{c \wedge c'}(\pi_{\vec{a}'}(R) \times (S \times T))))$$




$$\pi_{\vec{a}}(\sigma_c(\pi_{\vec{a}'}(R) \bowtie_{c'} (S \times T))))$$


Part 2: Optimization




Part 2: Optimization




Part 2: Optimization




Part 2: Optimization



Part 2: Optimization



Part 2: Optimization



Part 2: Optimization

- `PlanRewrite`: Apply a transformation to every node in a relational algebra tree.
- Find every node where a rewrite is applicable and apply it.
 - E.g., a `Select` node with a `NL join` as its child.
 - Also push-down selects, push-down projects.
- Implement Hybrid Hash or Sort/Merge Join.

Part 2: Optimization

- Implement an -explain flag in your project's main function. Should print out:
 - Before: The original RA query.
 - After: The optimized RA query.

Part 3: Indexing

- Look at the query workload:
 - TPC-H Queries 1, 3, 5, 6, 10 and 19
- Determine what indexes can/will help you.
 - Use -examine!
- Add an -index flag to build these indexes.

Part 3: Indexing

- Implement Index Scan.
- Implement Index-Nested Loop Join.
- Modify your optimizer to create these where appropriate.

Extra Credit

- Implement Sort/Limit
 - Can combine sort with merge-sort
- Track which subtrees are sorted, and on what attribute.
- Avoid re-sorting an already sorted dataset.

Racing Implementations

- Bi-Weekly (Wednesday night, Sunday night)
- All submitted projects will be timed and measured for memory use.
- Your score is based on how well your implementation does relative to the median performance of all groups.
- Top 5 teams each run win 15 bonus pts (up to 30)

Project 4

- Optional Project
- List of project ideas in project 3 writeup.
 - ... or pick your own!
- Add a feature to your query processor.
- Evaluate your added feature.
- Document your findings.

Project 4

- Step 1: Determine which project you want to do.
- Step 2: Decide on implementation and evaluation strategies.
- Step 3: Discuss the project with me.
- Step 4: Go!