

# Data Warehousing

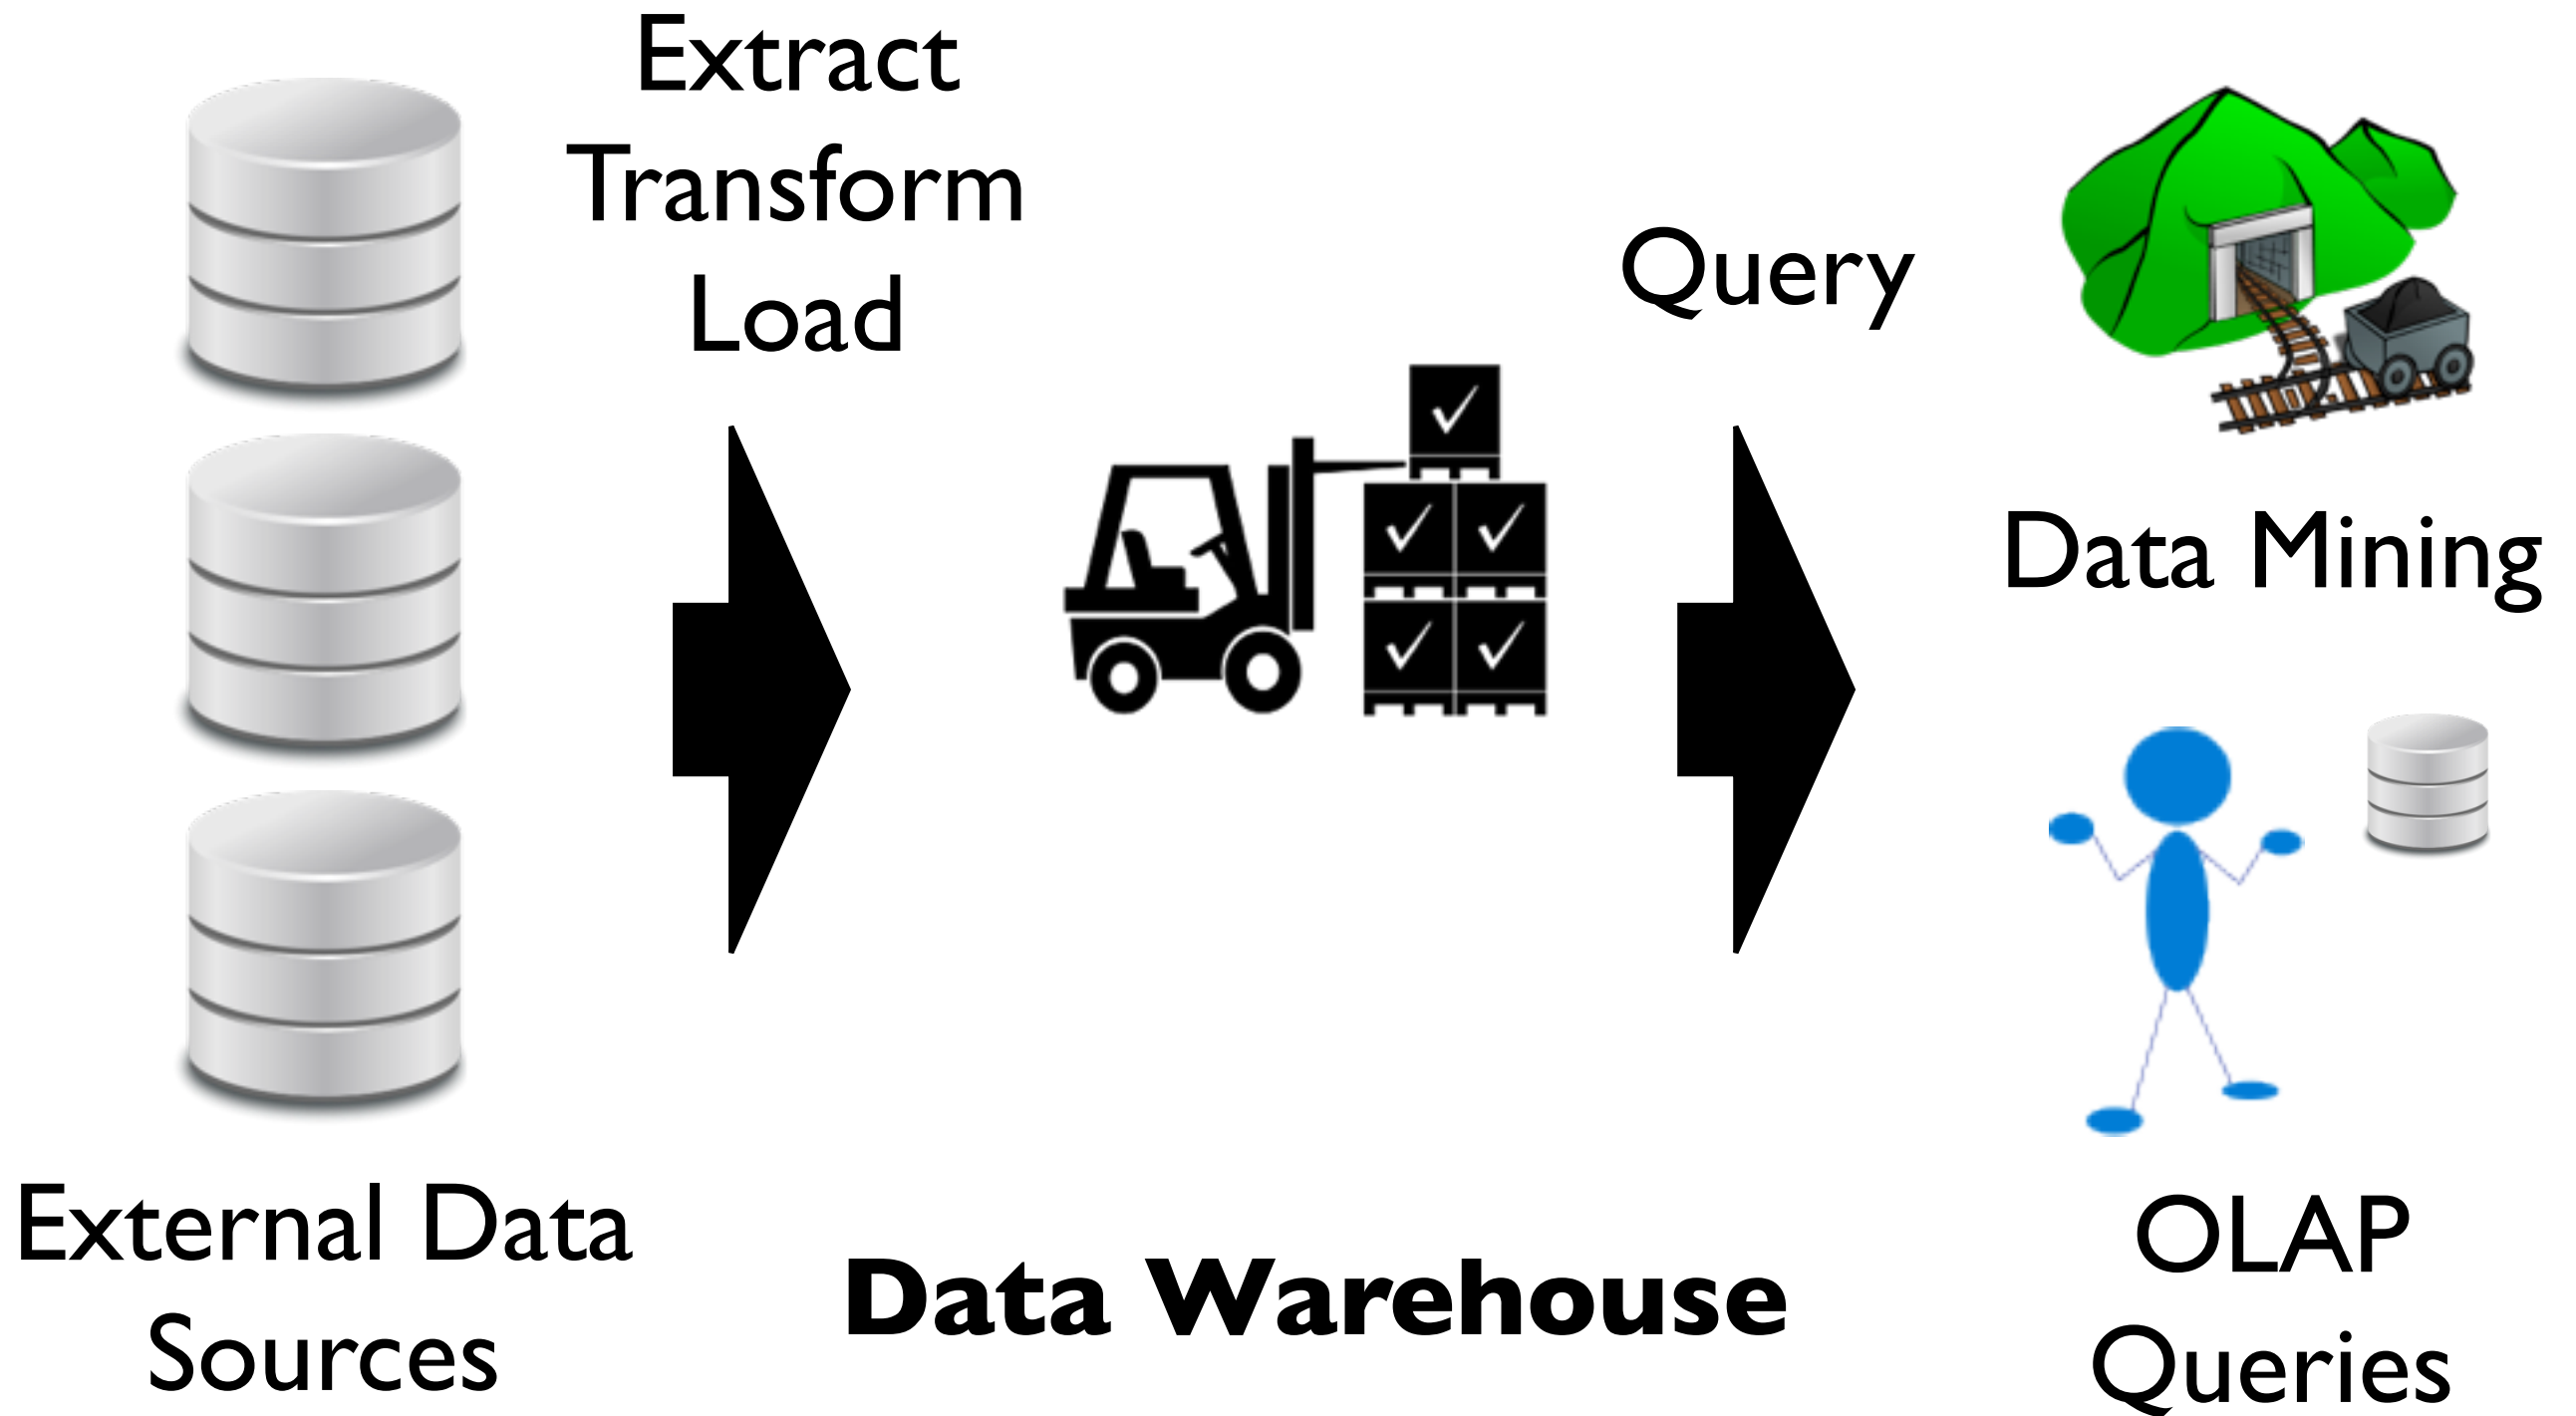
## R&G Chapter 25

(slides adapted from content by J.Gehrke, J.Shanmugasundaram, and/or C.Koch)

# Data Warehousing

- Big companies gather lots of data.
- This data can be exploited to...
  - ...identify interesting patterns.
  - ...correlate different data sources.
  - ...support hypotheses/what if questions.

# Data Warehousing



# A Data Warehouse


- Incorporates data spanning long time periods.
  - Merges data from many sources.
- Varies widely in scale: GB to TB to PB
- OLAP-centric workload:
  - Interactive query response times desirable.
  - Ad-hoc updates rare.

# Warehousing Challenges

- **Semantic Integration:** Data coming from different sources has different schemas, identifiers for the same entities.
- **Heterogeneous Sources:** Data arrives/is stored in in a variety of different formats.
- **Load/Refresh:** Must load/organize data, and keep it up to date with the outside world.
- **Metadata Management:** Increasingly important to track provenance of source data.

# Multidimensional Data Model (MOLAP)

- A Data Warehouse stores:
  - A collection of numeric measures...
  - ... that depend on a set of dimensions.
- For example:
  - Sales by Product, Location, Time



PID	TimeID	LocID	Sales
11	1	1	25
11	2	1	8
11	3	1	15
12	1	1	30
12	2	1	20
12	3	1	50
13	1	1	8
13	2	1	10
13	3	1	10
11	1	2	35
...			

# Relational Data Model (ROLAP)

- Central fact table relates multiple dimensions together.
- Extremely large (nearly all data here)
- Each dimension can have additional attributes stored in a dimension table.
- E.g. Products(pid, pname, category, price)

# Relational Data Model (ROLAP)

Times

timeid	date	week	month	quarter	year	holiday?
--------	------	------	-------	---------	------	----------

Sales  
(fact table)

pid	timeid	locid	sales
-----	--------	-------	-------

Products

pid	pname	category	price
-----	-------	----------	-------

Locations

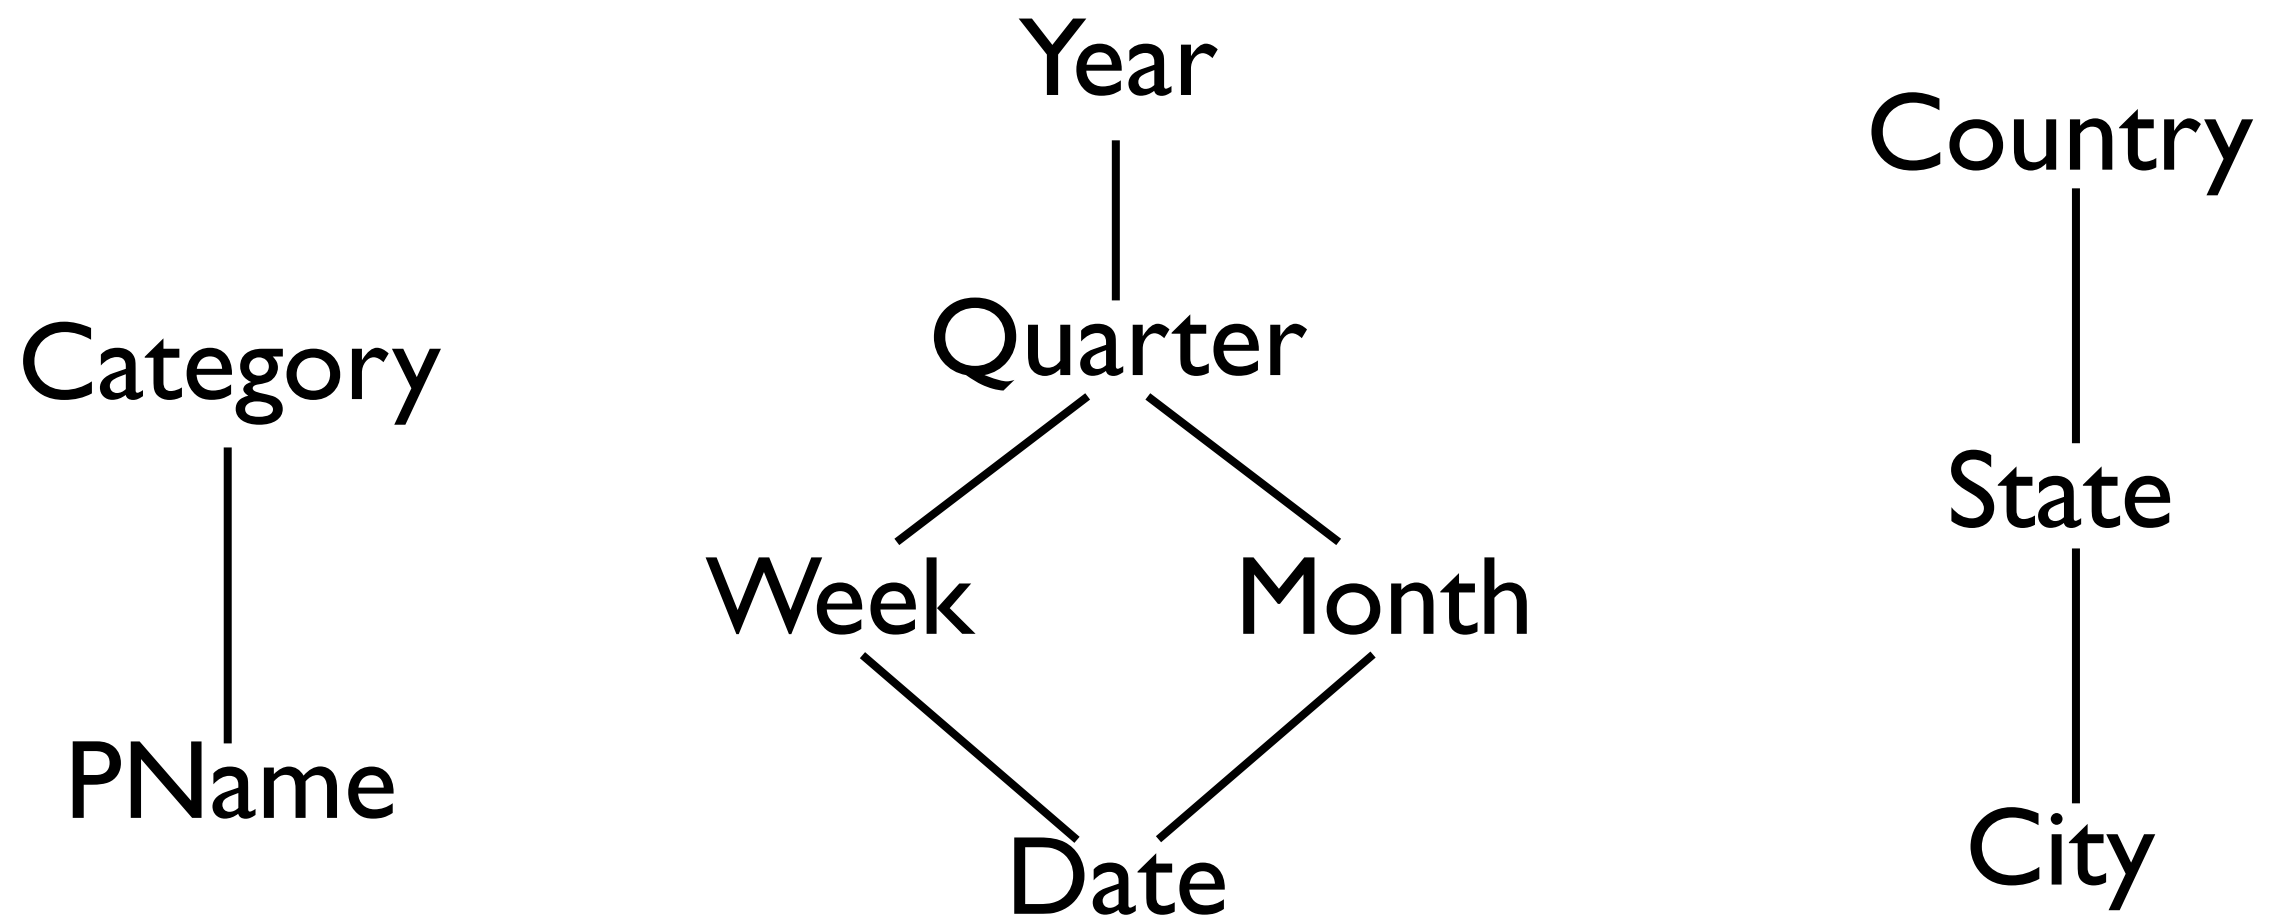
locid	city	state	country
-------	------	-------	---------

- Dimension Tables stored Un-Normalized
  - Updates/Inserts; Unlikely to cause issues
- Common schema in OLAP: The Star Schema



# Relational Data Model (ROLAP)

Dimension tables establish hierarchies of data



# OLAP Queries

- Influenced by SQL and Spreadsheets
- Most Common Operation: Aggregate over one or more dimensions.
  - Find total sales
  - Find total sales for each city, for each state.
  - Find top five products by sales

# OLAP Queries

- **Roll-Up:** Move to a coarser grained view of the data (more values aggregated)
  - “Moving up” in one dimension’s hierarchy
- **Drill-Down:** Move to a finer grained view of the data (more aggregate values shown)
  - “Moving down” in one dimension’s hierarchy

# OLAP Queries

- **Pivoting:** Aggregate over selected dimensions to get a cross-tabulation.
- **Slice and Dice:** Equality or Range Selections on one or more dimensions.

	NY	WA	<b>Total</b>
2012	63	81	144
2011	38	107	145
2010	75	35	110
<b>Total</b>	176	223	339

# OLAP Queries vs SQL Queries

# OLAP Queries vs SQL Queries

- Cross-tabulations can be obtained from SQL Queries. How?

# OLAP Queries vs SQL Queries

- Cross-tabulations can be obtained from SQL Queries. How?
- Generalizing, how many different group-by queries can we create over a fact table with  $k$  dimensions?

# OLAP Queries vs SQL Queries

- Cross-tabulations can be obtained from SQL Queries. How?
- Generalizing, how many different group-by queries can we create over a fact table with  $k$  dimensions?
- The CUBE operator creates all of these.



# OLAP Queries vs SQL Queries

- Cross-tabulations can be obtained from SQL Queries. How?
- Generalizing, how many different group-by queries can we create over a fact table with  $k$  dimensions?
- The CUBE operator creates all of these.
  - Prebuilds cube for faster OLAP queries.

# The CUBE Operator

- CUBE pid, locid, timeid BY SUM Sales
  - Contains results of all 8 possible group-by queries.
  - How would this be represented in SQL?
- This is an expensive operation, how might it be implemented efficiently?

# Bitmap Indexes

- Slice & Dice requires Equality Tests.
- Testing for equality is expensive (pipeline stalls)
  - Can we aggregate using purely arithmetic operations?
- Store enumerations as bitmaps
  - One bit per possible value: Each row sets one bit.

# Bitmap Indexes

Data				Bitmaps	
CustID	Name	Gender	Rating	Gender	Rating
112	Alice	F	3	01	00100
115	Bob	M	5	10	00001
119	Carol	F	5	01	00001
113	Dave	M	4	10	00010

```
SELECT SUM(Sales)
WHERE Gender = 'M'
AND Rating = 5
```

```
Tot += Sales * (
    ((G & 0x2) >> 1)
    & ((R & 0x1) >> 0)
)
```

# Join Index

- Consider the Join of Sales, Products, Times, Locs with the predicate Country = 'USA'
- A join index helps speed this process up.
- $\langle s, p, t, l \rangle$  in the index if SaleID  $s$  (respectively  $p, t, l$ ) matches the predicate.
- Problem: Too many, too big join indexes.
  - Solution: Index by value and SaleID (fact table rowid)
  - To match multiple predicates, intersect columns.

# Sequences

- Trend analysis is difficult (in SQL 92)
  - Find the % change in monthly sales
  - Find the daily top-5 product by sales in the last week
  - Find the trailing n-day moving average of sales.
- The first two examples are hard to express.
  - The third can not be expressed if n is a parameter
- SQL 99's WINDOW clause analyzes sequences.

# WINDOW

- Define a sequence (by sorting a relation)
- Generate all subsequences of size  $N$ .
- Compute aggregates for each subsequence.

# WINDOW

```
SELECT L.state, T.month,  
       AVG(S.sales) OVER W as movavg  
FROM   Sales S, Times T, Locations L  
WHERE  S.timeid = T.timeid  
       AND S.locid = L.locid  
WINDOW W AS (  
    PARTITION BY L.state  
    ORDER BY T.month  
    RANGE BETWEEN INTERVAL '1' MONTH PRECEDING  
           AND INTERVAL '1' MONTH FOLLOWING  
)
```