Approximation Techniques

Supplemental Reading
(papers posted on Piazza)

- Project 2 Graded.
  - Results will be posted tonight.
- If anyone has something to say...

The Leader Board

1. <for rent>
2. <for rent>
3. <for rent>
4. <for rent>
5. <for rent>

Query Approximation

- Summarize the data with a Sketching Algorithm.
  - Bloom Filters (Set Containment)
  - Flajolet & Martin Count-Distinct Sketch
  - Count Sketch (Frequent Items/Top-K)
- Sample the data and estimate the error
  - … or better yet, keep generating samples!
    - **Online Aggregation** & Ripple Joins

Online Aggregation

- Give (some) results immediately.
  - (Imprecise) results obtained by sampling.
- More time, more samples, more accuracy.
  - Query takes longer to complete.
  - … but result estimate available sooner.

Online Aggregation

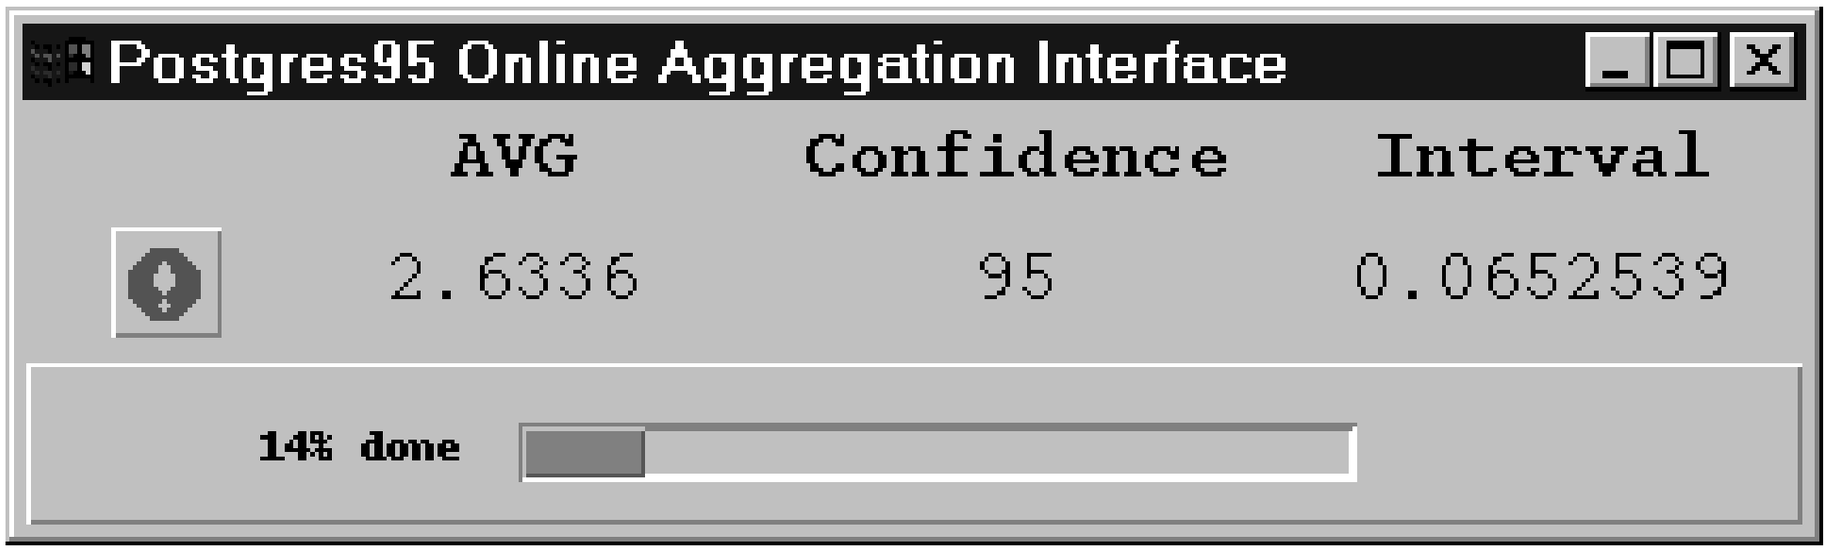

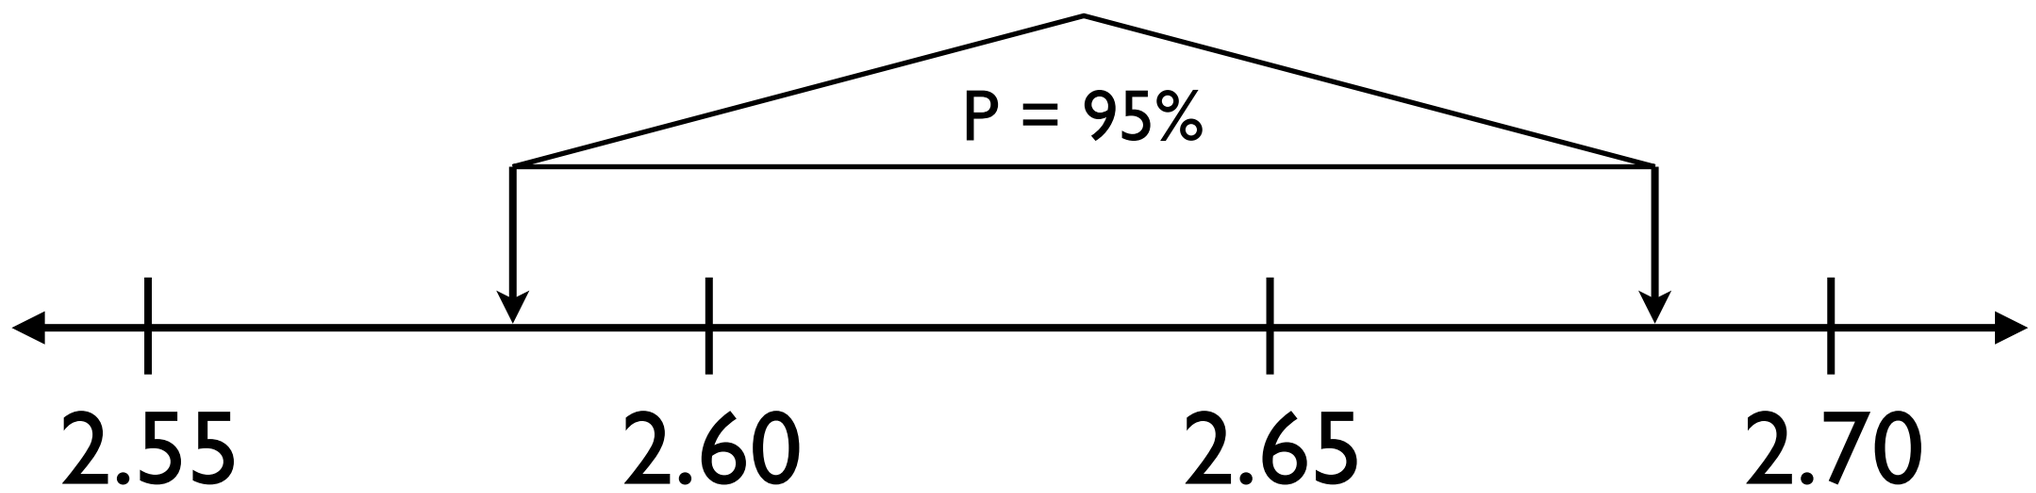

Online Aggregation

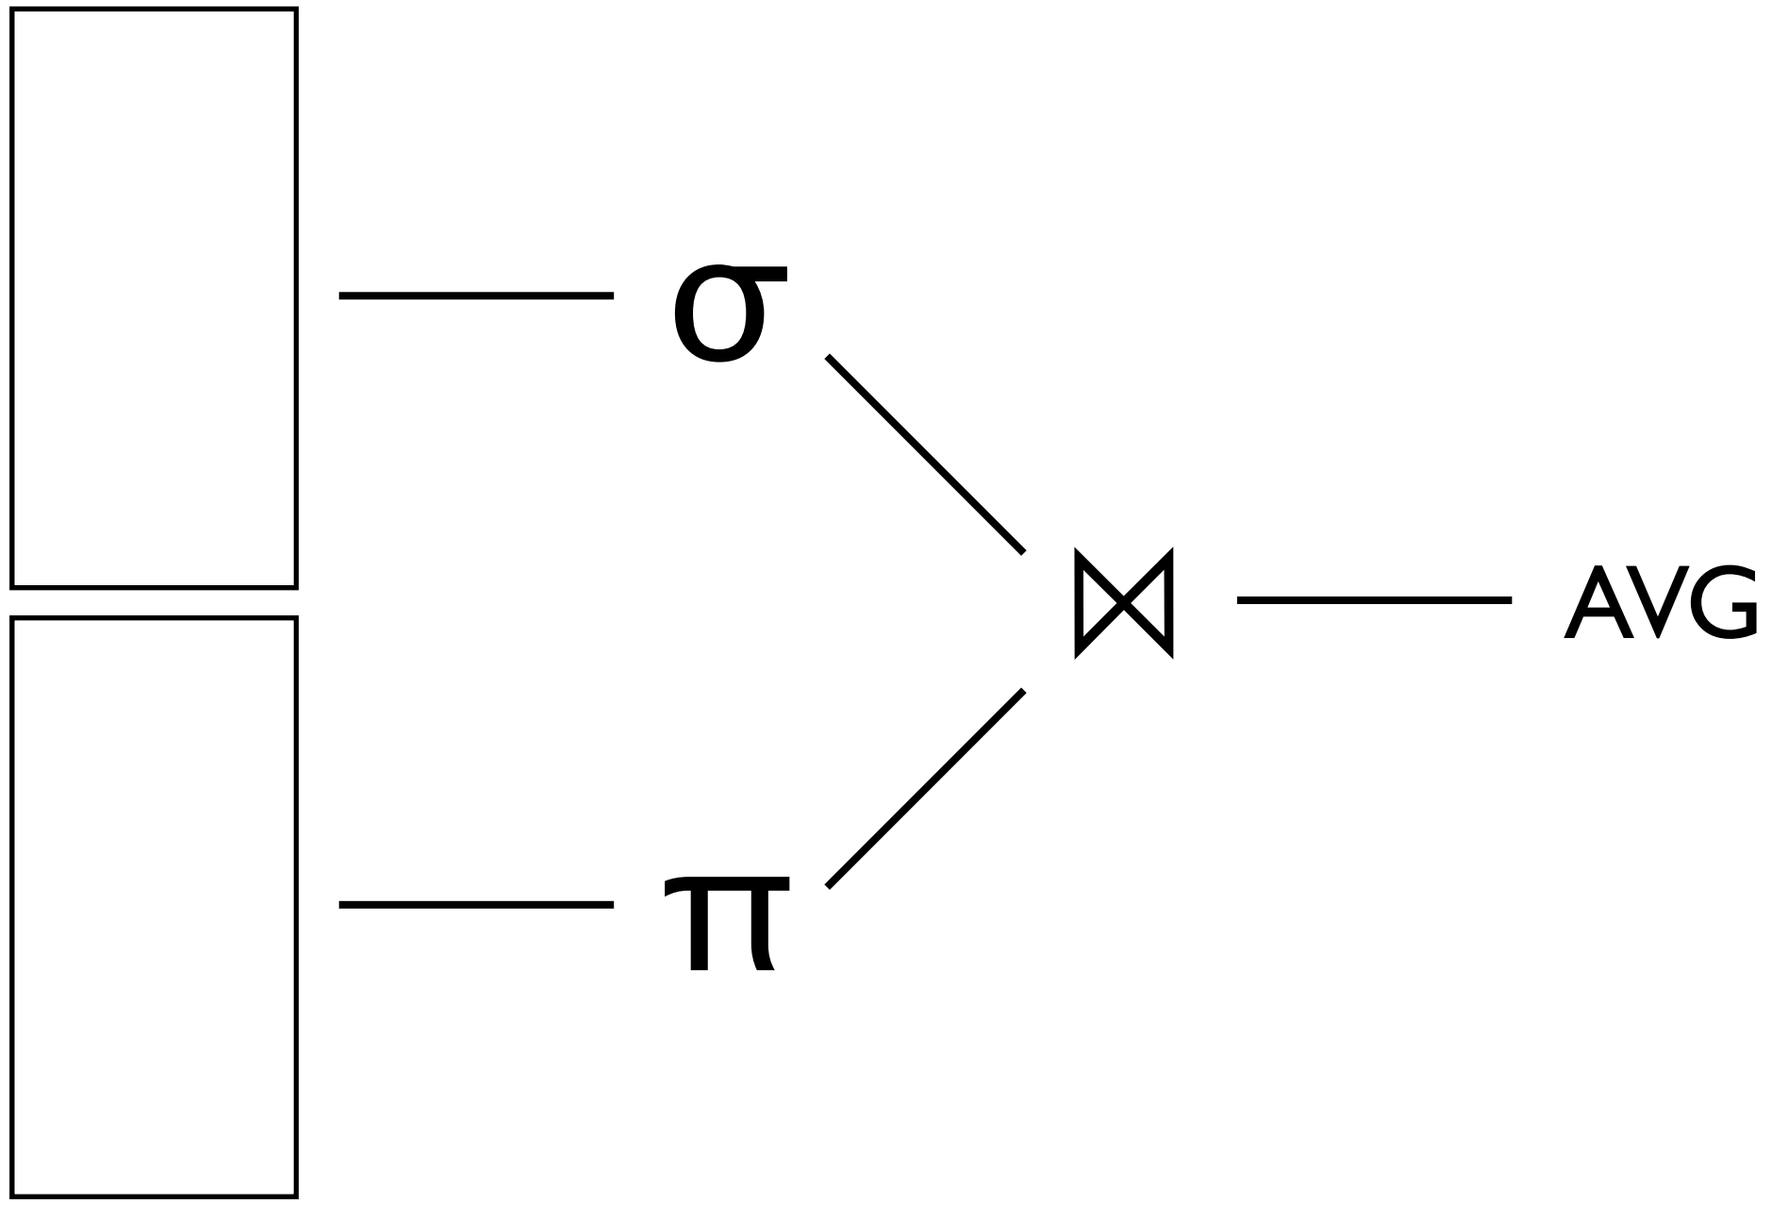

Online Aggregation

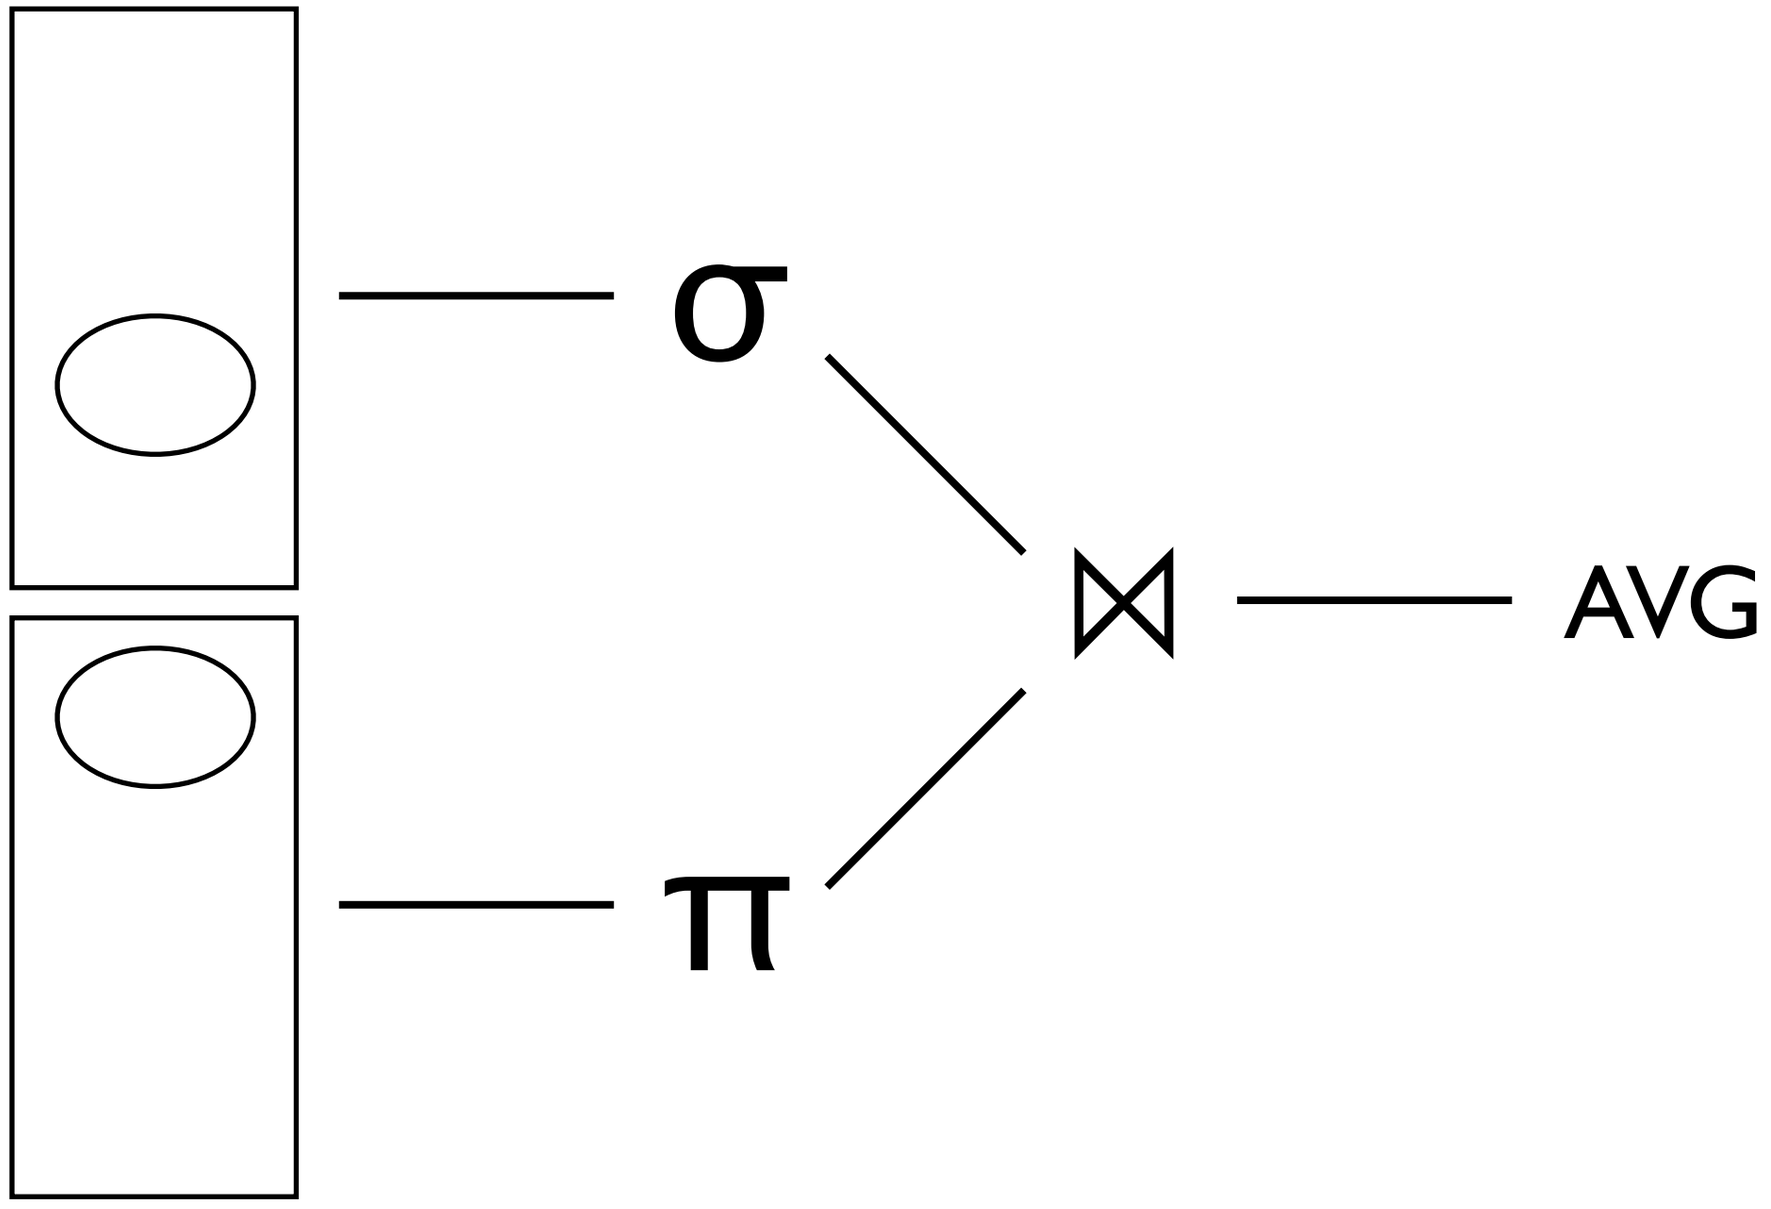

Online Aggregation

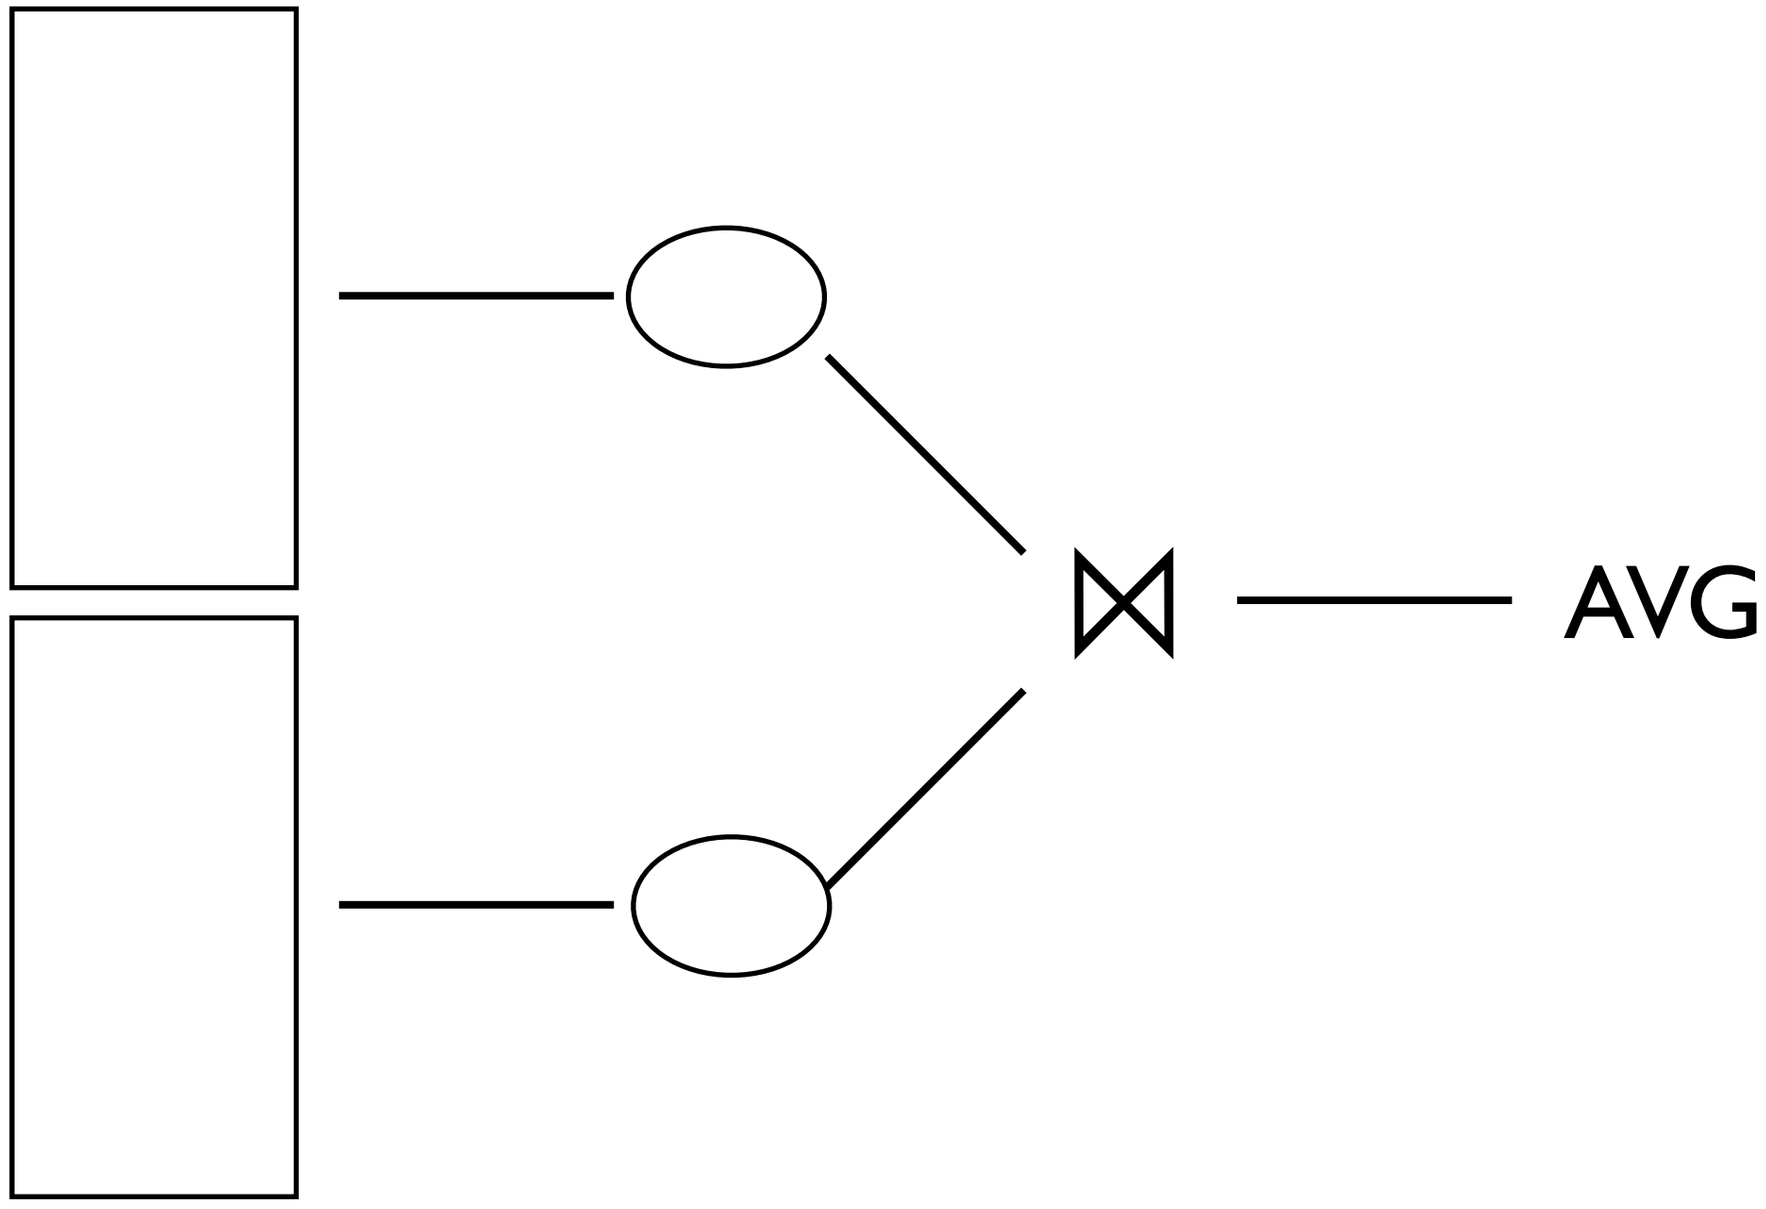

Online Aggregation

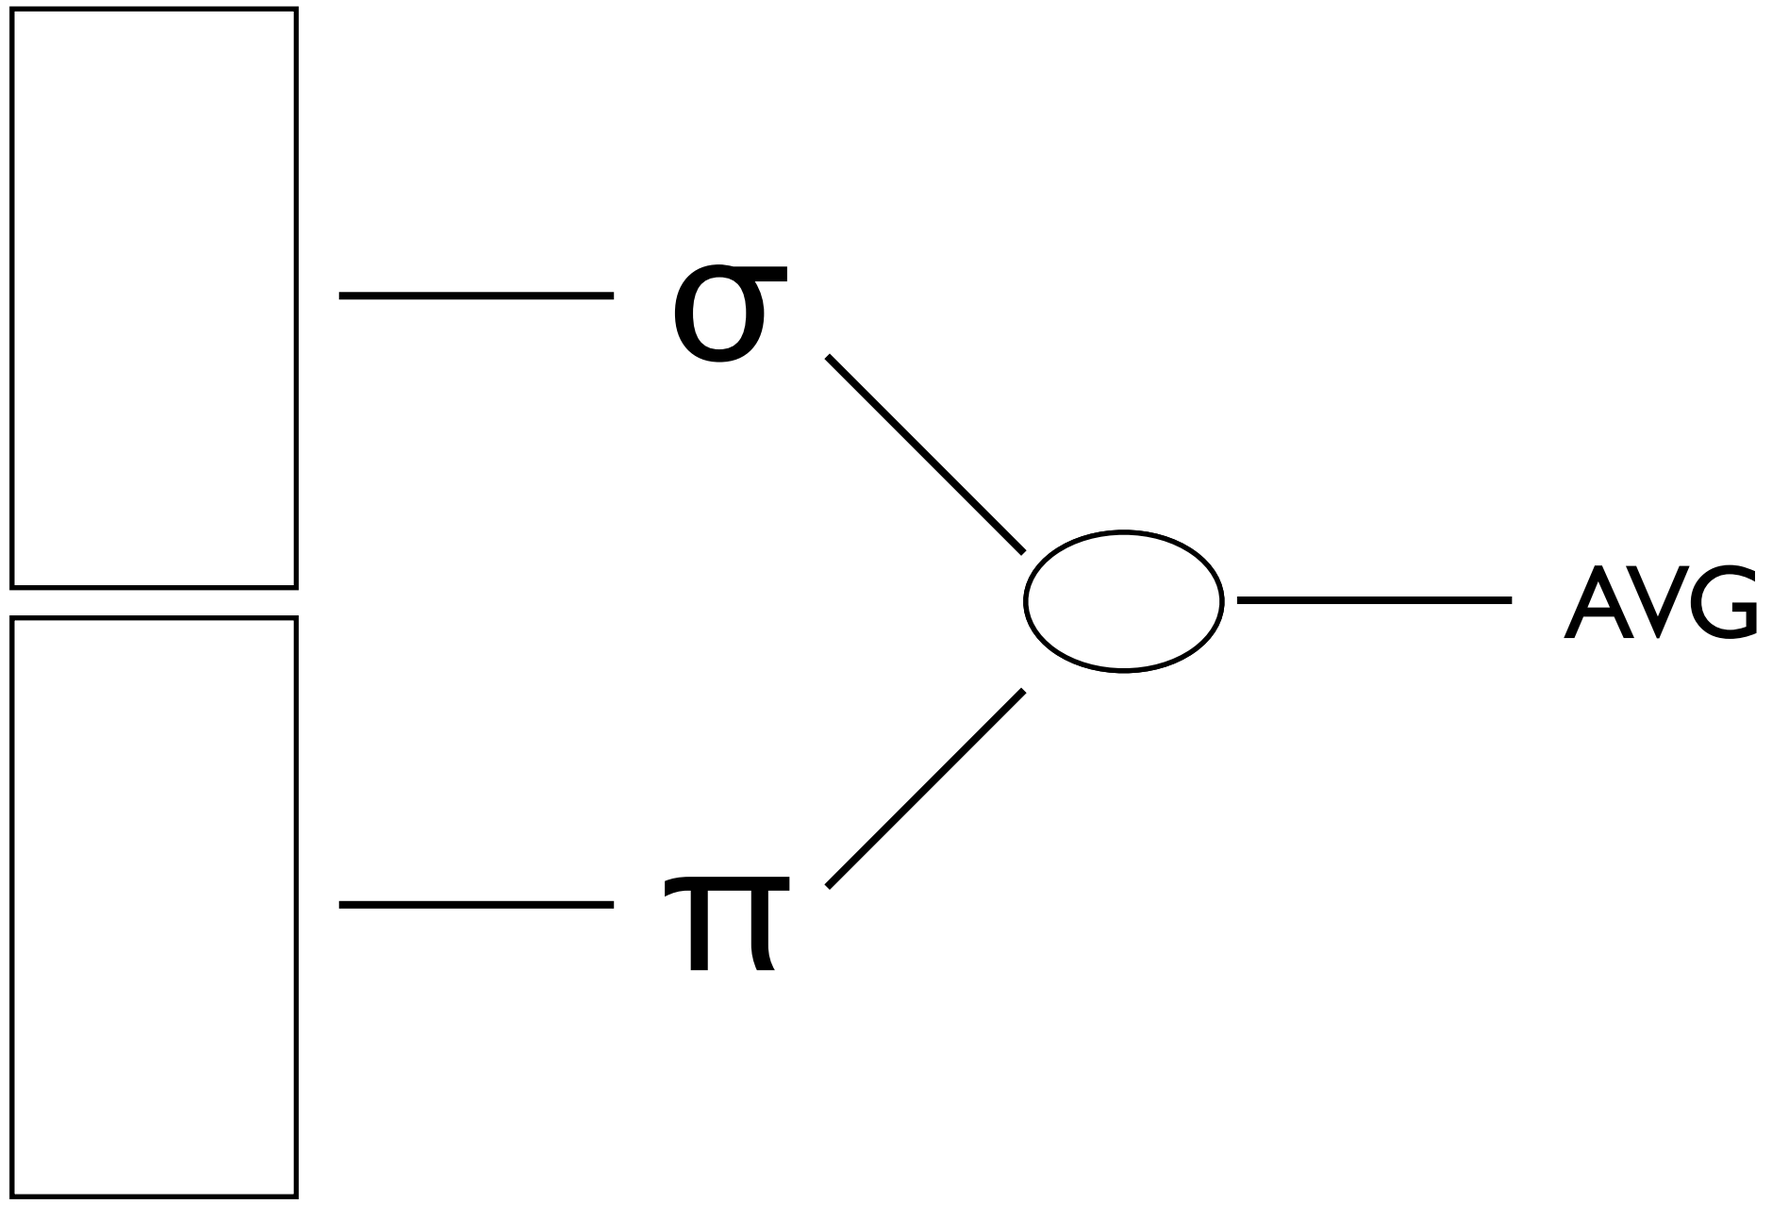

Online Aggregation

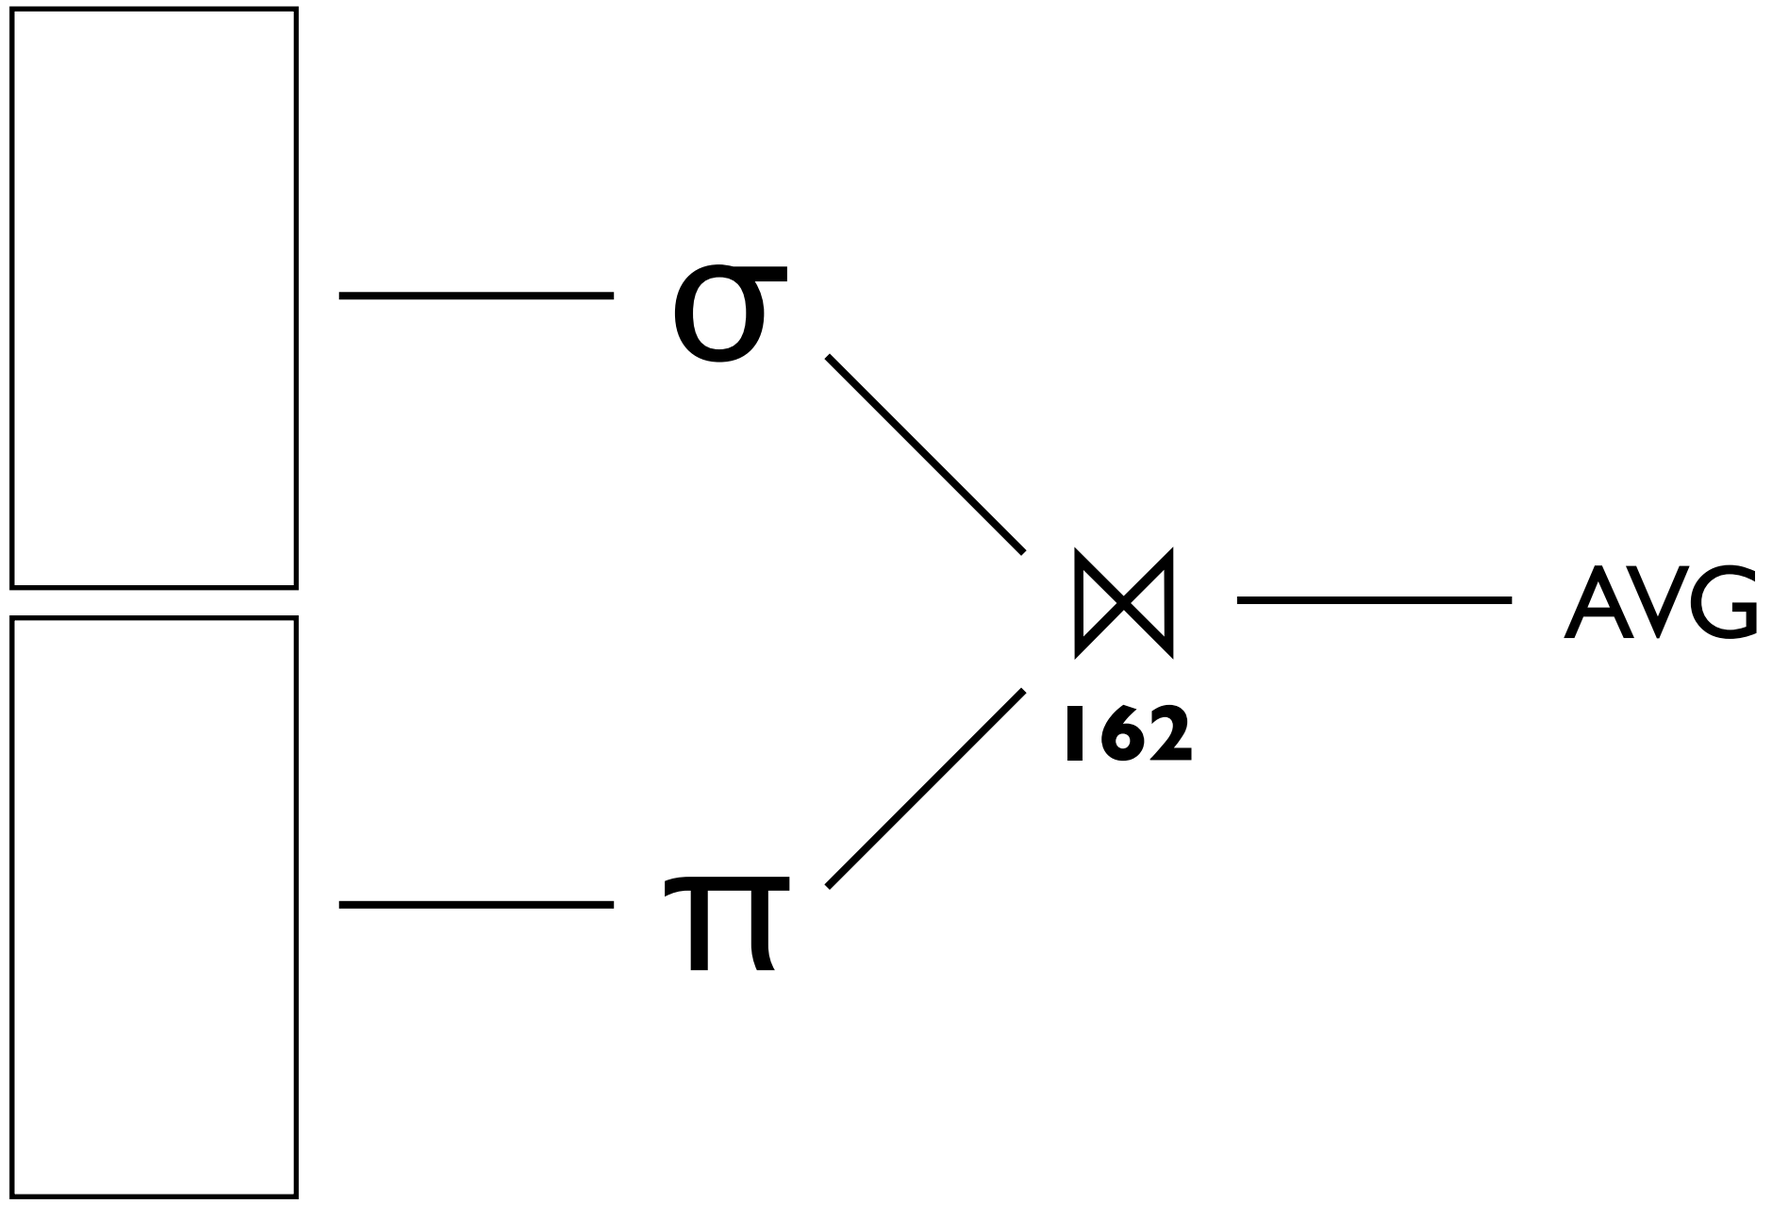

Online Aggregation

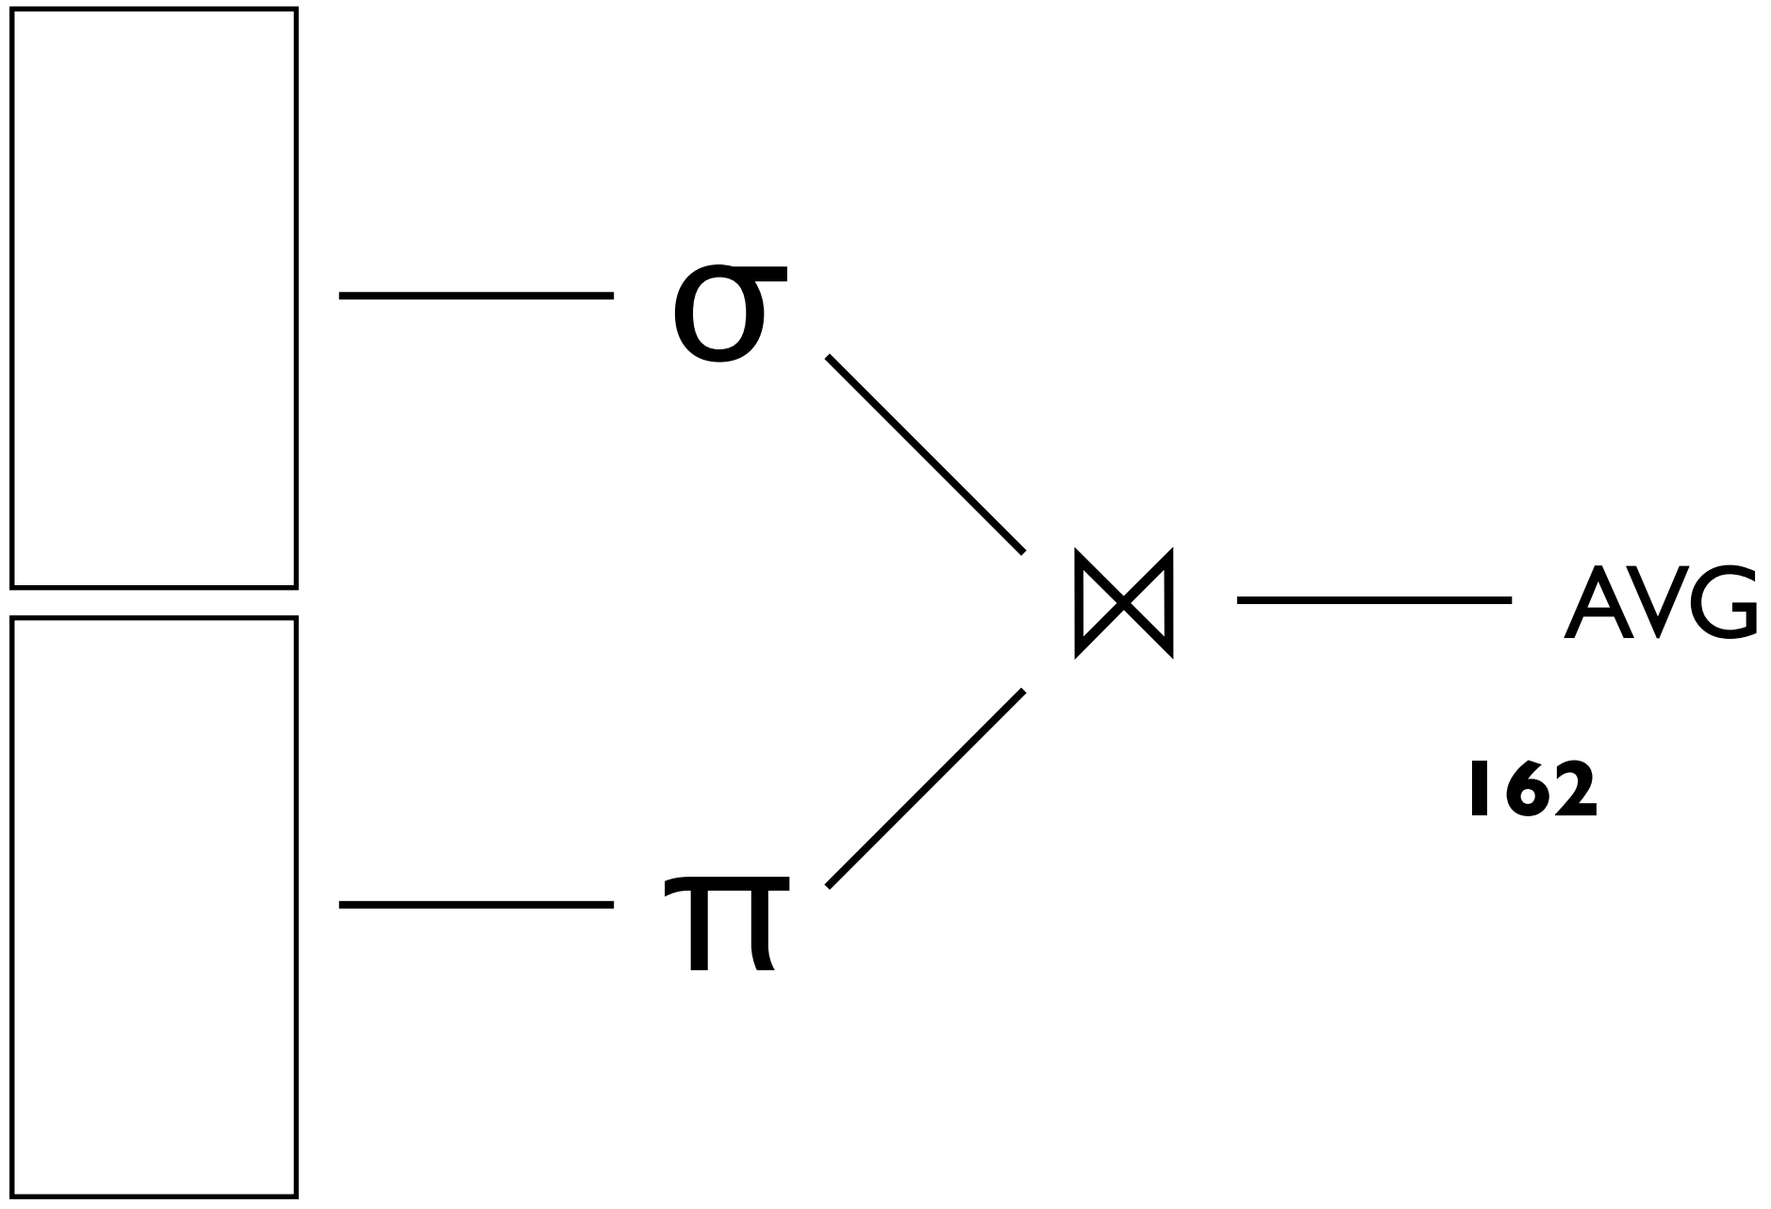

Online Aggregation

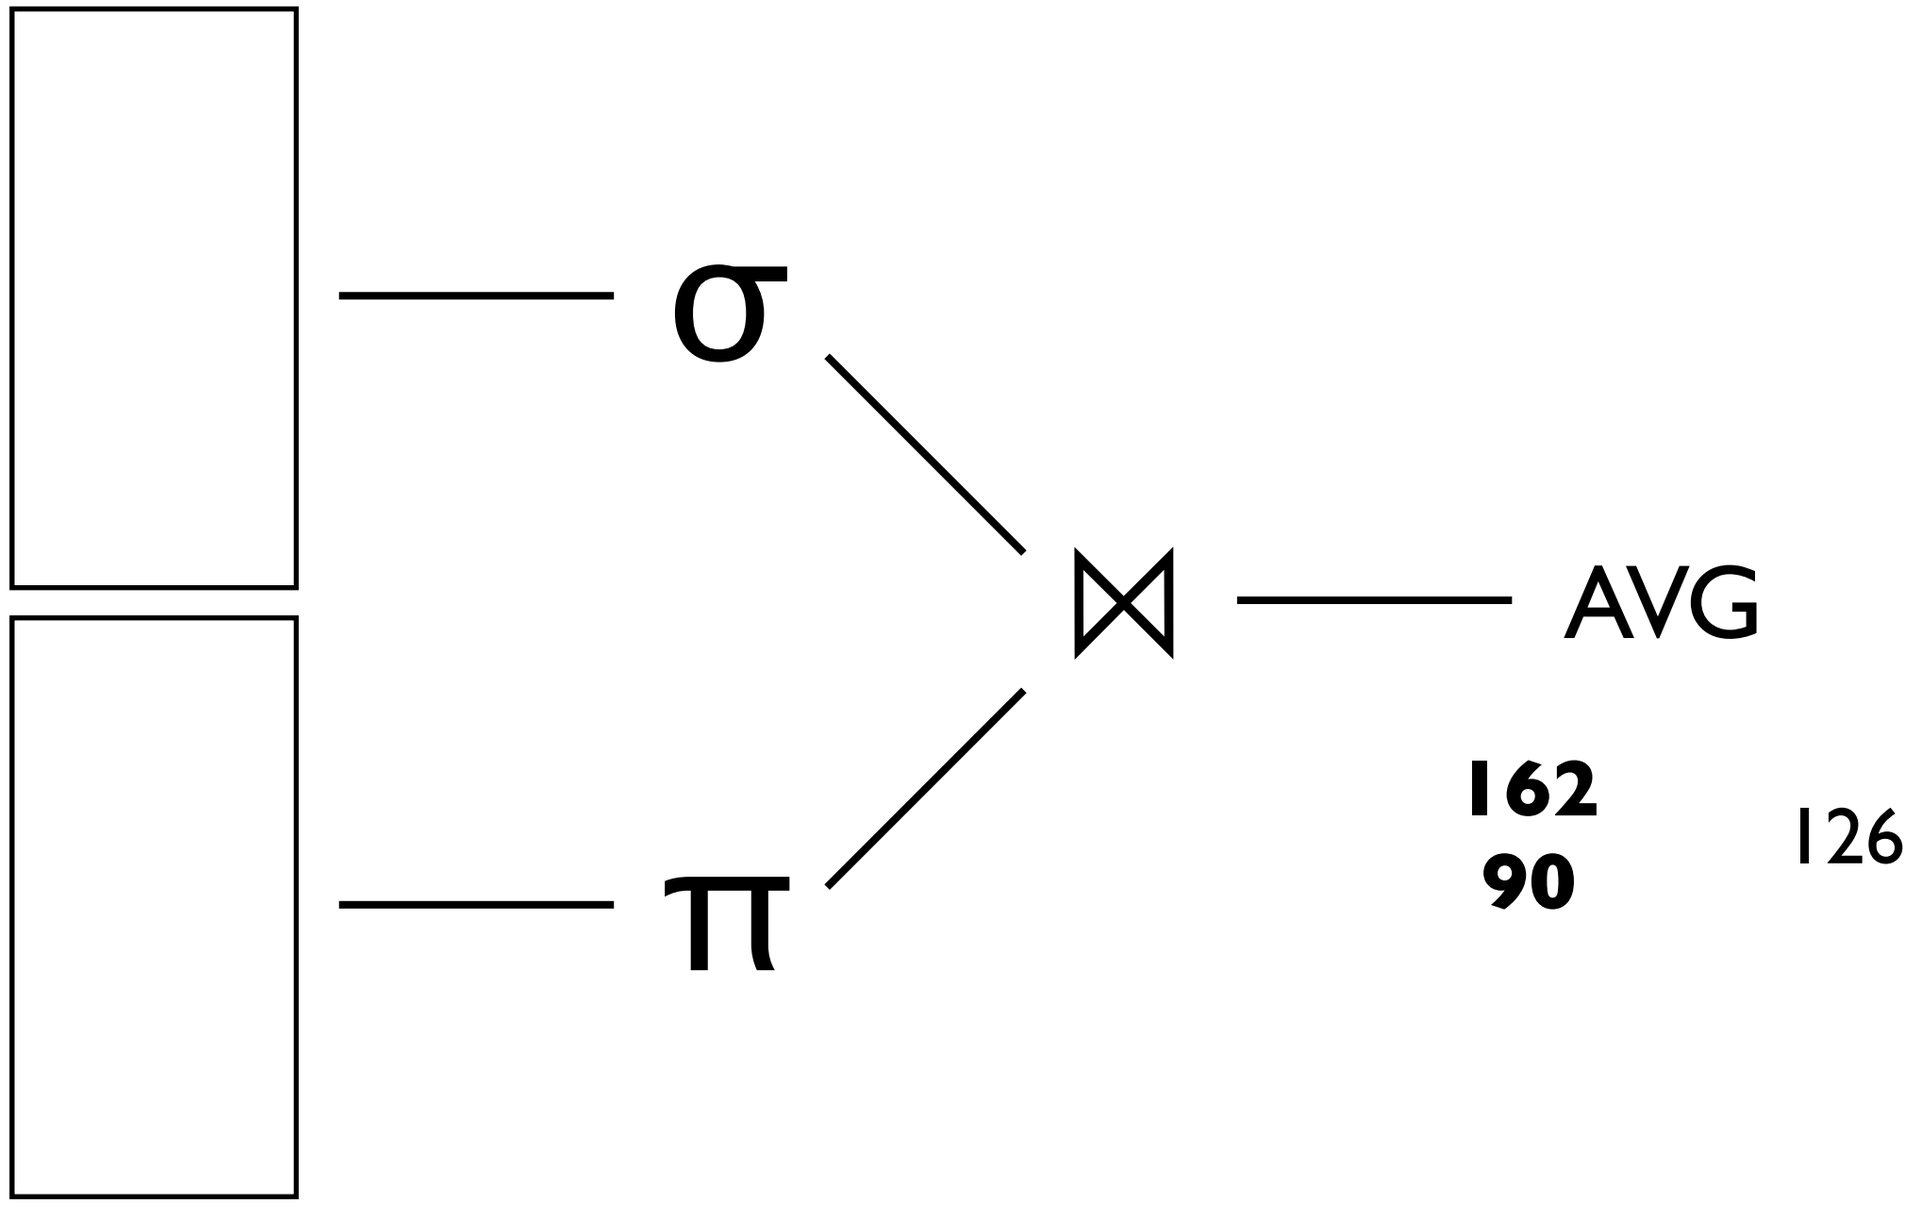

Online Aggregation

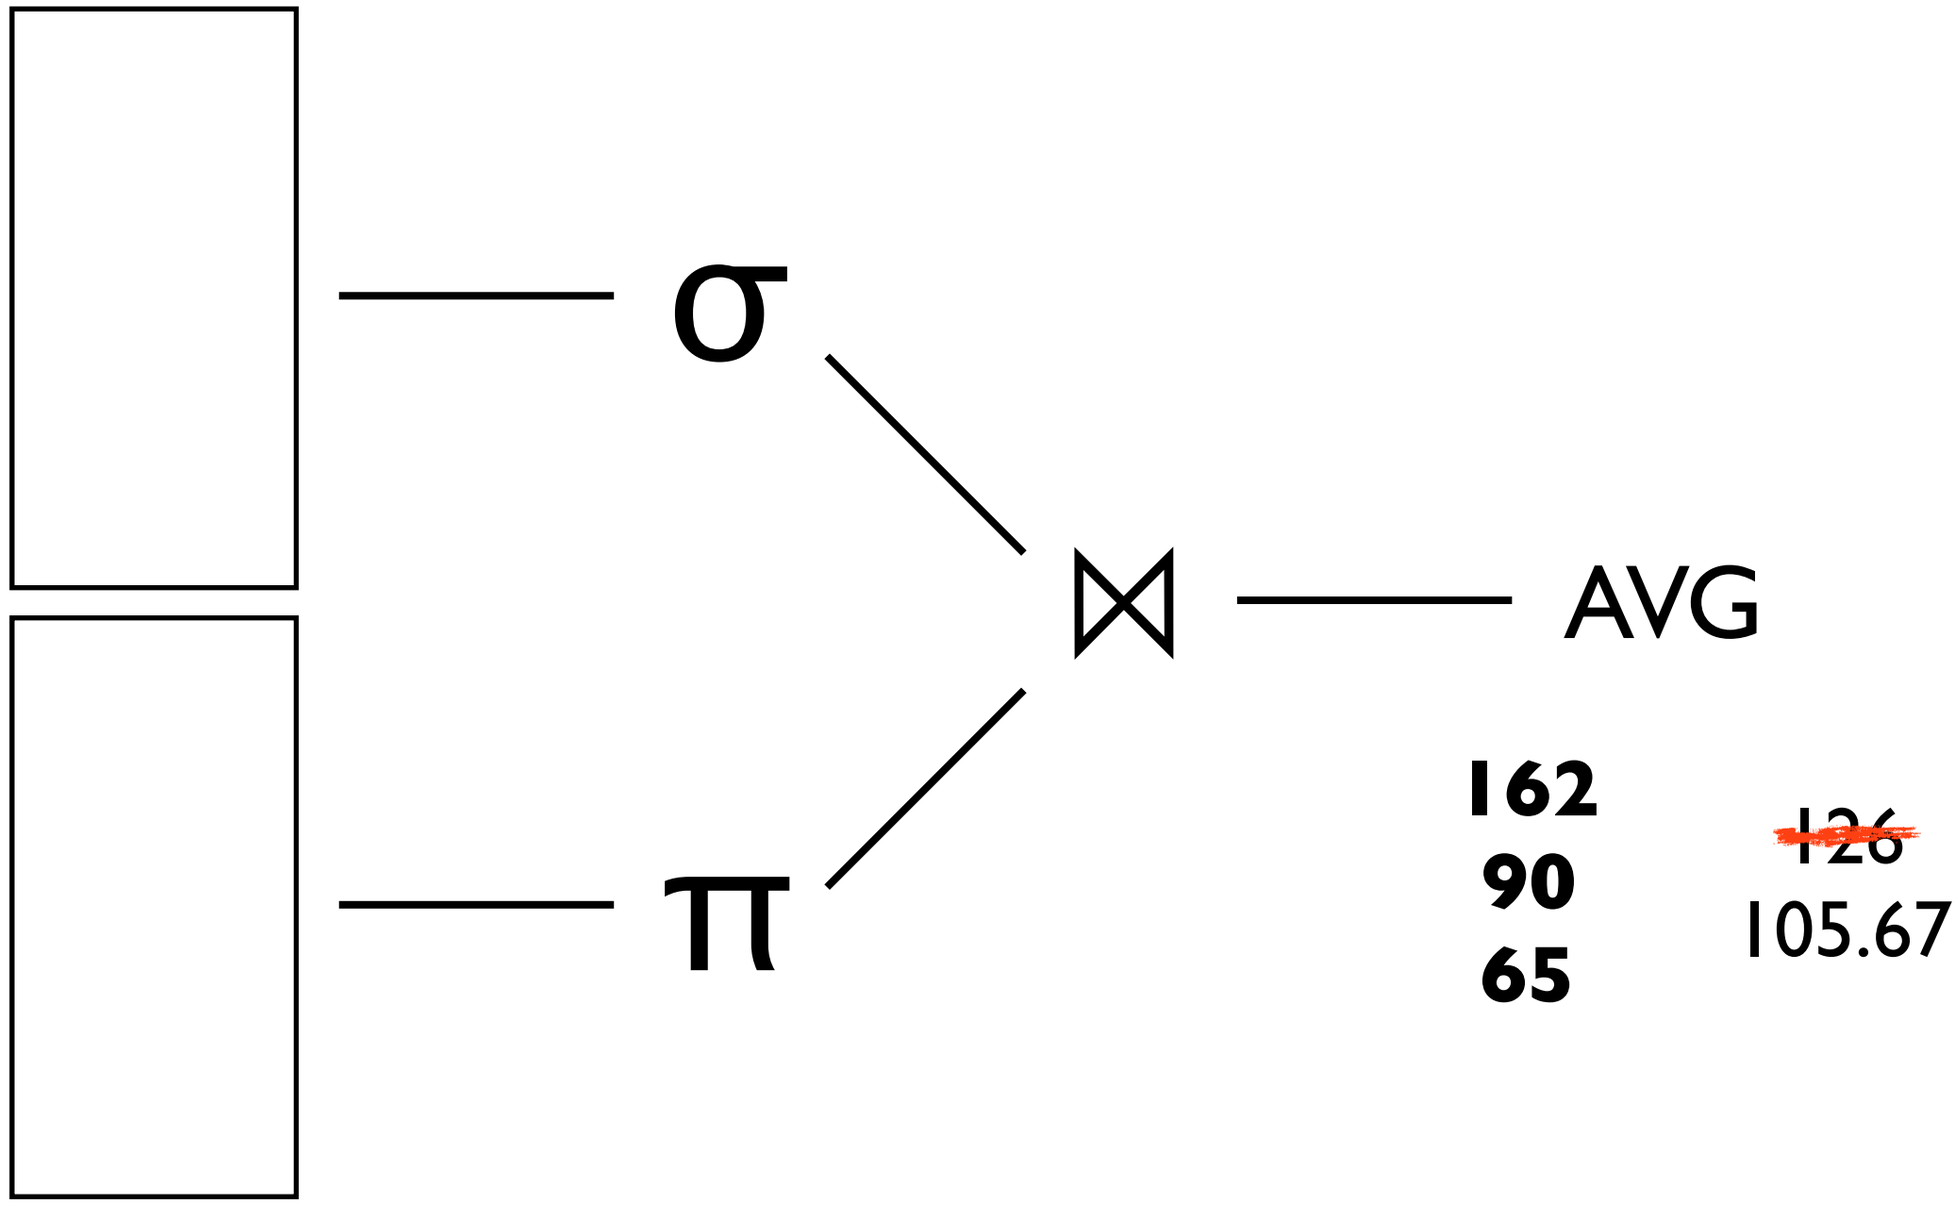

Online Aggregation

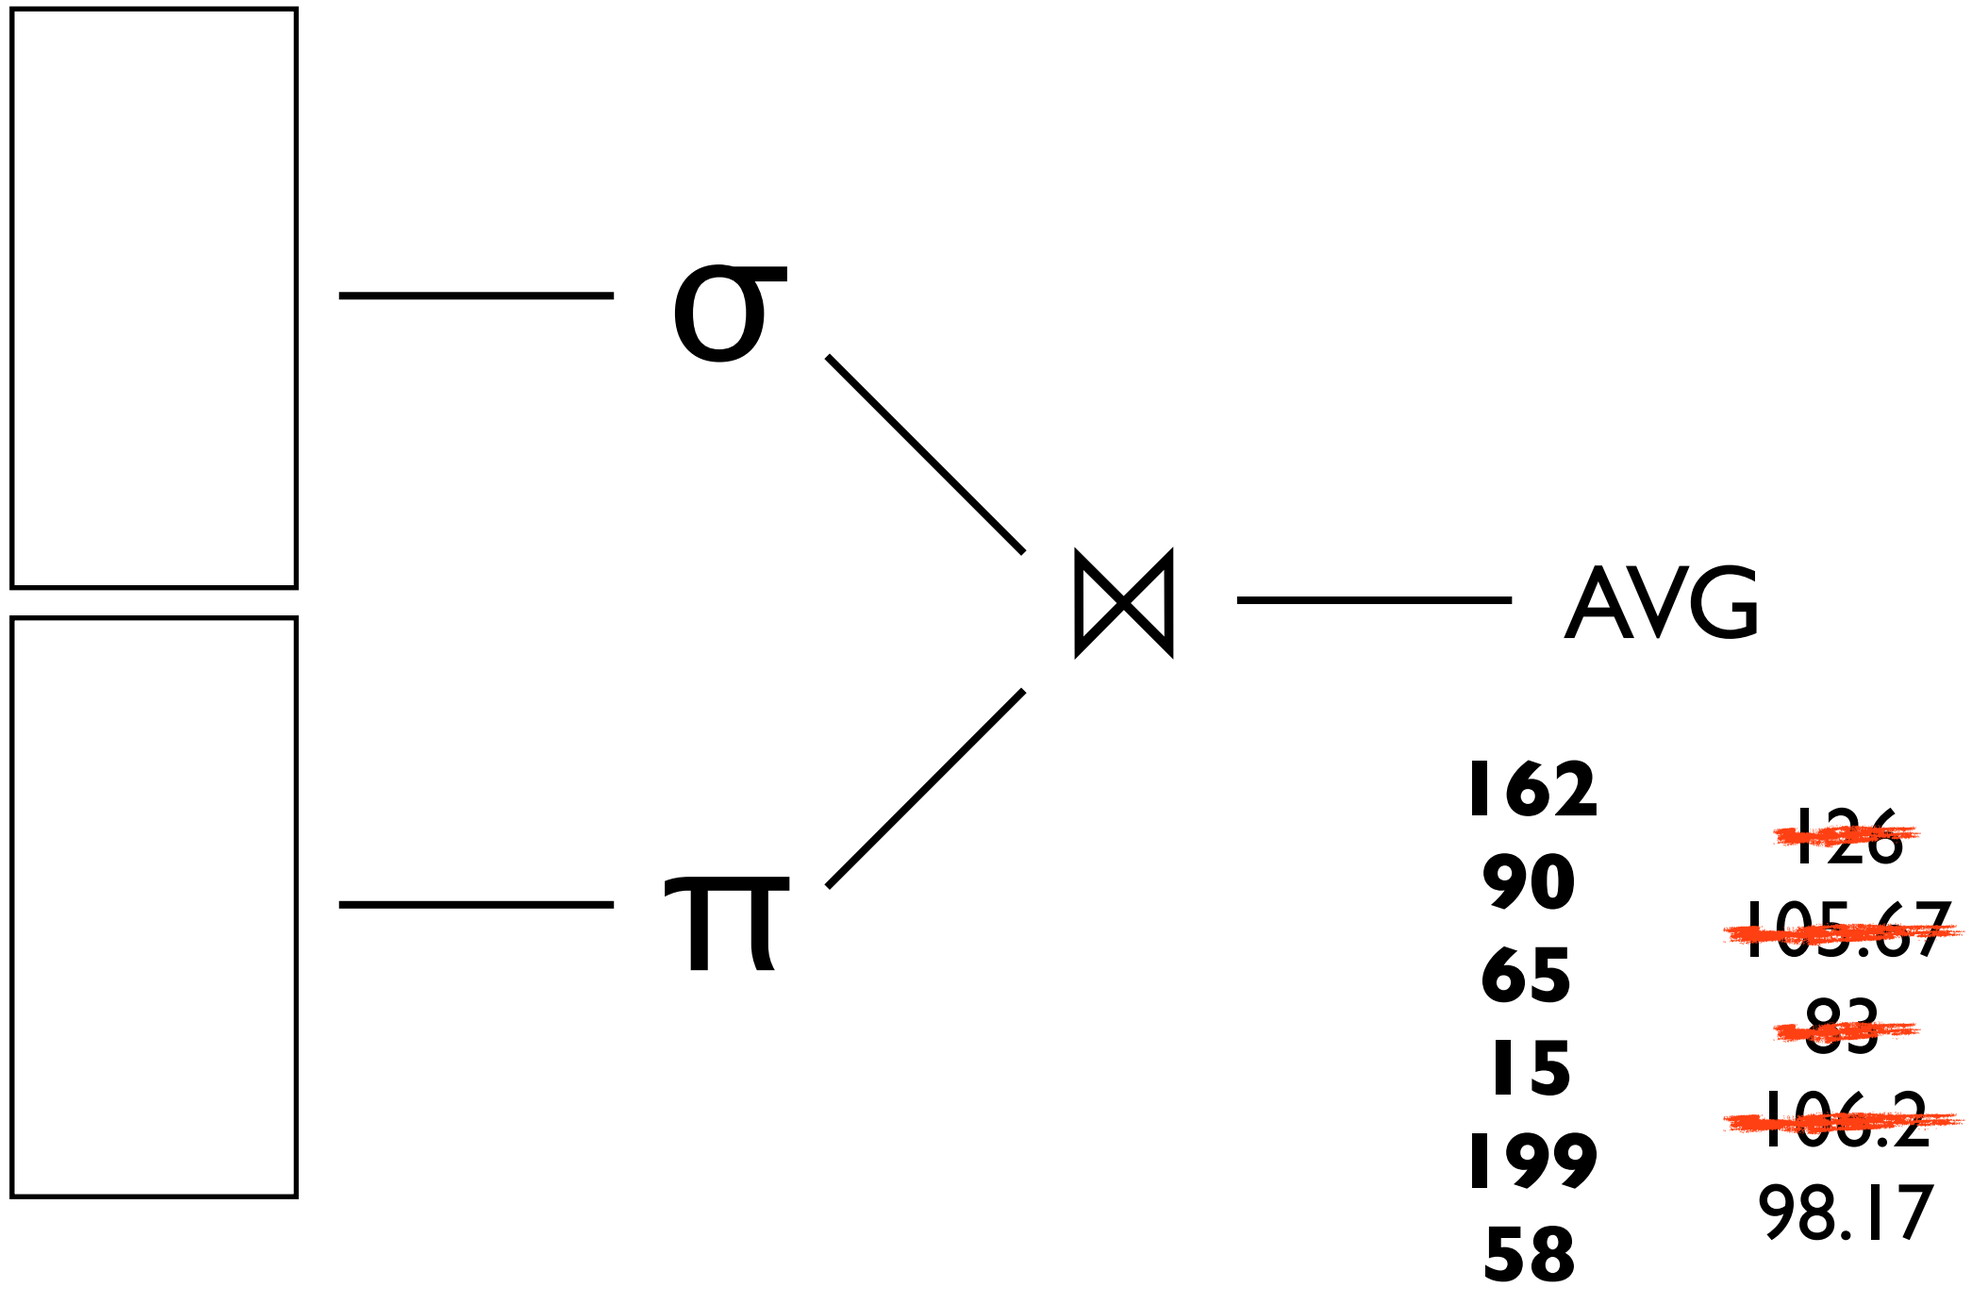

Online Challenges

- Sampling: Need Random Access to Data.
  - Heap (Unsorted) Files, Flash Drives
- Fairness: Sampling For “Rare” Group-By Columns.
  - Index Striding
- Blocking: Joins Must Be Streamed.
  - Ripple Join

Fairness

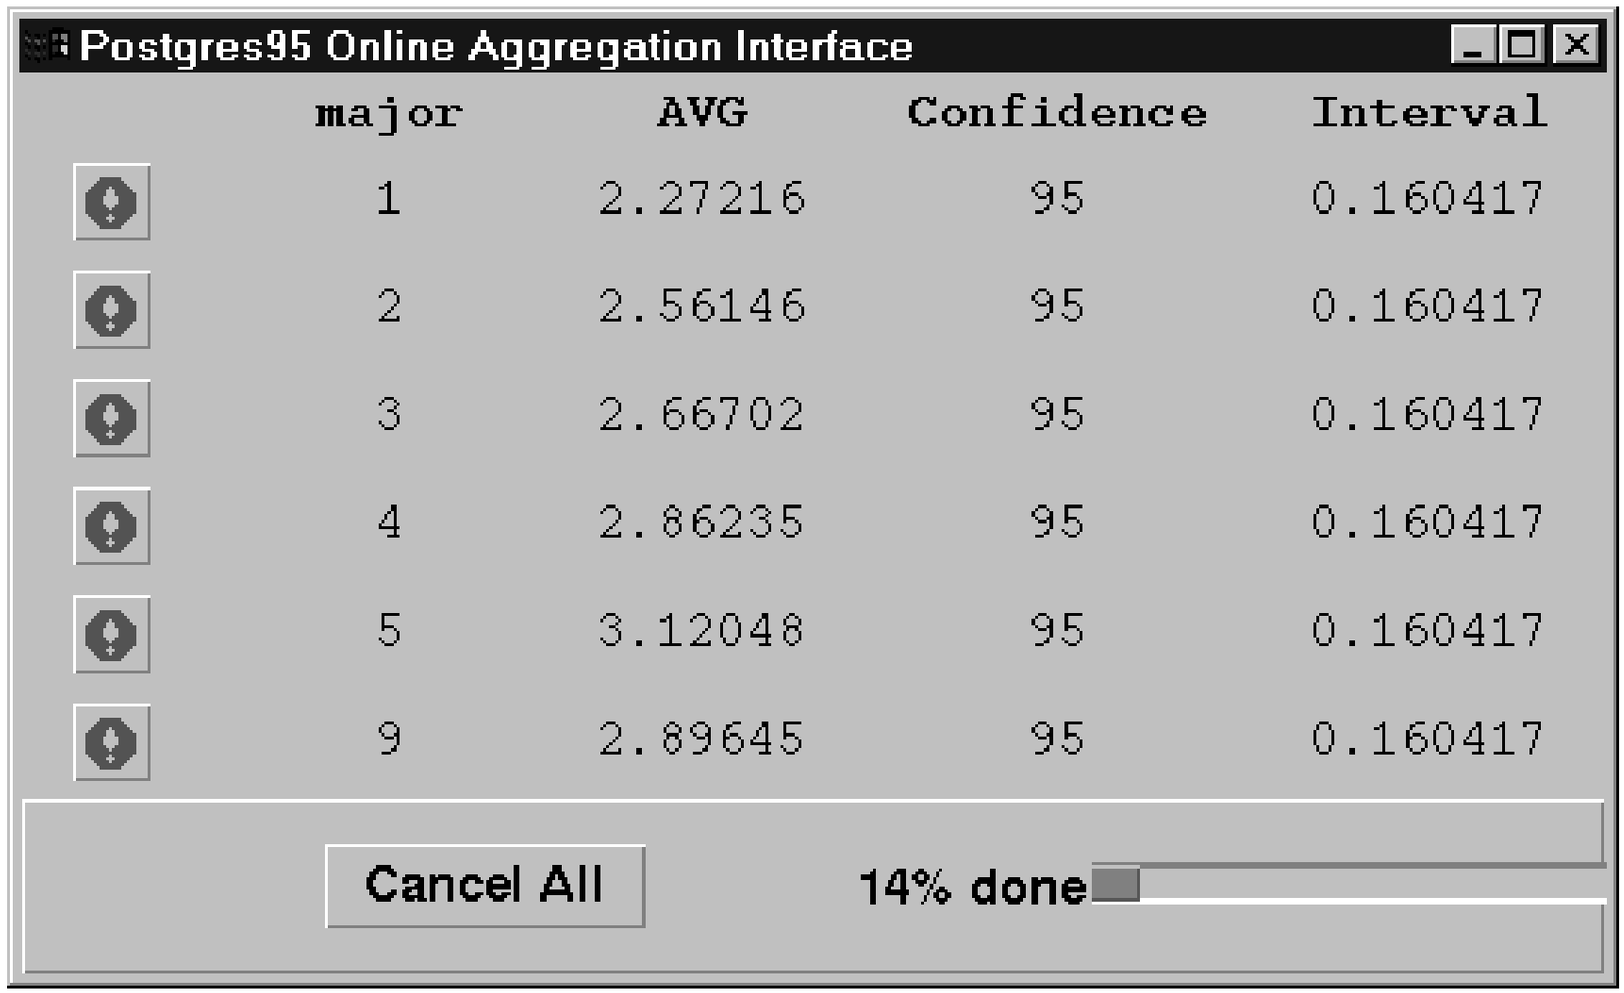

Group by aggregates produce many results.
Each sample contributes to one result.

Fairness

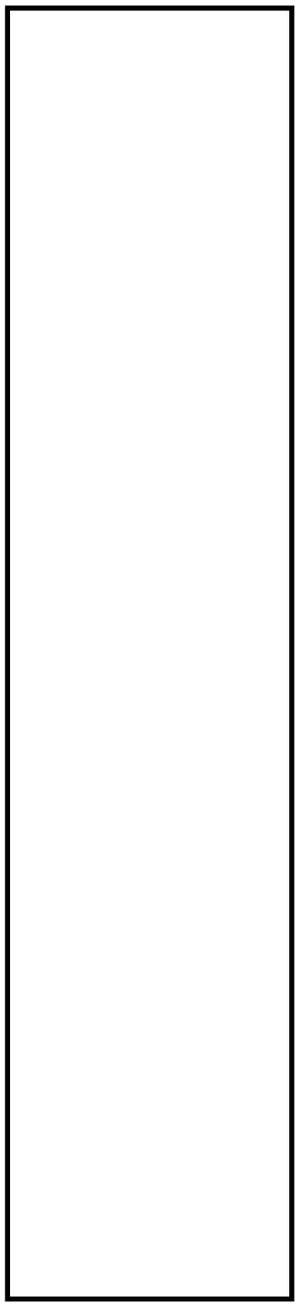

GB Key 1

GB Key 2

GB Key 3

GB Key 4

GB Key 5

Fairness

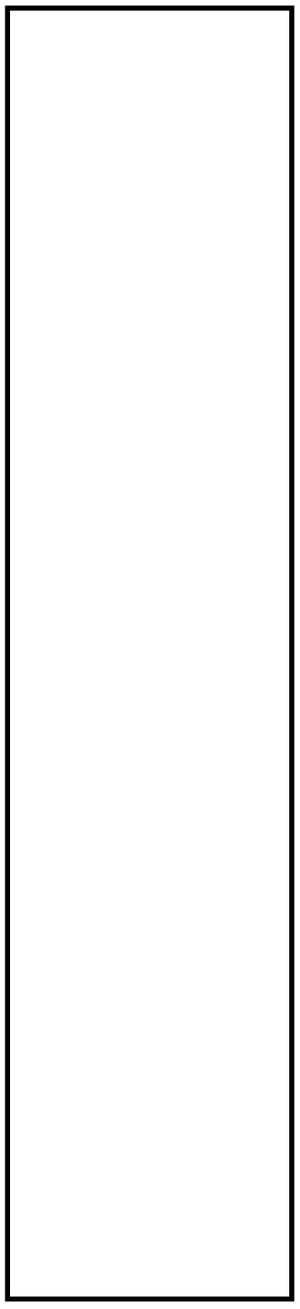

GB Key 1

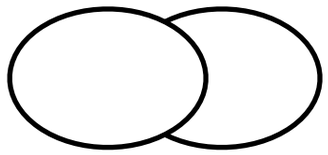

GB Key 2

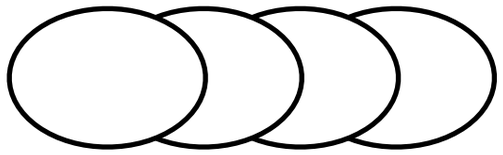

GB Key 3

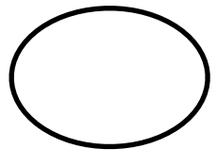

GB Key 4

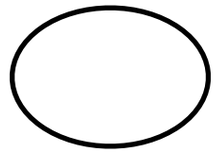

GB Key 5

Fairness

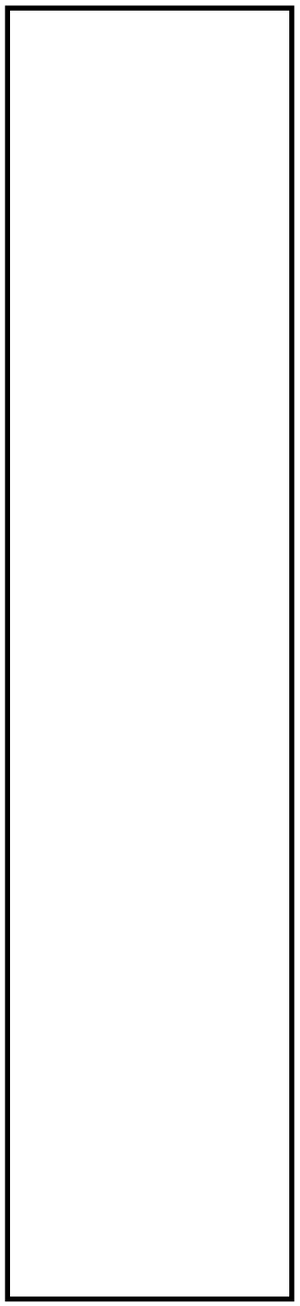

No/Few Samples
(Not Accurate)

Lots of Samples
(Very Accurate)

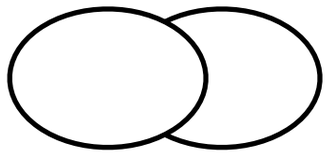

GB Key 1

GB Key 2

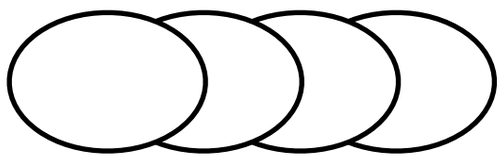

GB Key 3

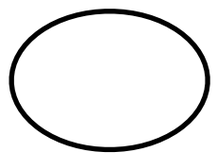

GB Key 4

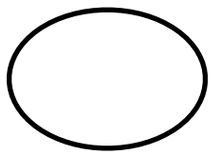

GB Key 5

Fairness

Problem: Group-By Key Distribution in Data May Be Skewed.

Solution: Scan All Group-By Keys In Parallel

Index Striding

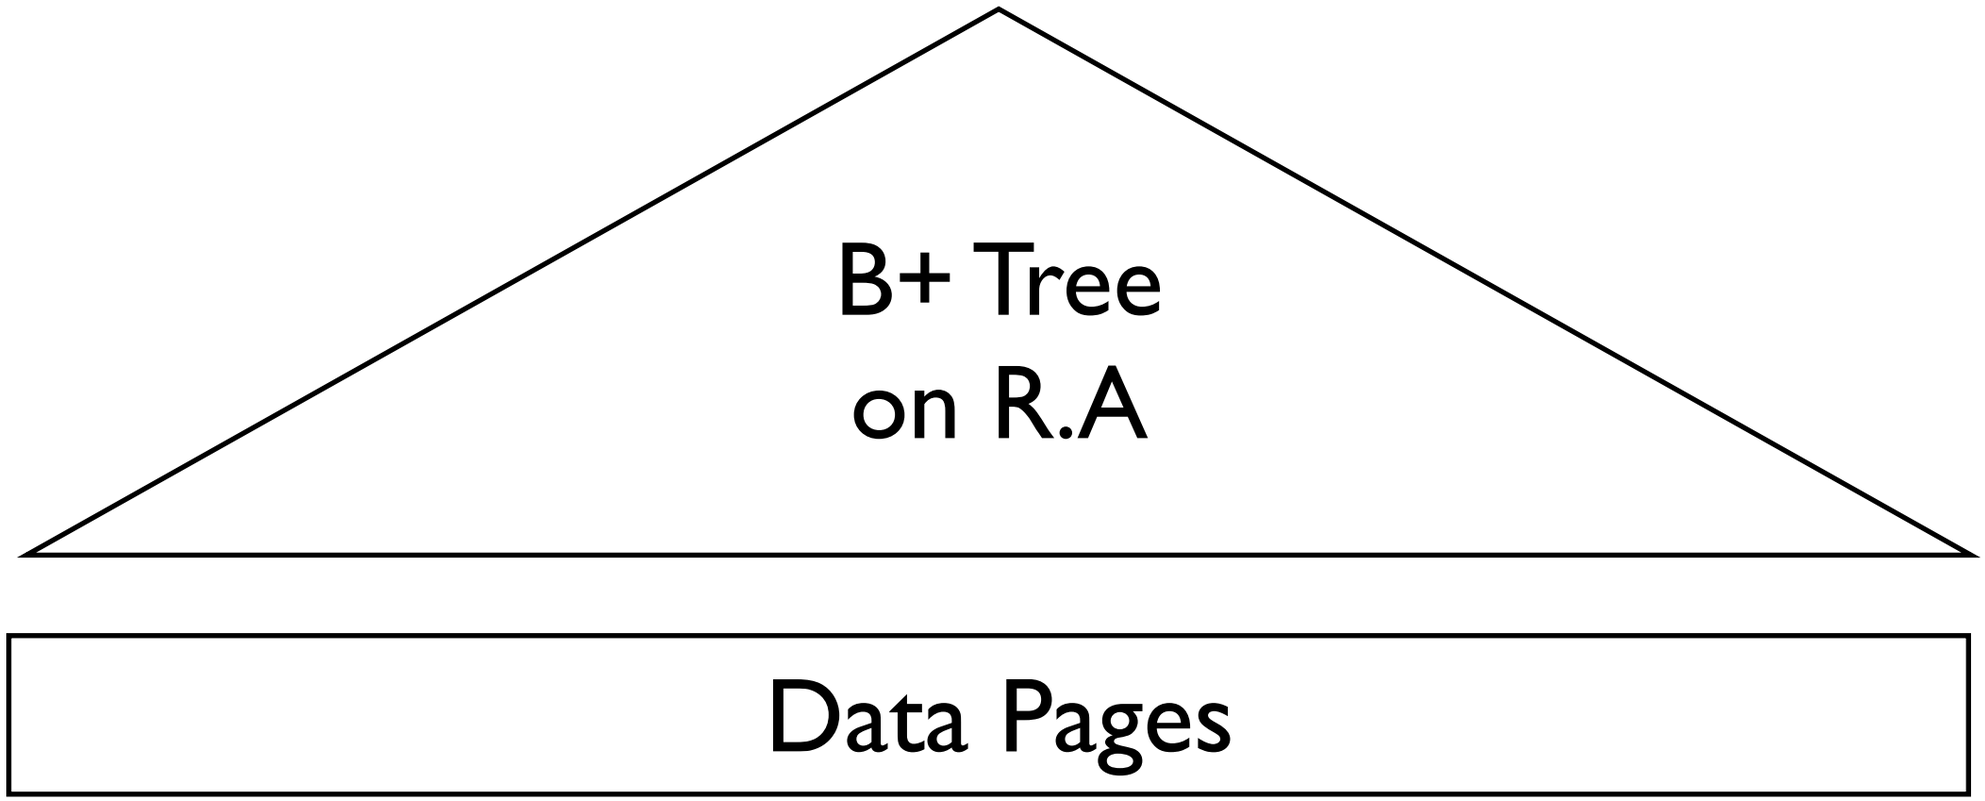

Index Striding

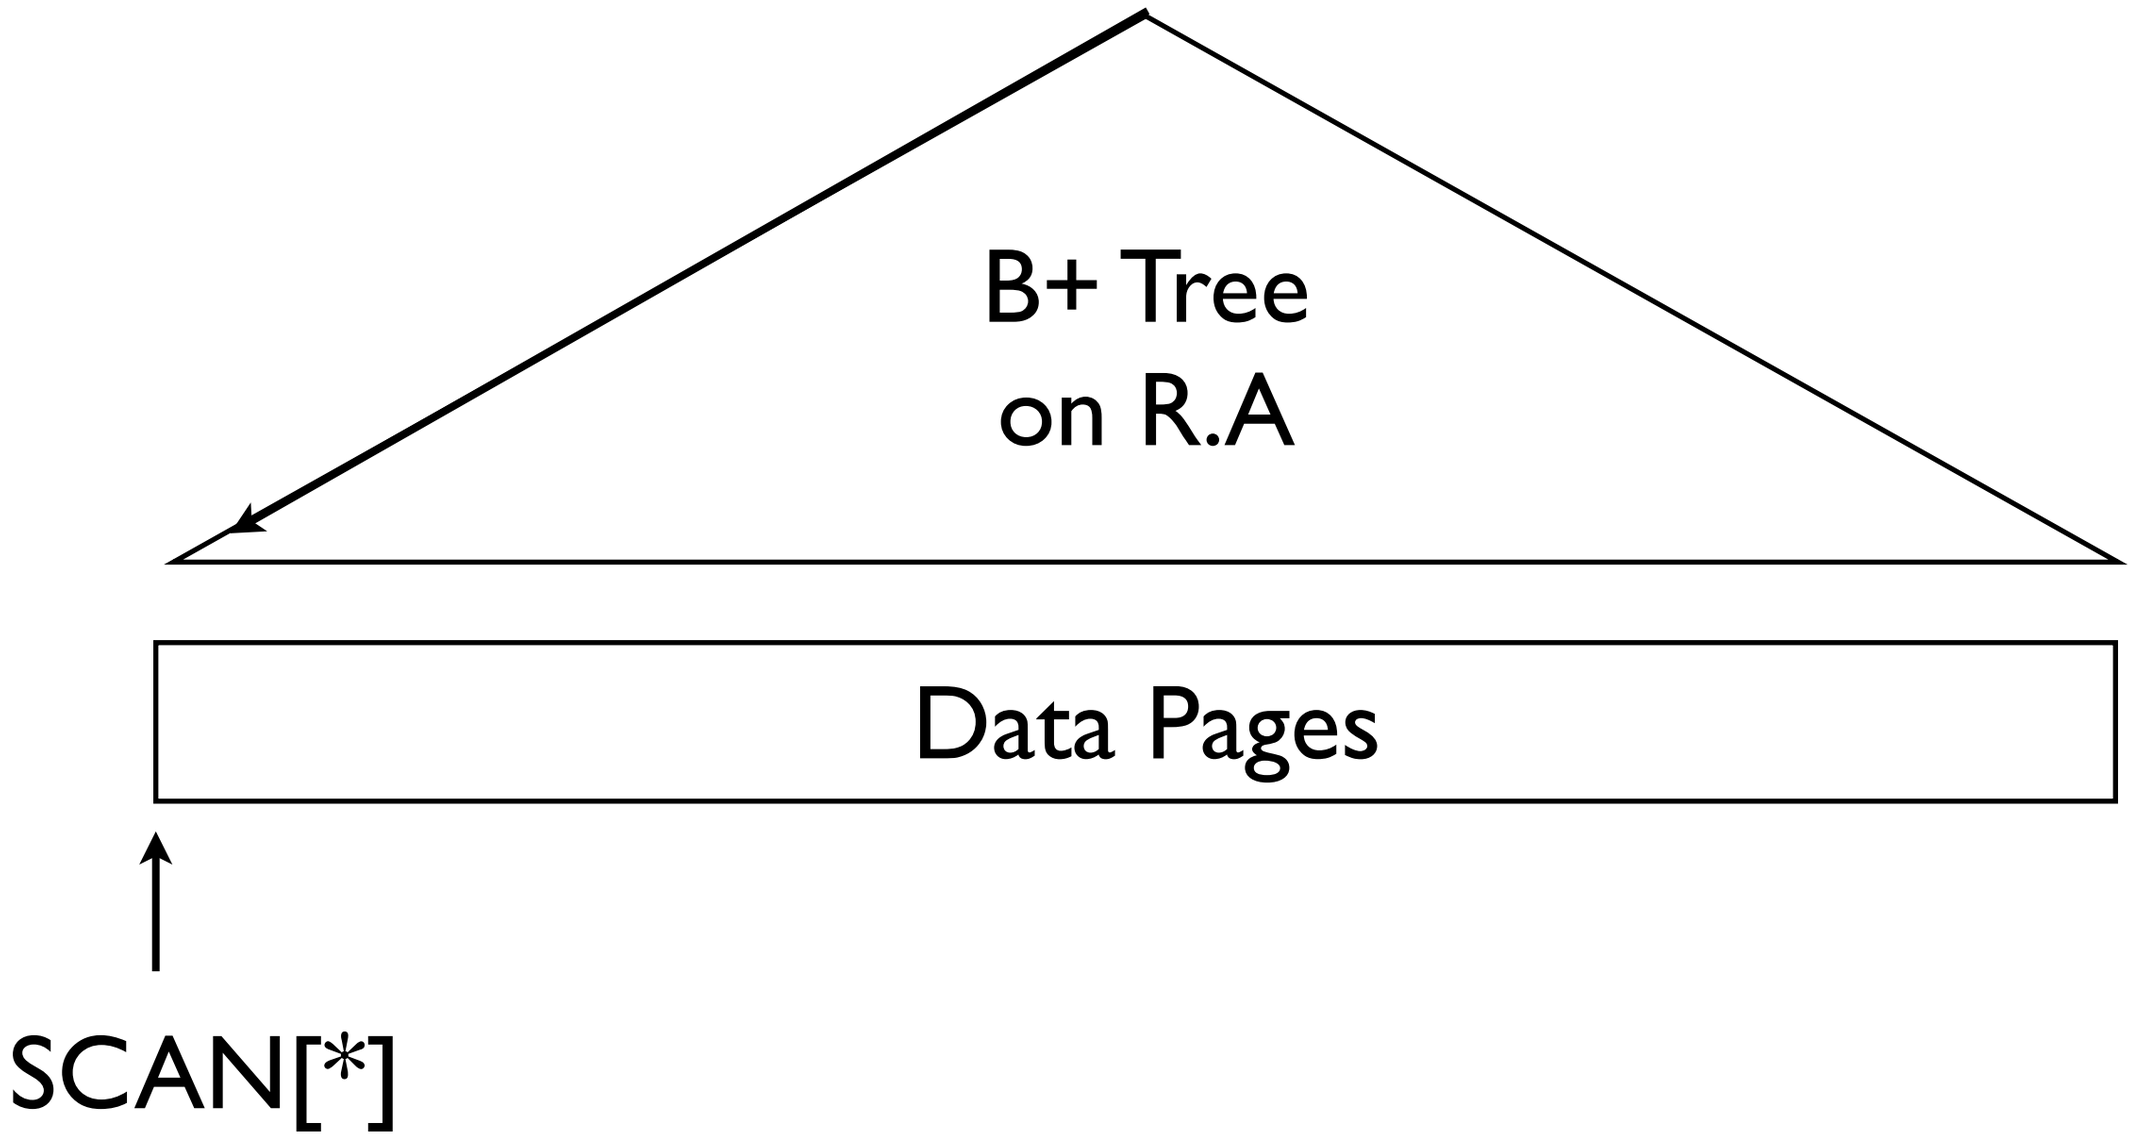

Index Striding

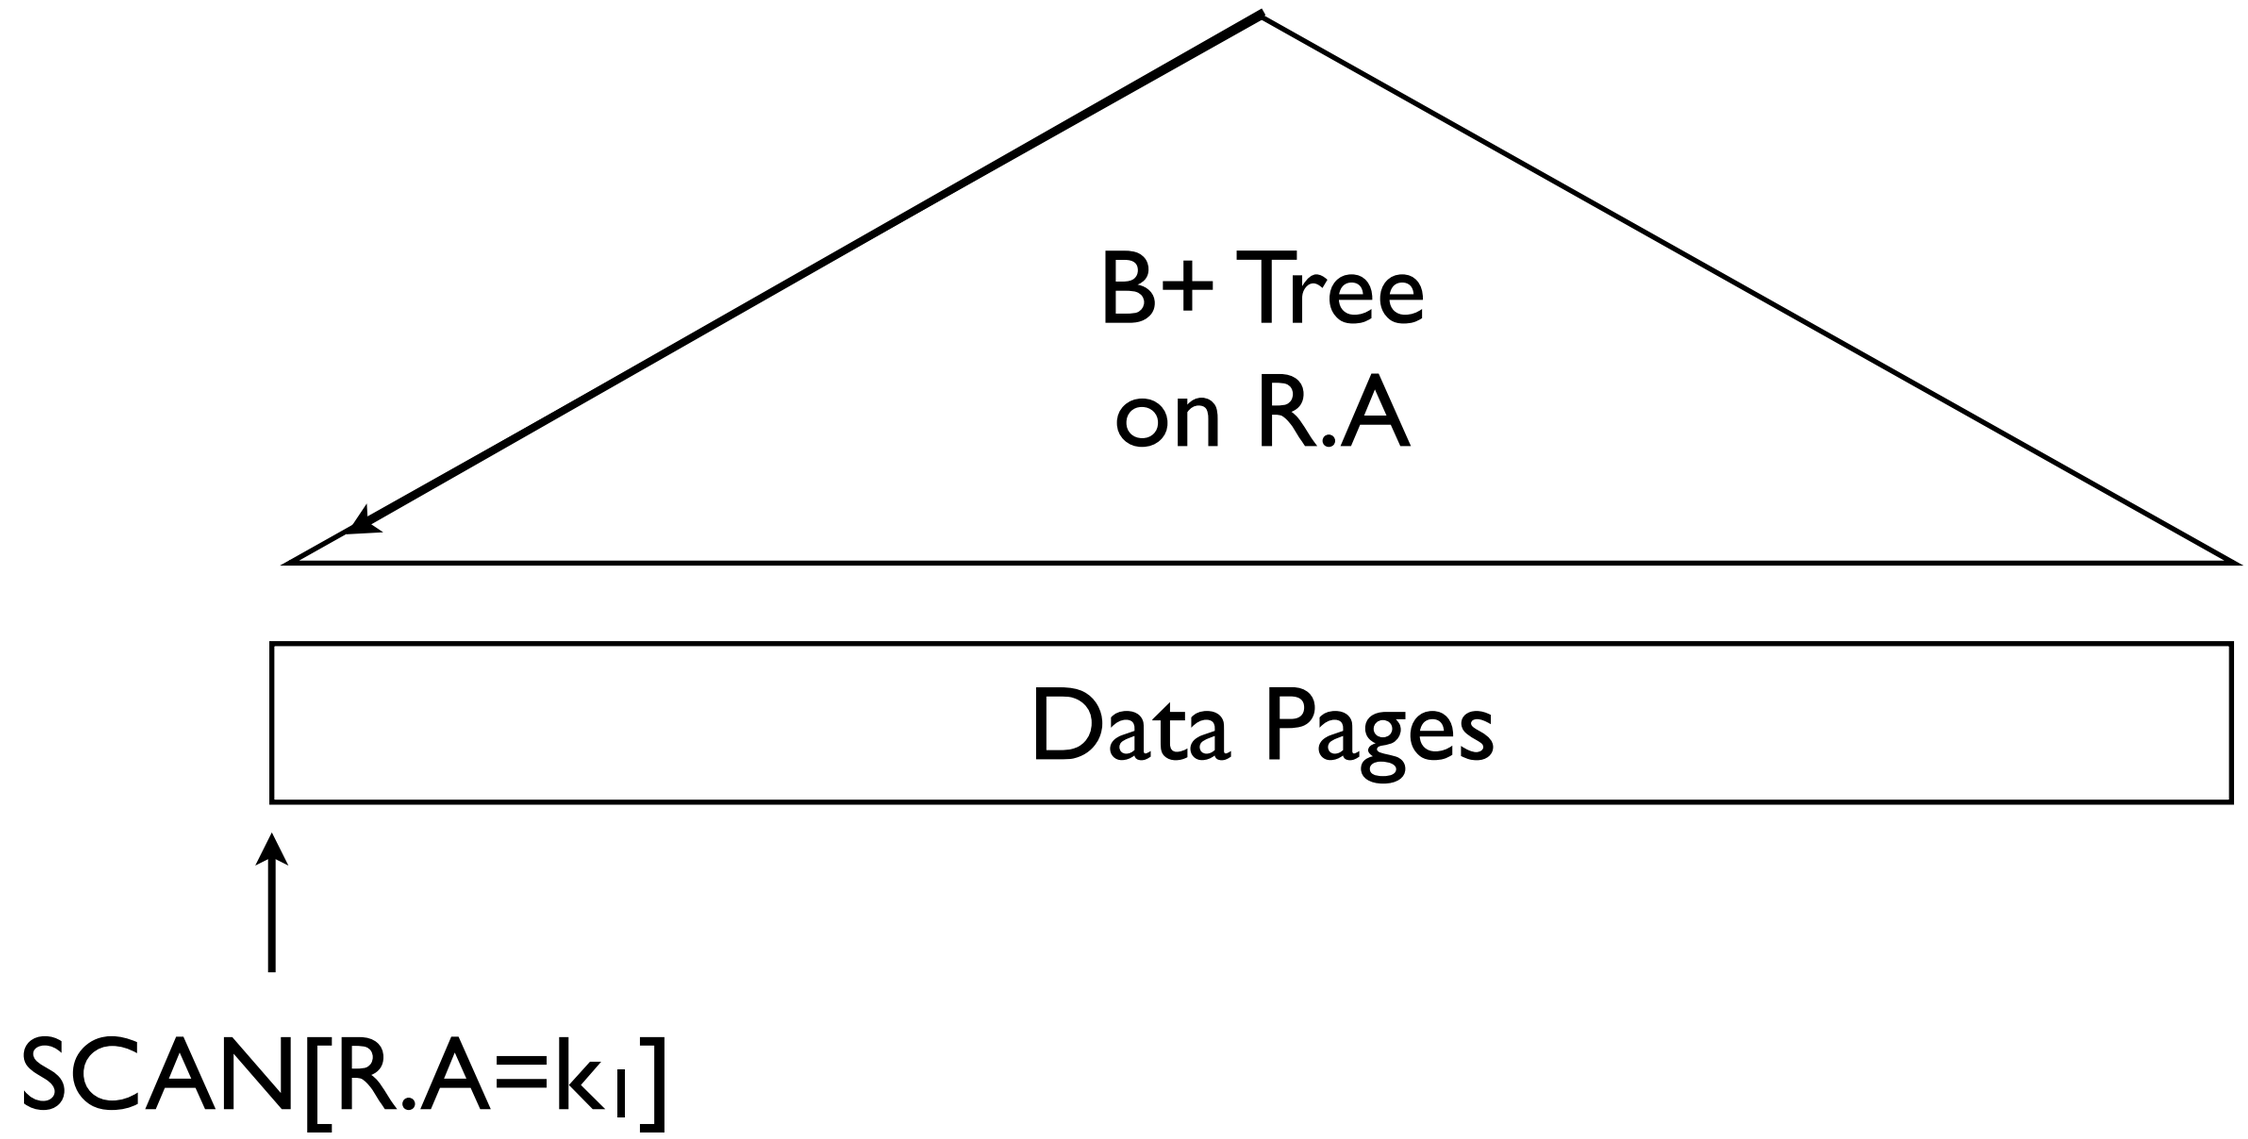

Index Striding

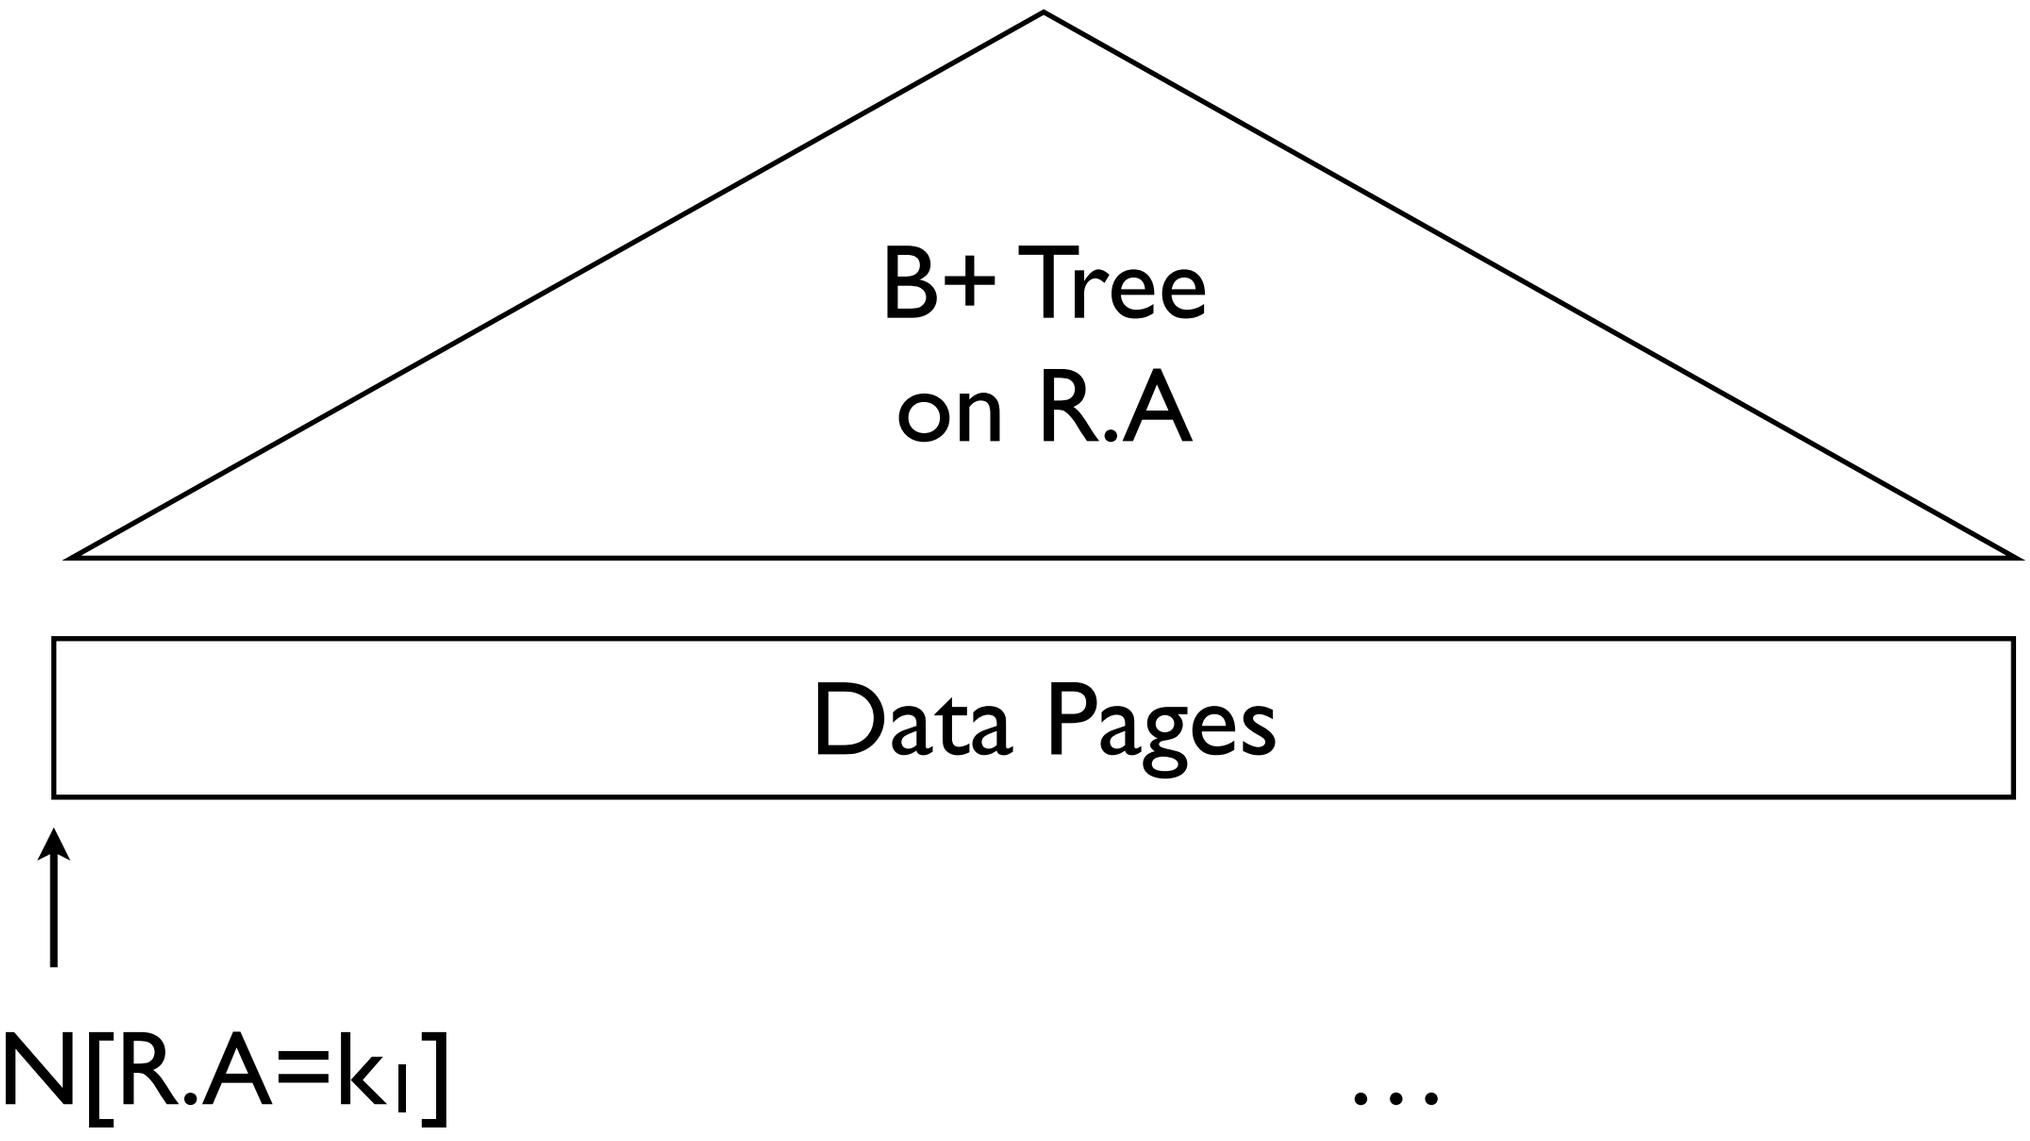

Index Striding

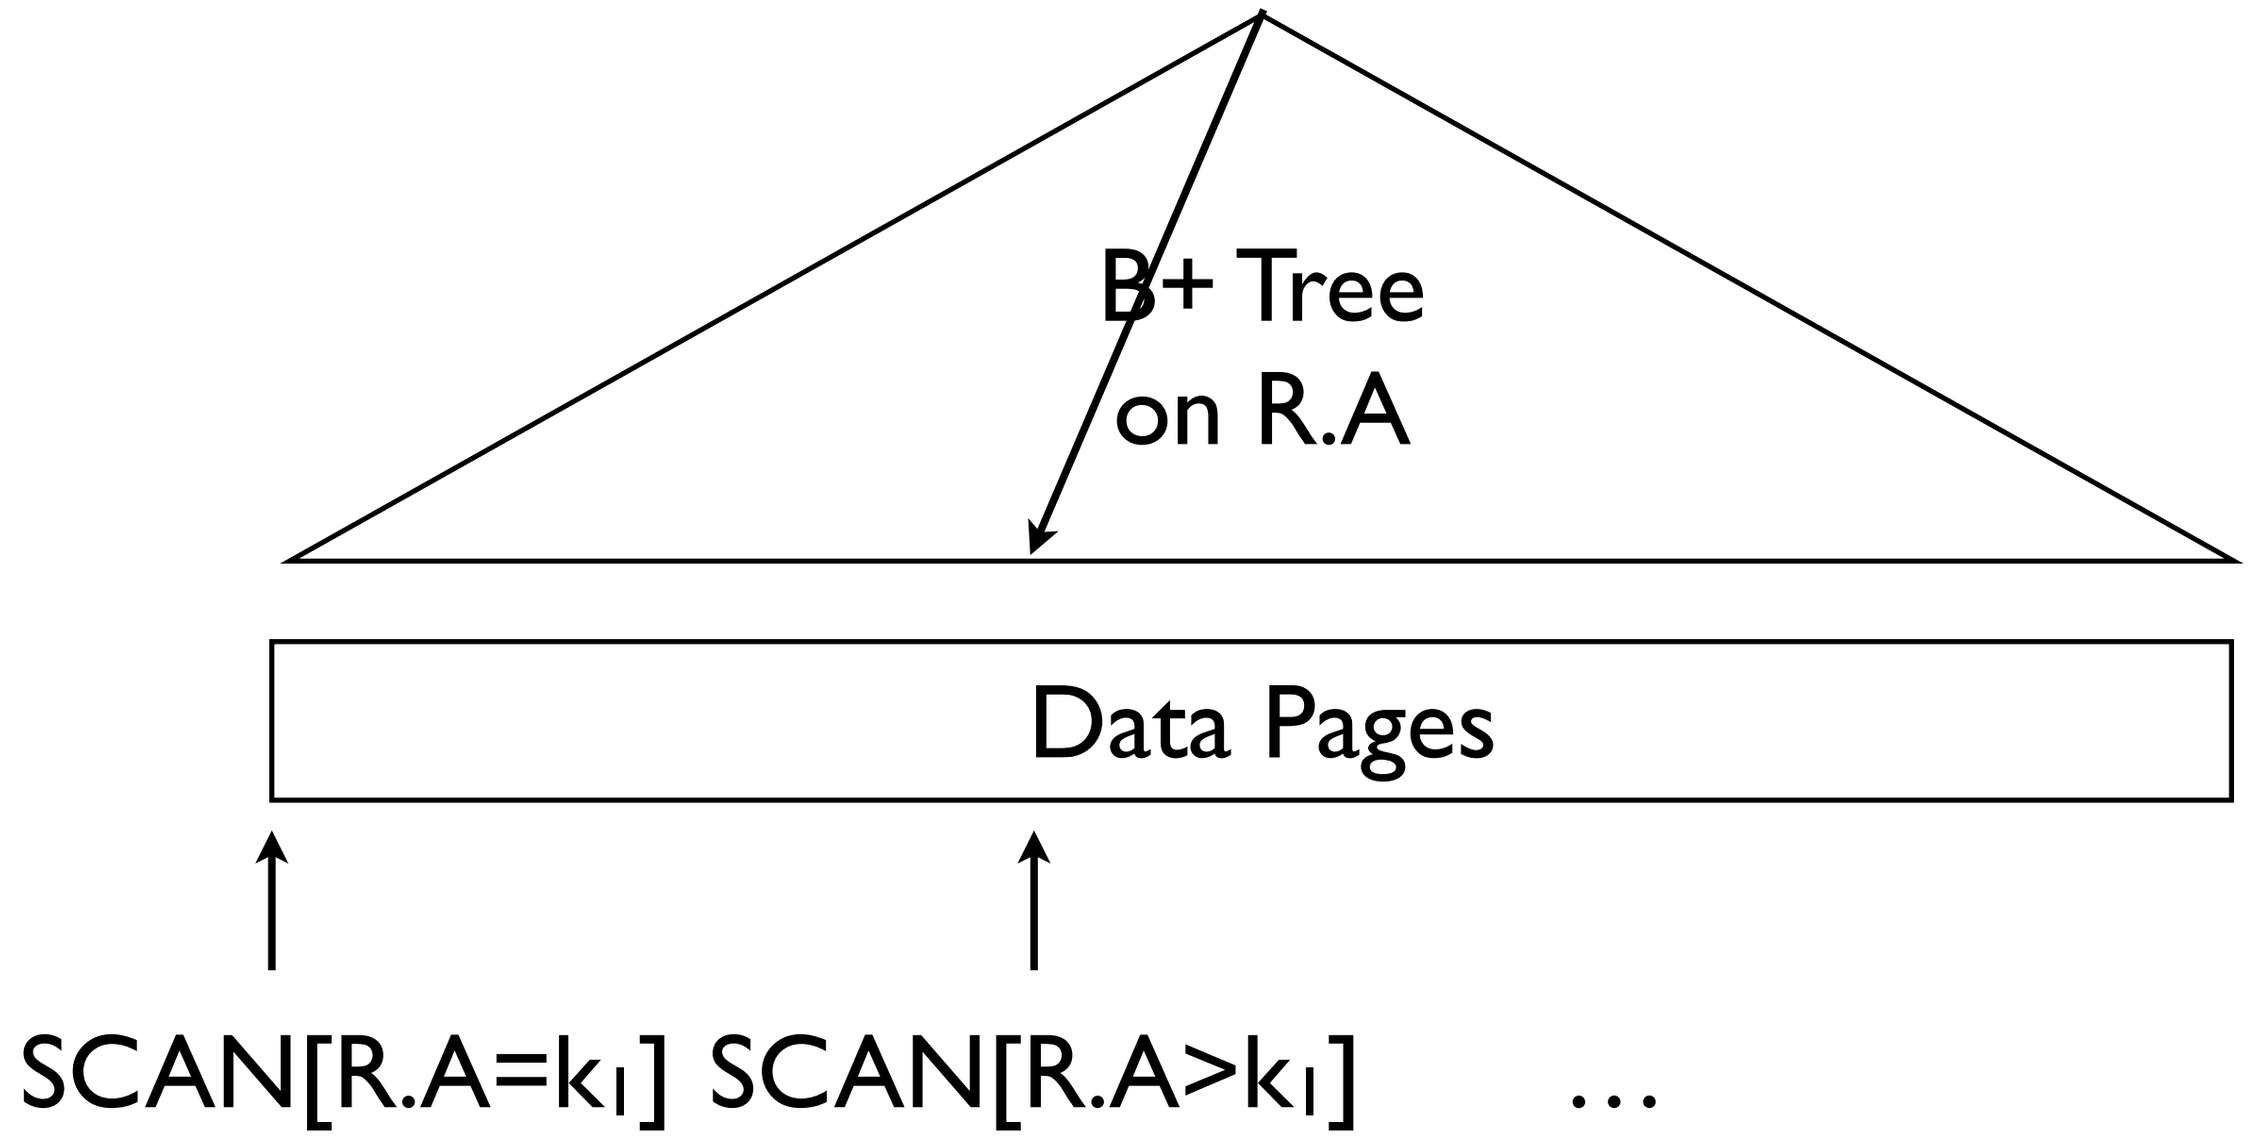

Index Striding

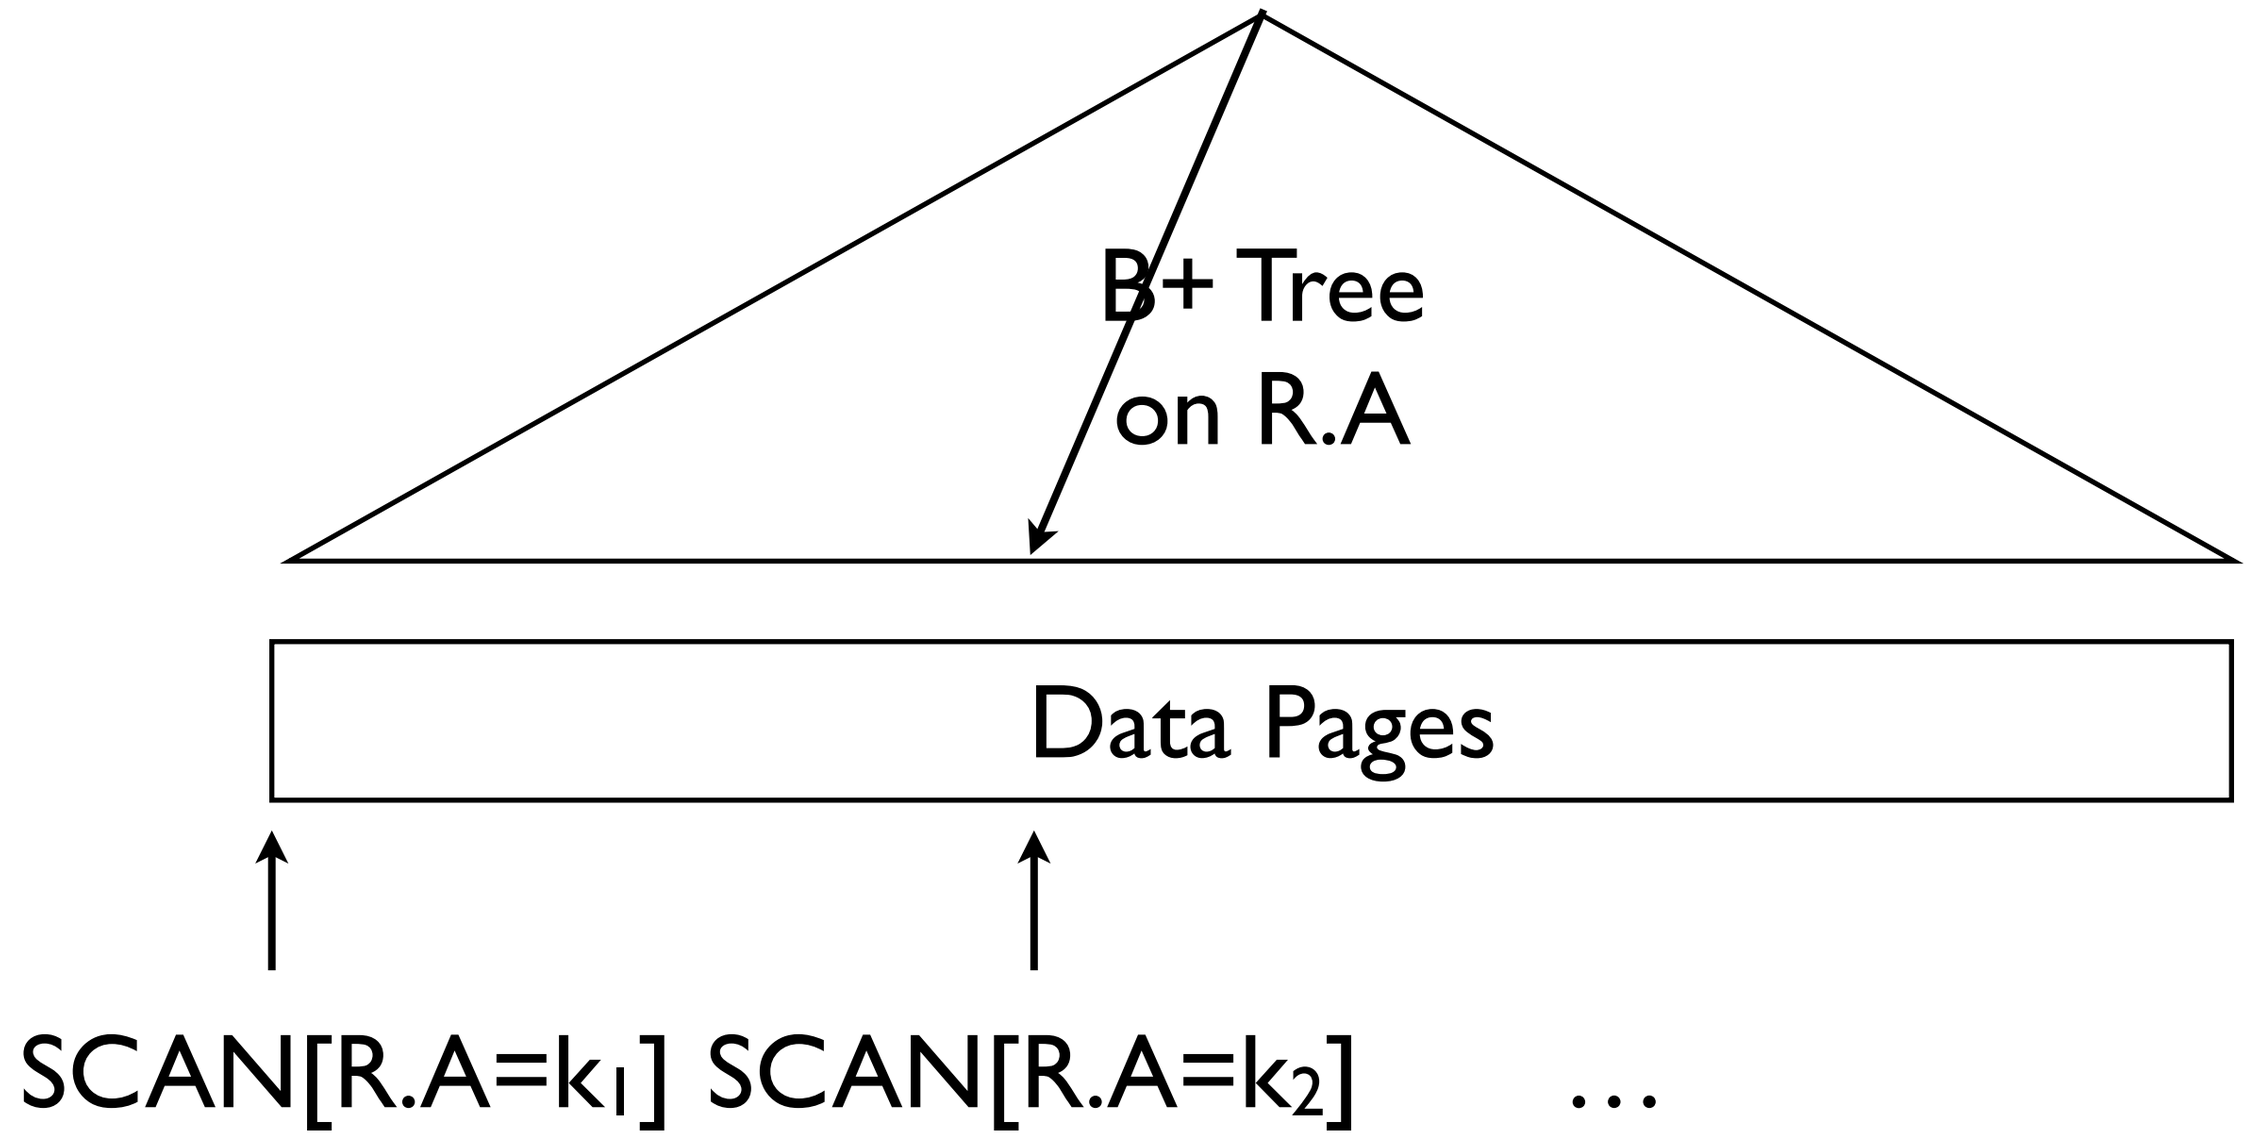

Index Striding

SCAN[R.A=k_1]

SCAN[R.A=k_2]

SCAN[R.A=k_3]

SCAN[R.A=k_4]

SCAN[R.A=k_5]

…

One Scan for Each GB Key

Each Scan is now a Heap Scan

Split resources evenly between each scan created.

Blocking

Problem: Can't get "Online" results with blocking operations

Solution: Only required blocking op: Join. Use nonblocking joins.

Non-Blocking Joins

- Sort/Merge Join
  - We want the data unsorted
- Index-Nested Loop Join
  - Could work if only few tuples matched.
- Hybrid Hash Join
  - Could work if one table is small.

Blocking

Problem: Join Algos are Blocking

Solution: Ripple Joins

Ripple Join

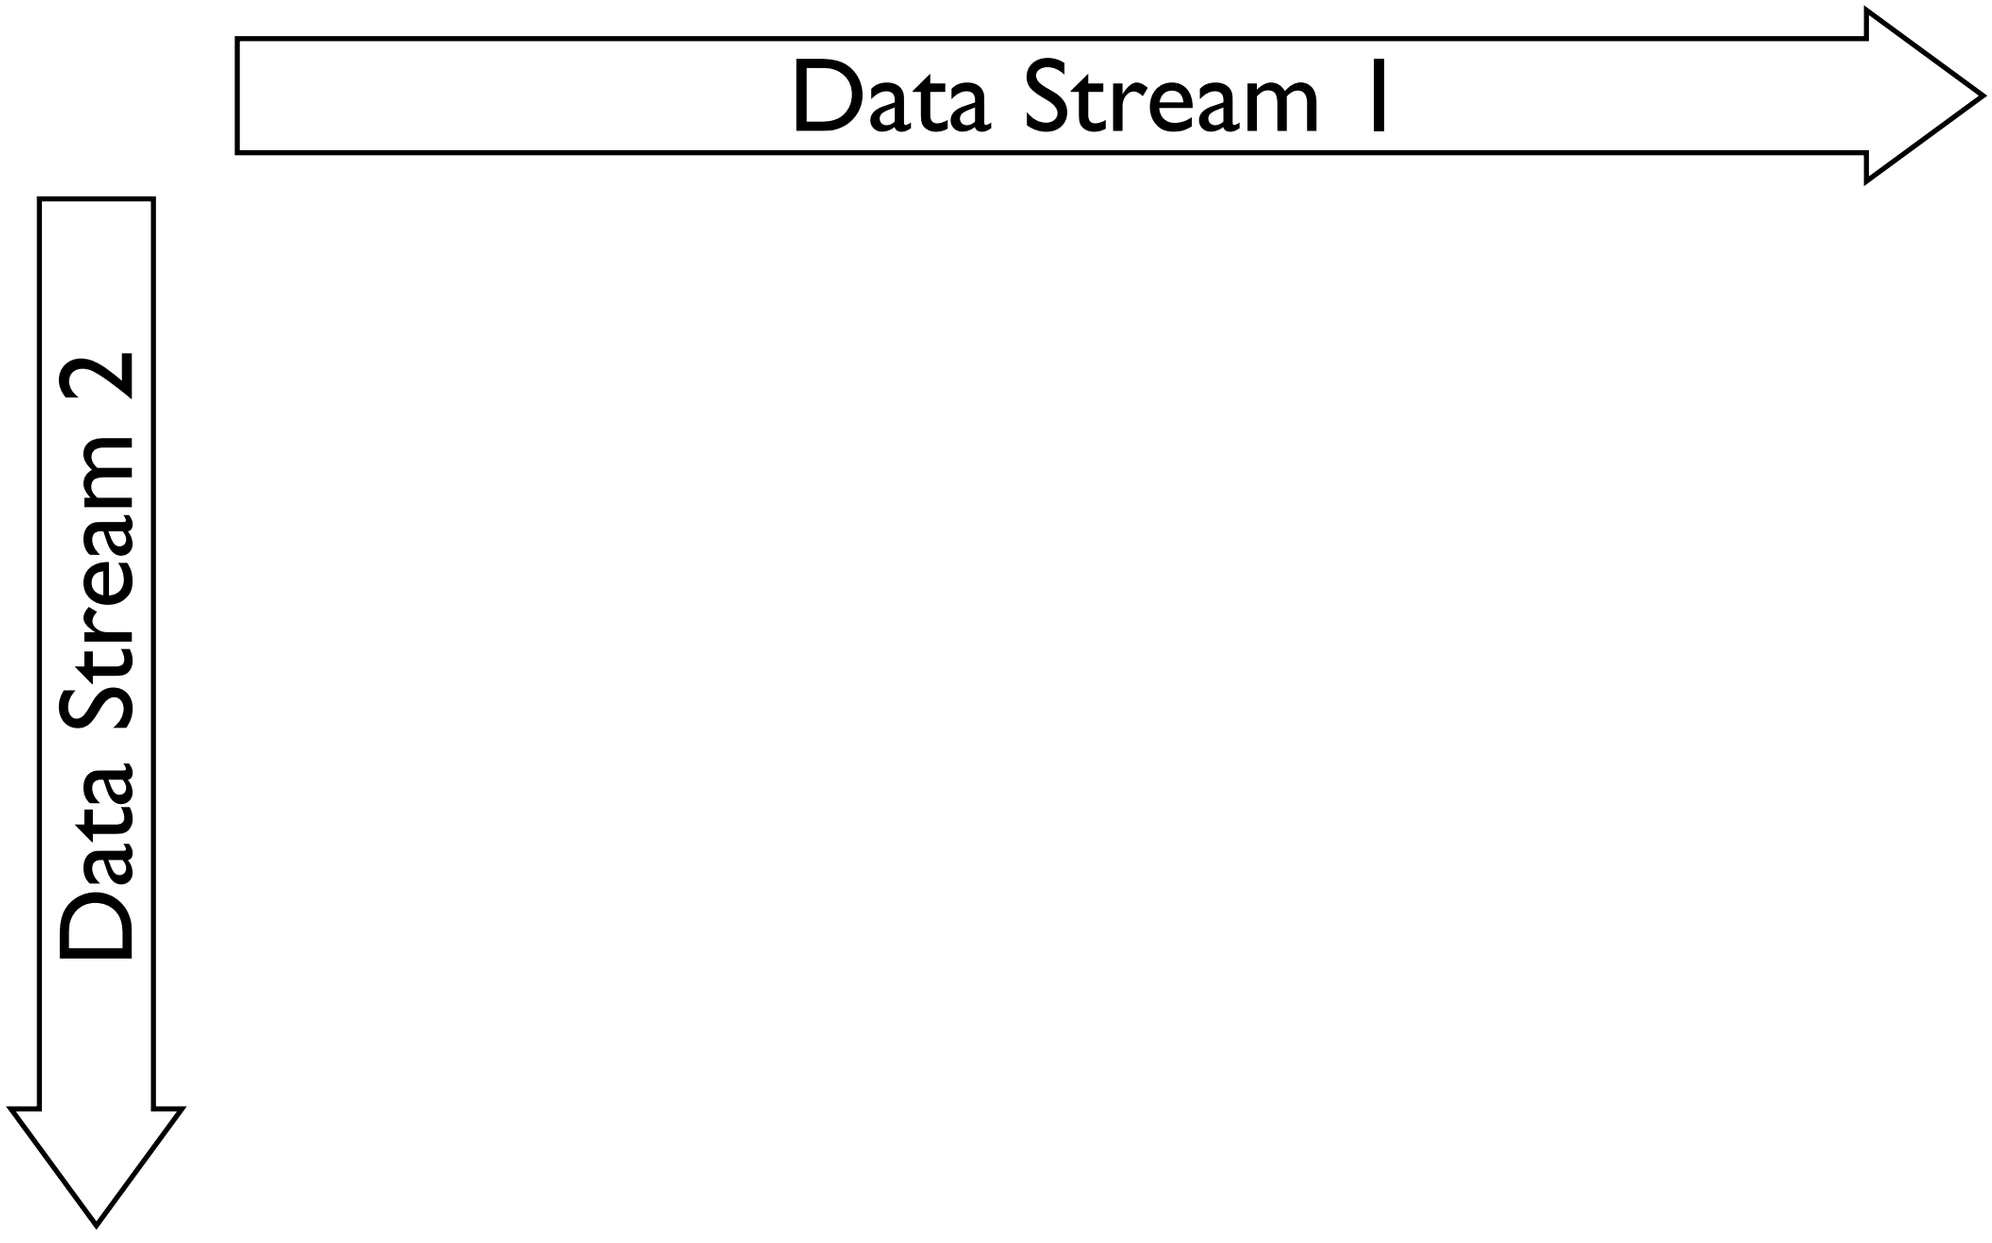

Ripple Join

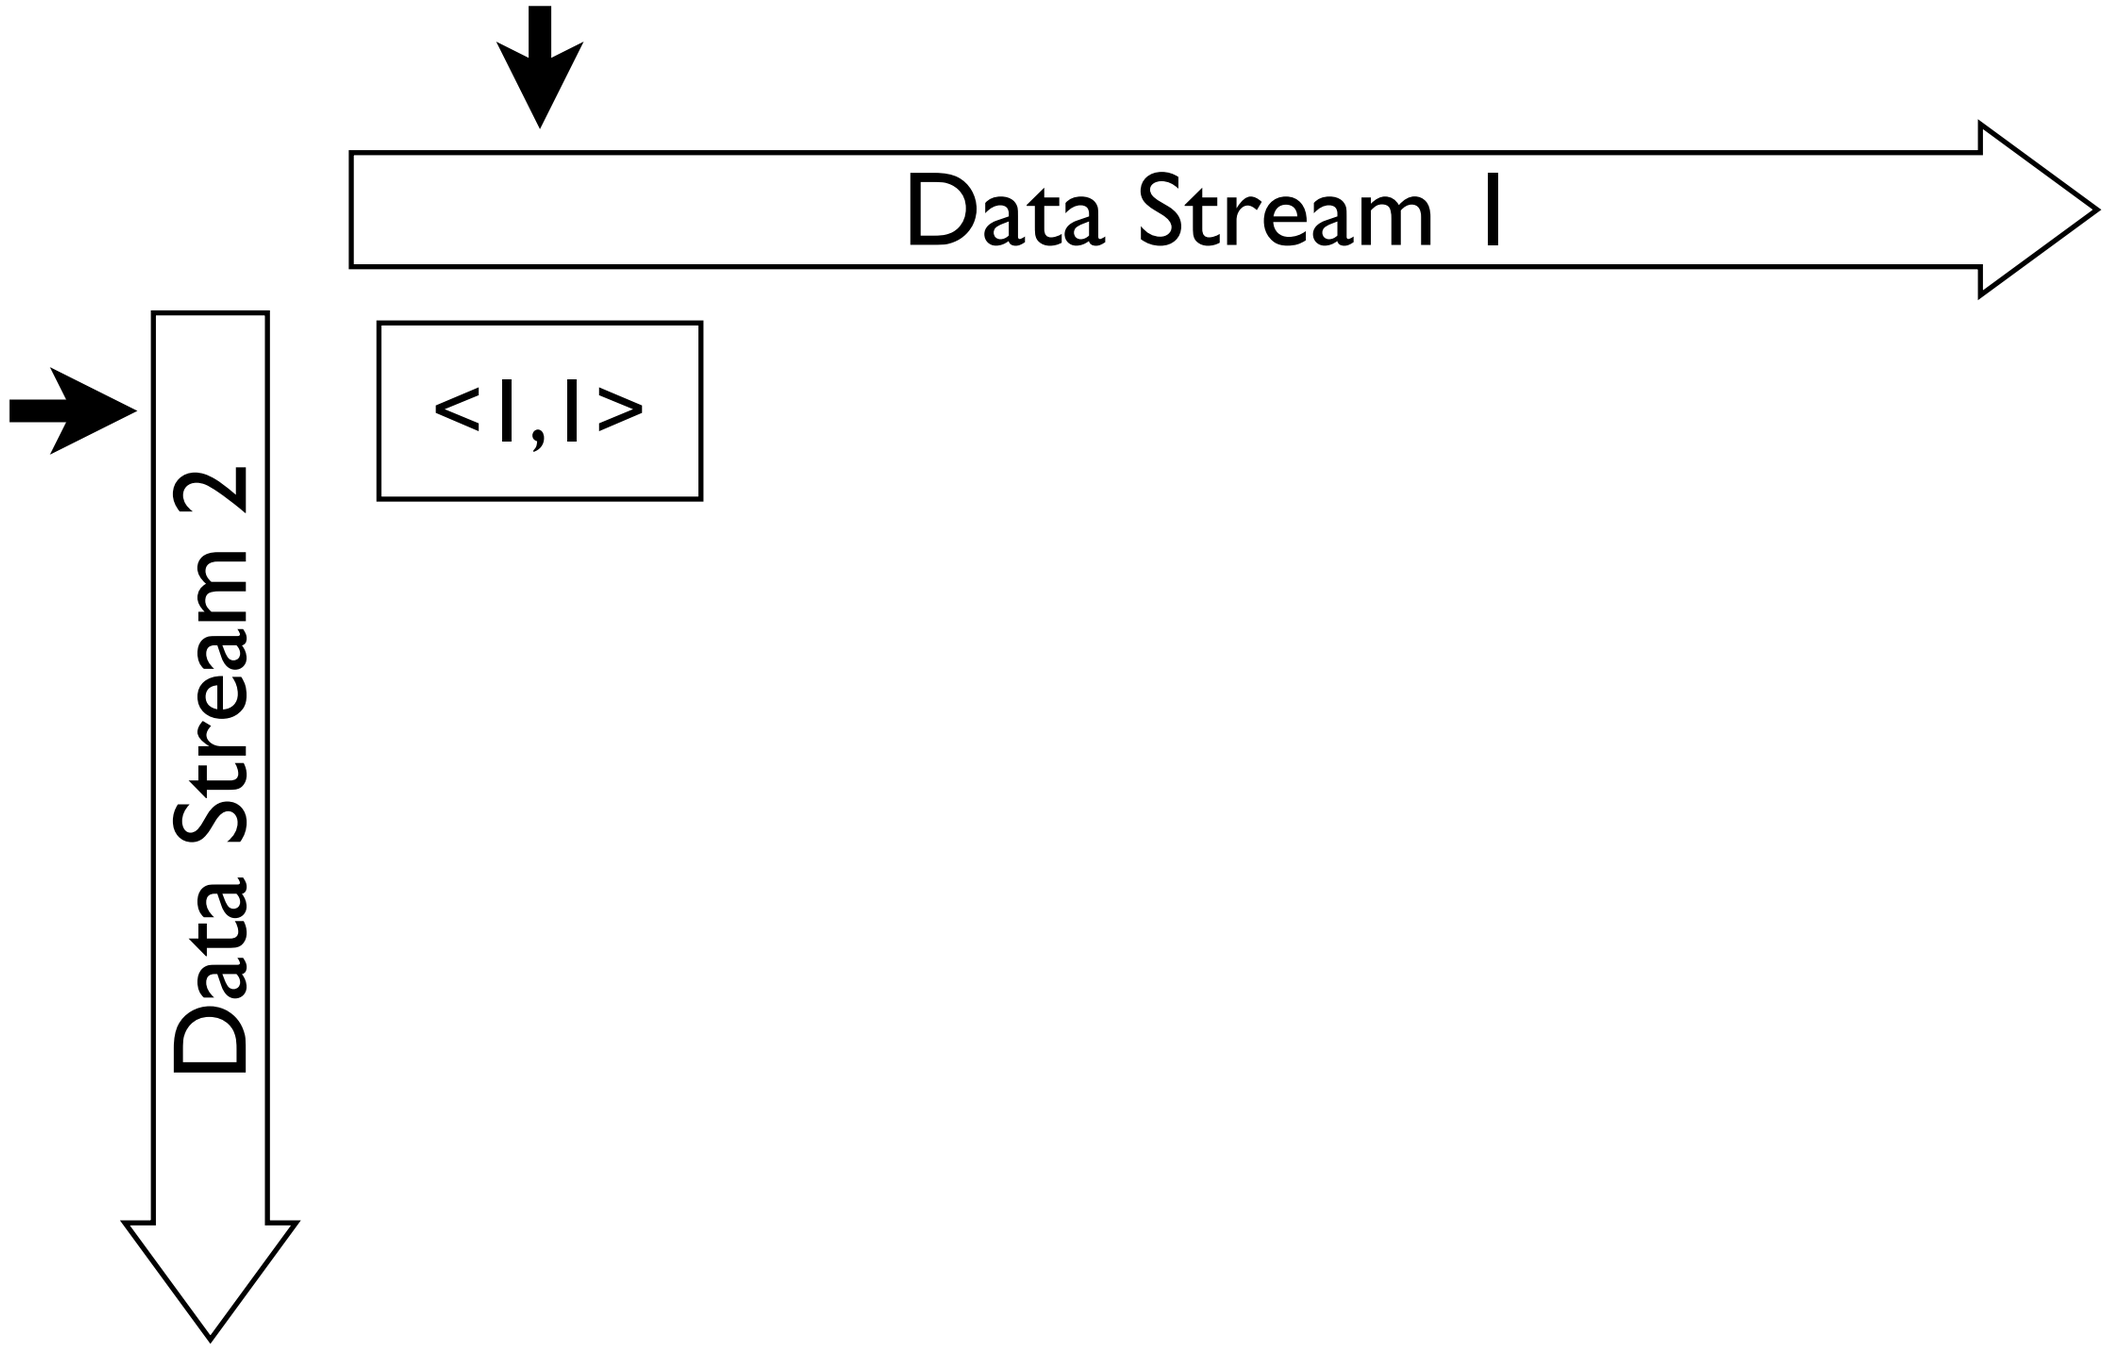

Ripple Join

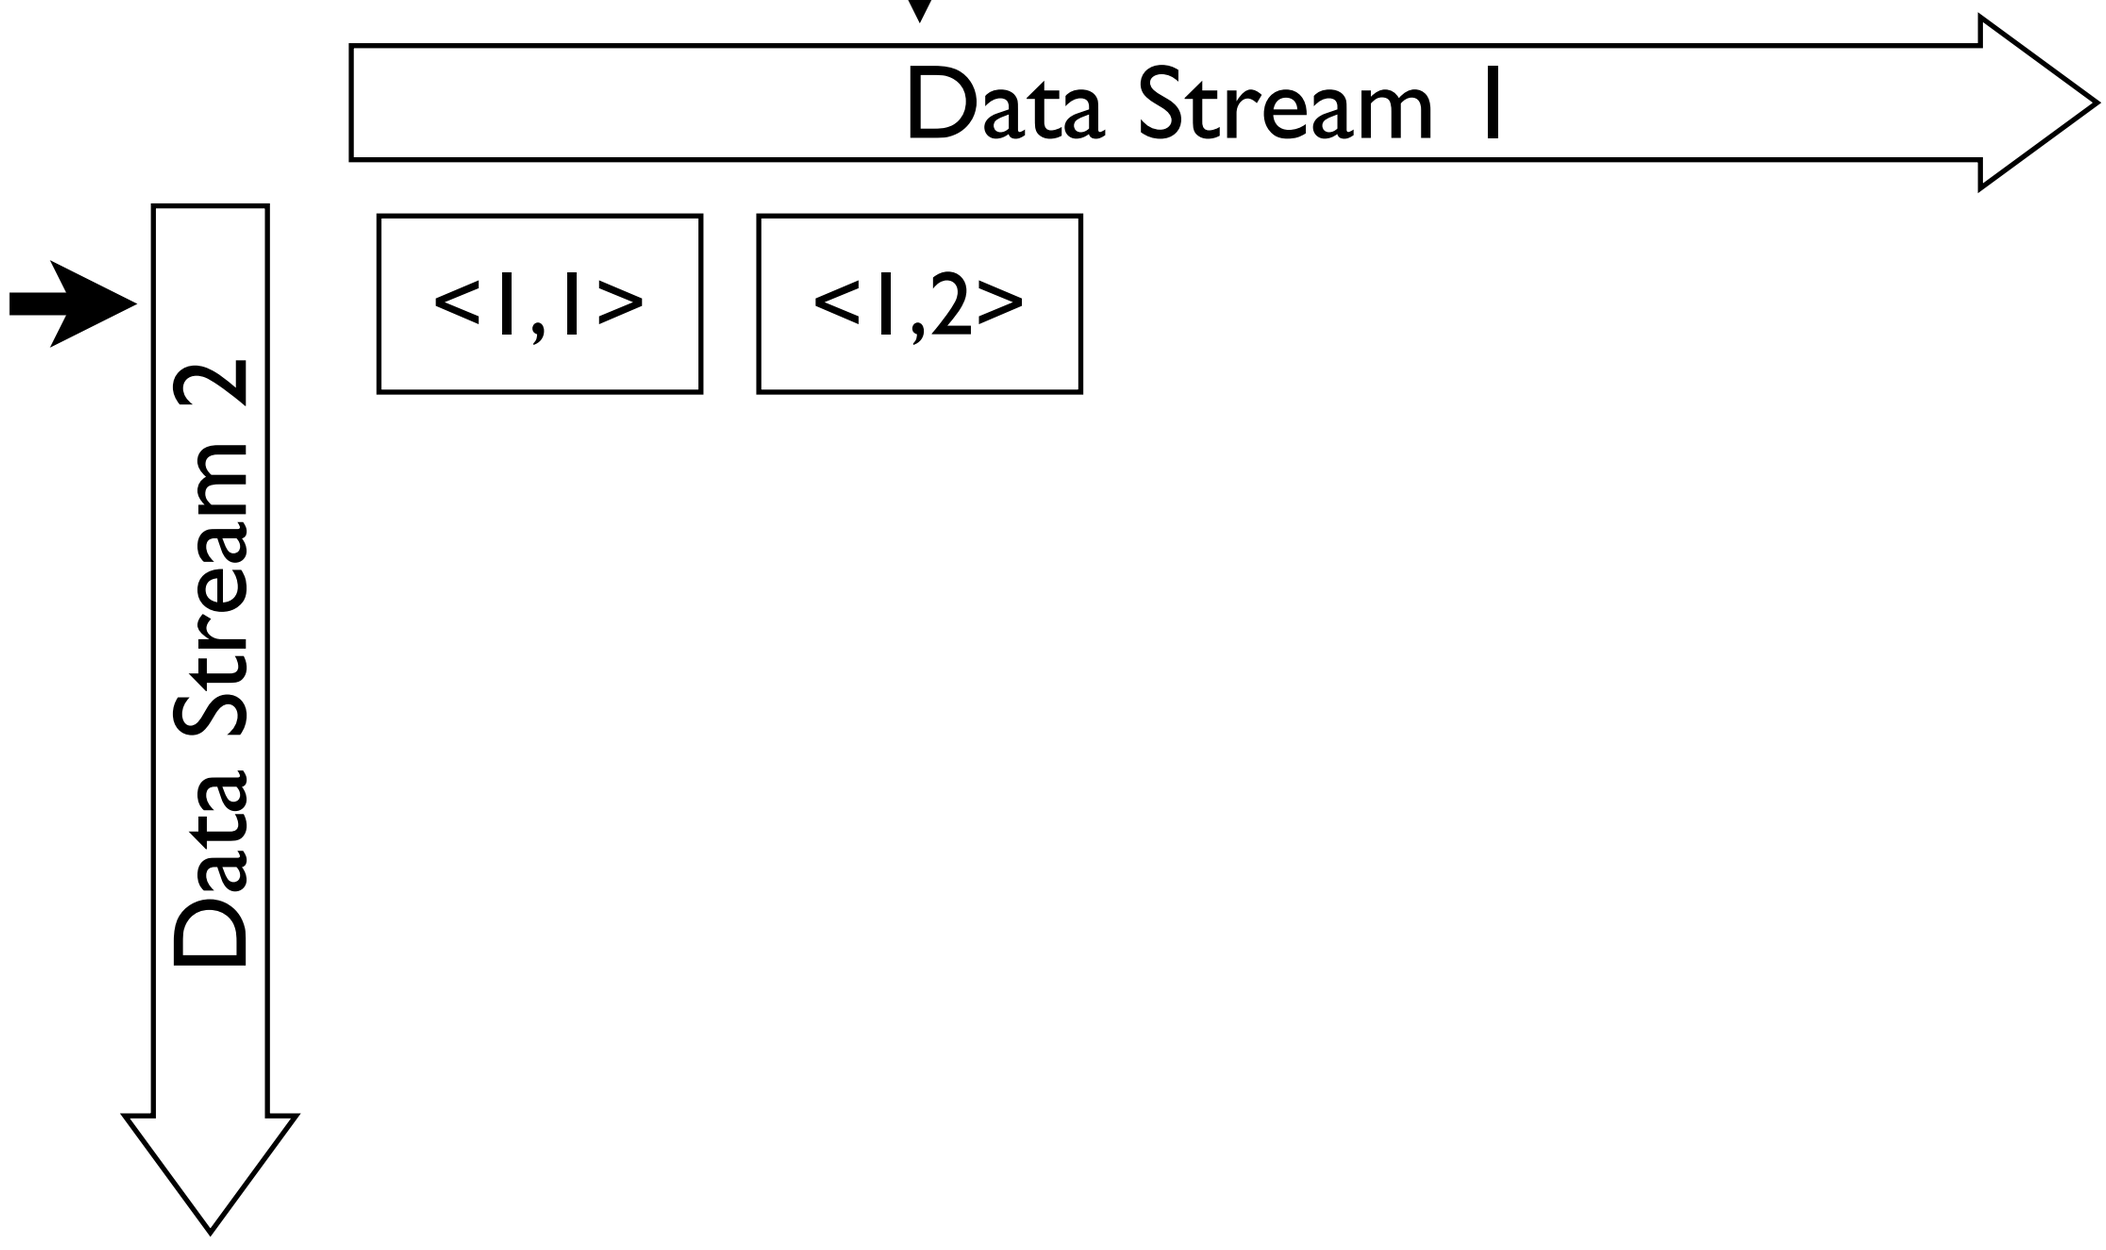

Ripple Join

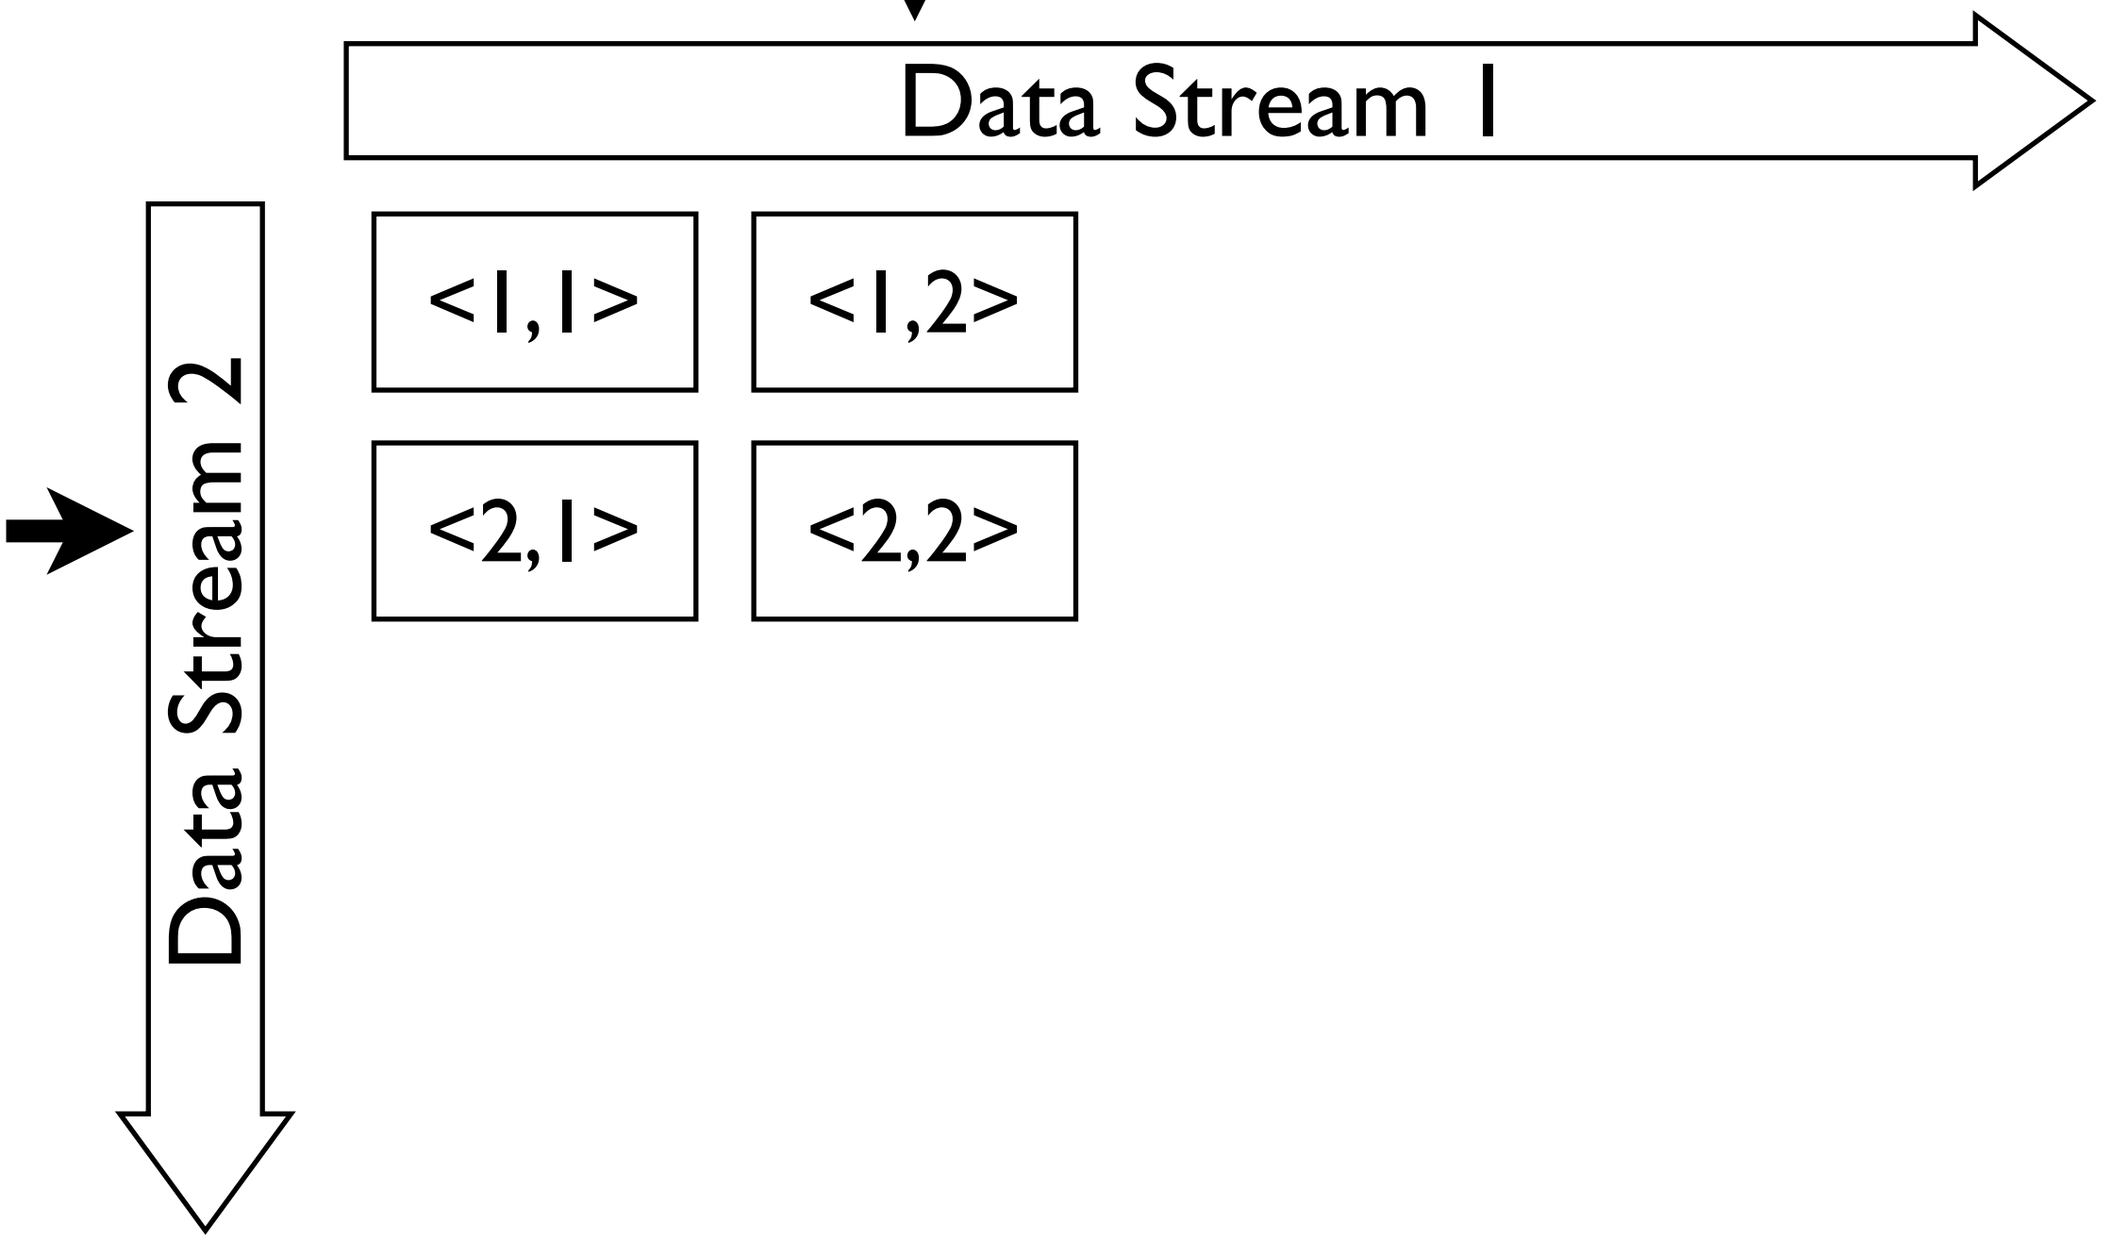

Ripple Join

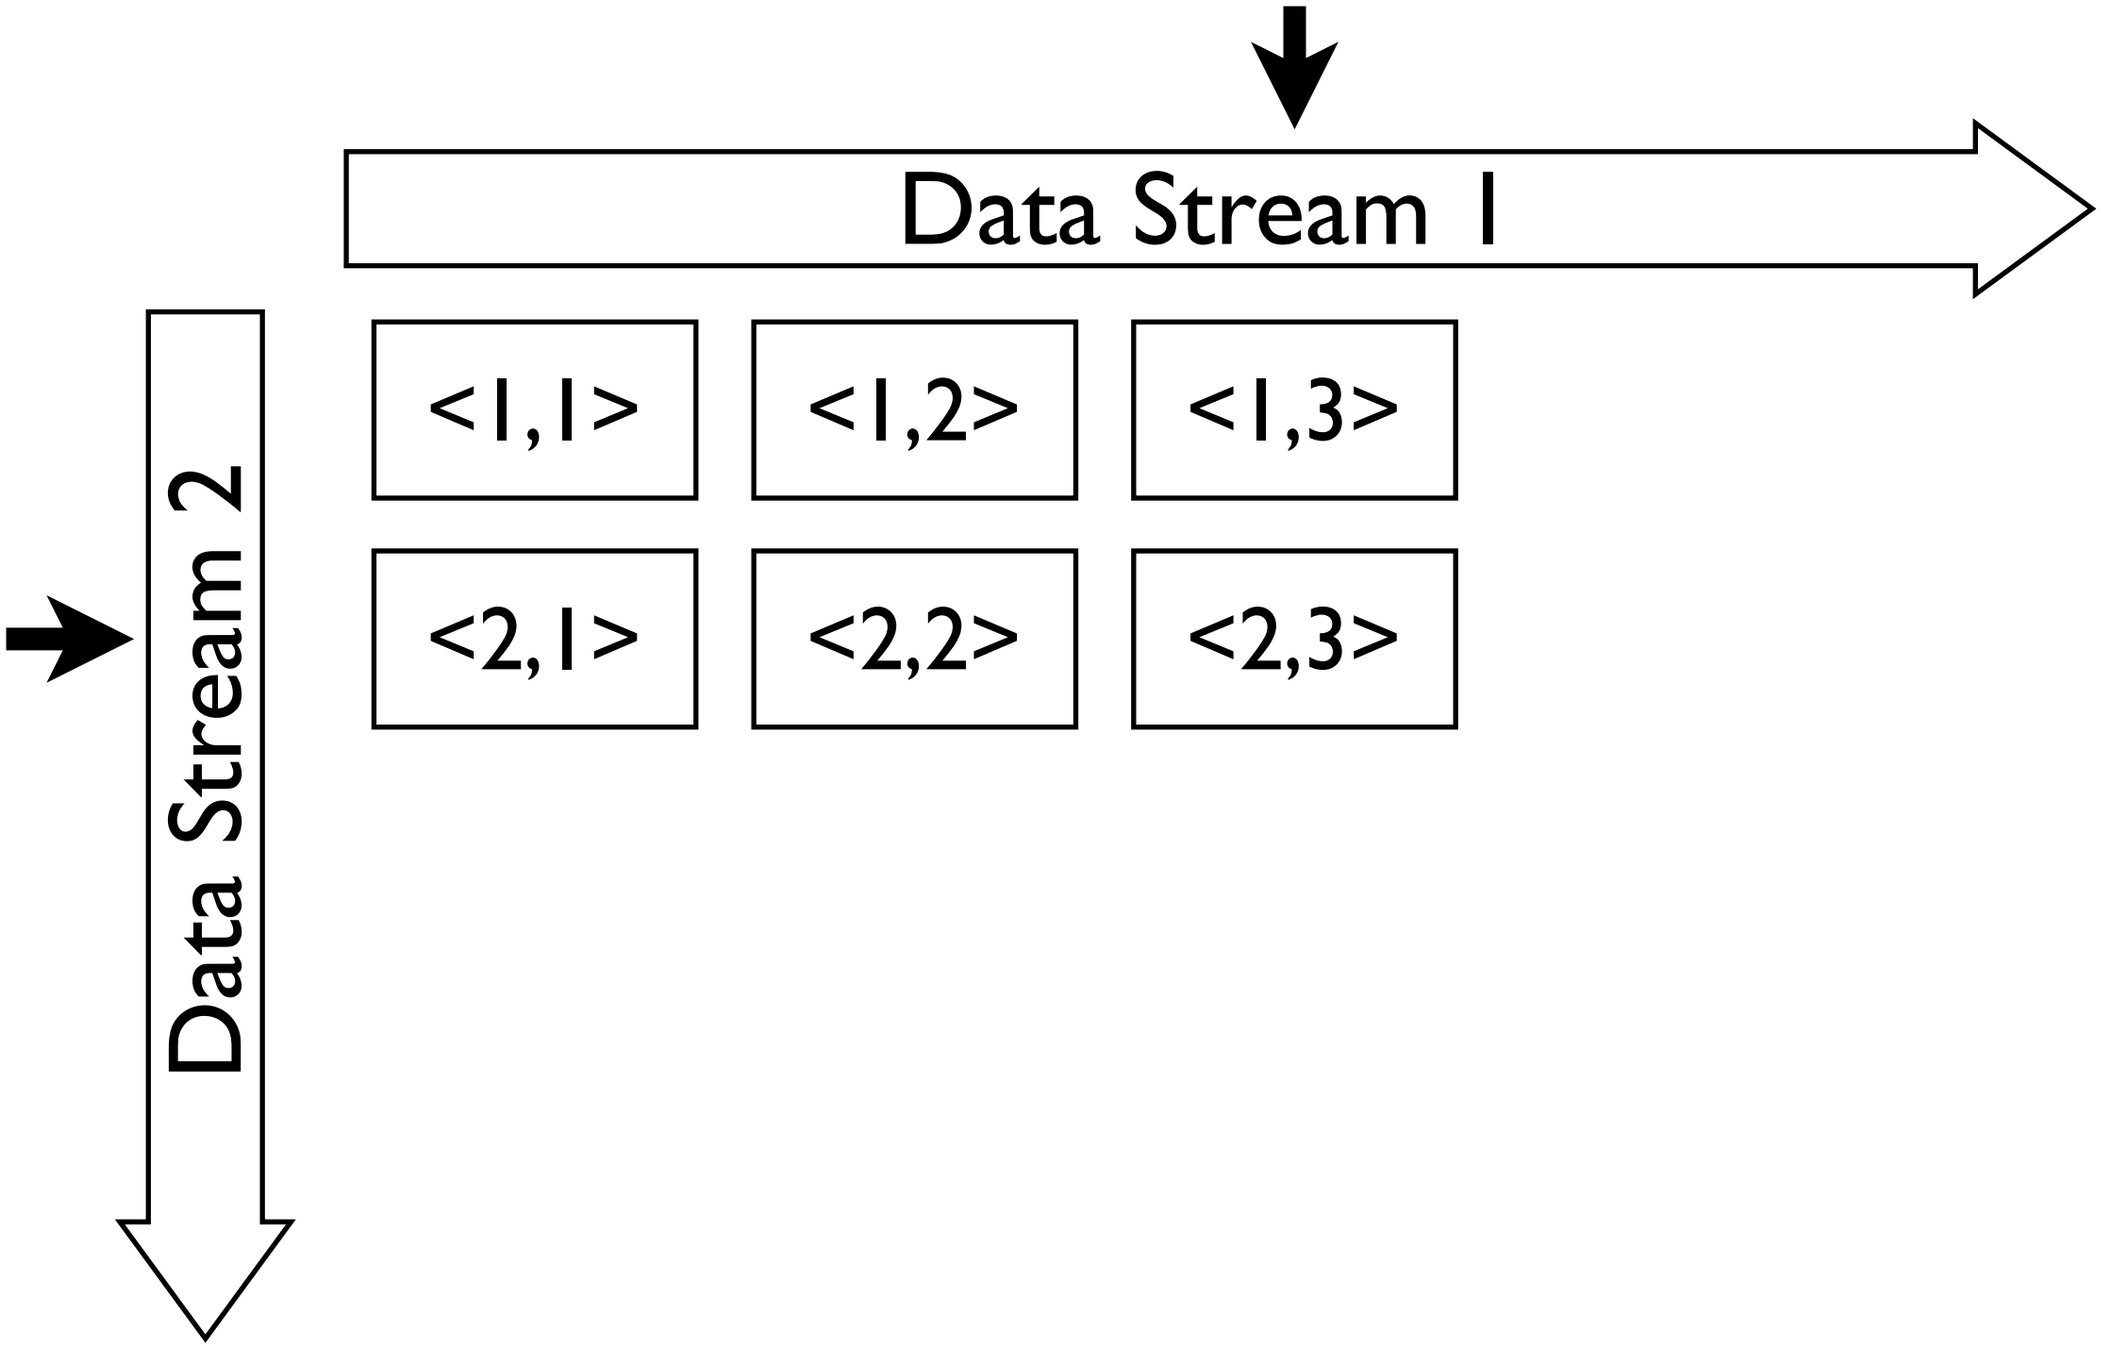

Ripple Join

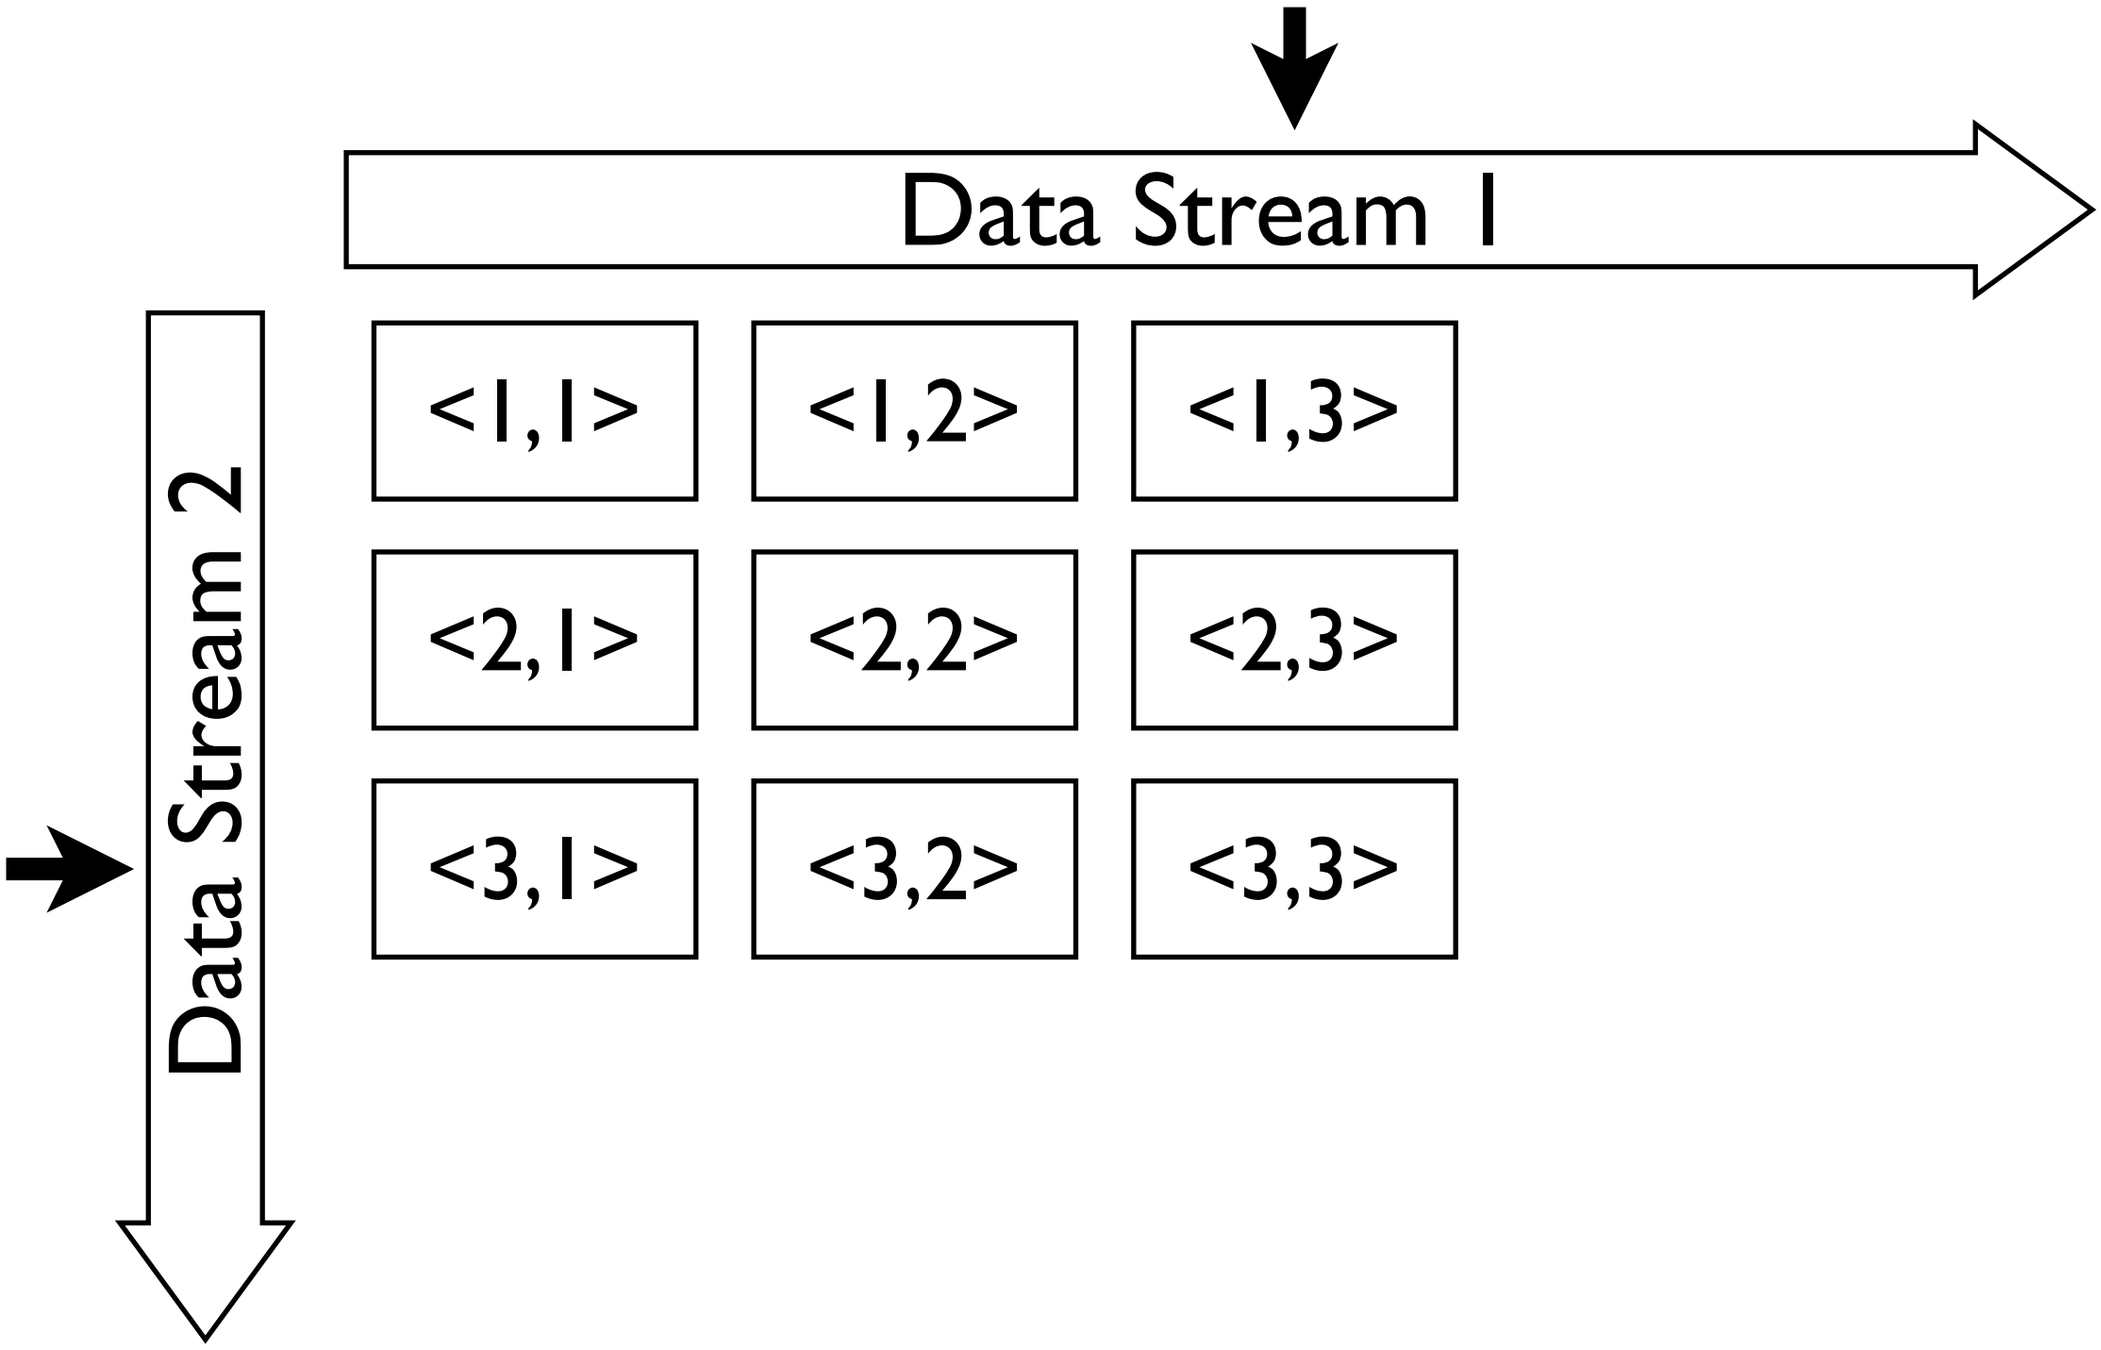

Ripple Join

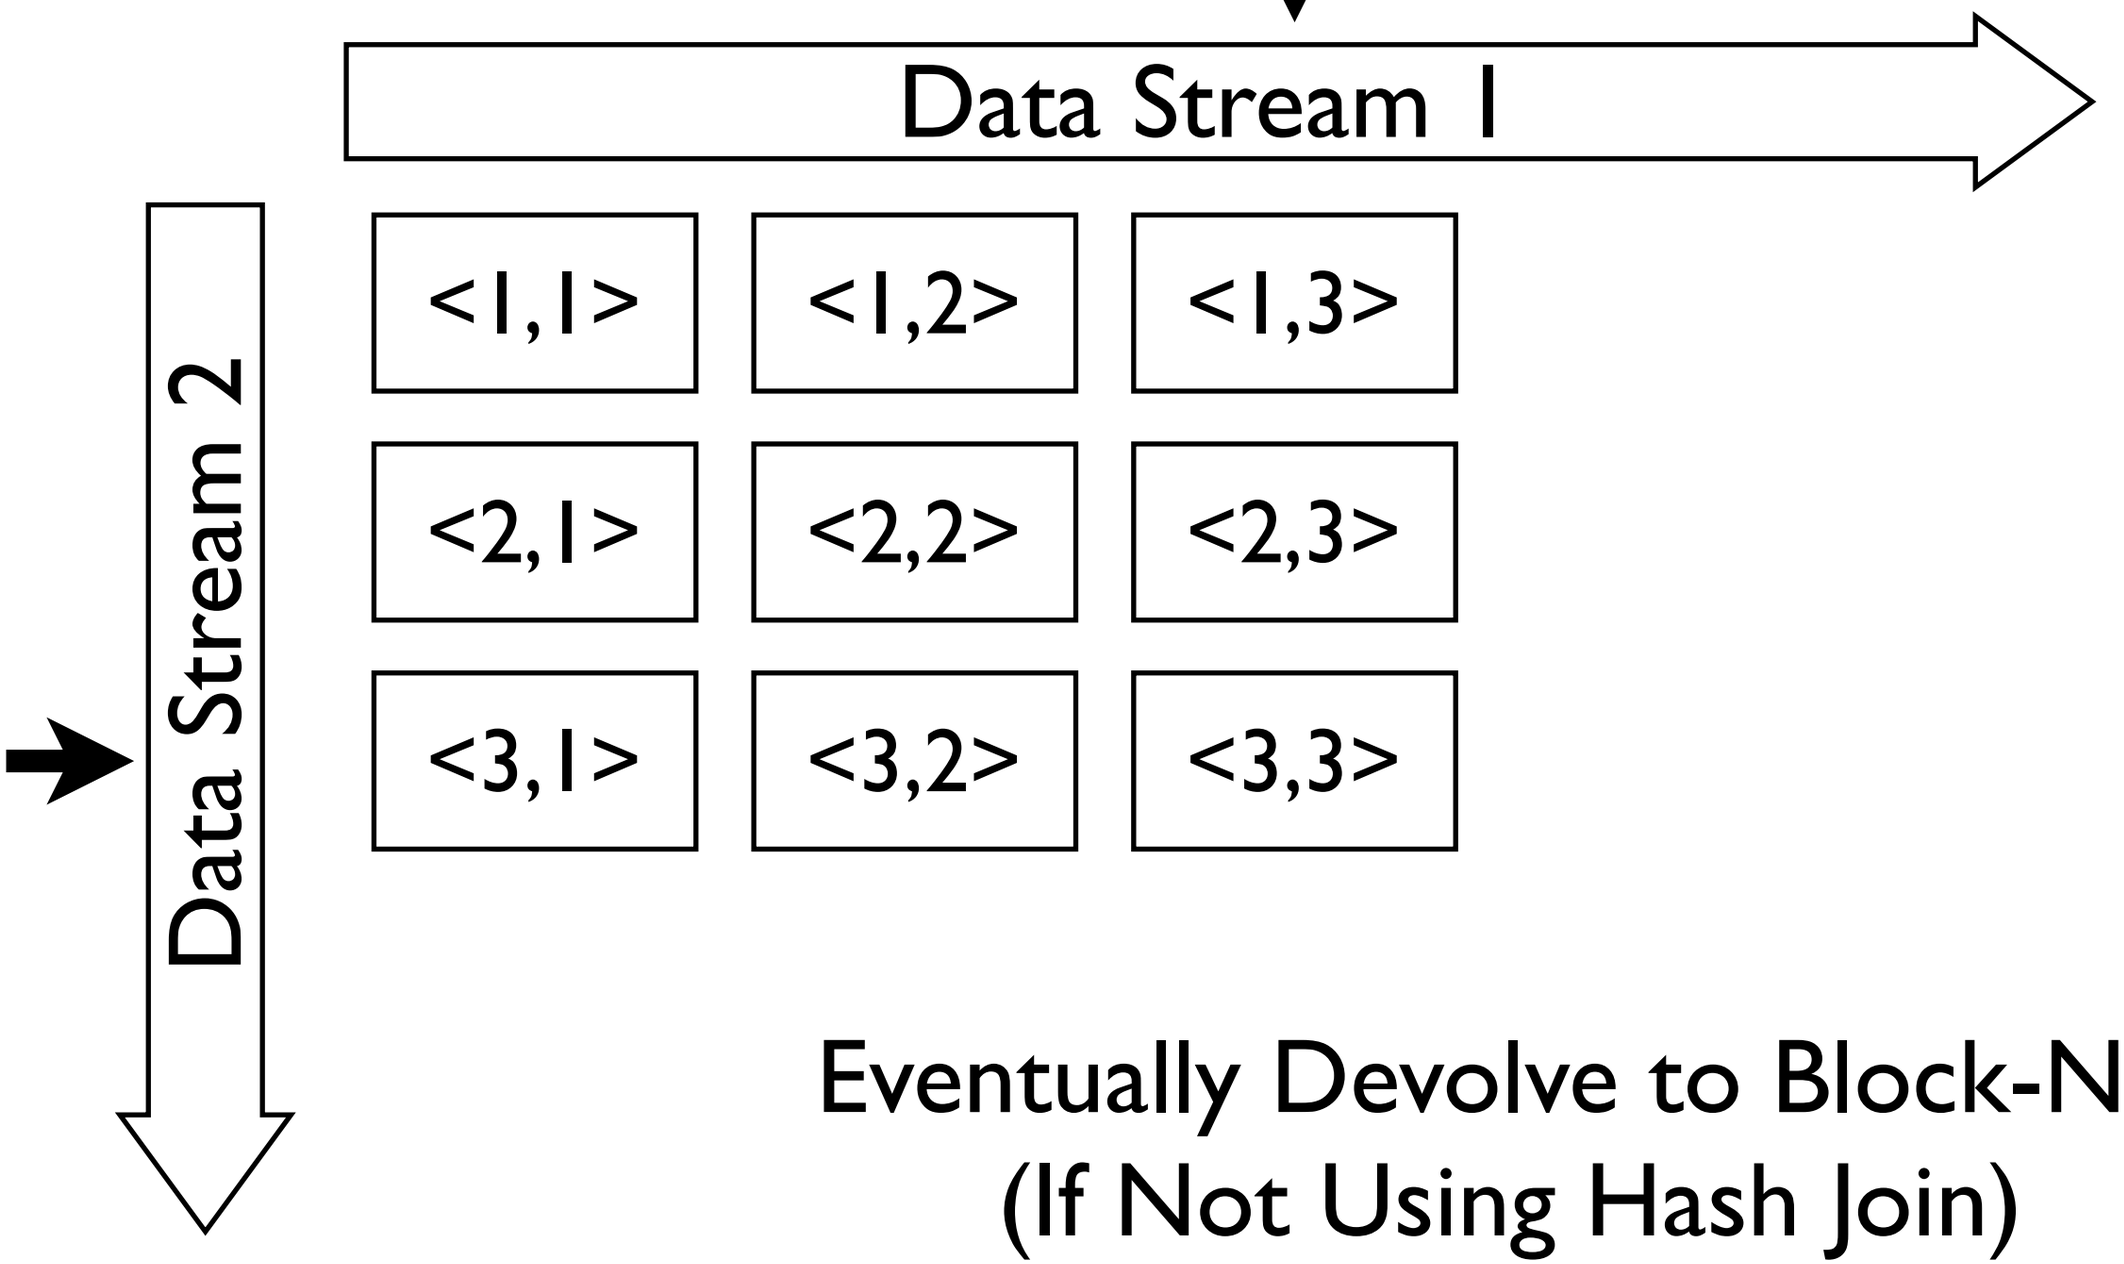

Eventually Devolve to Block-NLJ
(If Not Using Hash Join)

Online Aggregation

- Estimate Aggregate Results by Sampling.
- Continuously Sample and Update Results.
  - Should we sample with replacement?
- Online Aggregation Still an Open Challenge:
  - Random Access Expensive, Even on Flash.
  - Data is Changing Rapidly