

# Stream Processing and Incremental View Maintenance

Supplemental Reading  
(papers posted on Piazza)  
R&G Ch 3.6, Ch 25.8 - 25.10

# Streaming Data

- **Databases:** Repeatedly process different queries on the same data.
- **Streaming/IVM:** Repeatedly process different data with the same queries.

# Streaming Data Challenges


- **Blocking:** Streams are infinite, can't block until they finish.
  - Use Half Joins (Similar to Ripple Joins)
- **Scaling:** Insertion cost must be constant in the size (ripple joins are linear)
  - Use Window Joins
- **Performance:** Streaming data can arrive faster than the machine can handle.
  - Drop tuples as needed to maximize output rate.

# WINDOW

```
SELECT L.state, T.month,  
       AVG(S.sales) OVER W as movavg  
FROM   Sales S, Times T, Locations L  
WHERE  S.timeid = T.timeid  
       AND S.locid = L.locid  
WINDOW W AS (  
    PARTITION BY L.state  
    ORDER BY T.month  
    RANGE BETWEEN INTERVAL '1' MONTH PRECEDING  
           AND INTERVAL '1' MONTH FOLLOWING  
)
```


# Window Joins

$$\mathbf{R} \bowtie \mathbf{S}$$



# Window Joins


$$\mathbf{R} \bowtie \mathbf{S}$$



Do the full join for all tuples in the windows

# Window Joins


$$\mathbf{R} \bowtie \mathbf{S}$$



Slide the window and repeat

# Window Joins

$$\mathbf{R} \bowtie \mathbf{S}$$




**Deleted Results:** Tuples of R That Join With Expiring S Tuple

**New Results:** Tuples of R That Join With New S Tuple



# Window Joins

$$\mathbf{R} \bowtie \mathbf{S}$$



~~**Deleted Results:** Tuples of  
R That Join With Expiring S Tuple~~

**New Results:** Tuples of  
R That Join With New S Tuple

# Streaming Window Joins

- Stream joins produce streams of new join outputs.
- **Problem:** We must support insertions into either stream, in any order.
- **Solution:** Use a Ripple Join
- **Problem:** Joins are still slow
- **Solution:** Maintain an index over the contents of each window.

# Streaming Window Joins

- When a new tuple **t** arrives from R:
  - Use Index on S's Window to Join vs **t**.
  - Insert New Tuple into R's Window.
  - Delete Tuples Expiring From R's Window.

# Streaming Window Joins

- Use Index on  $S$ 's Window to Join vs  $t$ .
  - Hash Index ( $O(\sim 1)$  lookup,  $==$  only)
  - B+Tree Index ( $O(\log(n))$  lookup,  $==$  or  $>$ )
  - No Index ( $O(n)$  lookup, any predicate)

# Streaming Window Joins


- Insert New Tuple into R's Window.
  - Easy. Standard Index Insertion.
- Delete Tuples Expiring From R's Window.
  - Hard(er). Need to Find Expiring Tuples.

# Streaming Window Joins

- Insert New Tuple into R's ~~Window~~ <sup>Index</sup>.
- Easy. Standard Index Insertion.
- Delete Tuples Expiring From R's Window.
- Hard(er). Need to Find Expiring Tuples.


# Invalidation-NLJ

**Easy:** Preserve Order of Tuples in Window



Only ever delete the last tuple(s)


# Invalidation-Tree



**Problem:** Tuples sorted in tree order




# Invalidation-Tree




**Problem:** Tuples sorted in tree order


**Solution:** Maintain an array (like NLJ) in insertion order



# Invalidation-Hash




# Invalidation-Hash




# Invalidation-Hash

**2: Keep Tuples In Each Bucket Sorted**




**1: Keep Sorted List of Tuples in Window  
(used to find hash bucket to delete from)**




# Invalidation-Hash


**2:** Keep Tuples In Each Bucket Sorted



**3:** Even If A Bucket Has  $>1$  Item,  $O(1)$  deletion time for each bucket.



**1:** Keep Sorted List of Tuples in Window  
(used to find hash bucket to delete from)



# Streaming Window Joins

- The algorithms presented are sufficient to process half of a join (hence half join).
- The half join used for joins from R to S is independent of that for joins from S to R.
- Two half joins = A streaming join
- Choice of which half-join to use depends on *relative* insertion rates of R and S.

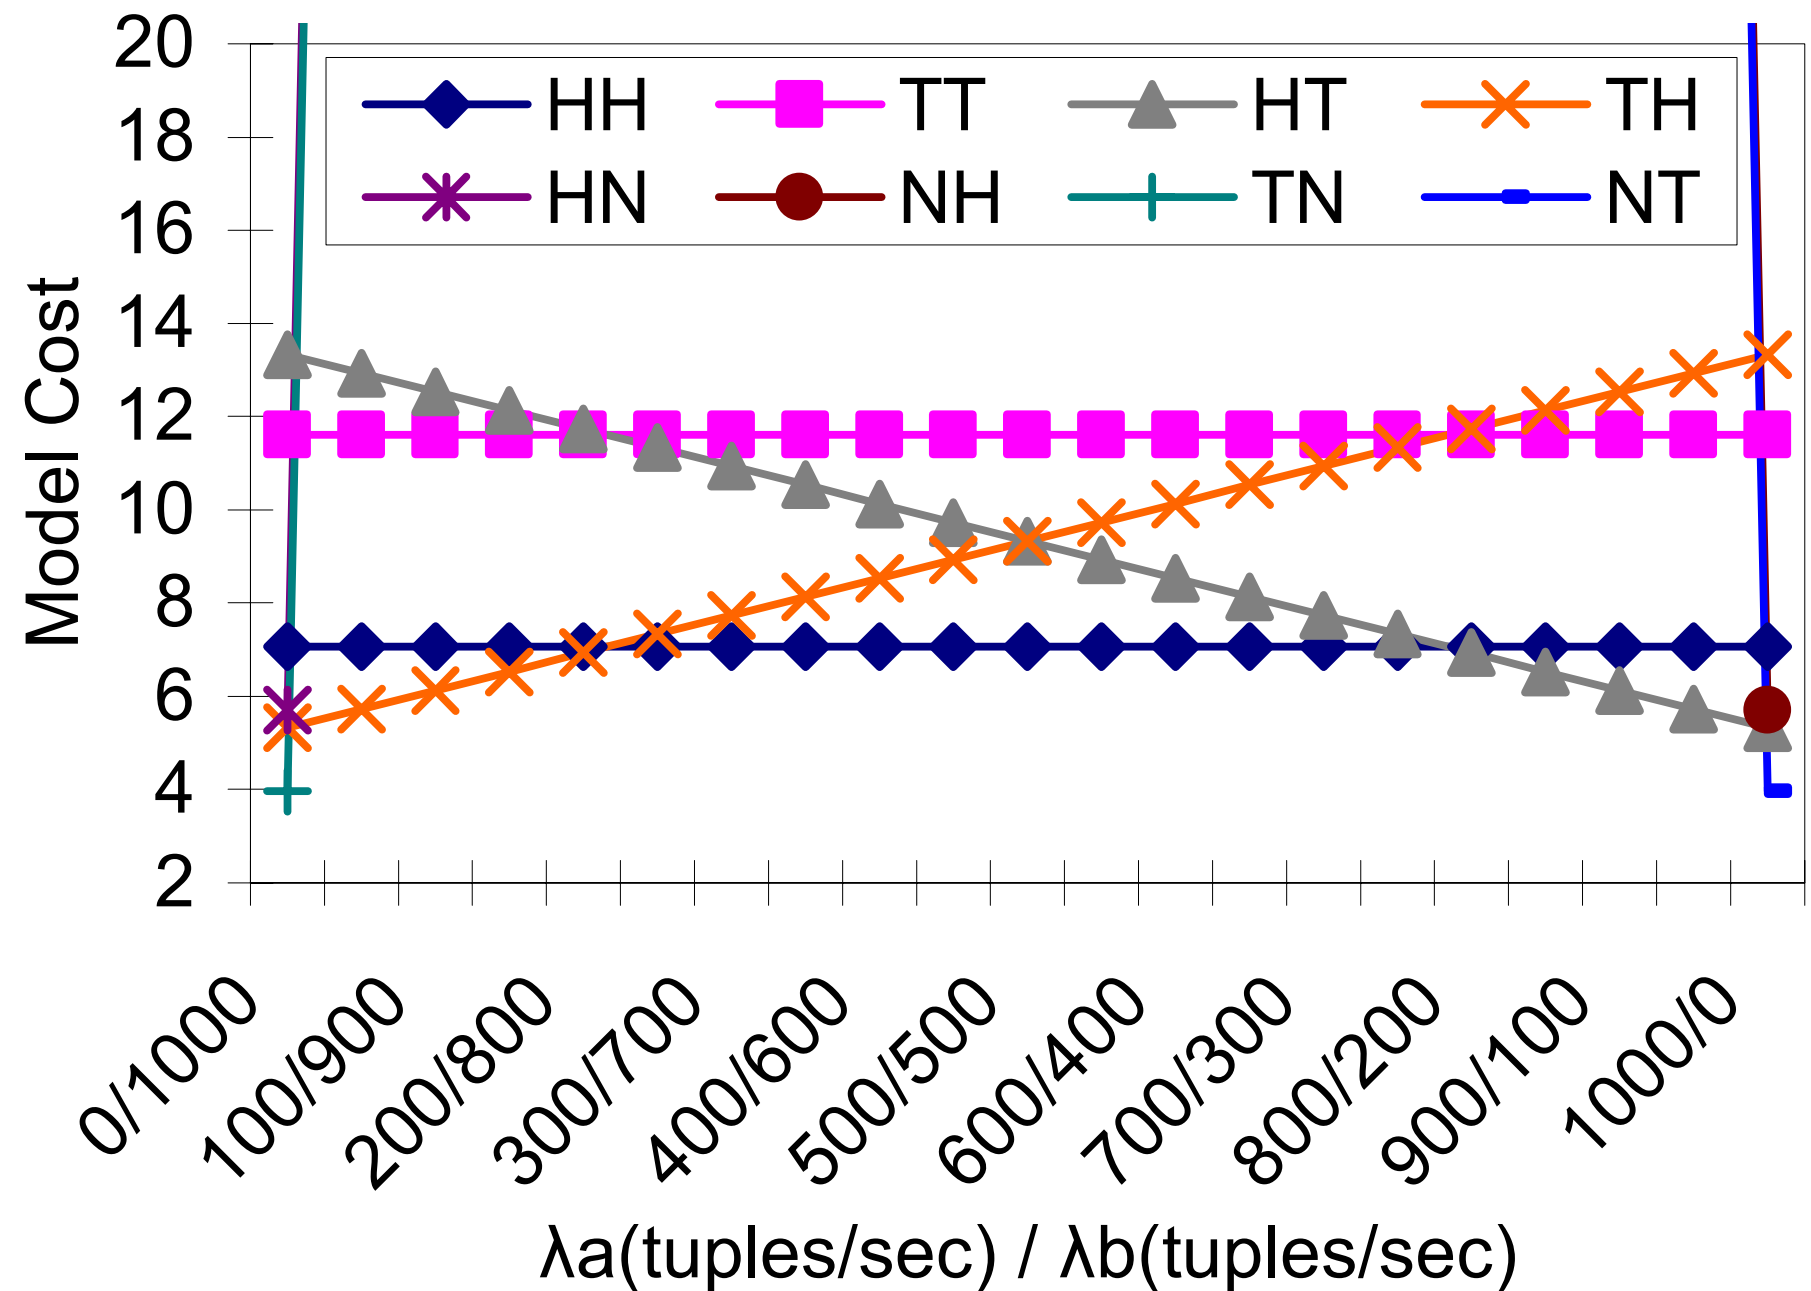
# Streaming Window Joins

- **Recap:** When a new tuple  $\mathbf{t}$  arrives from R:
  - Use Index on S's Window to Join vs  $\mathbf{t}$ .
  - Insert New Tuple into R's Window.
  - Delete Tuples Expiring From R's Window.
- When would you ever use an NLJ?

# Streaming Window Joins

H = Hash  
 N = NLJ  
 T = T-Tree  
 (Similar To B+)

Cost Model  
 In Paper

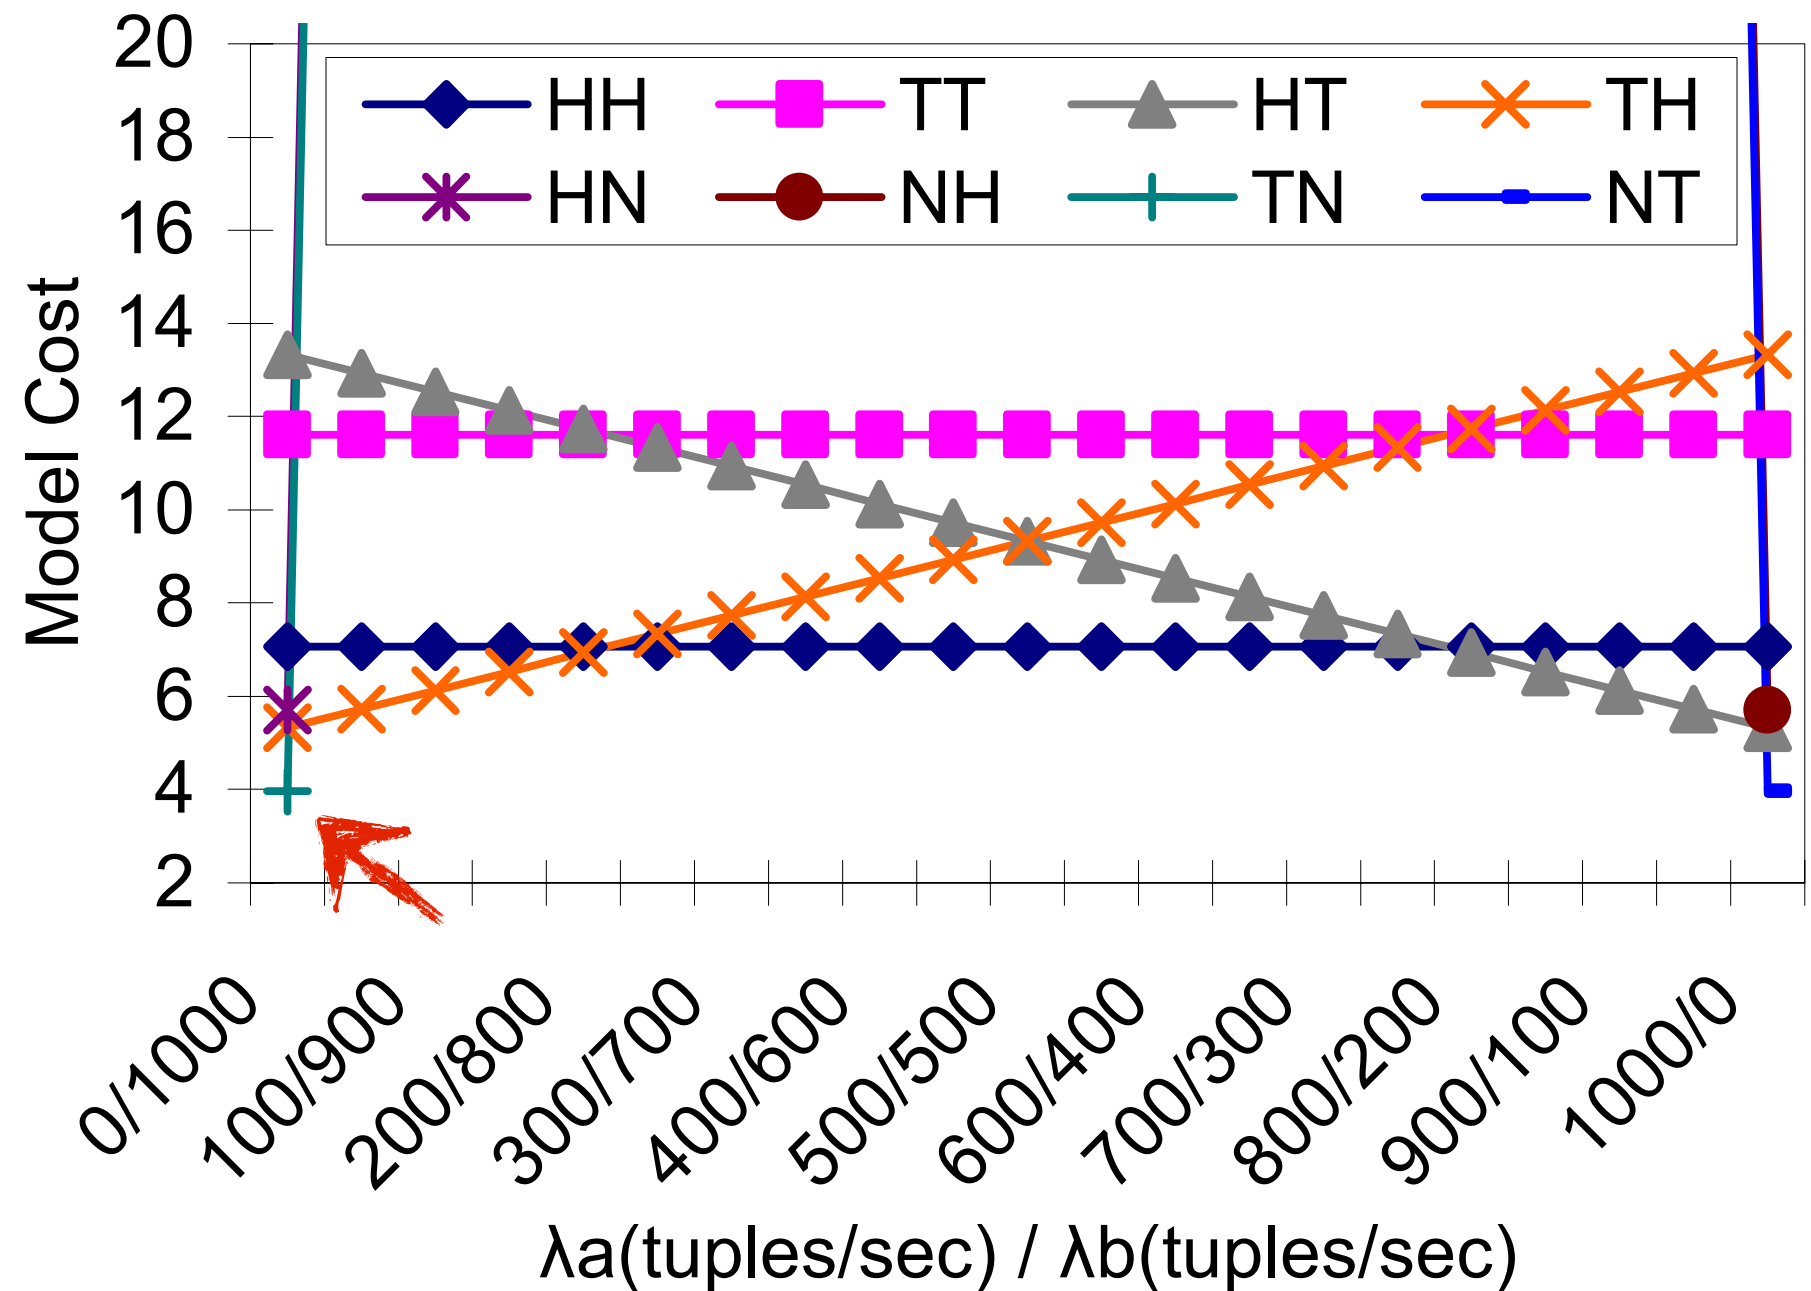




# Streaming Window Joins

H = Hash  
 N = NLJ  
 T = T-Tree  
 (Similar To B+)

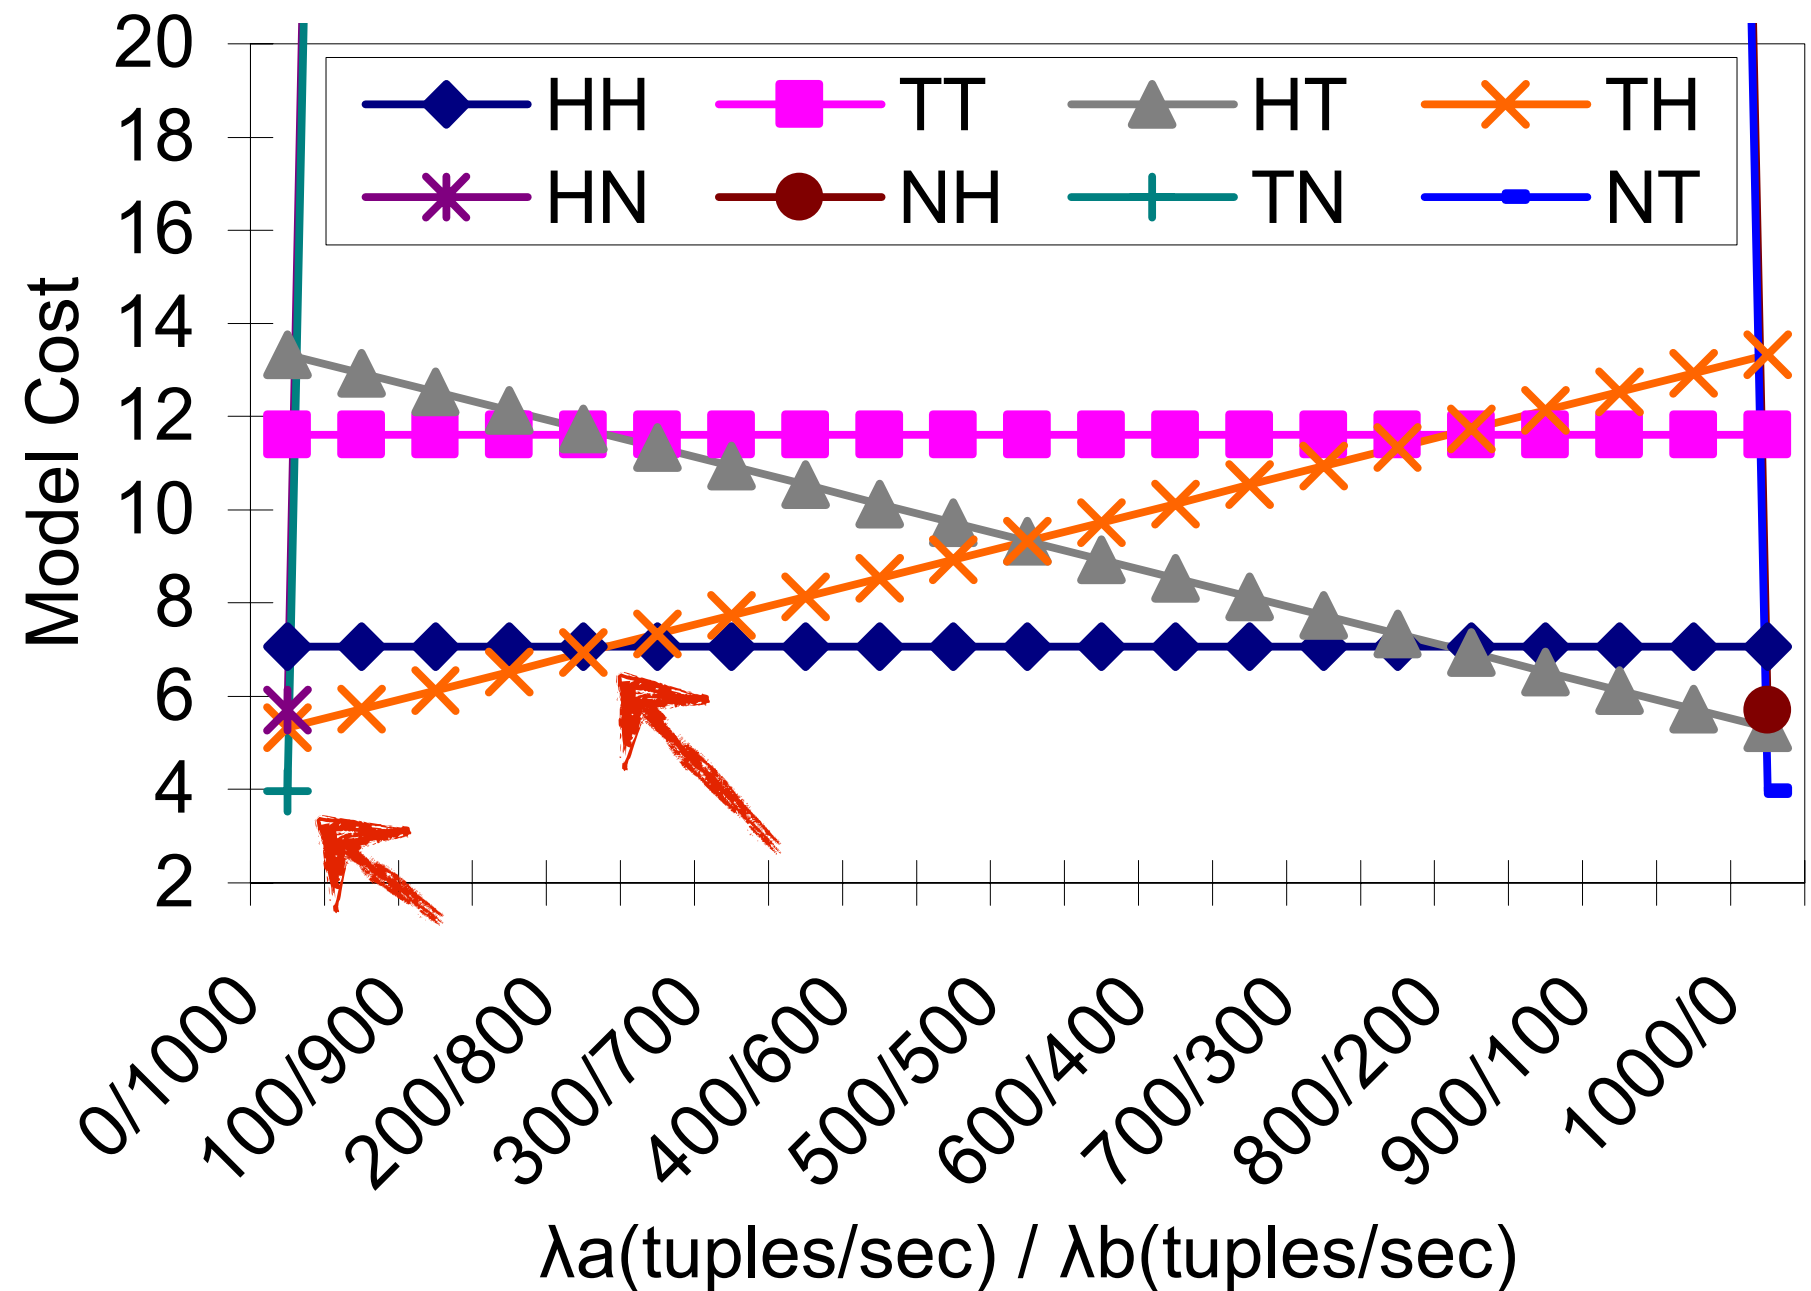
Cost Model  
 In Paper



# Streaming Window Joins

H = Hash  
 N = NLJ  
 T = T-Tree  
 (Similar To B+)

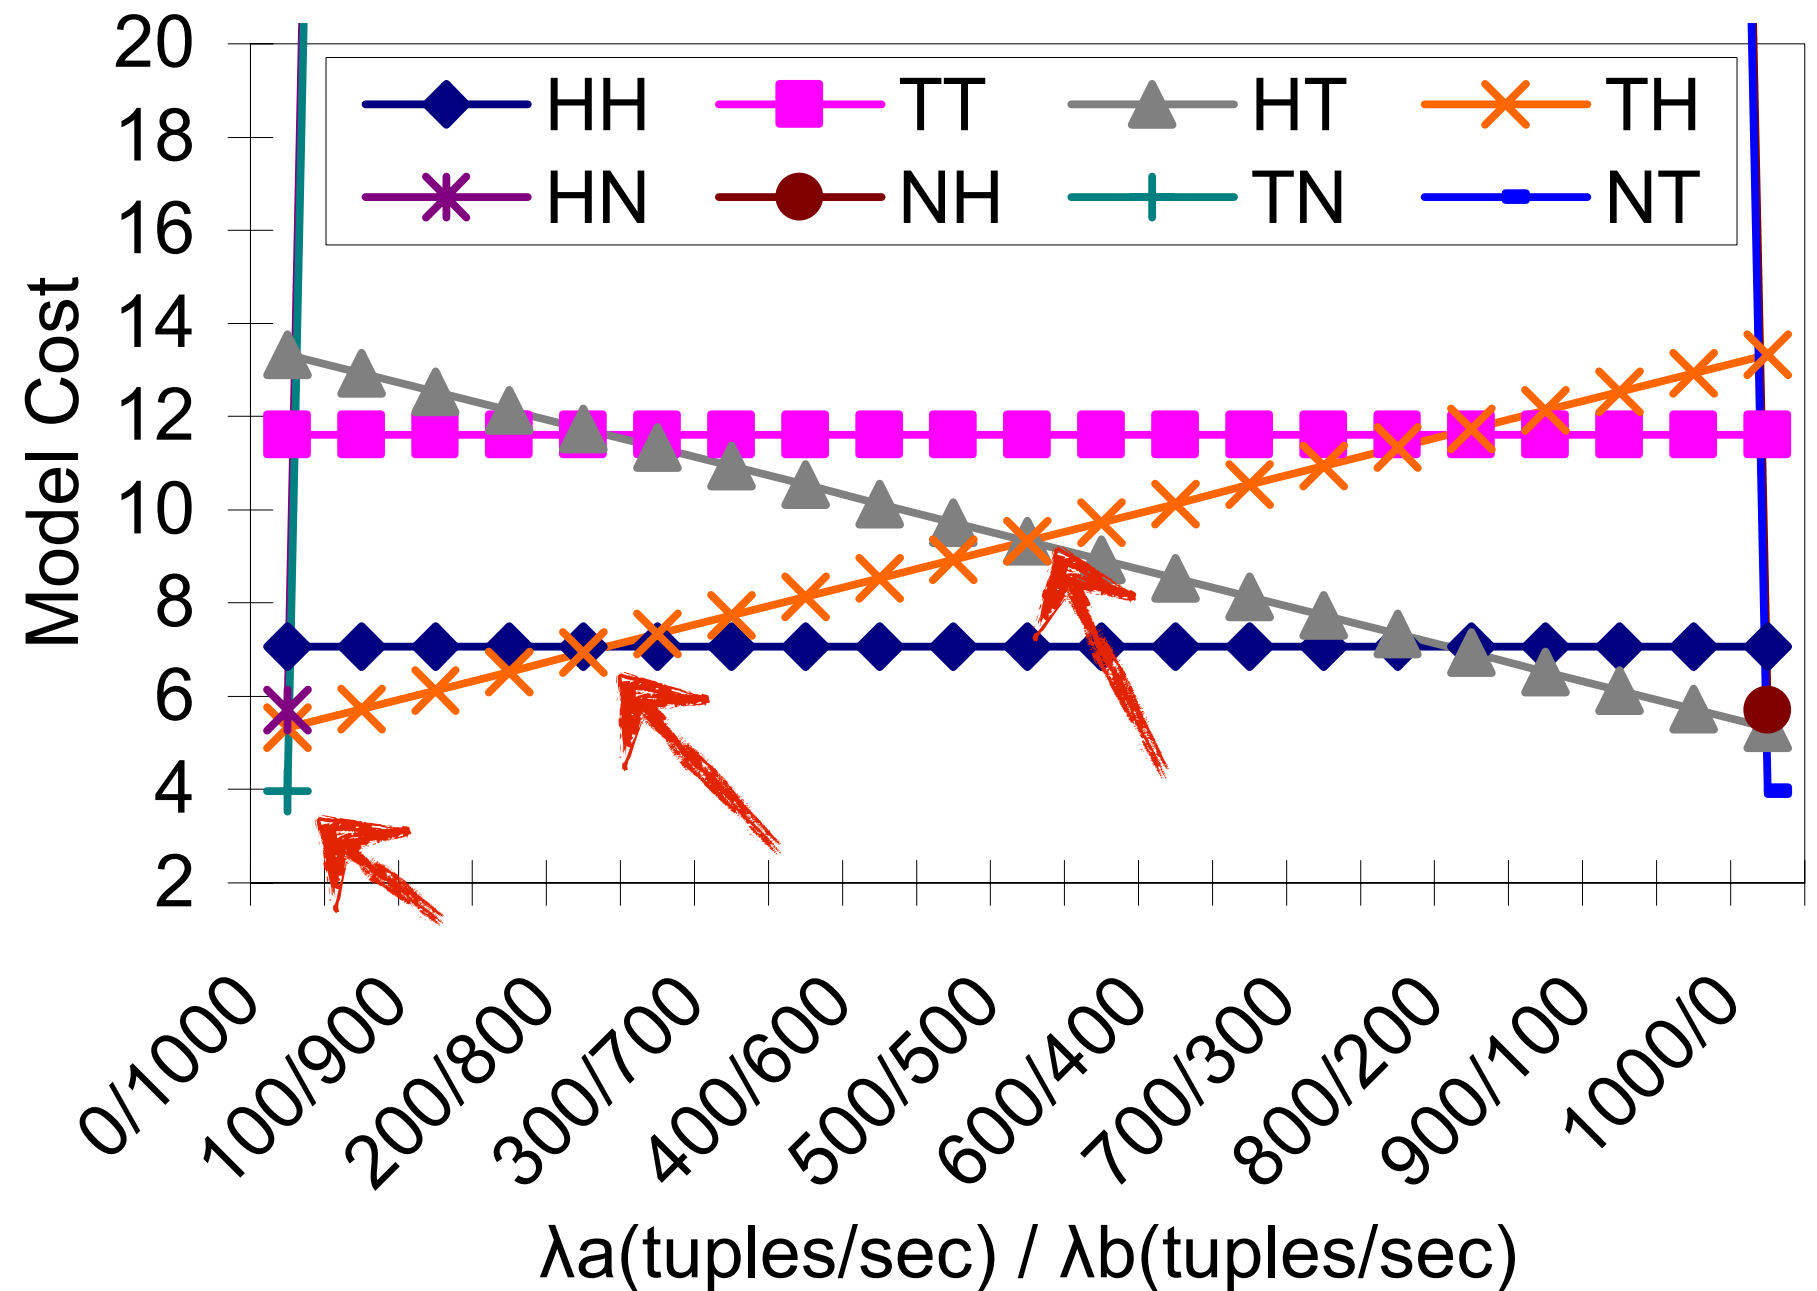
Cost Model  
 In Paper



# Streaming Window Joins

H = Hash  
 N = NLJ  
 T = T-Tree  
 (Similar To B+)

Cost Model  
 In Paper



# Views

```
CREATE VIEW EnterpriseOfficer(oid,name)  
  AS SELECT o.oid, o.name  
        FROM Officer o, Ship s  
        WHERE o.ship = s.sid;
```

# Views

```
CREATE VIEW EnterpriseOfficer(oid,name)  
  AS SELECT o.oid, o.name  
        FROM Officer o, Ship s  
        WHERE o.ship = s.sid;
```

```
SELECT eo.name  
FROM EnterpriseOfficer eo;
```

# Views

```
CREATE VIEW EnterpriseOfficer(oid,name)
  AS SELECT o.oid, o.name
      FROM Officer o, Ship s
      WHERE o.ship = s.sid;
```

```
SELECT eo.name
FROM (SELECT o.oid, o.name
      FROM Officer o, Ship s
      WHERE o.ship = s.sid
    ) eo;
```

# Materialized Views

- **Problem:** View Queries Can be Expensive
- **Solution:** Save the View Query Result as a normal table.
- **Problem:** What if the Data Changes
  - Re-evaluate the full query?
  - **Solution:** Delta Queries

# Delta Queries

- If  $D$  is your Database and  $Q$  is your Query:
  - $Q(D)$  is the result of your query on the database.
- Let's say you make a change  $\Delta D$  (Insert tuple)
  - $Q(D+\Delta D)$  is the new result
- If we have  $Q(D)$ , can we get  $Q(D+\Delta D)$  faster?
  - Analogy to Sum:  $\{34, 29, 10, 15\} + \{12\}$  ( $88+12$ )



# Delta Queries

$\sigma$   
|  
R

# Delta Queries

$\sigma$

|

R

R

$\Delta R$

Original R

Inserted/Deleted  
Tuples of R

# Delta Queries

$\sigma$   
|  
R

Original R

R

$\sigma$   
|  
 $\Delta R$

Inserted/Deleted  
Tuples of R

# Delta Queries

$\pi$

|

R

$\pi$

|

R

$\Delta R$

# Delta Queries

$\pi$   
|  
R

R

$\pi$   
|  
 $\Delta R$

**Problem:** What about Set-Semantics Projection?

# Delta Queries

$\pi$

|

R

$\pi$

|


R

$\Delta R$

**Problem:** What about Set-Semantics Projection?

**Problem:** What about Deletion?

# Delta Queries




$R$

$\Delta R$

$S$


# Delta Queries



**Problem:** What about Deletion?



# Delta Queries



**Problem:** What about Deletion?

# Delta Queries

What about aggregates?