CSE 486/586 Distributed Systems
## Gossiping

Steve Ko
Computer Sciences and Engineering
University at Buffalo

Recall: Passive Replication


- Request Communication: the request is issued to the primary RM and carries a unique request id.
- Coordination: Primary takes requests atomically, in order, checks id (resends response if not new id.)
- Execution: Primary executes & stores the response
- Agreement: If update, primary sends updated state/ result, req-id and response to all backup RMs (1-phase commit enough).
- Response: primary sends result to the front end

Recall: Active Replication

- Request Communication: The request contains a unique identifier and is multicast to all by a reliable totally-ordered multicast.
- Coordination: Group communication ensures that requests are delivered to each RM in the same order (but may be at different physical times!).
- Execution: Each replica executes the request. (Correct replicas return same result since they are running the same program, i.e., they are replicated protocols or replicated state machines)
- Agreement: No agreement phase is needed, because of multicast delivery semantics of requests
- Response: Each replica sends response directly to FE


Eager vs. Lazy

- Eager replication, e.g., B-multicast, R-multicast, etc. (previously in the course)
  - Multicast request to all RMs immediately in active replication
  - Multicast results to all RMs immediately in passive replication
- Alternative: Lazy replication
  - Allow replicas to converge eventually and lazily
  - Propagate updates and queries lazily, e.g., when network bandwidth available
  - FEs need to wait for reply from only one RM
  - Allow other RMs to be disconnected/unavailable
  - May provide weaker consistency than sequential consistency, but improves performance
- Lazy replication can be provided by using the gossiping

Revisiting Multicast

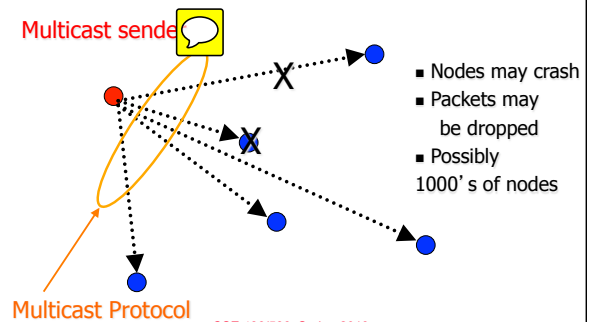Node with a piece of information to be communicated to everyone

Distributed Group of "Nodes"= Processes at Internet-based hosts

---

Fault-Tolerance and Scalability

Multicast sender

- Nodes may crash
- Packets may be dropped
- Possibly 1000's of nodes

Multicast Protocol

C

B-Multicast

UDP/TCP packets

- Simplest implementation
- Problems?

7

R-Multicast

UDP/TCP packets

- Stronger guarantees
- Overhead is quadratic in N

8

Any Other?

- E.g., tree-based multicast

UDP/TCP packets

- e.g., IPmulticast, SRM RMTP, TRAM,TMTP
- Tree setup and maintenance
- Problems?
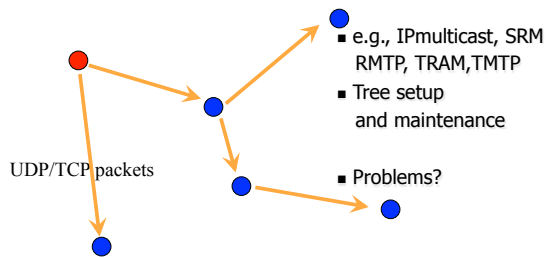
9

Another Approach

Multicast sender

10

Another Approach

Periodically, transmit to b random targets

→ Gossip messages (UDP)

11

Another Approach

Other nodes do same after receiving multicast

→ Gossip messages (UDP)

12

C

2

Another Approach

"Gossip" (or "Epidemic") Multicast

🔴 Infected

Protocol *rounds* (local clock)
b random targets per round

Gossip Message (UDP)

🔵 Uninfected

Properties

- Lightweight
- Quick spread
- Highly fault-tolerant
- Analysis from old mathematical branch of *Epidemiology* [Bailey 75]
- Parameters *c,b*:
  - *c* for determining rounds: (*c*log(n)), b*: # of nodes to contact
  - Can be small numbers independent of *n, e.g., c=2; b=2;*
- Within *c*log(n)* rounds, [low latency]
  - all but $\dfrac{1}{n^{cb-2}}$ of nodes receive the multicast [reliability]

  - each node has transmitted no more than *c*b*log(n)* gossip messages [lightweight]

Fault-Tolerance

- Packet loss
  - 50% packet loss: analyze with *b* replaced with *b/2*
  - To achieve same reliability as 0% packet loss, takes twice as many rounds
- Node failure
  - 50% of nodes fail: analyze with *n* replaced with *n/2* and *b* replaced with *b/2*
  - Same as above

Fault-Tolerance

- With failures, is it possible that the epidemic might die out quickly?
- Possible, but improbable:
  - Once a few nodes are infected, with high probability, the epidemic will not die out
  - So the analysis we saw in the previous slides is actually behavior *with high probability*
  - [Galey and Dani 98]
- The same applicable to:
  - Rumors
  - Infectious diseases
  - A worm such as Blaster
- Some implementations
  - Amazon Web Services EC2/S3 (rumored)
 - Usenet NNTP (Network News Transport Protocol)

Gossiping Architecture

- The RMs exchange "gossip" messages
  - Periodically and amongst each other.
  - Gossip messages convey updates they have each received from clients, and serve to achieve convergence of all RMs.
- Objective: provisioning of highly available service. Guarantee:
  - Each client obtains a consistent service over time: in response to a query, an RM may have to wait until it receives "required" updates from other RMs. The RM then provides client with data that at least reflects the updates that the client has observed so far.
  - Relaxed consistency among replicas: RMs may be inconsistent at any given point of time. Yet all RMs <u>eventually</u> receive all updates and they apply updates with ordering guarantees. Can be used to provide sequential consistency.

Gossip Architecture

Service

RM

RM — gossip — RM

Query, prev | Val, new | Update, prev | Update id

FE | FE

Query | Val | Update

Clients

Using Gossip for Failure Detection: Gossip-style Heartbeating

pi

☺ Using gossip to spread heartbeats gives better accuracy

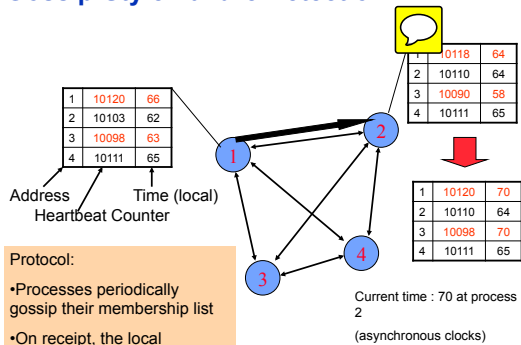**All-to-all heartbeating**

- Each process sends out heartbeats to every other process
- Con: Slow process/link causes false positives

Gossip-Style Failure Detection

1	10120	66
2	10103	62
3	10098	63
4	10111	65

Address   Time (local)

Heartbeat Counter

1	0118	64
2	10110	64
3	10090	58
4	10111	65

1	10120	70
2	10110	64
3	10098	70
4	10111	65

Current time : 70 at process 2

(asynchronous clocks)

Protocol:

- Processes periodically gossip their membership list
- On receipt, the local membership list is updated


Gossip-Style Failure Detection

- If the heartbeat has not increased for more than T_{fail} seconds (according to local time), the member is considered failed
- But don't delete it right away
- Wait another $T_{cleanup}$ seconds, then delete the member from the list

---

Gossip-Style Failure Detection

- What if an entry pointing to a failed process is deleted right after T_{fail} seconds?

1	10120	66
2	10103	62
3	10098	55
4	10111	65

1	10120	66
2	10110	64
3	10098	55
4	10111	65

Current time : 75 at process 2

- Fix: remember for another T_{fail}
- Ignore gossips for failed members
  – Don't include failed members in gossip messages

Summary

- Eager replication vs. lazy replication
  – Lazy replication propagates updates in the background
- Gossiping
  – One strategy for lazy replication
  – High-level of fault-tolerance & quick spread
- Another use case for gossiping
  – Failure detection

Acknowledgements

- These slides contain material developed and copyrighted by Indranil Gupta (UIUC).

25

C

5